

alzheimer-detection-part-3

October 28, 2025

```
[1]: import os
import matplotlib.pyplot as plt
from PIL import Image
import numpy as np

# Keep your original directories and subfolders
test_dir = "/kaggle/input/alzheimer-disease/test-20251010T094612Z-1-001/test"
train_dir = "/kaggle/input/alzheimer-disease/train-20251010T183510Z-1-001/train"
subfolders = ["AD", "CN", "MCI"]

# Function to show 3 images from each subfolder (unchanged signature)
def show_images_from_dir(path, n=3):
    # List only image files (same behavior; expanded to be robust if needed)
    files = [f for f in os.listdir(path) if f.lower().endswith('.png', '.jpg', '.jpeg', '.bmp', '.tif', '.tiff')]]
    files = files[:n] # Take only the first n images

    plt.figure(figsize=(15, 5))
    for i, file in enumerate(files):
        img_path = os.path.join(path, file)

        # Open with PIL
        img = Image.open(img_path)

        # --- CRITICAL FIX ---
        # Force grayscale to ensure single-channel display (no colorization)
        # Even if the source is already grayscale, this guarantees mode 'L'
        img = img.convert('L')

        # Convert to numpy for safe imshow with explicit bounds
        arr = np.asarray(img)

        # Display as true grayscale with fixed intensity bounds
        plt.subplot(1, len(files), i + 1)
        plt.imshow(arr, cmap='gray', vmin=0, vmax=255)
        plt.axis("off")
        plt.title(file[:10])
```

```

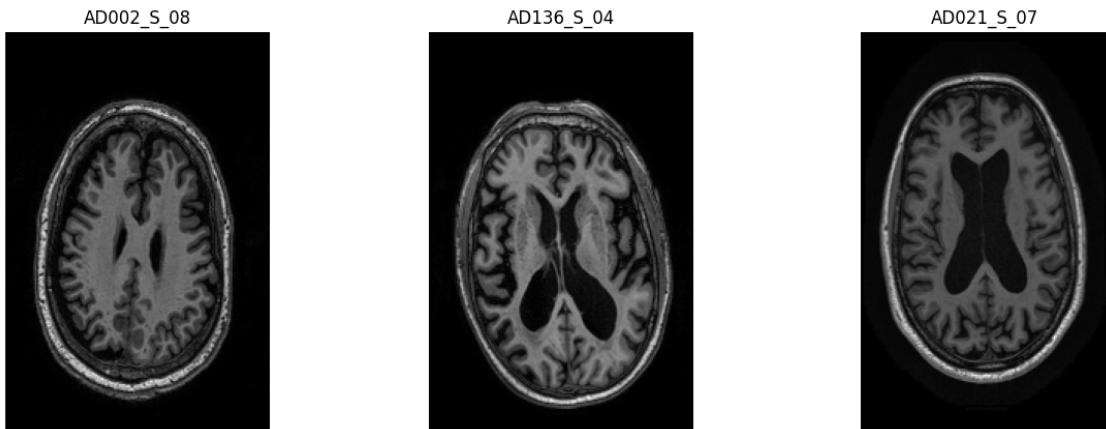
plt.show()

# Show images for Test set
for subfolder in subfolders:
    print(f"{subfolder} (Test)")
    show_images_from_dir(os.path.join(test_dir, subfolder))

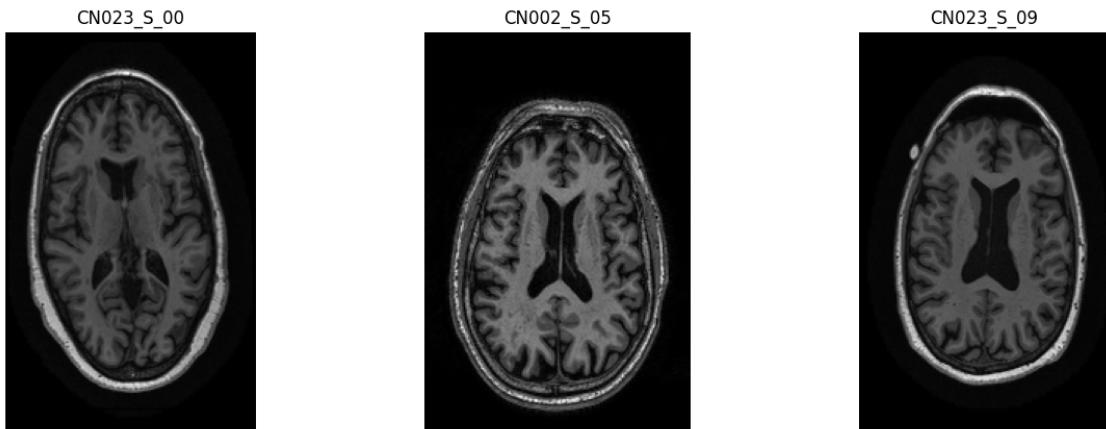
# Show images for Train set
for subfolder in subfolders:
    print(f"{subfolder} (Train)")
    show_images_from_dir(os.path.join(train_dir, subfolder))

```

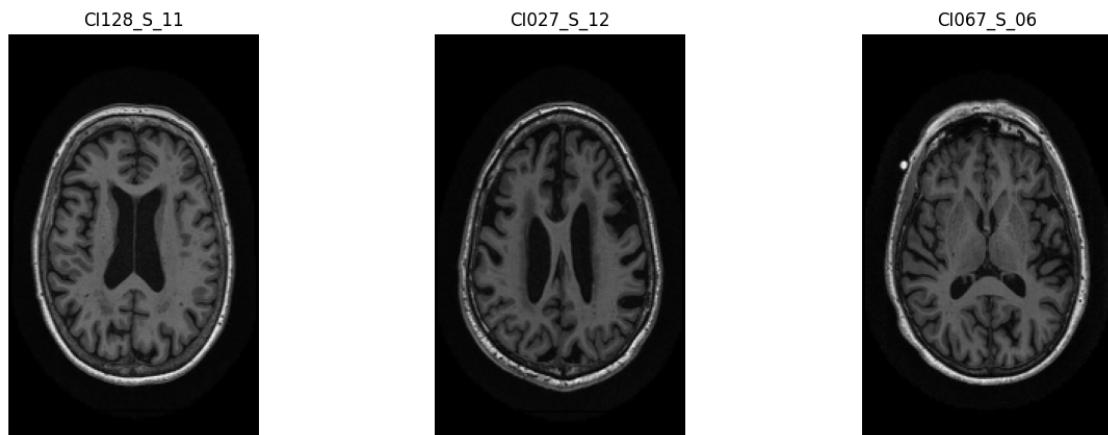
AD (Test)



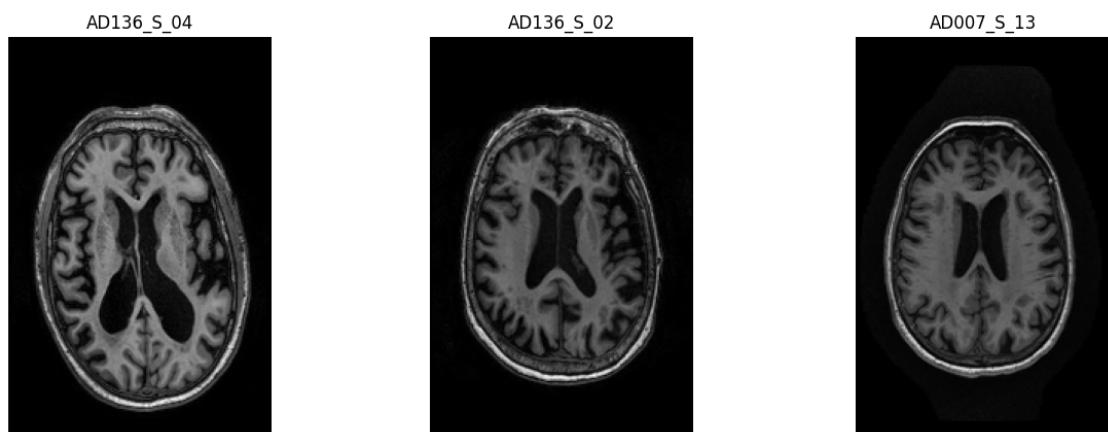
CN (Test)



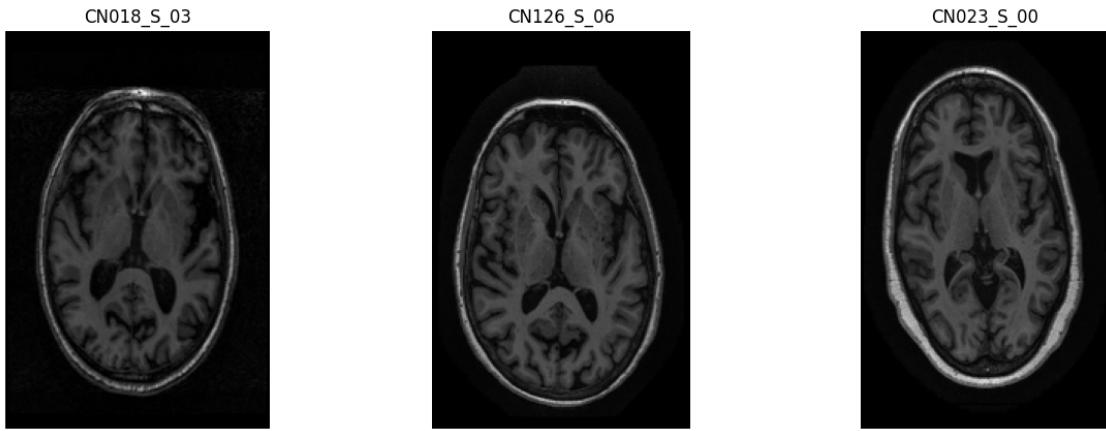
MCI (Test)



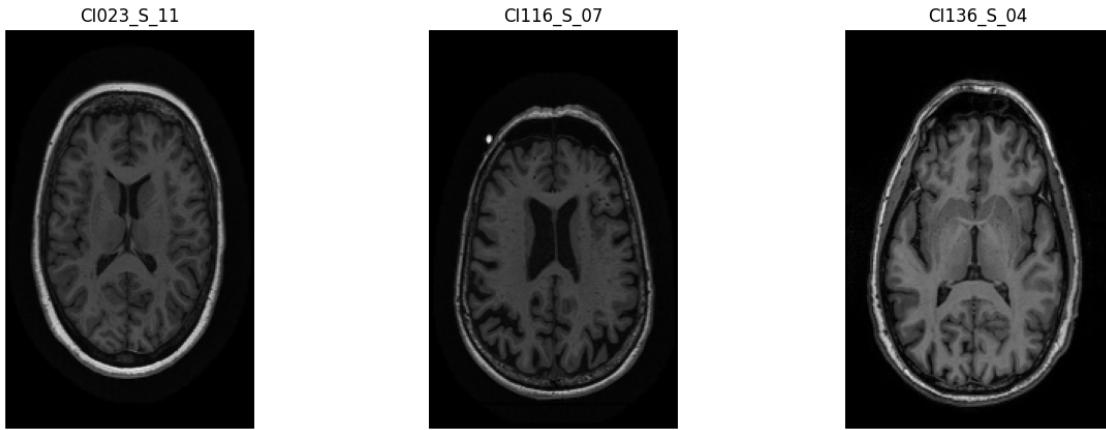
AD (Train)



CN (Train)



MCI (Train)



```
[2]: import os
import sys
import numpy as np
from PIL import Image, UnidentifiedImageError
import matplotlib.pyplot as plt

# SciPy for morphology / connected components (Kaggle preinstalled)
from scipy.ndimage import (
    gaussian_filter,
    binary_opening,
    binary_closing,
    binary_fill_holes,
    label
)
```

```

# Try skimage for CLAHE; fall back gracefully if missing
try:
    from skimage.exposure import equalize_adapthist
    _HAS_SKIMAGE = True
except Exception:
    _HAS_SKIMAGE = False

# GPU via PyTorch (for homomorphic filtering)

def _ensure(pkg):
    import importlib
    try:
        importlib.import_module(pkg)
    except Exception:
        import subprocess
        subprocess.check_call([sys.executable, "-m", "pip", "install", pkg, ↴"--quiet"])

_ensure("torch")
import torch
import torch.nn.functional as F

def torch_device():
    return torch.device("cuda" if torch.cuda.is_available() else "cpu")

def gaussian_kernel_2d(sigma: float, radius_factor: float = 3.0, device=None, ↴
    dtype=torch.float32):
    """Create a 2D Gaussian kernel tensor for conv2d (normalized)."""
    device = device or torch_device()
    rad = max(1, int(radius_factor * sigma))
    xs = torch.arange(-rad, rad + 1, device=device, dtype=dtype)
    g1 = torch.exp(-0.5 * (xs / sigma) ** 2)
    g1 = g1 / g1.sum()
    g2 = g1[:, None] @ g1[None, :]
    g2 = g2 / g2.sum()
    return g2 # (K, K)

def homomorphic_filter_gpu_u8(arr_u8: np.ndarray, sigma: float = 50.0) -> np. ↴
ndarray:
    """
    GPU homomorphic filtering (log + low-pass via conv2d + exp + robust ↴
    rescale).

    Input uint8 [0..255], output uint8 [0..255].
    """
    dev = torch_device()
    # [B=1,C=1,H,W] float32

```

```

    img = torch.from_numpy(arr_u8.astype(np.float32)).to(dev) + 1.0 # avoid underflow at log(0)
    img = img[None, None, :, :] # NCHW
    loga = torch.log(img)
    # 2D Gaussian low-pass via conv2d (reflect padding to avoid border artifacts)
    k = gaussian_kernel_2d(sigma=sigma, device=dev)
    k = k[None, None, :, :] # (out_c,in_c,H,W)
    pad = k.shape[-1] // 2
    low = F.conv2d(F.pad(loga, (pad, pad, pad, pad), mode="reflect"), k)
    high = loga - low
    corr = torch.exp(high) - 1.0 # back to linear domain

    # Robust percentile rescale to uint8
    a = corr.squeeze(0).squeeze(0) # HxW
    # Compute percentiles on CPU for simplicity
    a_np = a.detach().cpu().numpy().astype(np.float32)
    lo, hi = np.percentile(a_np, [0.5, 99.5])
    if hi - lo < 1e-6:
        a_np = (a_np - a_np.min()) / (a_np.ptp() + 1e-8)
    else:
        a_np = np.clip((a_np - lo) / (hi - lo), 0.0, 1.0)
    out_u8 = (a_np * 255.0 + 0.5).astype(np.uint8)
    return out_u8

test_dir = "/kaggle/input/alzheimer-disease/test-20251010T094612Z-1-001/test"
train_dir = "/kaggle/input/alzheimer-disease/train-20251010T183510Z-1-001/train"
subfolders = ["AD", "CN", "MCI"]

# Output (mirrors the structure)

preprocessed_root = "/kaggle/working/alzheimer-preprocessed"
preprocessed_test = os.path.join(preprocessed_root, "test")
preprocessed_train = os.path.join(preprocessed_root, "train")

IMG_EXT = (".png", ".jpg", ".jpeg", ".bmp", ".tif", ".tiff")

def ensure_dir(p: str):
    os.makedirs(p, exist_ok=True)

# Utilities

def load_gray(path: str) -> np.ndarray:
    """Load as uint8 grayscale [0..255]."""

```

```

    return np.asarray(Image.open(path).convert("L"), dtype=np.uint8)

def save_gray(arr_u8: np.ndarray, path: str):
    Image.fromarray(arr_u8, mode="L").save(path)

def percentile_rescale_u8(arr: np.ndarray, p_low=1.0, p_high=99.0) -> np.
    ndarray:
    """Robust contrast stretching to uint8."""
    a = arr.astype(np.float32)
    lo, hi = np.percentile(a, [p_low, p_high])
    if hi - lo < 1e-6:
        a = (a - a.min()) / (a.ptp() + 1e-8)
    else:
        a = np.clip((a - lo) / (hi - lo), 0.0, 1.0)
    return (a * 255.0 + 0.5).astype(np.uint8)

def otsu_threshold_u8(arr_u8: np.ndarray) -> int:
    """Pure NumPy Otsu threshold (returns 0..255)."""
    hist = np.bincount(arr_u8.ravel(), minlength=256).astype(np.float64)
    prob = hist / (arr_u8.size + 1e-12)
    omega = np.cumsum(prob)
    mu = np.cumsum(prob * np.arange(256))
    mu_t = mu[-1]
    sigma_b2 = (mu_t * omega - mu) ** 2 / (omega * (1 - omega) + 1e-12)
    sigma_b2[~np.isfinite(sigma_b2)] = 0.0
    return int(np.argmax(sigma_b2))

def largest_cc(mask: np.ndarray) -> np.ndarray:
    """Keep only largest connected component of a binary mask."""
    lbl, n = label(mask.astype(np.uint8))
    if n <= 1:
        return mask.astype(bool)
    counts = np.bincount(lbl.ravel())
    counts[0] = 0 # background
    keep = counts.argmax()
    return (lbl == keep)

def clahe_u8(arr_u8: np.ndarray) -> np.ndarray:
    """CLAHE if skimage is available; otherwise identity."""
    if _HAS_SKIMAGE:
        arr01 = arr_u8.astype(np.float32) / 255.0
        out01 = equalize_adapthist(arr01, clip_limit=0.01)
        return (np.clip(out01, 0.0, 1.0) * 255.0 + 0.5).astype(np.uint8)
    else:
        return arr_u8

```

```

# Core: one-image preprocessing

def preprocess_single_image(
    arr_u8: np.ndarray,
    do_skull_strip: bool = True,
    sigma_homomorphic: float = 50.0,
    min_area: int = 500
):
    """
    Steps:
        1) Robust pre-stretch (percentile)
        2) GPU homomorphic filter (PyTorch)
        3) (Optional) 2D skull/background stripping via Otsu + morphology + LCC + hole fill
        4) Gentle CLAHE (if available), else percentile rescale
    Returns processed uint8 image, plus an optional mask (uint8).
    """

    # 1) Robust stretch
    a1 = percentile_rescale_u8(arr_u8, 1.0, 99.0)

    # 2) Homomorphic filtering on GPU
    a2 = homomorphic_filter_gpu_u8(a1, sigma=sigma_homomorphic)

    brain_mask = None
    if do_skull_strip:
        # 3) Otsu thresholding (foreground bright)
        th = otsu_threshold_u8(a2)
        mask = (a2 >= th).astype(np.uint8)

        # Morphological clean-up
        mask = binary_opening(mask, structure=np.ones((3,3), dtype=np.uint8))
        mask = binary_closing(mask, structure=np.ones((5,5), dtype=np.uint8))
        mask = binary_fill_holes(mask)
        mask = largest_cc(mask)

        # Remove tiny masks (safety)
        if mask.sum() < min_area:
            mask = np.ones_like(mask, dtype=bool) # fallback: keep as-is

        a3 = (a2 * mask).astype(np.uint8)
        brain_mask = (mask.astype(np.uint8) * 255)
    else:
        a3 = a2

    # 4) CLAHE or robust stretch
    a4 = clahe_u8(a3)
    a4 = percentile_rescale_u8(a4, 0.5, 99.5)

```

```

    return a4, brain_mask

# Dataset-level processing & visualization

def preprocess_dataset(src_root: str, dst_root: str, n_preview: int = 3,
                      ↪do_skull_strip=True):
    """
    Applies the pipeline to all images under src_root/{AD/CN/MCI}
    and writes to dst_root/{AD/CN/MCI}. Silent on per-file issues.
    """
    ensure_dir(dst_root)
    summary = {}

    # Report GPU/CPU
    dev = torch_device()
    print(f"Device for homomorphic filtering: {dev}")

    for cls in subfolders:
        src_cls = os.path.join(src_root, cls)
        dst_cls = os.path.join(dst_root, cls)
        ensure_dir(dst_cls)

        processed = skipped = 0
        if not os.path.isdir(src_cls):
            summary[cls] = (0, 0)
            continue

        files = sorted([f for f in os.listdir(src_cls) if f.lower().
        ↪endswith(IMG_EXT)])

        # Process & save
        for fname in files:
            spath = os.path.join(src_cls, fname)
            dpath = os.path.join(dst_cls, fname)
            try:
                arr = load_gray(spath)
                proc, _ = preprocess_single_image(arr, ↪
                    ↪do_skull_strip=do_skull_strip)
                save_gray(proc, dpath)
                processed += 1
            except (UnidentifiedImageError, OSError, RuntimeError, ValueError):
                skipped += 1
                continue

        summary[cls] = (processed, skipped)

```

```

# Preview few examples
preview = files[:n_preview]
if preview:
    fig, axs = plt.subplots(len(preview), 3, figsize=(10, 3*len(preview)))
    if len(preview) == 1:
        axs = np.expand_dims(axs, 0)
    for i, fname in enumerate(preview):
        sp = os.path.join(src_cls, fname)
        dp = os.path.join(dst_cls, fname)
        try:
            orig = load_gray(sp)
            proc = load_gray(dp)
            diff = np.abs(proc.astype(np.int16) - orig.astype(np.int16)).astype(np.uint8)

            axs[i, 0].imshow(orig, cmap="gray", vmin=0, vmax=255); axs[i, 0].axis("off"); axs[i, 0].set_title(f"{cls}: Original")
            axs[i, 1].imshow(proc, cmap="gray", vmin=0, vmax=255); axs[i, 1].axis("off"); axs[i, 1].set_title("Processed")
            axs[i, 2].imshow(diff, cmap="gray", vmin=0, vmax=255); axs[i, 2].axis("off"); axs[i, 2].set_title("|Diff|")
        except Exception:
            continue
    plt.tight_layout()
    plt.show()

return summary

print("== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TEST ==")
sum_test = preprocess_dataset(test_dir, preprocessed_test, n_preview=3, do_skull_strip=True)

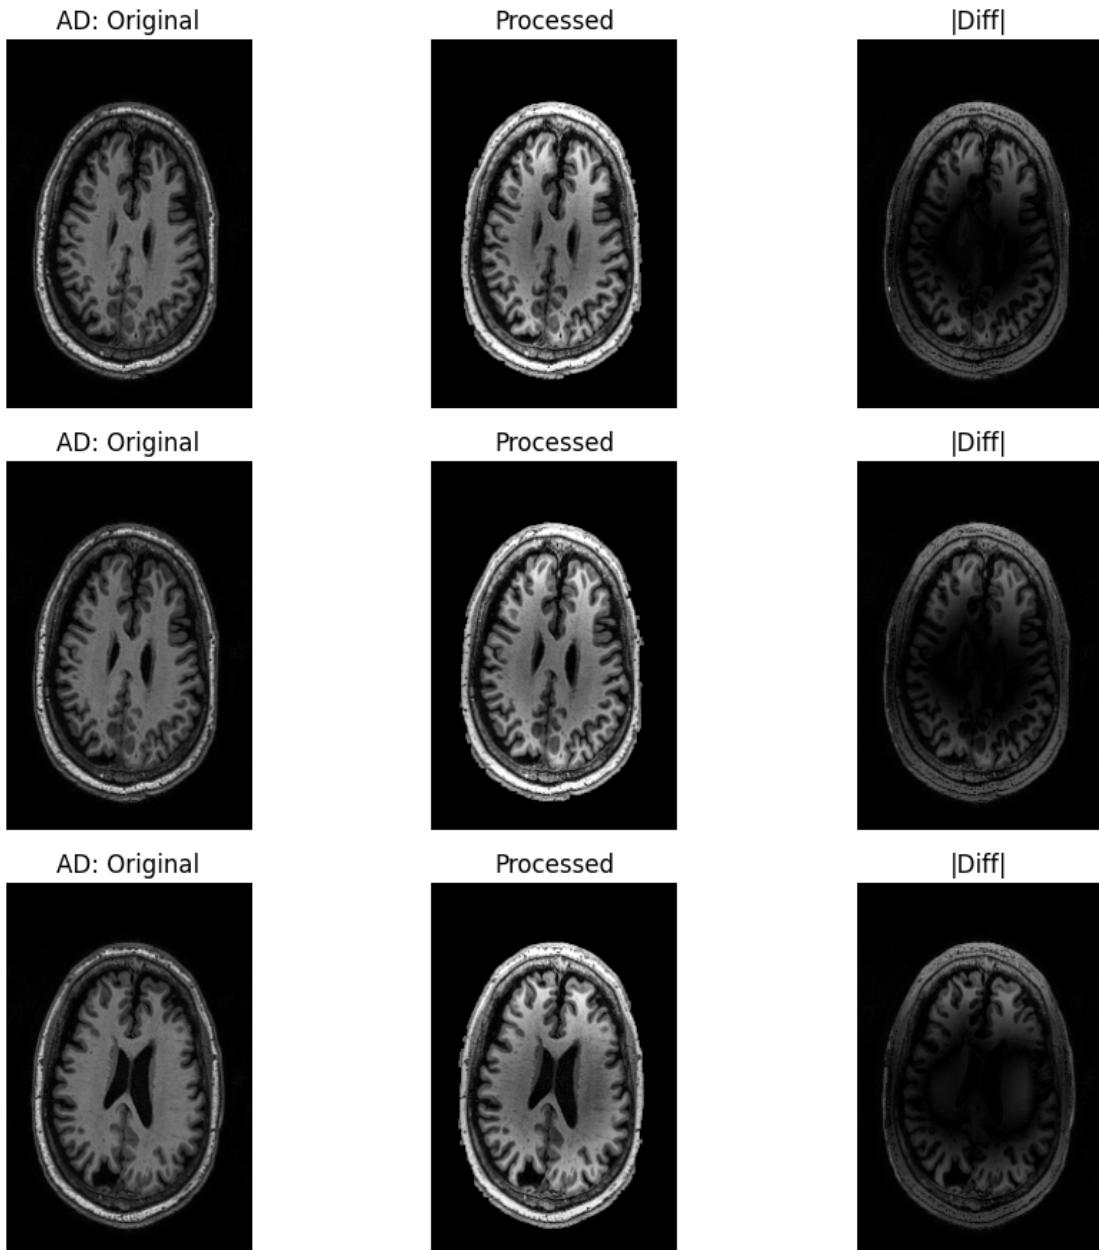
print("== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TRAIN ==")
sum_train = preprocess_dataset(train_dir, preprocessed_train, n_preview=3, do_skull_strip=True)

def _fmt(s): return ", ".join([f"{k}: {v[0]} ok / {v[1]} skipped" for k, v in s.items()])
print("Preprocessed images saved to:")
print(f" • Test : {preprocessed_test}")
print(f" • Train: {preprocessed_train}")
print("Summary TEST :", _fmt(sum_test))
print("Summary TRAIN:", _fmt(sum_train))

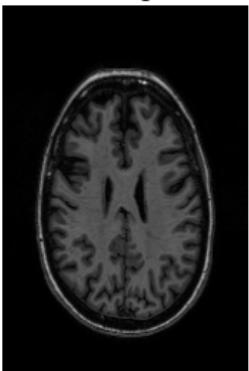
```

```
==== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TEST ====  
Device for homomorphic filtering: cuda
```

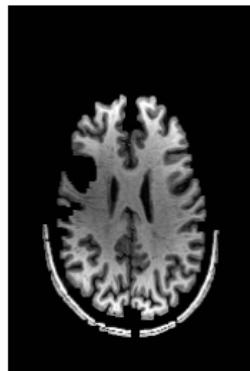
```
/tmp/ipykernel_19/2569239085.py:107: DeprecationWarning: 'mode' parameter is  
deprecated and will be removed in Pillow 13 (2026-10-15)  
Image.fromarray(arr_u8, mode="L").save(path)
```



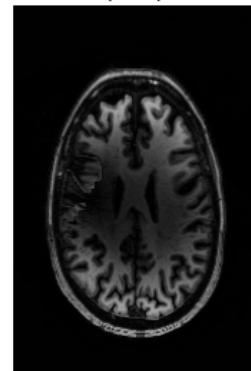
CN: Original



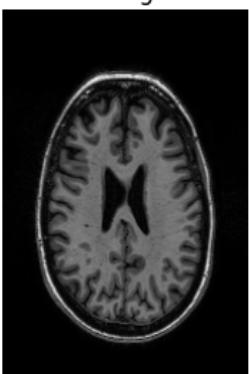
Processed



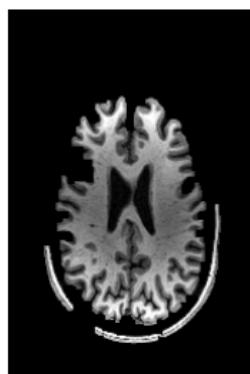
|Diff|



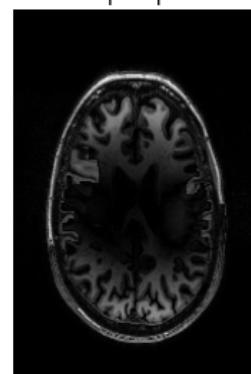
CN: Original



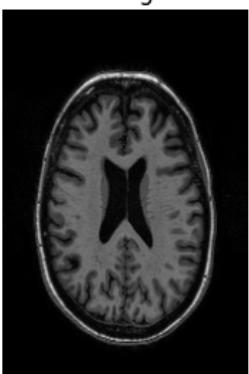
Processed



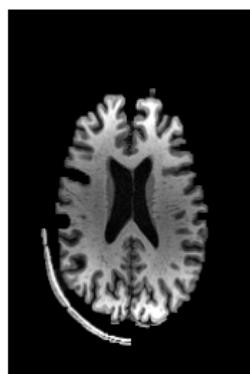
|Diff|



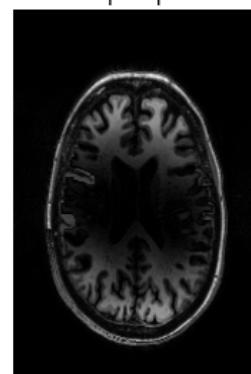
CN: Original

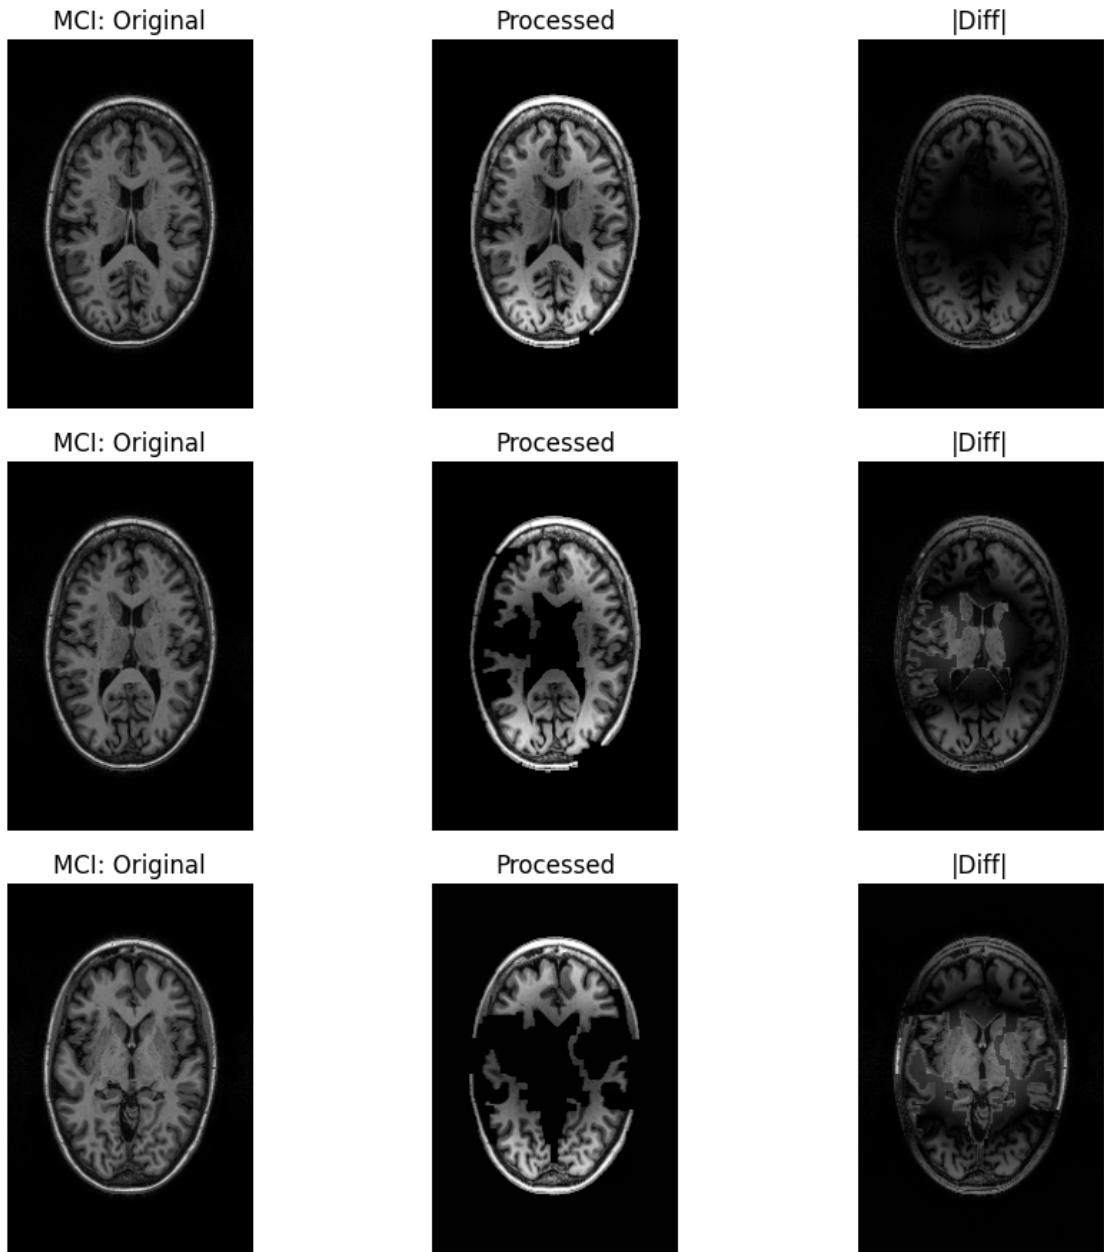


Processed



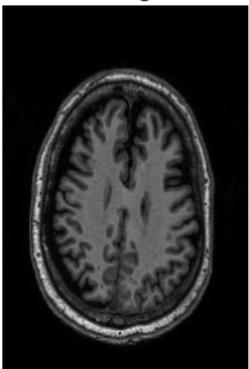
|Diff|



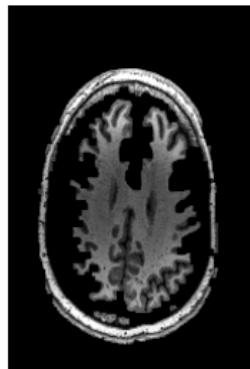


==== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TRAIN ====
Device for homomorphic filtering: cuda

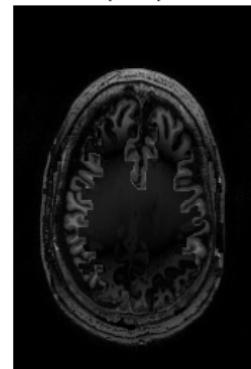
AD: Original



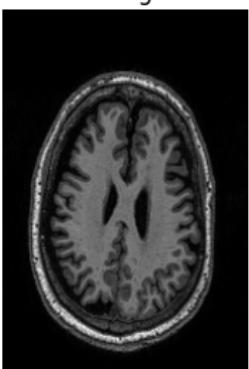
Processed



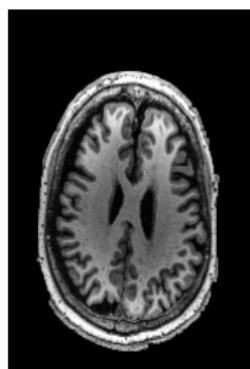
|Diff|



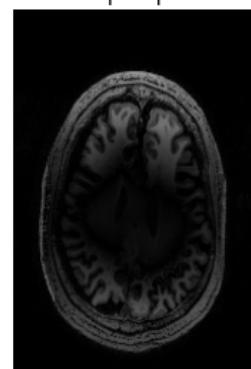
AD: Original



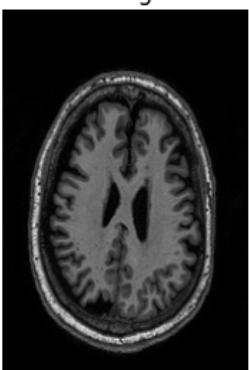
Processed



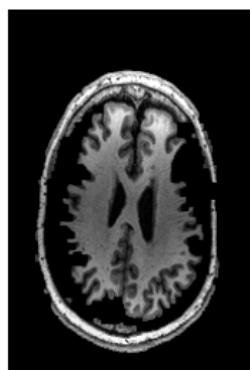
|Diff|



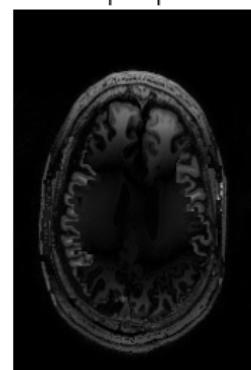
AD: Original



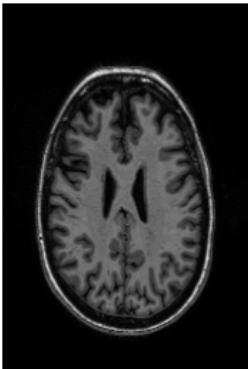
Processed



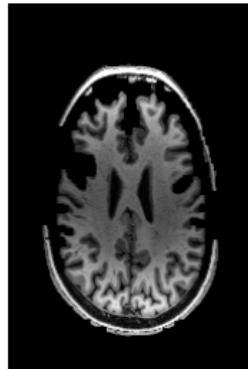
|Diff|



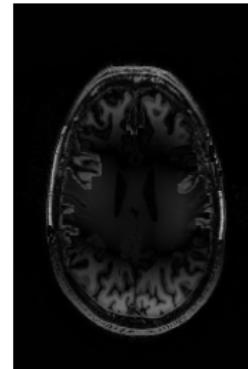
CN: Original



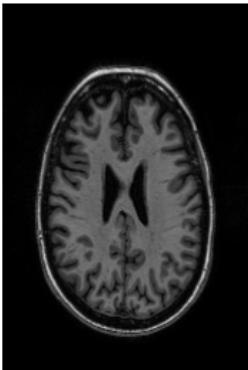
Processed



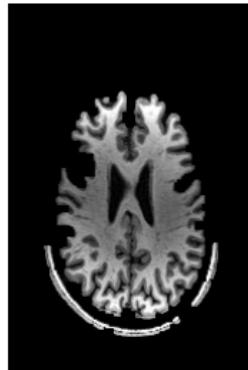
|Diff|



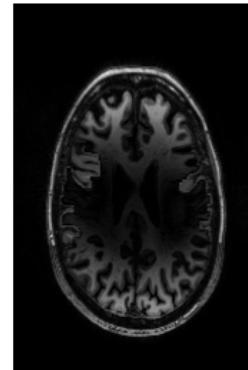
CN: Original



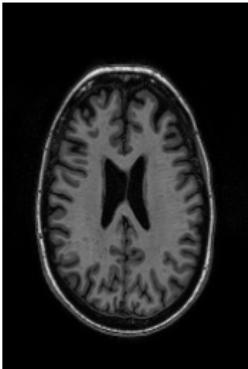
Processed



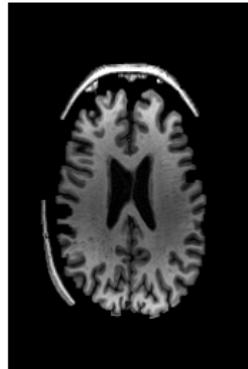
|Diff|



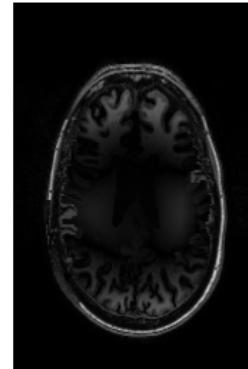
CN: Original

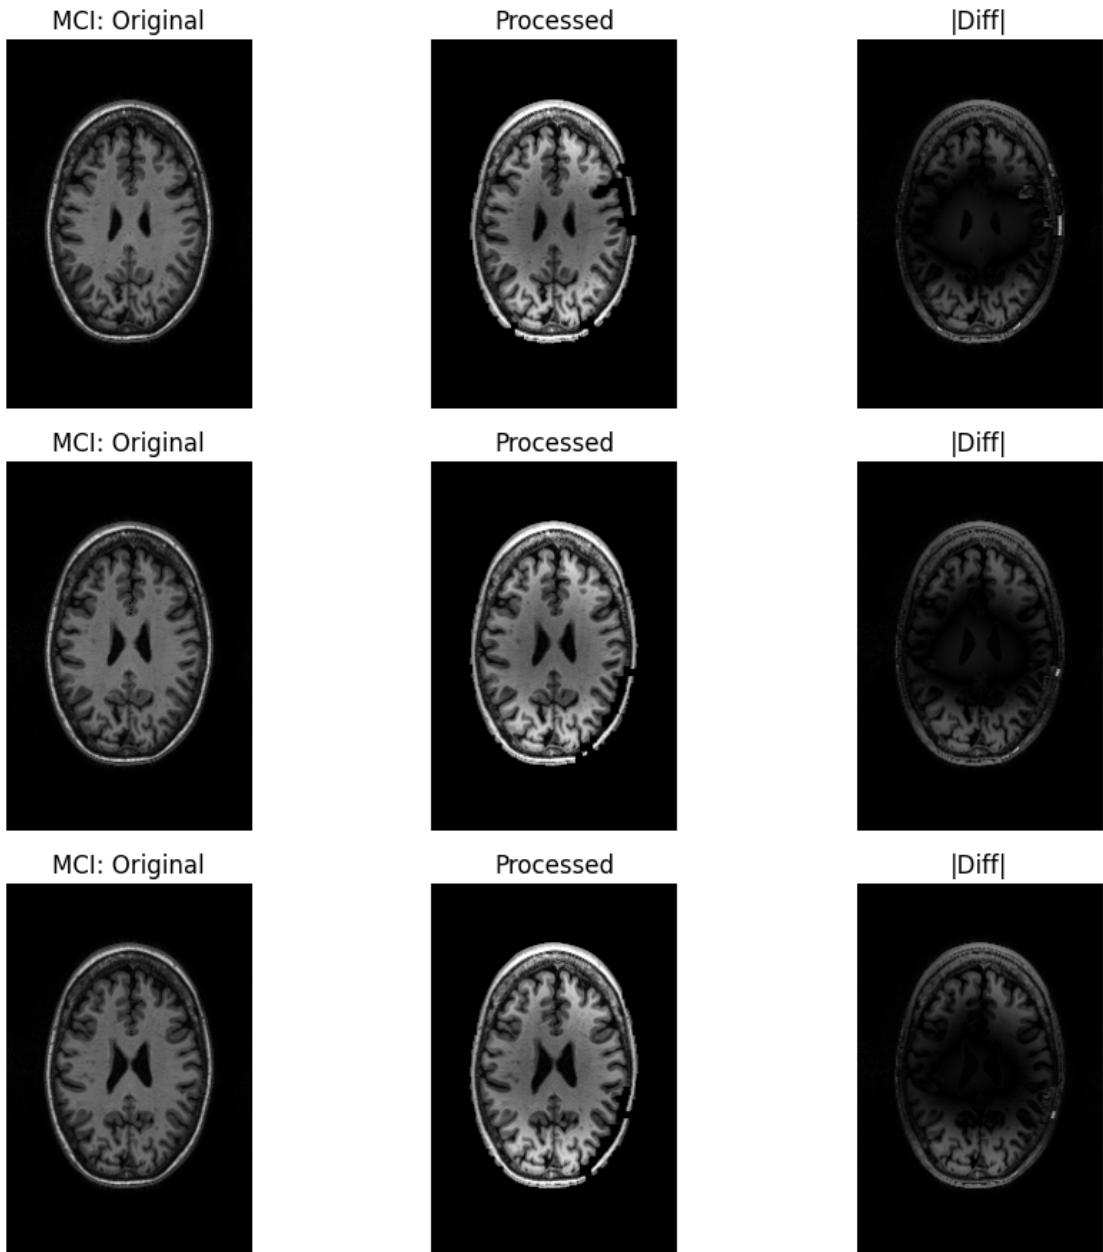


Processed



|Diff|





Preprocessed images saved to:

- Test : /kaggle/working/alzheimer-preprocessed/test
- Train: /kaggle/working/alzheimer-preprocessed/train

Summary TEST : AD: 225 ok / 0 skipped, CN: 288 ok / 0 skipped, MCI: 518 ok / 0 skipped

Summary TRAIN: AD: 899 ok / 0 skipped, CN: 1152 ok / 0 skipped, MCI: 2072 ok / 0 skipped

```
[3]: import os
import numpy as np
from PIL import Image, UnidentifiedImageError
import matplotlib.pyplot as plt

test_dir = "/kaggle/input/alzheimer-disease/test-20251010T094612Z-1-001/test"
train_dir = "/kaggle/input/alzheimer-disease/train-20251010T183510Z-1-001/train"
subfolders = ["AD", "CN", "MCI"]

preprocessed_test = "/kaggle/working/alzheimer-preprocessed/test"
preprocessed_train = "/kaggle/working/alzheimer-preprocessed/train"
resized_root = "/kaggle/working/alzheimer-resized-224"
resized_test = os.path.join(resized_root, "test")
resized_train = os.path.join(resized_root, "train")

IMG_EXT = (".png", ".jpg", ".jpeg", ".bmp", ".tif", ".tiff")
TARGET_SIZE = (224, 224)

def ensure_dir(p: str):
    os.makedirs(p, exist_ok=True)

def load_gray(path: str) -> np.ndarray:
    return np.asarray(Image.open(path).convert("L"), dtype=np.uint8)

def resize_and_save_gray(arr_u8: np.ndarray, path: str):
    img = Image.fromarray(arr_u8, mode="L")
    img_resized = img.resize(TARGET_SIZE, Image.Resampling.LANCZOS)
    img_resized.save(path)

def resize_dataset(src_root: str, dst_root: str, n_preview: int = 3):
    ensure_dir(dst_root)
    summary = {}

    for cls in subfolders:
        src_cls = os.path.join(src_root, cls)
        dst_cls = os.path.join(dst_root, cls)
        ensure_dir(dst_cls)

        processed = skipped = 0
        if not os.path.isdir(src_cls):
            summary[cls] = (0, 0)
            continue

        files = sorted([f for f in os.listdir(src_cls) if f.lower().
        ↪endswith(IMG_EXT)])
        for fname in files:
```

```

        spath = os.path.join(src_cls, fname)
        dpath = os.path.join(dst_cls, fname)
        try:
            arr = load_gray(spath)
            resize_and_save_gray(arr, dpath)
            processed += 1
        except (UnidentifiedImageError, OSError, RuntimeError, ValueError):
            skipped += 1
            continue

    summary[cls] = (processed, skipped)

    preview = files[:n_preview]
    if preview:
        fig, axs = plt.subplots(len(preview), 2, figsize=(8, 4*len(preview)))
        if len(preview) == 1:
            axs = np.expand_dims(axs, 0)
        for i, fname in enumerate(preview):
            sp = os.path.join(src_cls, fname)
            dp = os.path.join(dst_cls, fname)
            try:
                orig = load_gray(sp)
                resized = load_gray(dp)

                axs[i, 0].imshow(orig, cmap="gray", vmin=0, vmax=255)
                axs[i, 0].axis("off")
                axs[i, 0].set_title(f"{cls}: Original {orig.shape}")

                axs[i, 1].imshow(resized, cmap="gray", vmin=0, vmax=255)
                axs[i, 1].axis("off")
                axs[i, 1].set_title(f"Resized {resized.shape}")
            except Exception:
                continue
        plt.tight_layout()
        plt.show()

    return summary

print("== Resizing Preprocessed Images: TEST ==")
sum_test = resize_dataset(preprocessed_test, resized_test, n_preview=3)

print("== Resizing Preprocessed Images: TRAIN ==")
sum_train = resize_dataset(preprocessed_train, resized_train, n_preview=3)

def _fmt(s): return ", ".join([f"{k}: {v[0]} ok / {v[1]} skipped" for k, v in s.items()])

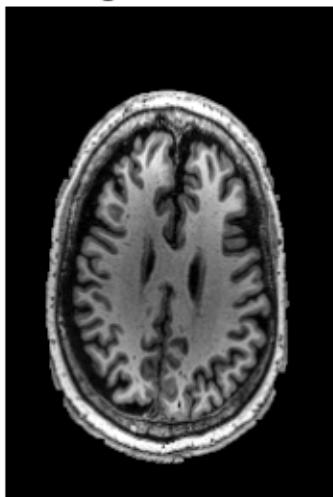
```

```
print("Resized images saved to:")
print(f" • Resized Test : {resized_test}")
print(f" • Resized Train: {resized_train}")
print("Summary TEST :", _fmt(sum_test))
print("Summary TRAIN:", _fmt(sum_train))
```

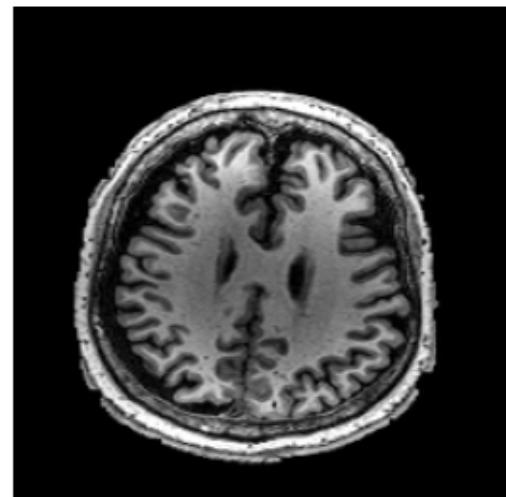
==== Resizing Preprocessed Images: TEST ===

```
/tmp/ipykernel_19/502029755.py:26: DeprecationWarning: 'mode' parameter is
deprecated and will be removed in Pillow 13 (2026-10-15)
    img = Image.fromarray(arr_u8, mode="L")
```

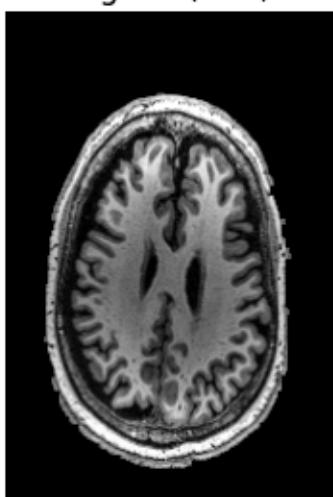
AD: Original (256, 170)



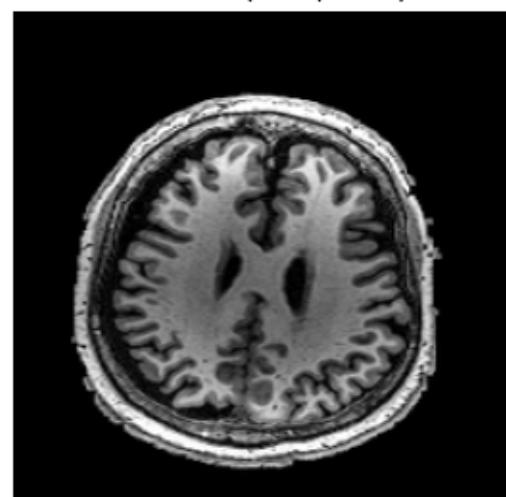
Resized (224, 224)



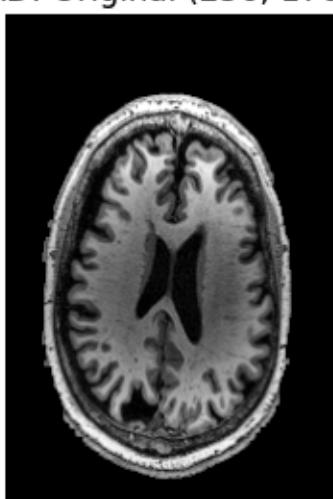
AD: Original (256, 170)



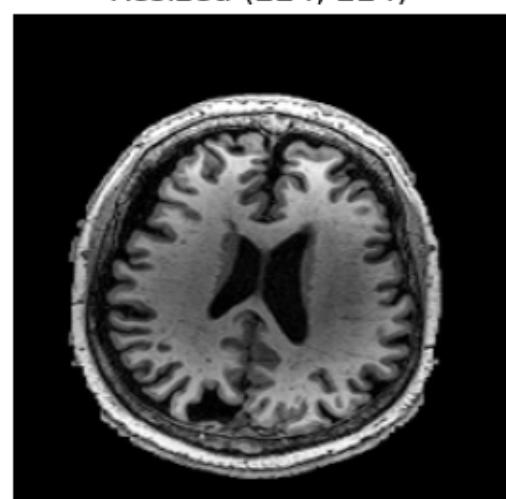
Resized (224, 224)



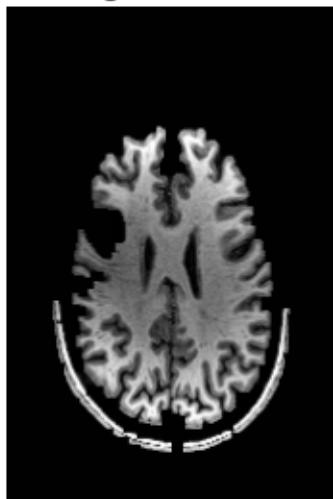
AD: Original (256, 170)



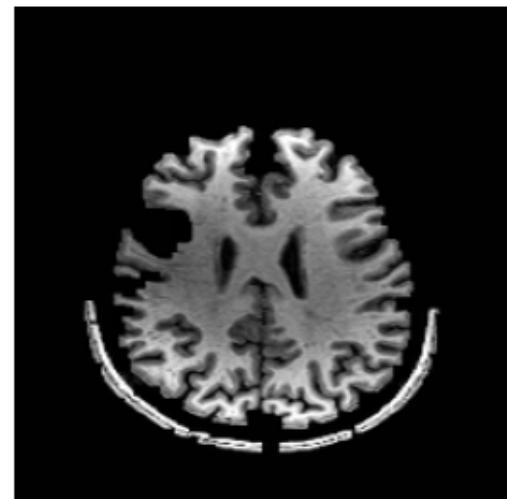
Resized (224, 224)



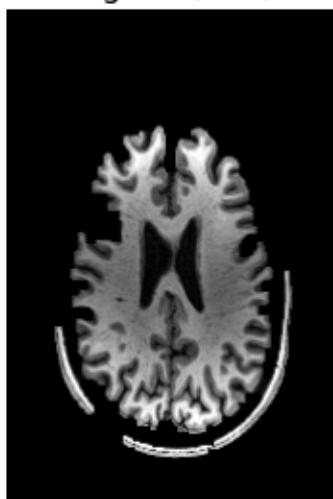
CN: Original (256, 170)



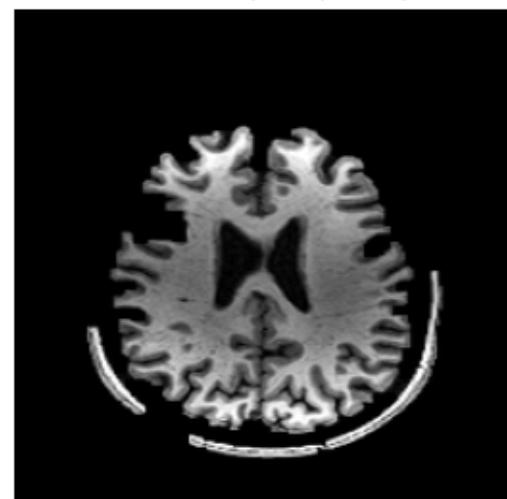
Resized (224, 224)



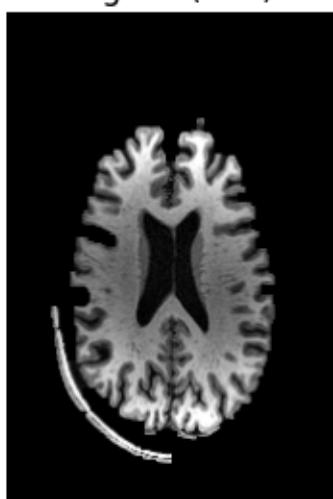
CN: Original (256, 170)



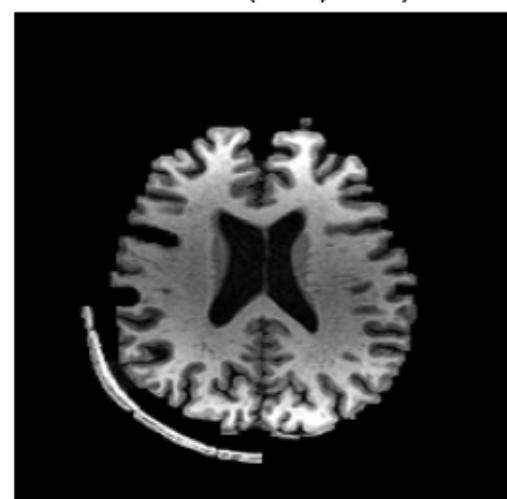
Resized (224, 224)



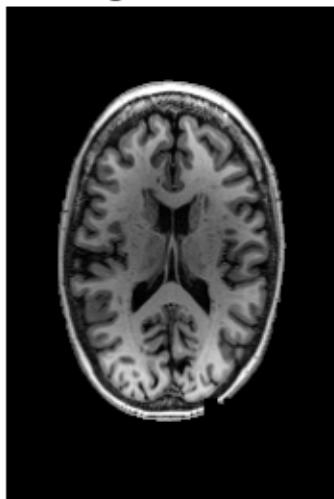
CN: Original (256, 170)



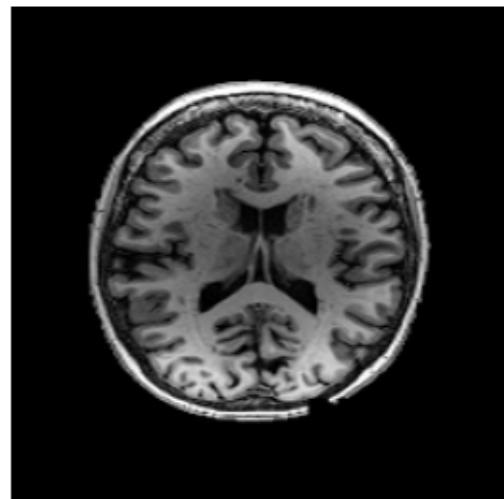
Resized (224, 224)



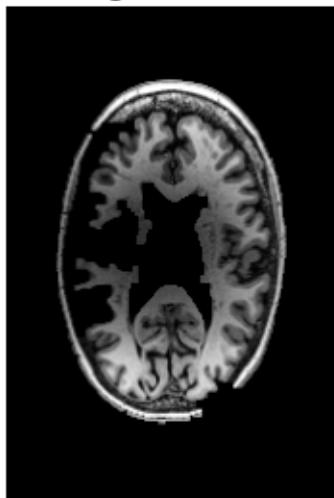
MCI: Original (256, 170)



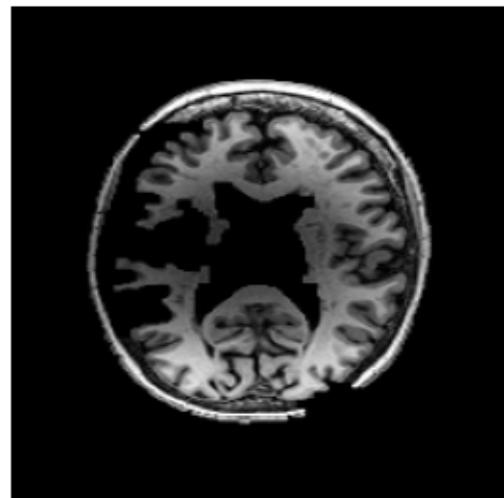
Resized (224, 224)



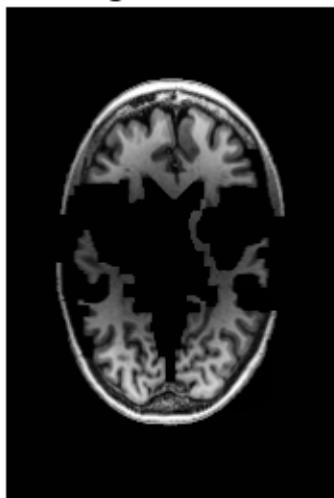
MCI: Original (256, 170)



Resized (224, 224)



MCI: Original (256, 170)

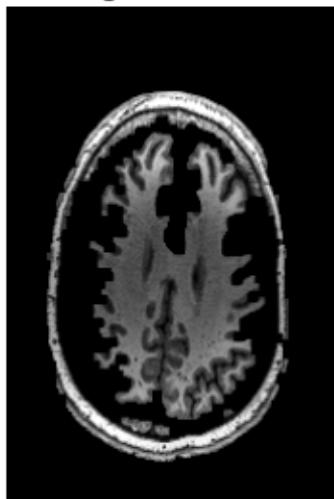


Resized (224, 224)

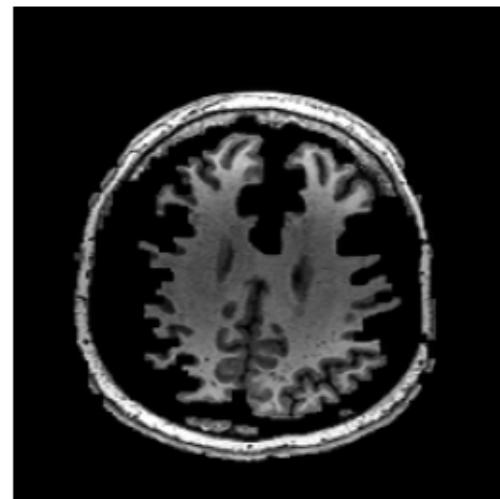


==== Resizing Preprocessed Images: TRAIN ===

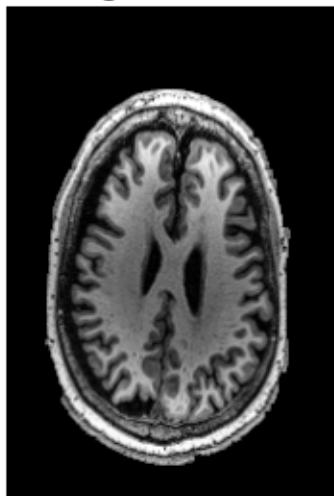
AD: Original (256, 170)



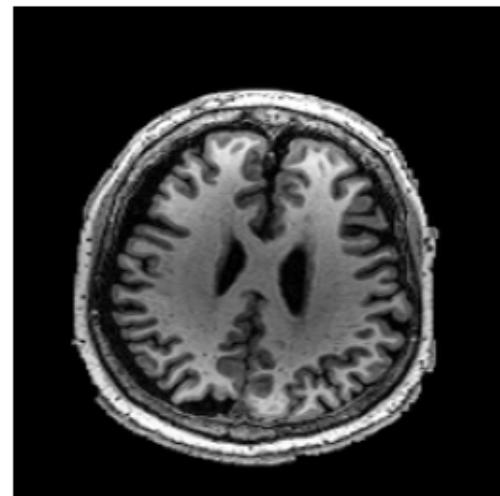
Resized (224, 224)



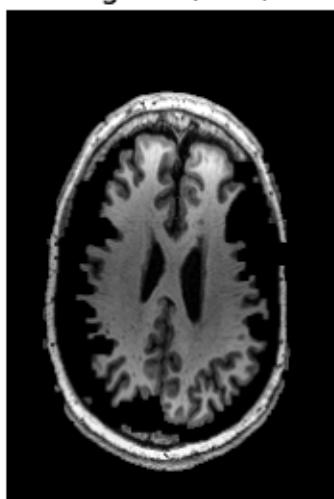
AD: Original (256, 170)



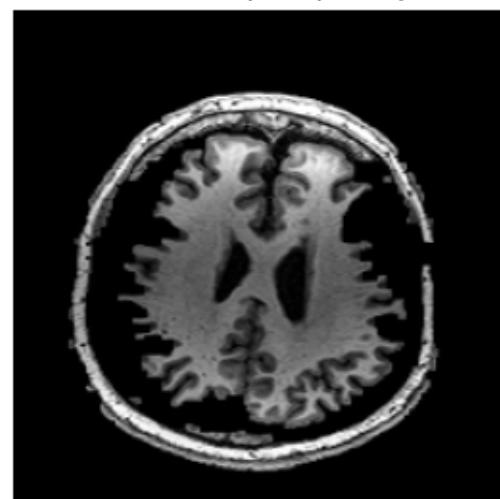
Resized (224, 224)



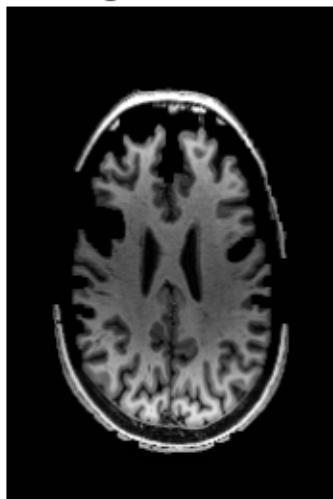
AD: Original (256, 170)



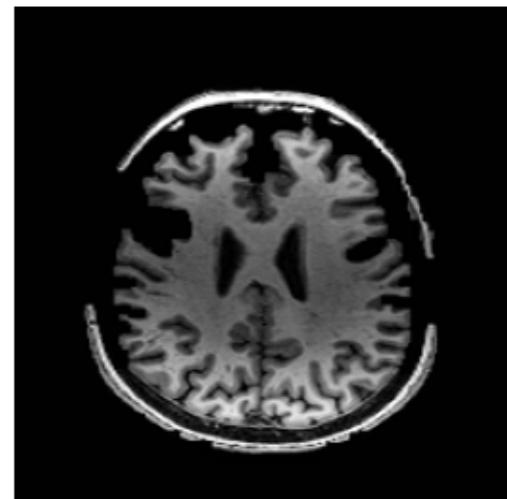
Resized (224, 224)



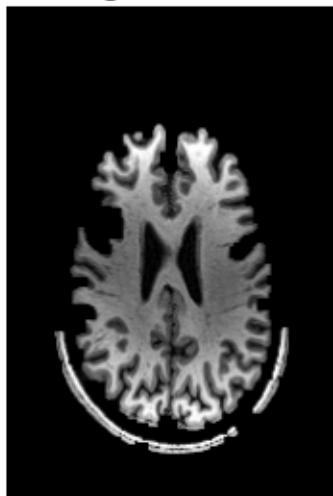
CN: Original (256, 170)



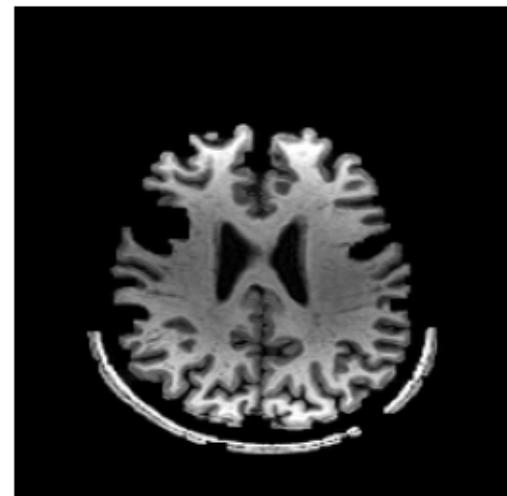
Resized (224, 224)



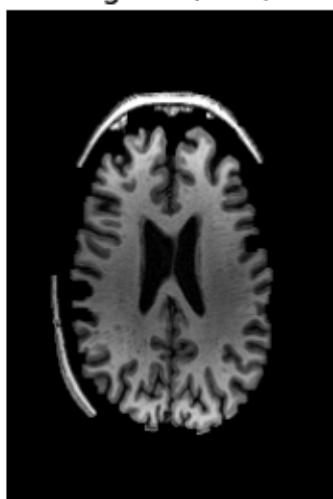
CN: Original (256, 170)



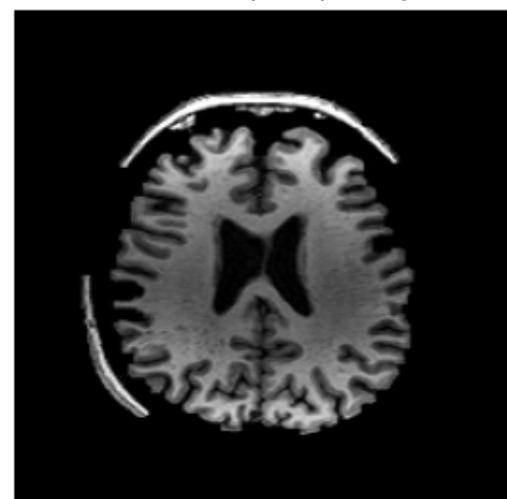
Resized (224, 224)



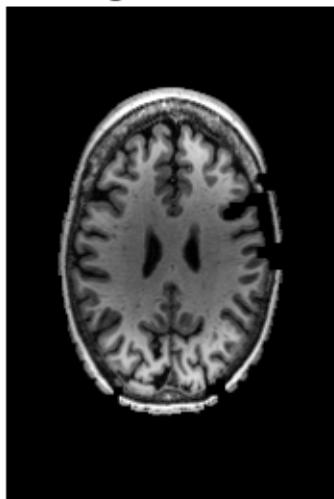
CN: Original (256, 170)



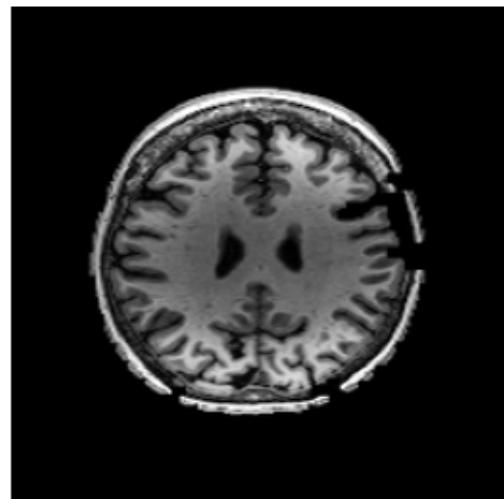
Resized (224, 224)



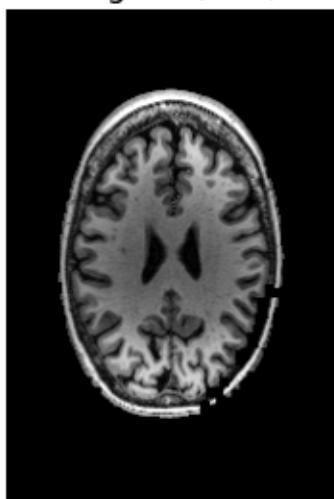
MCI: Original (256, 170)



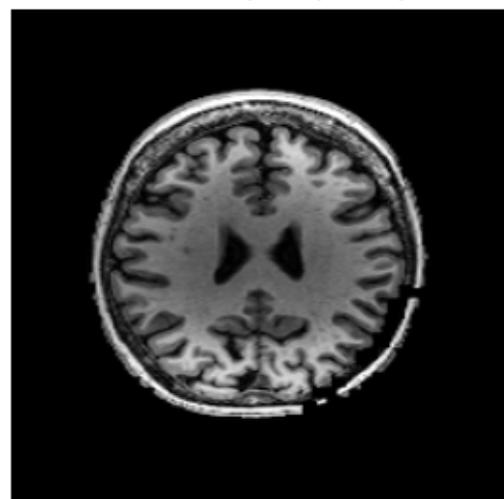
Resized (224, 224)



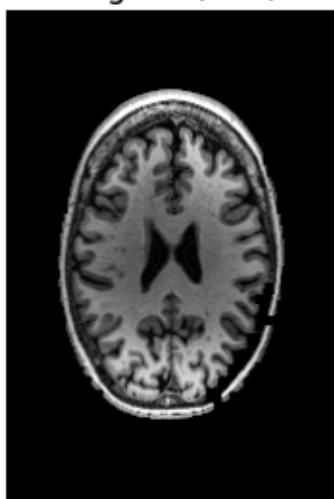
MCI: Original (256, 170)



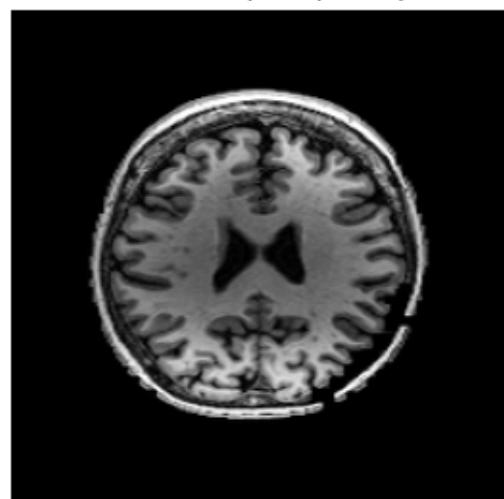
Resized (224, 224)



MCI: Original (256, 170)



Resized (224, 224)



```

Resized images saved to:
  • Resized Test : /kaggle/working/alzheimer-resized-224/test
  • Resized Train: /kaggle/working/alzheimer-resized-224/train
Summary TEST : AD: 225 ok / 0 skipped, CN: 288 ok / 0 skipped, MCI: 518 ok / 0
skipped
Summary TRAIN: AD: 899 ok / 0 skipped, CN: 1152 ok / 0 skipped, MCI: 2072 ok / 0
skipped

```

Train And Test Split

```
[4]: import os
import sys
import csv
import random
import shutil
from pathlib import Path
from typing import Dict, List, Tuple

# -----
# Configuration (modified for resized images)
# -----
# Source resized dataset (224x224)
RESIZED_ROOT = os.environ.get("RESIZED_ROOT", "/kaggle/working/
    ↪alzheimer-resized-224")
SRC_TRAIN = os.path.join(RESIZED_ROOT, "train")
SRC_TEST = os.path.join(RESIZED_ROOT, "test")

# Classes (same as before)
CLASSES = ["AD", "CN", "MCI"]

# Where to write splits
SPLITS_ROOT = "/kaggle/working/alzheimer-resized-224_splits"
os.makedirs(SPLITS_ROOT, exist_ok=True)

# The split ratios (train:test) you requested
RATIO_LIST = [
    (0.90, 0.10),
    (0.80, 0.20),
    (0.70, 0.30),
    (0.60, 0.40),
    (0.50, 0.50),
    (0.40, 0.60),
    (0.30, 0.70),
    (0.20, 0.80),
    (0.10, 0.90)
```

```

]

# Validation share taken from the training portion
VAL_FRACTION = 0.10

# Base random seed (deterministic builds). Each ratio derives its own seed.
BASE_SEED = 2025

# -----
# Helpers
# -----
IMG_EXTS = (".png", ".jpg", ".jpeg", ".bmp", ".tif", ".tiff")

def list_images_in_class(class_dir: str) -> List[str]:
    """Return absolute paths of all images in a class directory."""
    if not os.path.isdir(class_dir):
        return []
    files = sorted([
        str(Path(class_dir) / f) for f in os.listdir(class_dir)
        if f.lower().endswith(IMG_EXTS)
    ])
    return files

def gather_all_images() -> Dict[str, List[str]]:
    """
    Gather all resized images per class, from both 'train' and 'test'
    to form a single full pool for stratified splitting by ratio.
    """
    all_by_class = {c: [] for c in CLASSES}
    for c in CLASSES:
        # From resized train
        all_by_class[c].extend(list_images_in_class(os.path.join(SRC_TRAIN, c)))
        # From resized test
        all_by_class[c].extend(list_images_in_class(os.path.join(SRC_TEST, c)))
    return all_by_class

def ensure_dirs(*paths: str):
    for p in paths:
        os.makedirs(p, exist_ok=True)

def link_or_copy(src: str, dst: str):
    """Create a symlink; if not permitted, copy the file."""
    try:
        # Remove dst if exists
        if os.path.lexists(dst):
            os.unlink(dst)
        os.symlink(src, dst)
    except OSError:
        shutil.copy(src, dst)

```

```

except OSError:
    shutil.copy2(src, dst)

def write_manifest(split_dir: str, split_name: str, rows: List[Tuple[str, str, str, str]]):
    """
    Write both CSV manifest and plain path list.
    rows: list of (split, cls, filename, src_path)
    """
    # CSV
    csv_path = os.path.join(split_dir, f"{split_name}.csv")
    with open(csv_path, "w", newline="") as f:
        w = csv.writer(f)
        w.writerow(["split", "class", "filename", "path"])
        for r in rows:
            w.writerow(r)

    # TXT list
    txt_path = os.path.join(split_dir, f"{split_name}.txt")
    with open(txt_path, "w") as f:
        for _, _, _, p in rows:
            f.write(p + "\n")

def summarize_counts(counts: Dict[str, Dict[str, int]]):
    """
    Print counts per split (train/val/test) and per class for quick sanity
    check.
    """
    for split in ["train", "val", "test"]:
        info = counts.get(split, {})
        total = sum(info.values())
        detail = ", ".join([f"{k}: {v}" for k, v in info.items()])
        print(f"{split.capitalize():5s} => total {total:5d} | {detail}")

# -----
# Split builder
# -----
def build_splits():
    # 1) Pool all images by class (union of resized/train and resized/test)
    all_by_class = gather_all_images()

    # Optional: quick report of total availability
    print("Total images by class (union of resized/train + resized/test):")
    for c in CLASSES:
        print(f" {c}: {len(all_by_class[c])}")

    # 2) For each (train_ratio, test_ratio), create a split folder and populate

```

```

for tr_ratio, te_ratio in RATIO_LIST:
    # sanity: ratios sum approx 1
    assert abs(tr_ratio + te_ratio - 1.0) < 1e-6, "Train+Test ratio must sum to 1"

    # deterministic seed per ratio
    seed = BASE_SEED + int(round(te_ratio * 100))
    rng = random.Random(seed)

    # Split name and dirs
    split_name = f"split_{int(round(tr_ratio*100))}_{int(round(te_ratio*100))}"
    split_root = os.path.join(SPLITS_ROOT, split_name)
    train_root = os.path.join(split_root, "train")
    val_root = os.path.join(split_root, "val")
    test_root = os.path.join(split_root, "test")
    ensure_dirs(split_root, train_root, val_root, test_root)
    for c in CLASSES:
        ensure_dirs(os.path.join(train_root, c), os.path.join(val_root, c), os.path.join(test_root, c))

    rows_train, rows_val, rows_test = [], [], []
    counts = {"train": {}, "val": {}, "test": {}}

    # 3) Per-class stratified splitting
    for c in CLASSES:
        full_list = list(all_by_class[c]) # copy
        rng.shuffle(full_list) # deterministic shuffle

        n_total = len(full_list)
        n_test = max(0, int(round(n_total * te_ratio)))
        n_test = min(n_test, n_total) # guard

        test_list = full_list[:n_test]
        train_pool = full_list[n_test:]

        # Validation from training portion (10%)
        n_val = max(0, int(round(len(train_pool) * VAL_FRACTION)))
        val_list = train_pool[:n_val]
        train_list = train_pool[n_val:]

    # 4) Materialize (symlink/copy) into folders and write manifests
    # test
    for src_path in test_list:
        fname = os.path.basename(src_path)
        dst_path = os.path.join(test_root, c, fname)
        link_or_copy(src_path, dst_path)

```

```

        rows_test.append(("test", c, fname, dst_path))

    # val
    for src_path in val_list:
        fname = os.path.basename(src_path)
        dst_path = os.path.join(val_root, c, fname)
        link_or_copy(src_path, dst_path)
        rows_val.append(("val", c, fname, dst_path))

    # train
    for src_path in train_list:
        fname = os.path.basename(src_path)
        dst_path = os.path.join(train_root, c, fname)
        link_or_copy(src_path, dst_path)
        rows_train.append(("train", c, fname, dst_path))

    # counts
    counts["test"][c] = len(test_list)
    counts["val"][c] = len(val_list)
    counts["train"][c] = len(train_list)

    # 5) Write manifest files for this split
    write_manifest(split_root, "train", rows_train)
    write_manifest(split_root, "val", rows_val)
    write_manifest(split_root, "test", rows_test)

    # 6) Summary printout
    print(f"\n==== {split_name} ====")
    summarize_counts(counts)
    print(f"Paths:\n  Train: {train_root}\n  Val : {val_root}\n  Test : {test_root}\n")

    # 7) Export env var for convenient access in later cells
    os.environ["RESIZED_SPLITS_ROOT"] = SPLITS_ROOT
    print(f"All splits created under: {SPLITS_ROOT}")

# -----
# Execute
# -----
build_splits()

```

Total images by class (union of resized/train + resized/test):

AD: 1124
CN: 1440
MCI: 2590

==== split_90_10 ===

```

Train => total 4175 | AD: 911, CN: 1166, MCI: 2098
Val   => total 464 | AD: 101, CN: 130, MCI: 233
Test  => total 515 | AD: 112, CN: 144, MCI: 259
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_90_10/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_90_10/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_90_10/test

==== split_80_20 ====
Train => total 3711 | AD: 809, CN: 1037, MCI: 1865
Val   => total 412 | AD: 90, CN: 115, MCI: 207
Test  => total 1031 | AD: 225, CN: 288, MCI: 518
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_80_20/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_80_20/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_80_20/test

==== split_70_30 ====
Train => total 3247 | AD: 708, CN: 907, MCI: 1632
Val   => total 361 | AD: 79, CN: 101, MCI: 181
Test  => total 1546 | AD: 337, CN: 432, MCI: 777
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_70_30/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_70_30/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_70_30/test

==== split_60_40 ====
Train => total 2784 | AD: 607, CN: 778, MCI: 1399
Val   => total 308 | AD: 67, CN: 86, MCI: 155
Test  => total 2062 | AD: 450, CN: 576, MCI: 1036
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_60_40/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_60_40/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_60_40/test

==== split_50_50 ====
Train => total 2319 | AD: 506, CN: 648, MCI: 1165
Val   => total 258 | AD: 56, CN: 72, MCI: 130
Test  => total 2577 | AD: 562, CN: 720, MCI: 1295
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_50_50/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_50_50/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_50_50/test

```

```

==== split_40_60 ====
Train => total 1855 | AD: 405, CN: 518, MCI: 932
Val   => total 207 | AD: 45, CN: 58, MCI: 104
Test  => total 3092 | AD: 674, CN: 864, MCI: 1554
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_40_60/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_40_60/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_40_60/test

==== split_30_70 ====
Train => total 1391 | AD: 303, CN: 389, MCI: 699
Val   => total 155 | AD: 34, CN: 43, MCI: 78
Test  => total 3608 | AD: 787, CN: 1008, MCI: 1813
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_30_70/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_30_70/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_30_70/test

==== split_20_80 ====
Train => total 928 | AD: 203, CN: 259, MCI: 466
Val   => total 103 | AD: 22, CN: 29, MCI: 52
Test  => total 4123 | AD: 899, CN: 1152, MCI: 2072
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_20_80/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_20_80/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_20_80/test

==== split_10_90 ====
Train => total 464 | AD: 101, CN: 130, MCI: 233
Val   => total 51 | AD: 11, CN: 14, MCI: 26
Test  => total 4639 | AD: 1012, CN: 1296, MCI: 2331
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_10_90/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_10_90/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_10_90/test

All splits created under: /kaggle/working/alzheimer-resized-224_splits

```

```
[5]: import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, models
```

```

import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import CosineAnnealingLR, ReduceLROnPlateau

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)
        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label

def get_data_transforms():
    train_transform = transforms.Compose([
        transforms.Resize((380, 380)),
        transforms.RandomHorizontalFlip(p=0.5),
        transforms.RandomRotation(10),
        transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
        transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
        transforms.RandomGrayscale(p=0.1),
        transforms.GaussianBlur(kernel_size=3, sigma=(0.1, 2.0)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
    ])

```

```

        transforms.RandomErasing(p=0.2, scale=(0.02, 0.2), ratio=(0.3, 3.3))
    ])

    val_transform = transforms.Compose([
        transforms.Resize((380, 380)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
        ↪225])
    ])

    return train_transform, val_transform

def train_model(model, train_loader, val_loader, criterion, optimizer, ↪
    ↪scheduler, num_epochs, device, warmup_epochs=3):
    best_val_acc = 0
    patience = 8
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        # Learning rate warmup
        if epoch < warmup_epochs:
            lr_scale = min(1.0, float(epoch + 1) / warmup_epochs)
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['initial_lr'] * lr_scale

        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()

            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
            optimizer.step()

            running_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        epoch_loss = running_loss / len(train_loader)

```

```

epoch_acc = 100 * correct / total

val_acc = evaluate_model(model, val_loader, device)

if scheduler and epoch >= warmup_epochs:
    if isinstance(scheduler, CosineAnnealingLR):
        scheduler.step()
    else:
        scheduler.step(val_acc)

current_lr = optimizer.param_groups[0]['lr']
print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, TrainAcc: {epoch_acc:.2f}%, Val Acc: {val_acc:.2f}%, LR: {current_lr:.6f}')
```

```

if val_acc > best_val_acc:
    best_val_acc = val_acc
    patience_counter = 0
else:
    patience_counter += 1

if patience_counter >= patience:
    print(f"Early stopping at epoch {epoch+1}")
    break
```

```

def evaluate_model(model, data_loader, device):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in data_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    return accuracy
```

```

def test_model(model, test_loader, device):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for images, labels in test_loader:
```

```

        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)

        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

    return all_preds, all_labels

def calculate_metrics(y_true, y_pred, split_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred, average='weighted', zero_division=0)
    recall = recall_score(y_true, y_pred, average='weighted', zero_division=0)
    f1 = f1_score(y_true, y_pred, average='weighted', zero_division=0)

    print(f"\n==== {split_name} Results ===")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-Score: {f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=['AD', 'CN', 'MCI'], zero_division=0))

    return {
        'split': split_name,
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1_score': f1
    }

def run_efficientnet_b4_on_splits():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    splits_root = "/kaggle/working/alzheimer-resized-224_splits"
    split_folders = [f for f in os.listdir(splits_root) if f.startswith('split_')]
    split_folders.sort()

    results = []

    train_transform, val_transform = get_data_transforms()

    for split_folder in split_folders:

```

```

print(f"\n{'='*60}")
print(f"Processing: {split_folder}")
print(f"{'='*60}")

split_path = os.path.join(splits_root, split_folder)

train_dataset = AlzheimerDataset(split_path, 'train', train_transform)
val_dataset = AlzheimerDataset(split_path, 'val', val_transform)
test_dataset = AlzheimerDataset(split_path, 'test', val_transform)

train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True,
                           num_workers=2, pin_memory=True)
val_loader = DataLoader(val_dataset, batch_size=8, shuffle=False,
                           num_workers=2, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=8, shuffle=False,
                           num_workers=2, pin_memory=True)

print(f"Train samples: {len(train_dataset)}")
print(f"Val samples: {len(val_dataset)}")
print(f"Test samples: {len(test_dataset)}")

model = models.efficientnet_b4(weights=models.EfficientNet_B4_Weights.
                           IMAGENET1K_V1)

# Freeze all layers initially
for param in model.parameters():
    param.requires_grad = False

# Unfreeze classifier
num_ftrs = model.classifier[1].in_features
model.classifier[1] = nn.Linear(num_ftrs, 3)

# Unfreeze last blocks (blocks 6-7) for feature adaptation
for param in model.features[6:].parameters():
    param.requires_grad = True
for param in model.features[5][-2:].parameters():
    param.requires_grad = True

model = model.to(device)

# Layer-wise learning rates for EfficientNet
optimizer = torch.optim.AdamW([
    {'params': model.classifier.parameters(), 'lr': 0.0001,
     'initial_lr': 0.0001},
    {'params': model.features[7].parameters(), 'lr': 0.00005,
     'initial_lr': 0.00005},
    ...
])

```

```

        {'params': model.features[6].parameters(), 'lr': 0.00003, ↴
↳'initial_lr': 0.00003},
        {'params': model.features[5][-2:].parameters(), 'lr': 0.00001, ↴
↳'initial_lr': 0.00001},
    ], weight_decay=1e-5)

criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
scheduler = CosineAnnealingLR(optimizer, T_max=15, eta_min=1e-6)

print("Starting phase 1 training (partial unfreeze)...")
start_time = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, ↴
↳scheduler, num_epochs=18, device=device)
phase1_time = time.time() - start_time

print("Starting phase 2 training (full fine-tuning)...")
for param in model.parameters():
    param.requires_grad = True

optimizer = torch.optim.AdamW(model.parameters(), lr=0.000005, ↴
↳weight_decay=1e-6)
scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5, ↴
↳patience=4, verbose=True)

start_time_phase2 = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, ↴
↳scheduler, num_epochs=12, device=device, warmup_epochs=0)
total_time = phase1_time + (time.time() - start_time_phase2)

print("Testing model...")
test_preds, test_labels = test_model(model, test_loader, device)

split_results = calculate_metrics(test_labels, test_preds, split_folder)
split_results['training_time'] = total_time
results.append(split_results)

torch.cuda.empty_cache()

results_df = pd.DataFrame(results)
print(f"\n{'='*80}")
print("EfficientNet-B4 - SUMMARY OF ALL SPLITS")
print(f"\n{'='*80}")
print(results_df.to_string(index=False))

results_csv_path = "/kaggle/working/efficientnet_b4_results.csv"
results_df.to_csv(results_csv_path, index=False)

```

```

    print(f"\nDetailed results saved to: {results_csv_path}")

    return results_df

if __name__ == "__main__":
    results = run_efficientnet_b4_on_splits()

```

Using device: cuda

```

=====
Processing: split_10_90
=====

Train samples: 464
Val samples: 51
Test samples: 4639

Downloading:
"https://download.pytorch.org/models/efficientnet_b4_rwightman-23ab8bcd.pth" to
/root/.cache/torch/hub/checkpoints/efficientnet_b4_rwightman-23ab8bcd.pth
100%| 74.5M/74.5M [00:00<00:00, 184MB/s]

Starting phase 1 training (partial unfreeze)...
Epoch [1/18], Loss: 1.0975, Train Acc: 33.84%, Val Acc: 33.33%, LR: 0.000033
Epoch [2/18], Loss: 1.0774, Train Acc: 44.61%, Val Acc: 50.98%, LR: 0.000067
Epoch [3/18], Loss: 1.0545, Train Acc: 49.78%, Val Acc: 50.98%, LR: 0.000100
Epoch [4/18], Loss: 1.0435, Train Acc: 50.43%, Val Acc: 50.98%, LR: 0.000099
Epoch [5/18], Loss: 1.0455, Train Acc: 50.22%, Val Acc: 50.98%, LR: 0.000096
Epoch [6/18], Loss: 1.0326, Train Acc: 50.00%, Val Acc: 50.98%, LR: 0.000091
Epoch [7/18], Loss: 1.0308, Train Acc: 50.43%, Val Acc: 50.98%, LR: 0.000084
Epoch [8/18], Loss: 1.0319, Train Acc: 50.22%, Val Acc: 50.98%, LR: 0.000075
Epoch [9/18], Loss: 1.0239, Train Acc: 50.43%, Val Acc: 50.98%, LR: 0.000066
Epoch [10/18], Loss: 1.0302, Train Acc: 49.57%, Val Acc: 50.98%, LR: 0.000056
Early stopping at epoch 10
Starting phase 2 training (full fine-tuning)...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/12], Loss: 1.0167, Train Acc: 50.65%, Val Acc: 50.98%, LR: 0.000005
Epoch [2/12], Loss: 1.0248, Train Acc: 50.00%, Val Acc: 50.98%, LR: 0.000005
Epoch [3/12], Loss: 1.0250, Train Acc: 51.29%, Val Acc: 50.98%, LR: 0.000005
Epoch [4/12], Loss: 1.0238, Train Acc: 50.86%, Val Acc: 50.98%, LR: 0.000005
Epoch [5/12], Loss: 1.0152, Train Acc: 51.29%, Val Acc: 50.98%, LR: 0.000005
Epoch [6/12], Loss: 1.0222, Train Acc: 50.86%, Val Acc: 50.98%, LR: 0.000003
Epoch [7/12], Loss: 1.0234, Train Acc: 50.00%, Val Acc: 50.98%, LR: 0.000003
Epoch [8/12], Loss: 1.0161, Train Acc: 51.94%, Val Acc: 50.98%, LR: 0.000003
Epoch [9/12], Loss: 1.0193, Train Acc: 51.08%, Val Acc: 50.98%, LR: 0.000003

```

```
Early stopping at epoch 9
Testing model...
```

```
==== split_10_90 Results ====
Accuracy: 0.5031
Precision: 0.4124
Recall: 0.5031
F1-Score: 0.3380
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.00	0.00	0.00	1012
CN	0.57	0.00	0.01	1296
MCI	0.50	1.00	0.67	2331
accuracy			0.50	4639
macro avg	0.36	0.33	0.23	4639
weighted avg	0.41	0.50	0.34	4639

```
=====
Processing: split_20_80
=====
```

```
Train samples: 928
```

```
Val samples: 103
```

```
Test samples: 4123
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/18], Loss: 1.0934, Train Acc: 39.55%, Val Acc: 47.57%, LR: 0.000033
Epoch [2/18], Loss: 1.0604, Train Acc: 49.78%, Val Acc: 50.49%, LR: 0.000067
Epoch [3/18], Loss: 1.0485, Train Acc: 50.22%, Val Acc: 50.49%, LR: 0.000100
Epoch [4/18], Loss: 1.0394, Train Acc: 50.32%, Val Acc: 50.49%, LR: 0.000099
Epoch [5/18], Loss: 1.0389, Train Acc: 50.32%, Val Acc: 50.49%, LR: 0.000096
Epoch [6/18], Loss: 1.0383, Train Acc: 50.32%, Val Acc: 50.49%, LR: 0.000091
Epoch [7/18], Loss: 1.0330, Train Acc: 50.22%, Val Acc: 50.49%, LR: 0.000084
Epoch [8/18], Loss: 1.0290, Train Acc: 50.22%, Val Acc: 50.49%, LR: 0.000075
Epoch [9/18], Loss: 1.0279, Train Acc: 50.75%, Val Acc: 50.49%, LR: 0.000066
Epoch [10/18], Loss: 1.0264, Train Acc: 50.75%, Val Acc: 50.49%, LR: 0.000056
```

```
Early stopping at epoch 10
```

```
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/12], Loss: 1.0195, Train Acc: 50.32%, Val Acc: 50.49%, LR: 0.000005
Epoch [2/12], Loss: 1.0190, Train Acc: 50.54%, Val Acc: 50.49%, LR: 0.000005
Epoch [3/12], Loss: 1.0147, Train Acc: 50.54%, Val Acc: 50.49%, LR: 0.000005
```

```
Epoch [4/12], Loss: 1.0152, Train Acc: 50.86%, Val Acc: 50.49%, LR: 0.000005
Epoch [5/12], Loss: 1.0139, Train Acc: 50.75%, Val Acc: 50.49%, LR: 0.000005
Epoch [6/12], Loss: 1.0204, Train Acc: 50.54%, Val Acc: 50.49%, LR: 0.000003
Epoch [7/12], Loss: 1.0164, Train Acc: 51.29%, Val Acc: 50.49%, LR: 0.000003
Epoch [8/12], Loss: 1.0176, Train Acc: 50.32%, Val Acc: 50.49%, LR: 0.000003
Epoch [9/12], Loss: 1.0092, Train Acc: 50.75%, Val Acc: 50.49%, LR: 0.000003
Early stopping at epoch 9
Testing model...
```

```
==== split_20_80 Results ===
```

```
Accuracy: 0.5035
Precision: 0.4764
Recall: 0.5035
F1-Score: 0.3384
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.00	0.00	0.00	899
CN	0.80	0.00	0.01	1152
MCI	0.50	1.00	0.67	2072
accuracy			0.50	4123
macro avg	0.43	0.33	0.23	4123
weighted avg	0.48	0.50	0.34	4123

```
=====
Processing: split_30_70
=====
```

```
Train samples: 1391
```

```
Val samples: 155
```

```
Test samples: 3608
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/18], Loss: 1.0764, Train Acc: 47.09%, Val Acc: 50.32%, LR: 0.000033
Epoch [2/18], Loss: 1.0476, Train Acc: 50.11%, Val Acc: 50.32%, LR: 0.000067
Epoch [3/18], Loss: 1.0421, Train Acc: 50.32%, Val Acc: 50.32%, LR: 0.000100
Epoch [4/18], Loss: 1.0389, Train Acc: 50.61%, Val Acc: 50.32%, LR: 0.000099
Epoch [5/18], Loss: 1.0331, Train Acc: 50.40%, Val Acc: 50.32%, LR: 0.000096
Epoch [6/18], Loss: 1.0227, Train Acc: 50.47%, Val Acc: 50.97%, LR: 0.000091
Epoch [7/18], Loss: 1.0179, Train Acc: 52.34%, Val Acc: 52.26%, LR: 0.000084
Epoch [8/18], Loss: 1.0135, Train Acc: 51.83%, Val Acc: 52.90%, LR: 0.000075
Epoch [9/18], Loss: 1.0093, Train Acc: 52.05%, Val Acc: 52.26%, LR: 0.000066
Epoch [10/18], Loss: 1.0031, Train Acc: 52.70%, Val Acc: 52.26%, LR: 0.000056
Epoch [11/18], Loss: 0.9939, Train Acc: 53.06%, Val Acc: 53.55%, LR: 0.000045
Epoch [12/18], Loss: 0.9945, Train Acc: 52.19%, Val Acc: 53.55%, LR: 0.000035
Epoch [13/18], Loss: 0.9808, Train Acc: 53.99%, Val Acc: 54.84%, LR: 0.000026
Epoch [14/18], Loss: 0.9871, Train Acc: 53.41%, Val Acc: 52.90%, LR: 0.000017
```

```
Epoch [15/18], Loss: 0.9827, Train Acc: 53.41%, Val Acc: 52.26%, LR: 0.000010
Epoch [16/18], Loss: 0.9775, Train Acc: 54.78%, Val Acc: 56.77%, LR: 0.000005
Epoch [17/18], Loss: 0.9773, Train Acc: 54.35%, Val Acc: 57.42%, LR: 0.000002
Epoch [18/18], Loss: 0.9769, Train Acc: 54.78%, Val Acc: 56.13%, LR: 0.000001
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
warnings.warn(
Epoch [1/12], Loss: 0.9749, Train Acc: 54.92%, Val Acc: 56.13%, LR: 0.000005
Epoch [2/12], Loss: 0.9678, Train Acc: 55.28%, Val Acc: 56.13%, LR: 0.000005
Epoch [3/12], Loss: 0.9658, Train Acc: 55.14%, Val Acc: 59.35%, LR: 0.000005
Epoch [4/12], Loss: 0.9638, Train Acc: 56.29%, Val Acc: 58.71%, LR: 0.000005
Epoch [5/12], Loss: 0.9682, Train Acc: 55.36%, Val Acc: 59.35%, LR: 0.000005
Epoch [6/12], Loss: 0.9550, Train Acc: 56.58%, Val Acc: 60.00%, LR: 0.000005
Epoch [7/12], Loss: 0.9602, Train Acc: 54.57%, Val Acc: 61.29%, LR: 0.000005
Epoch [8/12], Loss: 0.9549, Train Acc: 56.22%, Val Acc: 63.23%, LR: 0.000005
Epoch [9/12], Loss: 0.9552, Train Acc: 55.86%, Val Acc: 63.23%, LR: 0.000005
Epoch [10/12], Loss: 0.9527, Train Acc: 55.64%, Val Acc: 61.94%, LR: 0.000005
Epoch [11/12], Loss: 0.9464, Train Acc: 56.00%, Val Acc: 61.94%, LR: 0.000005
Epoch [12/12], Loss: 0.9458, Train Acc: 58.23%, Val Acc: 62.58%, LR: 0.000005
Testing model...
```

```
==== split_30_70 Results ===
```

```
Accuracy: 0.5654
Precision: 0.5878
Recall: 0.5654
F1-Score: 0.4849
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.59	0.10	0.17	787
CN	0.65	0.24	0.34	1008
MCI	0.55	0.95	0.70	1813
accuracy			0.57	3608
macro avg	0.60	0.43	0.40	3608
weighted avg	0.59	0.57	0.48	3608

```
=====
=====
```

```
Processing: split_40_60
=====
```

```
Train samples: 1855
Val samples: 207
Test samples: 3092
```

```
Starting phase 1 training (partial unfreeze)...
Epoch [1/18], Loss: 1.0797, Train Acc: 44.47%, Val Acc: 50.24%, LR: 0.000033
Epoch [2/18], Loss: 1.0486, Train Acc: 50.30%, Val Acc: 50.24%, LR: 0.000067
Epoch [3/18], Loss: 1.0409, Train Acc: 50.35%, Val Acc: 50.24%, LR: 0.000100
Epoch [4/18], Loss: 1.0289, Train Acc: 50.94%, Val Acc: 52.17%, LR: 0.000099
Epoch [5/18], Loss: 1.0250, Train Acc: 51.64%, Val Acc: 53.62%, LR: 0.000096
Epoch [6/18], Loss: 1.0100, Train Acc: 52.61%, Val Acc: 52.66%, LR: 0.000091
Epoch [7/18], Loss: 0.9961, Train Acc: 52.88%, Val Acc: 53.14%, LR: 0.000084
Epoch [8/18], Loss: 0.9873, Train Acc: 53.64%, Val Acc: 53.14%, LR: 0.000075
Epoch [9/18], Loss: 0.9715, Train Acc: 55.63%, Val Acc: 56.04%, LR: 0.000066
Epoch [10/18], Loss: 0.9639, Train Acc: 57.04%, Val Acc: 56.52%, LR: 0.000056
Epoch [11/18], Loss: 0.9465, Train Acc: 58.76%, Val Acc: 57.97%, LR: 0.000045
Epoch [12/18], Loss: 0.9362, Train Acc: 58.87%, Val Acc: 59.90%, LR: 0.000035
Epoch [13/18], Loss: 0.9340, Train Acc: 58.38%, Val Acc: 56.52%, LR: 0.000026
Epoch [14/18], Loss: 0.9121, Train Acc: 61.29%, Val Acc: 56.52%, LR: 0.000017
Epoch [15/18], Loss: 0.9165, Train Acc: 60.49%, Val Acc: 57.97%, LR: 0.000010
Epoch [16/18], Loss: 0.9067, Train Acc: 61.83%, Val Acc: 60.39%, LR: 0.000005
Epoch [17/18], Loss: 0.9059, Train Acc: 60.86%, Val Acc: 61.35%, LR: 0.000002
Epoch [18/18], Loss: 0.9080, Train Acc: 60.49%, Val Acc: 58.45%, LR: 0.000001
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/12], Loss: 0.8994, Train Acc: 63.40%, Val Acc: 60.87%, LR: 0.000005
Epoch [2/12], Loss: 0.9134, Train Acc: 61.40%, Val Acc: 57.97%, LR: 0.000005
Epoch [3/12], Loss: 0.9025, Train Acc: 62.91%, Val Acc: 59.42%, LR: 0.000005
Epoch [4/12], Loss: 0.8985, Train Acc: 62.05%, Val Acc: 59.90%, LR: 0.000005
Epoch [5/12], Loss: 0.9022, Train Acc: 62.32%, Val Acc: 60.87%, LR: 0.000005
Epoch [6/12], Loss: 0.9060, Train Acc: 62.21%, Val Acc: 59.90%, LR: 0.000003
Epoch [7/12], Loss: 0.8898, Train Acc: 63.61%, Val Acc: 58.45%, LR: 0.000003
Epoch [8/12], Loss: 0.8981, Train Acc: 61.67%, Val Acc: 59.90%, LR: 0.000003
Epoch [9/12], Loss: 0.8890, Train Acc: 63.02%, Val Acc: 60.87%, LR: 0.000003
Early stopping at epoch 9
Testing model...
```

```
==== split_40_60 Results ====
Accuracy: 0.6239
Precision: 0.6387
Recall: 0.6239
F1-Score: 0.5858
```

Classification Report:

	precision	recall	f1-score	support
AD	0.74	0.19	0.30	674
CN	0.57	0.52	0.54	864

MCI	0.64	0.87	0.74	1554
accuracy			0.62	3092
macro avg	0.65	0.53	0.52	3092
weighted avg	0.64	0.62	0.59	3092

```
=====
Processing: split_50_50
=====
Train samples: 2319
Val samples: 258
Test samples: 2577
Starting phase 1 training (partial unfreeze)...
Epoch [1/18], Loss: 1.0744, Train Acc: 46.44%, Val Acc: 50.39%, LR: 0.000033
Epoch [2/18], Loss: 1.0458, Train Acc: 50.24%, Val Acc: 50.39%, LR: 0.000067
Epoch [3/18], Loss: 1.0324, Train Acc: 50.24%, Val Acc: 50.39%, LR: 0.000100
Epoch [4/18], Loss: 1.0265, Train Acc: 50.67%, Val Acc: 50.39%, LR: 0.000099
Epoch [5/18], Loss: 1.0176, Train Acc: 50.84%, Val Acc: 53.88%, LR: 0.000096
Epoch [6/18], Loss: 1.0034, Train Acc: 53.39%, Val Acc: 53.88%, LR: 0.000091
Epoch [7/18], Loss: 0.9831, Train Acc: 55.02%, Val Acc: 54.26%, LR: 0.000084
Epoch [8/18], Loss: 0.9658, Train Acc: 56.92%, Val Acc: 57.36%, LR: 0.000075
Epoch [9/18], Loss: 0.9356, Train Acc: 58.86%, Val Acc: 55.04%, LR: 0.000066
Epoch [10/18], Loss: 0.9329, Train Acc: 59.29%, Val Acc: 56.59%, LR: 0.000056
Epoch [11/18], Loss: 0.9146, Train Acc: 61.41%, Val Acc: 60.47%, LR: 0.000045
Epoch [12/18], Loss: 0.9034, Train Acc: 62.01%, Val Acc: 62.40%, LR: 0.000035
Epoch [13/18], Loss: 0.8995, Train Acc: 62.35%, Val Acc: 65.50%, LR: 0.000026
Epoch [14/18], Loss: 0.8950, Train Acc: 62.23%, Val Acc: 64.34%, LR: 0.000017
Epoch [15/18], Loss: 0.8815, Train Acc: 64.04%, Val Acc: 63.18%, LR: 0.000010
Epoch [16/18], Loss: 0.8909, Train Acc: 61.92%, Val Acc: 65.12%, LR: 0.000005
Epoch [17/18], Loss: 0.8784, Train Acc: 64.12%, Val Acc: 64.34%, LR: 0.000002
Epoch [18/18], Loss: 0.8814, Train Acc: 63.09%, Val Acc: 67.83%, LR: 0.000001
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/12], Loss: 0.8801, Train Acc: 62.83%, Val Acc: 64.73%, LR: 0.000005
Epoch [2/12], Loss: 0.8794, Train Acc: 63.04%, Val Acc: 64.34%, LR: 0.000005
Epoch [3/12], Loss: 0.8727, Train Acc: 62.70%, Val Acc: 65.50%, LR: 0.000005
Epoch [4/12], Loss: 0.8748, Train Acc: 64.08%, Val Acc: 67.83%, LR: 0.000005
Epoch [5/12], Loss: 0.8690, Train Acc: 64.42%, Val Acc: 69.38%, LR: 0.000005
Epoch [6/12], Loss: 0.8622, Train Acc: 64.98%, Val Acc: 69.38%, LR: 0.000005
Epoch [7/12], Loss: 0.8624, Train Acc: 65.33%, Val Acc: 68.60%, LR: 0.000005
Epoch [8/12], Loss: 0.8589, Train Acc: 64.90%, Val Acc: 67.83%, LR: 0.000005
Epoch [9/12], Loss: 0.8609, Train Acc: 65.24%, Val Acc: 68.22%, LR: 0.000005
Epoch [10/12], Loss: 0.8500, Train Acc: 64.98%, Val Acc: 67.05%, LR: 0.000003
```

```
Epoch [11/12], Loss: 0.8725, Train Acc: 64.21%, Val Acc: 67.44%, LR: 0.000003
Epoch [12/12], Loss: 0.8427, Train Acc: 67.31%, Val Acc: 70.54%, LR: 0.000003
Testing model...
```

```
==== split_50_50 Results ===
```

```
Accuracy: 0.6546
Precision: 0.6613
Recall: 0.6546
F1-Score: 0.6267
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.70	0.31	0.43	562
CN	0.65	0.47	0.55	720
MCI	0.65	0.91	0.76	1295
accuracy			0.65	2577
macro avg	0.67	0.56	0.58	2577
weighted avg	0.66	0.65	0.63	2577

```
=====
Processing: split_60_40
=====
```

```
Train samples: 2784
```

```
Val samples: 308
```

```
Test samples: 2062
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/18], Loss: 1.0683, Train Acc: 47.16%, Val Acc: 50.32%, LR: 0.000033
Epoch [2/18], Loss: 1.0404, Train Acc: 50.40%, Val Acc: 50.00%, LR: 0.000067
Epoch [3/18], Loss: 1.0331, Train Acc: 50.36%, Val Acc: 50.00%, LR: 0.000100
Epoch [4/18], Loss: 1.0201, Train Acc: 50.90%, Val Acc: 50.00%, LR: 0.000099
Epoch [5/18], Loss: 1.0122, Train Acc: 51.76%, Val Acc: 50.97%, LR: 0.000096
Epoch [6/18], Loss: 0.9950, Train Acc: 53.12%, Val Acc: 53.25%, LR: 0.000091
Epoch [7/18], Loss: 0.9775, Train Acc: 54.45%, Val Acc: 54.22%, LR: 0.000084
Epoch [8/18], Loss: 0.9535, Train Acc: 57.15%, Val Acc: 58.44%, LR: 0.000075
Epoch [9/18], Loss: 0.9353, Train Acc: 58.98%, Val Acc: 61.04%, LR: 0.000066
Epoch [10/18], Loss: 0.9212, Train Acc: 60.70%, Val Acc: 62.66%, LR: 0.000056
Epoch [11/18], Loss: 0.9014, Train Acc: 61.93%, Val Acc: 63.64%, LR: 0.000045
Epoch [12/18], Loss: 0.9095, Train Acc: 60.78%, Val Acc: 65.58%, LR: 0.000035
Epoch [13/18], Loss: 0.8870, Train Acc: 61.60%, Val Acc: 68.51%, LR: 0.000026
Epoch [14/18], Loss: 0.8795, Train Acc: 63.58%, Val Acc: 65.91%, LR: 0.000017
Epoch [15/18], Loss: 0.8831, Train Acc: 63.00%, Val Acc: 65.91%, LR: 0.000010
Epoch [16/18], Loss: 0.8798, Train Acc: 63.79%, Val Acc: 67.86%, LR: 0.000005
Epoch [17/18], Loss: 0.8861, Train Acc: 63.33%, Val Acc: 65.58%, LR: 0.000002
Epoch [18/18], Loss: 0.8630, Train Acc: 64.73%, Val Acc: 64.94%, LR: 0.000001
```

```
Starting phase 2 training (full fine-tuning)...
```

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.

    warnings.warn(
Epoch [1/12], Loss: 0.8718, Train Acc: 63.79%, Val Acc: 67.21%, LR: 0.000005
Epoch [2/12], Loss: 0.8690, Train Acc: 64.19%, Val Acc: 66.88%, LR: 0.000005
Epoch [3/12], Loss: 0.8549, Train Acc: 65.37%, Val Acc: 67.53%, LR: 0.000005
Epoch [4/12], Loss: 0.8663, Train Acc: 63.90%, Val Acc: 71.10%, LR: 0.000005
Epoch [5/12], Loss: 0.8625, Train Acc: 64.37%, Val Acc: 69.81%, LR: 0.000005
Epoch [6/12], Loss: 0.8556, Train Acc: 65.95%, Val Acc: 70.78%, LR: 0.000005
Epoch [7/12], Loss: 0.8582, Train Acc: 66.06%, Val Acc: 72.73%, LR: 0.000005
Epoch [8/12], Loss: 0.8441, Train Acc: 66.77%, Val Acc: 69.16%, LR: 0.000005
Epoch [9/12], Loss: 0.8324, Train Acc: 67.06%, Val Acc: 70.45%, LR: 0.000005
Epoch [10/12], Loss: 0.8455, Train Acc: 66.06%, Val Acc: 69.48%, LR: 0.000005
Epoch [11/12], Loss: 0.8322, Train Acc: 67.03%, Val Acc: 67.86%, LR: 0.000005
Epoch [12/12], Loss: 0.8299, Train Acc: 66.88%, Val Acc: 70.78%, LR: 0.000003
Testing model...

==== split_60_40 Results ===
Accuracy: 0.6848
Precision: 0.6808
Recall: 0.6848
F1-Score: 0.6704

```

	precision	recall	f1-score	support
AD	0.64	0.42	0.51	450
CN	0.69	0.55	0.61	576
MCI	0.69	0.87	0.77	1036
accuracy			0.68	2062
macro avg	0.67	0.62	0.63	2062
weighted avg	0.68	0.68	0.67	2062

```

=====
Processing: split_70_30
=====
Train samples: 3247
Val samples: 361
Test samples: 1546
Starting phase 1 training (partial unfreeze)...
Epoch [1/18], Loss: 1.0636, Train Acc: 48.01%, Val Acc: 50.14%, LR: 0.000033
Epoch [2/18], Loss: 1.0439, Train Acc: 50.32%, Val Acc: 50.14%, LR: 0.000067
Epoch [3/18], Loss: 1.0398, Train Acc: 50.20%, Val Acc: 50.14%, LR: 0.000100
Epoch [4/18], Loss: 1.0278, Train Acc: 50.79%, Val Acc: 50.14%, LR: 0.000099

```

```
Epoch [5/18], Loss: 1.0066, Train Acc: 51.59%, Val Acc: 52.35%, LR: 0.000096
Epoch [6/18], Loss: 0.9865, Train Acc: 54.17%, Val Acc: 57.89%, LR: 0.000091
Epoch [7/18], Loss: 0.9532, Train Acc: 56.76%, Val Acc: 58.17%, LR: 0.000084
Epoch [8/18], Loss: 0.9453, Train Acc: 57.13%, Val Acc: 60.94%, LR: 0.000075
Epoch [9/18], Loss: 0.9217, Train Acc: 60.33%, Val Acc: 63.99%, LR: 0.000066
Epoch [10/18], Loss: 0.9149, Train Acc: 61.04%, Val Acc: 66.20%, LR: 0.000056
Epoch [11/18], Loss: 0.8992, Train Acc: 61.60%, Val Acc: 64.82%, LR: 0.000045
Epoch [12/18], Loss: 0.8853, Train Acc: 62.49%, Val Acc: 67.31%, LR: 0.000035
Epoch [13/18], Loss: 0.8660, Train Acc: 63.60%, Val Acc: 66.20%, LR: 0.000026
Epoch [14/18], Loss: 0.8750, Train Acc: 63.66%, Val Acc: 67.59%, LR: 0.000017
Epoch [15/18], Loss: 0.8664, Train Acc: 64.64%, Val Acc: 67.59%, LR: 0.000010
Epoch [16/18], Loss: 0.8495, Train Acc: 65.29%, Val Acc: 67.59%, LR: 0.000005
Epoch [17/18], Loss: 0.8575, Train Acc: 65.26%, Val Acc: 67.04%, LR: 0.000002
Epoch [18/18], Loss: 0.8598, Train Acc: 65.60%, Val Acc: 66.76%, LR: 0.000001
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/12], Loss: 0.8516, Train Acc: 66.40%, Val Acc: 66.76%, LR: 0.000005
Epoch [2/12], Loss: 0.8378, Train Acc: 66.34%, Val Acc: 66.48%, LR: 0.000005
Epoch [3/12], Loss: 0.8473, Train Acc: 66.28%, Val Acc: 66.76%, LR: 0.000005
Epoch [4/12], Loss: 0.8387, Train Acc: 65.66%, Val Acc: 67.31%, LR: 0.000005
Epoch [5/12], Loss: 0.8481, Train Acc: 66.83%, Val Acc: 67.04%, LR: 0.000005
Epoch [6/12], Loss: 0.8397, Train Acc: 66.15%, Val Acc: 68.42%, LR: 0.000005
Epoch [7/12], Loss: 0.8273, Train Acc: 67.23%, Val Acc: 68.14%, LR: 0.000005
Epoch [8/12], Loss: 0.8343, Train Acc: 66.68%, Val Acc: 68.14%, LR: 0.000005
Epoch [9/12], Loss: 0.8207, Train Acc: 67.79%, Val Acc: 68.14%, LR: 0.000005
Epoch [10/12], Loss: 0.8247, Train Acc: 67.69%, Val Acc: 68.42%, LR: 0.000005
Epoch [11/12], Loss: 0.8221, Train Acc: 68.37%, Val Acc: 67.59%, LR: 0.000003
Epoch [12/12], Loss: 0.8088, Train Acc: 69.11%, Val Acc: 68.98%, LR: 0.000003
Testing model...
```

```
==== split_70_30 Results ===
```

```
Accuracy: 0.7160
Precision: 0.7100
Recall: 0.7160
F1-Score: 0.7059
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.66	0.45	0.53	337
CN	0.71	0.66	0.68	432
MCI	0.74	0.86	0.79	777
accuracy			0.72	1546

macro avg	0.70	0.66	0.67	1546
weighted avg	0.71	0.72	0.71	1546

```
=====
Processing: split_80_20
=====
Train samples: 3711
Val samples: 412
Test samples: 1031
Starting phase 1 training (partial unfreeze)...
Epoch [1/18], Loss: 1.0625, Train Acc: 47.45%, Val Acc: 50.24%, LR: 0.000033
Epoch [2/18], Loss: 1.0434, Train Acc: 50.42%, Val Acc: 50.24%, LR: 0.000067
Epoch [3/18], Loss: 1.0309, Train Acc: 50.47%, Val Acc: 50.97%, LR: 0.000100
Epoch [4/18], Loss: 1.0209, Train Acc: 51.09%, Val Acc: 52.67%, LR: 0.000099
Epoch [5/18], Loss: 1.0002, Train Acc: 53.27%, Val Acc: 54.13%, LR: 0.000096
Epoch [6/18], Loss: 0.9680, Train Acc: 55.56%, Val Acc: 59.71%, LR: 0.000091
Epoch [7/18], Loss: 0.9493, Train Acc: 57.80%, Val Acc: 58.74%, LR: 0.000084
Epoch [8/18], Loss: 0.9205, Train Acc: 59.71%, Val Acc: 59.95%, LR: 0.000075
Epoch [9/18], Loss: 0.8889, Train Acc: 62.38%, Val Acc: 65.29%, LR: 0.000066
Epoch [10/18], Loss: 0.8732, Train Acc: 63.86%, Val Acc: 65.78%, LR: 0.000056
Epoch [11/18], Loss: 0.8732, Train Acc: 63.97%, Val Acc: 69.90%, LR: 0.000045
Epoch [12/18], Loss: 0.8527, Train Acc: 64.67%, Val Acc: 68.45%, LR: 0.000035
Epoch [13/18], Loss: 0.8482, Train Acc: 65.94%, Val Acc: 68.45%, LR: 0.000026
Epoch [14/18], Loss: 0.8382, Train Acc: 66.83%, Val Acc: 68.69%, LR: 0.000017
Epoch [15/18], Loss: 0.8374, Train Acc: 66.88%, Val Acc: 68.45%, LR: 0.000010
Epoch [16/18], Loss: 0.8316, Train Acc: 66.61%, Val Acc: 70.15%, LR: 0.000005
Epoch [17/18], Loss: 0.8325, Train Acc: 67.34%, Val Acc: 69.42%, LR: 0.000002
Epoch [18/18], Loss: 0.8248, Train Acc: 67.56%, Val Acc: 69.17%, LR: 0.000001
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/12], Loss: 0.8266, Train Acc: 67.80%, Val Acc: 72.09%, LR: 0.000005
Epoch [2/12], Loss: 0.8251, Train Acc: 66.91%, Val Acc: 71.36%, LR: 0.000005
Epoch [3/12], Loss: 0.8192, Train Acc: 67.48%, Val Acc: 71.60%, LR: 0.000005
Epoch [4/12], Loss: 0.8228, Train Acc: 67.91%, Val Acc: 70.39%, LR: 0.000005
Epoch [5/12], Loss: 0.8112, Train Acc: 68.96%, Val Acc: 70.87%, LR: 0.000005
Epoch [6/12], Loss: 0.7917, Train Acc: 69.82%, Val Acc: 71.12%, LR: 0.000003
Epoch [7/12], Loss: 0.7940, Train Acc: 69.44%, Val Acc: 72.33%, LR: 0.000003
Epoch [8/12], Loss: 0.7859, Train Acc: 70.22%, Val Acc: 71.36%, LR: 0.000003
Epoch [9/12], Loss: 0.8089, Train Acc: 67.74%, Val Acc: 71.60%, LR: 0.000003
Epoch [10/12], Loss: 0.7979, Train Acc: 69.31%, Val Acc: 72.09%, LR: 0.000003
Epoch [11/12], Loss: 0.7857, Train Acc: 69.60%, Val Acc: 72.09%, LR: 0.000003
Epoch [12/12], Loss: 0.7869, Train Acc: 70.52%, Val Acc: 71.84%, LR: 0.000001
Testing model...
```

```
==== split_80_20 Results ====
Accuracy: 0.7255
Precision: 0.7198
Recall: 0.7255
F1-Score: 0.7083
```

Classification Report:

	precision	recall	f1-score	support
AD	0.69	0.38	0.49	225
CN	0.71	0.66	0.68	288
MCI	0.74	0.91	0.82	518
accuracy			0.73	1031
macro avg	0.71	0.65	0.66	1031
weighted avg	0.72	0.73	0.71	1031

```
=====
Processing: split_90_10
=====
```

Train samples: 4175

Val samples: 464

Test samples: 515

Starting phase 1 training (partial unfreeze)...

```
Epoch [1/18], Loss: 1.0567, Train Acc: 49.75%, Val Acc: 50.22%, LR: 0.000033
Epoch [2/18], Loss: 1.0379, Train Acc: 50.35%, Val Acc: 50.65%, LR: 0.000067
Epoch [3/18], Loss: 1.0292, Train Acc: 50.49%, Val Acc: 51.29%, LR: 0.000100
Epoch [4/18], Loss: 1.0097, Train Acc: 51.93%, Val Acc: 52.16%, LR: 0.000099
Epoch [5/18], Loss: 0.9791, Train Acc: 54.85%, Val Acc: 54.96%, LR: 0.000096
Epoch [6/18], Loss: 0.9514, Train Acc: 56.77%, Val Acc: 59.05%, LR: 0.000091
Epoch [7/18], Loss: 0.9312, Train Acc: 59.21%, Val Acc: 60.56%, LR: 0.000084
Epoch [8/18], Loss: 0.9063, Train Acc: 61.58%, Val Acc: 61.42%, LR: 0.000075
Epoch [9/18], Loss: 0.8784, Train Acc: 63.52%, Val Acc: 65.09%, LR: 0.000066
Epoch [10/18], Loss: 0.8629, Train Acc: 66.06%, Val Acc: 67.03%, LR: 0.000056
Epoch [11/18], Loss: 0.8647, Train Acc: 65.03%, Val Acc: 66.81%, LR: 0.000045
Epoch [12/18], Loss: 0.8421, Train Acc: 66.71%, Val Acc: 66.81%, LR: 0.000035
Epoch [13/18], Loss: 0.8373, Train Acc: 66.78%, Val Acc: 65.95%, LR: 0.000026
Epoch [14/18], Loss: 0.8241, Train Acc: 67.31%, Val Acc: 70.47%, LR: 0.000017
Epoch [15/18], Loss: 0.8256, Train Acc: 67.81%, Val Acc: 68.10%, LR: 0.000010
Epoch [16/18], Loss: 0.8182, Train Acc: 68.17%, Val Acc: 66.81%, LR: 0.000005
Epoch [17/18], Loss: 0.8128, Train Acc: 69.34%, Val Acc: 68.53%, LR: 0.000002
Epoch [18/18], Loss: 0.8140, Train Acc: 69.17%, Val Acc: 70.47%, LR: 0.000001
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
```

```

access the learning rate.
warnings.warn(
Epoch [1/12], Loss: 0.8114, Train Acc: 68.91%, Val Acc: 72.84%, LR: 0.000005
Epoch [2/12], Loss: 0.8168, Train Acc: 68.43%, Val Acc: 70.69%, LR: 0.000005
Epoch [3/12], Loss: 0.7996, Train Acc: 69.87%, Val Acc: 70.47%, LR: 0.000005
Epoch [4/12], Loss: 0.7956, Train Acc: 69.56%, Val Acc: 71.55%, LR: 0.000005
Epoch [5/12], Loss: 0.7989, Train Acc: 69.89%, Val Acc: 70.69%, LR: 0.000005
Epoch [6/12], Loss: 0.7967, Train Acc: 70.49%, Val Acc: 71.55%, LR: 0.000003
Epoch [7/12], Loss: 0.7897, Train Acc: 69.44%, Val Acc: 71.12%, LR: 0.000003
Epoch [8/12], Loss: 0.7809, Train Acc: 70.66%, Val Acc: 72.20%, LR: 0.000003
Epoch [9/12], Loss: 0.7743, Train Acc: 70.95%, Val Acc: 71.55%, LR: 0.000003
Early stopping at epoch 9
Testing model...

```

```

==== split_90_10 Results ===
Accuracy: 0.7262
Precision: 0.7293
Recall: 0.7262
F1-Score: 0.7158

```

Classification Report:

	precision	recall	f1-score	support
AD	0.75	0.42	0.54	112
CN	0.65	0.75	0.70	144
MCI	0.77	0.85	0.80	259
accuracy			0.73	515
macro avg	0.72	0.67	0.68	515
weighted avg	0.73	0.73	0.72	515

=====
EfficientNet-B4 - SUMMARY OF ALL SPLITS
=====

split	accuracy	precision	recall	f1_score	training_time
split_10_90	0.503126	0.412399	0.503126	0.338000	209.842631
split_20_80	0.503517	0.476386	0.503517	0.338370	363.769990
split_30_70	0.565410	0.587784	0.565410	0.484915	840.035301
split_40_60	0.623868	0.638671	0.623868	0.585815	950.033990
split_50_50	0.654637	0.661264	0.654637	0.626693	1494.124134
split_60_40	0.684772	0.680800	0.684772	0.670393	1763.249976
split_70_30	0.716041	0.710036	0.716041	0.705865	2104.602802
split_80_20	0.725509	0.719785	0.725509	0.708299	2544.441258
split_90_10	0.726214	0.729257	0.726214	0.715817	2554.996111

Detailed results saved to: /kaggle/working/efficientnet_b4_results.csv

```
[6]: import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, models
import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import CosineAnnealingLR, ReduceLROnPlateau

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)
        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label

    def get_data_transforms():
        train_transform = transforms.Compose([
            transforms.Resize((384, 384)),
            transforms.RandomHorizontalFlip(p=0.5),
            transforms.RandomRotation(10),
            transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
            transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
```

```

        transforms.RandomGrayscale(p=0.1),
        transforms.GaussianBlur(kernel_size=3, sigma=(0.1, 2.0)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
        ↪225]),
        transforms.RandomErasing(p=0.2, scale=(0.02, 0.2), ratio=(0.3, 3.3))
    ])

    val_transform = transforms.Compose([
        transforms.Resize((384, 384)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
        ↪225])
])

return train_transform, val_transform
}

def train_model(model, train_loader, val_loader, criterion, optimizer, ↪
    ↪scheduler, num_epochs, device, warmup_epochs=3):
    best_val_acc = 0
    patience = 8
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        # Learning rate warmup
        if epoch < warmup_epochs:
            lr_scale = min(1.0, float(epoch + 1) / warmup_epochs)
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['initial_lr'] * lr_scale

        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()

            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
            optimizer.step()

            running_loss += loss.item()
    
```

```

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    epoch_loss = running_loss / len(train_loader)
    epoch_acc = 100 * correct / total

    val_acc = evaluate_model(model, val_loader, device)

    if scheduler and epoch >= warmup_epochs:
        if isinstance(scheduler, CosineAnnealingLR):
            scheduler.step()
        else:
            scheduler.step(val_acc)

    current_lr = optimizer.param_groups[0]['lr']
    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train_U
→Acc: {epoch_acc:.2f}%, Val Acc: {val_acc:.2f}%, LR: {current_lr:.6f}')
```

```

    if val_acc > best_val_acc:
        best_val_acc = val_acc
        patience_counter = 0
    else:
        patience_counter += 1

    if patience_counter >= patience:
        print("Early stopping at epoch {epoch+1}")
        break

def evaluate_model(model, data_loader, device):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in data_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    return accuracy

def test_model(model, test_loader, device):
    model.eval()
```

```

all_preds = []
all_labels = []

with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)

        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

return all_preds, all_labels

def calculate_metrics(y_true, y_pred, split_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred, average='weighted', zero_division=0)
    recall = recall_score(y_true, y_pred, average='weighted', zero_division=0)
    f1 = f1_score(y_true, y_pred, average='weighted', zero_division=0)

    print(f"\n==== {split_name} Results ===")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-Score: {f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=['AD', 'CN', 'MCI'], zero_division=0))

    return {
        'split': split_name,
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1_score': f1
    }

def run_efficientnetv2_s_on_splits():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    splits_root = "/kaggle/working/alzheimer-resized-224_splits"
    split_folders = [f for f in os.listdir(splits_root) if f.startswith('split_')]
    split_folders.sort()

```

```

results = []

train_transform, val_transform = get_data_transforms()

for split_folder in split_folders:
    print(f"\n{'='*60}")
    print(f"Processing: {split_folder}")
    print(f"{'='*60}")

    split_path = os.path.join(splits_root, split_folder)

    train_dataset = AlzheimerDataset(split_path, 'train', train_transform)
    val_dataset = AlzheimerDataset(split_path, 'val', val_transform)
    test_dataset = AlzheimerDataset(split_path, 'test', val_transform)

    train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True,
                             num_workers=2, pin_memory=True)
    val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False,
                           num_workers=2, pin_memory=True)
    test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False,
                            num_workers=2, pin_memory=True)

    print(f"Train samples: {len(train_dataset)}")
    print(f"Val samples: {len(val_dataset)}")
    print(f"Test samples: {len(test_dataset)}")

model = models.efficientnet_v2_s(weights=models.
EfficientNet_V2_S_Weights.IMAGENET1K_V1)

# Freeze all layers initially
for param in model.parameters():
    param.requires_grad = False

# Unfreeze classifier
num_ftrs = model.classifier[1].in_features
model.classifier[1] = nn.Linear(num_ftrs, 3)

# Unfreeze last blocks for feature adaptation
for param in model.features[-3:].parameters(): # Last 3 blocks
    param.requires_grad = True
for param in model.features[-4][-2:].parameters(): # Last part of 4th
    param.requires_grad = True

model = model.to(device)

# Layer-wise learning rates for EfficientNetV2

```

```

optimizer = torch.optim.AdamW([
    {'params': model.classifier.parameters(), 'lr': 0.0001, ↴
    'initial_lr': 0.0001},
    {'params': model.features[-1].parameters(), 'lr': 0.00005, ↴
    'initial_lr': 0.00005},
    {'params': model.features[-2].parameters(), 'lr': 0.00003, ↴
    'initial_lr': 0.00003},
    {'params': model.features[-3].parameters(), 'lr': 0.00002, ↴
    'initial_lr': 0.00002},
    {'params': model.features[-4][-2:].parameters(), 'lr': 0.00001, ↴
    'initial_lr': 0.00001},
], weight_decay=1e-5)

criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
scheduler = CosineAnnealingLR(optimizer, T_max=20, eta_min=1e-6)

print("Starting phase 1 training (partial unfreeze)...")
start_time = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, ↴
scheduler, num_epochs=20, device=device)
phase1_time = time.time() - start_time

print("Starting phase 2 training (full fine-tuning)...")
for param in model.parameters():
    param.requires_grad = True

optimizer = torch.optim.AdamW(model.parameters(), lr=0.000005, ↴
weight_decay=1e-6)
scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5, ↴
patience=4, verbose=True)

start_time_phase2 = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, ↴
scheduler, num_epochs=15, device=device, warmup_epochs=0)
total_time = phase1_time + (time.time() - start_time_phase2)

print("Testing model...")
test_preds, test_labels = test_model(model, test_loader, device)

split_results = calculate_metrics(test_labels, test_preds, split_folder)
split_results['training_time'] = total_time
results.append(split_results)

torch.cuda.empty_cache()

results_df = pd.DataFrame(results)

```

```

print(f"\n{'='*80}")
print("EfficientNetV2-S - SUMMARY OF ALL SPLITS")
print(f"{'='*80}")
print(results_df.to_string(index=False))

results_csv_path = "/kaggle/working/efficientnetv2_s_results.csv"
results_df.to_csv(results_csv_path, index=False)
print(f"\nDetailed results saved to: {results_csv_path}")

return results_df

if __name__ == "__main__":
    results = run_efficientnetv2_s_on_splits()

```

Using device: cuda

```

=====
Processing: split_10_90
=====
Train samples: 464
Val samples: 51
Test samples: 4639

Downloading: "https://download.pytorch.org/models/efficientnet_v2_s-
dd5fe13b.pth" to /root/.cache/torch/hub/checkpoints/efficientnet_v2_s-
dd5fe13b.pth
100%| 82.7M/82.7M [00:00<00:00, 206MB/s]

Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.1189, Train Acc: 31.25%, Val Acc: 25.49%, LR: 0.000033
Epoch [2/20], Loss: 1.0671, Train Acc: 42.89%, Val Acc: 47.06%, LR: 0.000067
Epoch [3/20], Loss: 1.0449, Train Acc: 49.14%, Val Acc: 50.98%, LR: 0.000100
Epoch [4/20], Loss: 1.0263, Train Acc: 50.65%, Val Acc: 49.02%, LR: 0.000099
Epoch [5/20], Loss: 1.0197, Train Acc: 50.22%, Val Acc: 50.98%, LR: 0.000098
Epoch [6/20], Loss: 1.0034, Train Acc: 51.29%, Val Acc: 52.94%, LR: 0.000095
Epoch [7/20], Loss: 1.0005, Train Acc: 50.86%, Val Acc: 50.98%, LR: 0.000091
Epoch [8/20], Loss: 0.9832, Train Acc: 53.02%, Val Acc: 41.18%, LR: 0.000086
Epoch [9/20], Loss: 0.9601, Train Acc: 55.39%, Val Acc: 52.94%, LR: 0.000080
Epoch [10/20], Loss: 0.9389, Train Acc: 58.62%, Val Acc: 52.94%, LR: 0.000073
Epoch [11/20], Loss: 0.9256, Train Acc: 60.13%, Val Acc: 50.98%, LR: 0.000066
Epoch [12/20], Loss: 0.8987, Train Acc: 61.21%, Val Acc: 50.98%, LR: 0.000058
Epoch [13/20], Loss: 0.8979, Train Acc: 60.99%, Val Acc: 52.94%, LR: 0.000051
Epoch [14/20], Loss: 0.8558, Train Acc: 66.38%, Val Acc: 52.94%, LR: 0.000043
Early stopping at epoch 14
Starting phase 2 training (full fine-tuning)...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```

warnings.warn(
Epoch [1/15], Loss: 0.8802, Train Acc: 64.22%, Val Acc: 54.90%, LR: 0.000005
Epoch [2/15], Loss: 0.8575, Train Acc: 67.24%, Val Acc: 56.86%, LR: 0.000005
Epoch [3/15], Loss: 0.8376, Train Acc: 67.67%, Val Acc: 62.75%, LR: 0.000005
Epoch [4/15], Loss: 0.8378, Train Acc: 67.24%, Val Acc: 56.86%, LR: 0.000005
Epoch [5/15], Loss: 0.8278, Train Acc: 69.18%, Val Acc: 58.82%, LR: 0.000005
Epoch [6/15], Loss: 0.8212, Train Acc: 69.83%, Val Acc: 54.90%, LR: 0.000005
Epoch [7/15], Loss: 0.8184, Train Acc: 69.83%, Val Acc: 56.86%, LR: 0.000005
Epoch [8/15], Loss: 0.8036, Train Acc: 70.26%, Val Acc: 52.94%, LR: 0.000003
Epoch [9/15], Loss: 0.7915, Train Acc: 70.91%, Val Acc: 56.86%, LR: 0.000003
Epoch [10/15], Loss: 0.8184, Train Acc: 70.04%, Val Acc: 49.02%, LR: 0.000003
Epoch [11/15], Loss: 0.8113, Train Acc: 70.26%, Val Acc: 54.90%, LR: 0.000003
Early stopping at epoch 11
Testing model...

```

```

==== split_10_90 Results ===
Accuracy: 0.5756
Precision: 0.5856
Recall: 0.5756
F1-Score: 0.5794

```

Classification Report:

	precision	recall	f1-score	support
AD	0.42	0.45	0.44	1012
CN	0.50	0.55	0.53	1296
MCI	0.70	0.64	0.67	2331
accuracy			0.58	4639
macro avg	0.54	0.55	0.55	4639
weighted avg	0.59	0.58	0.58	4639

```
=====
Processing: split_20_80
=====
```

```
Train samples: 928
```

```
Val samples: 103
```

```
Test samples: 4123
```

```
Starting phase 1 training (partial unfreeze)...
```

```

Epoch [1/20], Loss: 1.0824, Train Acc: 41.38%, Val Acc: 43.69%, LR: 0.000033
Epoch [2/20], Loss: 1.0547, Train Acc: 50.00%, Val Acc: 50.49%, LR: 0.000067
Epoch [3/20], Loss: 1.0410, Train Acc: 49.46%, Val Acc: 48.54%, LR: 0.000100
Epoch [4/20], Loss: 1.0322, Train Acc: 50.32%, Val Acc: 51.46%, LR: 0.000099
Epoch [5/20], Loss: 1.0167, Train Acc: 51.19%, Val Acc: 55.34%, LR: 0.000098
Epoch [6/20], Loss: 1.0000, Train Acc: 53.56%, Val Acc: 54.37%, LR: 0.000095
Epoch [7/20], Loss: 0.9780, Train Acc: 54.42%, Val Acc: 54.37%, LR: 0.000091

```

```
Epoch [8/20], Loss: 0.9630, Train Acc: 56.36%, Val Acc: 57.28%, LR: 0.000086
Epoch [9/20], Loss: 0.9191, Train Acc: 60.02%, Val Acc: 55.34%, LR: 0.000080
Epoch [10/20], Loss: 0.9157, Train Acc: 60.45%, Val Acc: 56.31%, LR: 0.000073
Epoch [11/20], Loss: 0.8950, Train Acc: 62.07%, Val Acc: 58.25%, LR: 0.000066
Epoch [12/20], Loss: 0.8671, Train Acc: 64.22%, Val Acc: 60.19%, LR: 0.000058
Epoch [13/20], Loss: 0.8614, Train Acc: 66.16%, Val Acc: 58.25%, LR: 0.000051
Epoch [14/20], Loss: 0.8357, Train Acc: 65.09%, Val Acc: 59.22%, LR: 0.000043
Epoch [15/20], Loss: 0.8248, Train Acc: 67.46%, Val Acc: 61.17%, LR: 0.000035
Epoch [16/20], Loss: 0.8003, Train Acc: 67.67%, Val Acc: 63.11%, LR: 0.000028
Epoch [17/20], Loss: 0.7831, Train Acc: 68.97%, Val Acc: 58.25%, LR: 0.000021
Epoch [18/20], Loss: 0.7808, Train Acc: 71.12%, Val Acc: 64.08%, LR: 0.000015
Epoch [19/20], Loss: 0.7726, Train Acc: 70.69%, Val Acc: 65.05%, LR: 0.000010
Epoch [20/20], Loss: 0.7500, Train Acc: 73.28%, Val Acc: 66.99%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.7805, Train Acc: 69.61%, Val Acc: 67.96%, LR: 0.000005
Epoch [2/15], Loss: 0.7589, Train Acc: 70.80%, Val Acc: 65.05%, LR: 0.000005
Epoch [3/15], Loss: 0.7349, Train Acc: 74.68%, Val Acc: 67.96%, LR: 0.000005
Epoch [4/15], Loss: 0.7561, Train Acc: 72.41%, Val Acc: 65.05%, LR: 0.000005
Epoch [5/15], Loss: 0.7438, Train Acc: 73.49%, Val Acc: 66.02%, LR: 0.000005
Epoch [6/15], Loss: 0.7365, Train Acc: 73.28%, Val Acc: 68.93%, LR: 0.000005
Epoch [7/15], Loss: 0.7202, Train Acc: 74.46%, Val Acc: 66.99%, LR: 0.000005
Epoch [8/15], Loss: 0.7097, Train Acc: 76.94%, Val Acc: 70.87%, LR: 0.000005
Epoch [9/15], Loss: 0.7127, Train Acc: 74.89%, Val Acc: 69.90%, LR: 0.000005
Epoch [10/15], Loss: 0.7247, Train Acc: 75.65%, Val Acc: 66.02%, LR: 0.000005
Epoch [11/15], Loss: 0.7204, Train Acc: 75.75%, Val Acc: 68.93%, LR: 0.000005
Epoch [12/15], Loss: 0.6925, Train Acc: 78.12%, Val Acc: 67.96%, LR: 0.000005
Epoch [13/15], Loss: 0.6980, Train Acc: 76.29%, Val Acc: 68.93%, LR: 0.000003
Epoch [14/15], Loss: 0.6736, Train Acc: 78.66%, Val Acc: 69.90%, LR: 0.000003
Epoch [15/15], Loss: 0.6872, Train Acc: 77.26%, Val Acc: 70.87%, LR: 0.000003
Testing model...
```

```
==== split_20_80 Results ===
Accuracy: 0.7080
Precision: 0.7014
Recall: 0.7080
F1-Score: 0.7024
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.61	0.55	0.58	899
CN	0.67	0.58	0.62	1152
MCI	0.76	0.85	0.80	2072

accuracy			0.71	4123
macro avg	0.68	0.66	0.67	4123
weighted avg	0.70	0.71	0.70	4123

```
=====
Processing: split_30_70
=====
```

Train samples: 1391

Val samples: 155

Test samples: 3608

Starting phase 1 training (partial unfreeze)...

Epoch [1/20], Loss: 1.0879, Train Acc: 38.46%, Val Acc: 46.45%, LR: 0.000033
 Epoch [2/20], Loss: 1.0427, Train Acc: 49.68%, Val Acc: 49.03%, LR: 0.000067
 Epoch [3/20], Loss: 1.0267, Train Acc: 50.25%, Val Acc: 53.55%, LR: 0.000100
 Epoch [4/20], Loss: 1.0070, Train Acc: 51.76%, Val Acc: 54.19%, LR: 0.000099
 Epoch [5/20], Loss: 0.9824, Train Acc: 54.64%, Val Acc: 59.35%, LR: 0.000098
 Epoch [6/20], Loss: 0.9626, Train Acc: 56.00%, Val Acc: 67.10%, LR: 0.000095
 Epoch [7/20], Loss: 0.9238, Train Acc: 59.74%, Val Acc: 59.35%, LR: 0.000091
 Epoch [8/20], Loss: 0.8766, Train Acc: 61.90%, Val Acc: 70.97%, LR: 0.000086
 Epoch [9/20], Loss: 0.8487, Train Acc: 66.14%, Val Acc: 70.97%, LR: 0.000080
 Epoch [10/20], Loss: 0.8314, Train Acc: 66.79%, Val Acc: 67.10%, LR: 0.000073
 Epoch [11/20], Loss: 0.7920, Train Acc: 69.73%, Val Acc: 71.61%, LR: 0.000066
 Epoch [12/20], Loss: 0.7884, Train Acc: 69.73%, Val Acc: 65.16%, LR: 0.000058
 Epoch [13/20], Loss: 0.7414, Train Acc: 73.69%, Val Acc: 74.84%, LR: 0.000051
 Epoch [14/20], Loss: 0.7347, Train Acc: 72.83%, Val Acc: 70.97%, LR: 0.000043
 Epoch [15/20], Loss: 0.7270, Train Acc: 74.55%, Val Acc: 77.42%, LR: 0.000035
 Epoch [16/20], Loss: 0.6959, Train Acc: 75.92%, Val Acc: 75.48%, LR: 0.000028
 Epoch [17/20], Loss: 0.6719, Train Acc: 77.43%, Val Acc: 75.48%, LR: 0.000021
 Epoch [18/20], Loss: 0.6736, Train Acc: 77.57%, Val Acc: 75.48%, LR: 0.000015
 Epoch [19/20], Loss: 0.6954, Train Acc: 75.05%, Val Acc: 74.19%, LR: 0.000010
 Epoch [20/20], Loss: 0.6706, Train Acc: 78.58%, Val Acc: 71.61%, LR: 0.000006

Starting phase 2 training (full fine-tuning)...

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:  

UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to  

access the learning rate.
```

warnings.warn(

Epoch [1/15], Loss: 0.6679, Train Acc: 77.71%, Val Acc: 73.55%, LR: 0.000005
 Epoch [2/15], Loss: 0.6471, Train Acc: 79.22%, Val Acc: 73.55%, LR: 0.000005
 Epoch [3/15], Loss: 0.6386, Train Acc: 80.73%, Val Acc: 76.77%, LR: 0.000005
 Epoch [4/15], Loss: 0.6575, Train Acc: 77.93%, Val Acc: 78.06%, LR: 0.000005
 Epoch [5/15], Loss: 0.6439, Train Acc: 80.16%, Val Acc: 73.55%, LR: 0.000005
 Epoch [6/15], Loss: 0.6273, Train Acc: 80.30%, Val Acc: 77.42%, LR: 0.000005
 Epoch [7/15], Loss: 0.6385, Train Acc: 79.65%, Val Acc: 76.77%, LR: 0.000005
 Epoch [8/15], Loss: 0.6246, Train Acc: 80.95%, Val Acc: 76.77%, LR: 0.000005
 Epoch [9/15], Loss: 0.6270, Train Acc: 80.52%, Val Acc: 78.06%, LR: 0.000003

```
Epoch [10/15], Loss: 0.6048, Train Acc: 82.31%, Val Acc: 75.48%, LR: 0.000003
Epoch [11/15], Loss: 0.6071, Train Acc: 82.31%, Val Acc: 78.71%, LR: 0.000003
Epoch [12/15], Loss: 0.6025, Train Acc: 82.67%, Val Acc: 81.29%, LR: 0.000003
Epoch [13/15], Loss: 0.5998, Train Acc: 83.90%, Val Acc: 80.00%, LR: 0.000003
Epoch [14/15], Loss: 0.6034, Train Acc: 81.60%, Val Acc: 80.00%, LR: 0.000003
Epoch [15/15], Loss: 0.6003, Train Acc: 81.38%, Val Acc: 76.13%, LR: 0.000003
Testing model...
```

```
==== split_30_70 Results ===
```

```
Accuracy: 0.8004
```

```
Precision: 0.7985
```

```
Recall: 0.8004
```

```
F1-Score: 0.7989
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.75	0.67	0.71	787
CN	0.76	0.76	0.76	1008
MCI	0.84	0.88	0.86	1813
accuracy			0.80	3608
macro avg	0.78	0.77	0.78	3608
weighted avg	0.80	0.80	0.80	3608

```
=====
Processing: split_40_60
=====
```

```
Train samples: 1855
```

```
Val samples: 207
```

```
Test samples: 3092
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/20], Loss: 1.0677, Train Acc: 45.44%, Val Acc: 49.76%, LR: 0.000033
Epoch [2/20], Loss: 1.0400, Train Acc: 49.76%, Val Acc: 52.17%, LR: 0.000067
Epoch [3/20], Loss: 1.0213, Train Acc: 51.32%, Val Acc: 53.14%, LR: 0.000100
Epoch [4/20], Loss: 0.9992, Train Acc: 53.15%, Val Acc: 54.59%, LR: 0.000099
Epoch [5/20], Loss: 0.9605, Train Acc: 56.33%, Val Acc: 60.87%, LR: 0.000098
Epoch [6/20], Loss: 0.9189, Train Acc: 60.97%, Val Acc: 64.73%, LR: 0.000095
Epoch [7/20], Loss: 0.8785, Train Acc: 62.43%, Val Acc: 67.63%, LR: 0.000091
Epoch [8/20], Loss: 0.8565, Train Acc: 65.39%, Val Acc: 63.77%, LR: 0.000086
Epoch [9/20], Loss: 0.8242, Train Acc: 66.52%, Val Acc: 65.70%, LR: 0.000080
Epoch [10/20], Loss: 0.7876, Train Acc: 70.78%, Val Acc: 68.60%, LR: 0.000073
Epoch [11/20], Loss: 0.7520, Train Acc: 72.56%, Val Acc: 72.46%, LR: 0.000066
Epoch [12/20], Loss: 0.7203, Train Acc: 74.23%, Val Acc: 74.88%, LR: 0.000058
Epoch [13/20], Loss: 0.7040, Train Acc: 76.06%, Val Acc: 75.85%, LR: 0.000051
Epoch [14/20], Loss: 0.6739, Train Acc: 78.44%, Val Acc: 74.88%, LR: 0.000043
Epoch [15/20], Loss: 0.6699, Train Acc: 78.01%, Val Acc: 77.78%, LR: 0.000035
```

```
Epoch [16/20], Loss: 0.6547, Train Acc: 78.76%, Val Acc: 76.33%, LR: 0.000028
Epoch [17/20], Loss: 0.6403, Train Acc: 81.08%, Val Acc: 79.23%, LR: 0.000021
Epoch [18/20], Loss: 0.6208, Train Acc: 82.05%, Val Acc: 79.23%, LR: 0.000015
Epoch [19/20], Loss: 0.6163, Train Acc: 81.99%, Val Acc: 79.71%, LR: 0.000010
Epoch [20/20], Loss: 0.6004, Train Acc: 83.07%, Val Acc: 81.16%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.6209, Train Acc: 81.40%, Val Acc: 80.68%, LR: 0.000005
Epoch [2/15], Loss: 0.5999, Train Acc: 82.86%, Val Acc: 79.71%, LR: 0.000005
Epoch [3/15], Loss: 0.6031, Train Acc: 82.32%, Val Acc: 81.16%, LR: 0.000005
Epoch [4/15], Loss: 0.5915, Train Acc: 83.56%, Val Acc: 82.13%, LR: 0.000005
Epoch [5/15], Loss: 0.5913, Train Acc: 83.72%, Val Acc: 82.13%, LR: 0.000005
Epoch [6/15], Loss: 0.5641, Train Acc: 85.82%, Val Acc: 81.64%, LR: 0.000005
Epoch [7/15], Loss: 0.5680, Train Acc: 84.96%, Val Acc: 81.16%, LR: 0.000005
Epoch [8/15], Loss: 0.5696, Train Acc: 84.69%, Val Acc: 84.06%, LR: 0.000005
Epoch [9/15], Loss: 0.5586, Train Acc: 86.20%, Val Acc: 83.09%, LR: 0.000005
Epoch [10/15], Loss: 0.5684, Train Acc: 85.12%, Val Acc: 84.54%, LR: 0.000005
Epoch [11/15], Loss: 0.5481, Train Acc: 85.61%, Val Acc: 83.09%, LR: 0.000005
Epoch [12/15], Loss: 0.5336, Train Acc: 86.68%, Val Acc: 85.51%, LR: 0.000005
Epoch [13/15], Loss: 0.5320, Train Acc: 87.33%, Val Acc: 83.57%, LR: 0.000005
Epoch [14/15], Loss: 0.5304, Train Acc: 87.28%, Val Acc: 83.09%, LR: 0.000005
Epoch [15/15], Loss: 0.5302, Train Acc: 86.85%, Val Acc: 84.54%, LR: 0.000005
Testing model...
```

```
==== split_40_60 Results ===
```

```
Accuracy: 0.8490
Precision: 0.8495
Recall: 0.8490
F1-Score: 0.8490
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.77	0.80	0.78	674
CN	0.85	0.81	0.83	864
MCI	0.88	0.89	0.89	1554
accuracy			0.85	3092
macro avg	0.84	0.83	0.83	3092
weighted avg	0.85	0.85	0.85	3092

```
=====
```

```
Processing: split_50_50
```

```
=====
Train samples: 2319
Val samples: 258
Test samples: 2577
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0559, Train Acc: 47.65%, Val Acc: 50.00%, LR: 0.000033
Epoch [2/20], Loss: 1.0336, Train Acc: 51.01%, Val Acc: 52.33%, LR: 0.000067
Epoch [3/20], Loss: 1.0009, Train Acc: 53.08%, Val Acc: 52.71%, LR: 0.000100
Epoch [4/20], Loss: 0.9538, Train Acc: 57.14%, Val Acc: 58.14%, LR: 0.000099
Epoch [5/20], Loss: 0.9149, Train Acc: 59.38%, Val Acc: 56.98%, LR: 0.000098
Epoch [6/20], Loss: 0.8768, Train Acc: 63.52%, Val Acc: 67.83%, LR: 0.000095
Epoch [7/20], Loss: 0.8191, Train Acc: 67.23%, Val Acc: 67.83%, LR: 0.000091
Epoch [8/20], Loss: 0.7878, Train Acc: 69.86%, Val Acc: 70.54%, LR: 0.000086
Epoch [9/20], Loss: 0.7598, Train Acc: 71.54%, Val Acc: 71.32%, LR: 0.000080
Epoch [10/20], Loss: 0.7317, Train Acc: 74.13%, Val Acc: 77.13%, LR: 0.000073
Epoch [11/20], Loss: 0.6930, Train Acc: 76.63%, Val Acc: 79.46%, LR: 0.000066
Epoch [12/20], Loss: 0.6722, Train Acc: 77.92%, Val Acc: 81.40%, LR: 0.000058
Epoch [13/20], Loss: 0.6335, Train Acc: 80.29%, Val Acc: 82.17%, LR: 0.000051
Epoch [14/20], Loss: 0.6300, Train Acc: 80.85%, Val Acc: 83.33%, LR: 0.000043
Epoch [15/20], Loss: 0.6317, Train Acc: 80.03%, Val Acc: 82.56%, LR: 0.000035
Epoch [16/20], Loss: 0.6053, Train Acc: 82.79%, Val Acc: 83.72%, LR: 0.000028
Epoch [17/20], Loss: 0.5836, Train Acc: 84.35%, Val Acc: 84.50%, LR: 0.000021
Epoch [18/20], Loss: 0.5889, Train Acc: 83.05%, Val Acc: 86.82%, LR: 0.000015
Epoch [19/20], Loss: 0.5771, Train Acc: 84.04%, Val Acc: 86.05%, LR: 0.000010
Epoch [20/20], Loss: 0.5870, Train Acc: 83.53%, Val Acc: 85.27%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/15], Loss: 0.5622, Train Acc: 86.07%, Val Acc: 85.27%, LR: 0.000005
Epoch [2/15], Loss: 0.5636, Train Acc: 85.21%, Val Acc: 87.21%, LR: 0.000005
Epoch [3/15], Loss: 0.5563, Train Acc: 85.81%, Val Acc: 85.66%, LR: 0.000005
Epoch [4/15], Loss: 0.5434, Train Acc: 86.76%, Val Acc: 89.53%, LR: 0.000005
Epoch [5/15], Loss: 0.5377, Train Acc: 86.93%, Val Acc: 85.66%, LR: 0.000005
Epoch [6/15], Loss: 0.5410, Train Acc: 86.80%, Val Acc: 87.98%, LR: 0.000005
Epoch [7/15], Loss: 0.5228, Train Acc: 87.02%, Val Acc: 87.60%, LR: 0.000005
Epoch [8/15], Loss: 0.5087, Train Acc: 88.96%, Val Acc: 87.60%, LR: 0.000005
Epoch [9/15], Loss: 0.5206, Train Acc: 88.06%, Val Acc: 85.27%, LR: 0.000003
Epoch [10/15], Loss: 0.5112, Train Acc: 87.54%, Val Acc: 87.98%, LR: 0.000003
Epoch [11/15], Loss: 0.4945, Train Acc: 88.75%, Val Acc: 87.98%, LR: 0.000003
Epoch [12/15], Loss: 0.5035, Train Acc: 89.09%, Val Acc: 88.37%, LR: 0.000003
Early stopping at epoch 12
Testing model...

==== split_50_50 Results ====
Accuracy: 0.8700
```

```
Precision: 0.8723  
Recall: 0.8700  
F1-Score: 0.8690
```

Classification Report:

	precision	recall	f1-score	support
AD	0.81	0.83	0.82	562
CN	0.92	0.78	0.84	720
MCI	0.87	0.94	0.90	1295
accuracy			0.87	2577
macro avg	0.87	0.85	0.86	2577
weighted avg	0.87	0.87	0.87	2577

```
=====  
Processing: split_60_40  
=====
```

```
Train samples: 2784
```

```
Val samples: 308
```

```
Test samples: 2062
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/20], Loss: 1.0576, Train Acc: 47.59%, Val Acc: 49.35%, LR: 0.000033  
Epoch [2/20], Loss: 1.0301, Train Acc: 50.57%, Val Acc: 52.27%, LR: 0.000067  
Epoch [3/20], Loss: 0.9987, Train Acc: 52.80%, Val Acc: 54.87%, LR: 0.000100  
Epoch [4/20], Loss: 0.9458, Train Acc: 57.29%, Val Acc: 59.42%, LR: 0.000099  
Epoch [5/20], Loss: 0.8996, Train Acc: 61.17%, Val Acc: 69.48%, LR: 0.000098  
Epoch [6/20], Loss: 0.8462, Train Acc: 65.16%, Val Acc: 67.21%, LR: 0.000095  
Epoch [7/20], Loss: 0.8012, Train Acc: 68.21%, Val Acc: 70.45%, LR: 0.000091  
Epoch [8/20], Loss: 0.7645, Train Acc: 71.95%, Val Acc: 75.00%, LR: 0.000086  
Epoch [9/20], Loss: 0.7284, Train Acc: 73.92%, Val Acc: 76.62%, LR: 0.000080  
Epoch [10/20], Loss: 0.6870, Train Acc: 77.08%, Val Acc: 76.62%, LR: 0.000073  
Epoch [11/20], Loss: 0.6672, Train Acc: 77.37%, Val Acc: 80.84%, LR: 0.000066  
Epoch [12/20], Loss: 0.6516, Train Acc: 78.99%, Val Acc: 81.82%, LR: 0.000058  
Epoch [13/20], Loss: 0.6241, Train Acc: 80.82%, Val Acc: 81.82%, LR: 0.000051  
Epoch [14/20], Loss: 0.6113, Train Acc: 81.50%, Val Acc: 81.82%, LR: 0.000043  
Epoch [15/20], Loss: 0.5832, Train Acc: 83.66%, Val Acc: 86.36%, LR: 0.000035  
Epoch [16/20], Loss: 0.5973, Train Acc: 83.48%, Val Acc: 85.71%, LR: 0.000028  
Epoch [17/20], Loss: 0.5864, Train Acc: 83.33%, Val Acc: 86.36%, LR: 0.000021  
Epoch [18/20], Loss: 0.5643, Train Acc: 84.23%, Val Acc: 86.36%, LR: 0.000015  
Epoch [19/20], Loss: 0.5447, Train Acc: 86.64%, Val Acc: 87.34%, LR: 0.000010  
Epoch [20/20], Loss: 0.5438, Train Acc: 86.67%, Val Acc: 87.66%, LR: 0.000006
```

```
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:  
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to  
access the learning rate.
```

```
warnings.warn(  
Epoch [1/15], Loss: 0.5545, Train Acc: 85.99%, Val Acc: 87.34%, LR: 0.000005  
Epoch [2/15], Loss: 0.5516, Train Acc: 86.31%, Val Acc: 88.31%, LR: 0.000005  
Epoch [3/15], Loss: 0.5400, Train Acc: 86.64%, Val Acc: 87.99%, LR: 0.000005  
Epoch [4/15], Loss: 0.5220, Train Acc: 87.68%, Val Acc: 88.96%, LR: 0.000005  
Epoch [5/15], Loss: 0.5127, Train Acc: 87.82%, Val Acc: 89.94%, LR: 0.000005  
Epoch [6/15], Loss: 0.5021, Train Acc: 88.51%, Val Acc: 88.64%, LR: 0.000005  
Epoch [7/15], Loss: 0.5032, Train Acc: 88.97%, Val Acc: 89.61%, LR: 0.000005  
Epoch [8/15], Loss: 0.5000, Train Acc: 89.40%, Val Acc: 89.61%, LR: 0.000005  
Epoch [9/15], Loss: 0.4942, Train Acc: 89.15%, Val Acc: 89.94%, LR: 0.000005  
Epoch [10/15], Loss: 0.4871, Train Acc: 90.09%, Val Acc: 89.61%, LR: 0.000003  
Epoch [11/15], Loss: 0.4751, Train Acc: 90.91%, Val Acc: 89.94%, LR: 0.000003  
Epoch [12/15], Loss: 0.4920, Train Acc: 89.62%, Val Acc: 90.58%, LR: 0.000003  
Epoch [13/15], Loss: 0.4838, Train Acc: 90.23%, Val Acc: 90.26%, LR: 0.000003  
Epoch [14/15], Loss: 0.4840, Train Acc: 89.58%, Val Acc: 90.58%, LR: 0.000003  
Epoch [15/15], Loss: 0.4759, Train Acc: 90.70%, Val Acc: 91.56%, LR: 0.000003  
Testing model...
```

```
==== split_60_40 Results ====  
Accuracy: 0.9234  
Precision: 0.9234  
Recall: 0.9234  
F1-Score: 0.9229
```

Classification Report:

	precision	recall	f1-score	support
AD	0.91	0.86	0.88	450
CN	0.93	0.90	0.92	576
MCI	0.92	0.96	0.94	1036
accuracy			0.92	2062
macro avg	0.92	0.91	0.91	2062
weighted avg	0.92	0.92	0.92	2062

```
=====  
Processing: split_70_30  
=====
```

```
Train samples: 3247  
Val samples: 361  
Test samples: 1546  
Starting phase 1 training (partial unfreeze)...  
Epoch [1/20], Loss: 1.0644, Train Acc: 45.24%, Val Acc: 50.14%, LR: 0.000033  
Epoch [2/20], Loss: 1.0274, Train Acc: 50.63%, Val Acc: 50.14%, LR: 0.000067  
Epoch [3/20], Loss: 0.9898, Train Acc: 54.11%, Val Acc: 57.62%, LR: 0.000100  
Epoch [4/20], Loss: 0.9299, Train Acc: 59.13%, Val Acc: 63.16%, LR: 0.000099
```

```
Epoch [5/20], Loss: 0.8717, Train Acc: 64.12%, Val Acc: 67.04%, LR: 0.000098
Epoch [6/20], Loss: 0.8072, Train Acc: 68.28%, Val Acc: 70.08%, LR: 0.000095
Epoch [7/20], Loss: 0.7711, Train Acc: 71.39%, Val Acc: 70.36%, LR: 0.000091
Epoch [8/20], Loss: 0.7358, Train Acc: 73.36%, Val Acc: 76.73%, LR: 0.000086
Epoch [9/20], Loss: 0.7001, Train Acc: 75.73%, Val Acc: 78.12%, LR: 0.000080
Epoch [10/20], Loss: 0.6685, Train Acc: 78.72%, Val Acc: 78.39%, LR: 0.000073
Epoch [11/20], Loss: 0.6354, Train Acc: 80.57%, Val Acc: 81.72%, LR: 0.000066
Epoch [12/20], Loss: 0.6232, Train Acc: 80.94%, Val Acc: 83.10%, LR: 0.000058
Epoch [13/20], Loss: 0.5979, Train Acc: 82.38%, Val Acc: 81.72%, LR: 0.000051
Epoch [14/20], Loss: 0.5794, Train Acc: 83.74%, Val Acc: 84.49%, LR: 0.000043
Epoch [15/20], Loss: 0.5667, Train Acc: 84.35%, Val Acc: 86.15%, LR: 0.000035
Epoch [16/20], Loss: 0.5676, Train Acc: 84.76%, Val Acc: 86.43%, LR: 0.000028
Epoch [17/20], Loss: 0.5418, Train Acc: 86.57%, Val Acc: 86.70%, LR: 0.000021
Epoch [18/20], Loss: 0.5262, Train Acc: 87.03%, Val Acc: 86.43%, LR: 0.000015
Epoch [19/20], Loss: 0.5261, Train Acc: 87.00%, Val Acc: 86.98%, LR: 0.000010
Epoch [20/20], Loss: 0.5061, Train Acc: 89.07%, Val Acc: 88.37%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.5227, Train Acc: 87.68%, Val Acc: 88.09%, LR: 0.000005
Epoch [2/15], Loss: 0.5041, Train Acc: 88.79%, Val Acc: 88.09%, LR: 0.000005
Epoch [3/15], Loss: 0.4929, Train Acc: 89.31%, Val Acc: 90.03%, LR: 0.000005
Epoch [4/15], Loss: 0.4933, Train Acc: 89.04%, Val Acc: 89.47%, LR: 0.000005
Epoch [5/15], Loss: 0.4894, Train Acc: 89.47%, Val Acc: 92.24%, LR: 0.000005
Epoch [6/15], Loss: 0.4702, Train Acc: 90.76%, Val Acc: 90.58%, LR: 0.000005
Epoch [7/15], Loss: 0.4705, Train Acc: 91.07%, Val Acc: 91.97%, LR: 0.000005
Epoch [8/15], Loss: 0.4696, Train Acc: 90.70%, Val Acc: 92.24%, LR: 0.000005
Epoch [9/15], Loss: 0.4718, Train Acc: 91.04%, Val Acc: 91.69%, LR: 0.000005
Epoch [10/15], Loss: 0.4614, Train Acc: 91.47%, Val Acc: 92.24%, LR: 0.000003
Epoch [11/15], Loss: 0.4476, Train Acc: 92.49%, Val Acc: 91.41%, LR: 0.000003
Epoch [12/15], Loss: 0.4566, Train Acc: 91.90%, Val Acc: 93.07%, LR: 0.000003
Epoch [13/15], Loss: 0.4308, Train Acc: 93.44%, Val Acc: 93.07%, LR: 0.000003
Epoch [14/15], Loss: 0.4460, Train Acc: 92.18%, Val Acc: 92.80%, LR: 0.000003
Epoch [15/15], Loss: 0.4410, Train Acc: 92.55%, Val Acc: 93.63%, LR: 0.000003
Testing model...
```

```
==== split_70_30 Results ====
Accuracy: 0.9282
Precision: 0.9285
Recall: 0.9282
F1-Score: 0.9278
```

```
Classification Report:
```

precision	recall	f1-score	support
-----------	--------	----------	---------

AD	0.89	0.88	0.88	337
CN	0.95	0.89	0.92	432
MCI	0.93	0.97	0.95	777
accuracy			0.93	1546
macro avg	0.93	0.91	0.92	1546
weighted avg	0.93	0.93	0.93	1546

=====

Processing: split_80_20

=====

Train samples: 3711

Val samples: 412

Test samples: 1031

Starting phase 1 training (partial unfreeze)...

Epoch [1/20], Loss: 1.0536, Train Acc: 48.69%, Val Acc: 49.51%, LR: 0.000033
 Epoch [2/20], Loss: 1.0205, Train Acc: 51.47%, Val Acc: 51.94%, LR: 0.000067
 Epoch [3/20], Loss: 0.9745, Train Acc: 55.38%, Val Acc: 59.71%, LR: 0.000100
 Epoch [4/20], Loss: 0.9052, Train Acc: 61.36%, Val Acc: 66.75%, LR: 0.000099
 Epoch [5/20], Loss: 0.8270, Train Acc: 67.50%, Val Acc: 71.36%, LR: 0.000098
 Epoch [6/20], Loss: 0.7884, Train Acc: 69.82%, Val Acc: 75.49%, LR: 0.000095
 Epoch [7/20], Loss: 0.7391, Train Acc: 73.30%, Val Acc: 79.85%, LR: 0.000091
 Epoch [8/20], Loss: 0.6958, Train Acc: 76.48%, Val Acc: 83.98%, LR: 0.000086
 Epoch [9/20], Loss: 0.6579, Train Acc: 79.12%, Val Acc: 84.71%, LR: 0.000080
 Epoch [10/20], Loss: 0.6385, Train Acc: 79.92%, Val Acc: 86.17%, LR: 0.000073
 Epoch [11/20], Loss: 0.6009, Train Acc: 82.16%, Val Acc: 88.59%, LR: 0.000066
 Epoch [12/20], Loss: 0.5866, Train Acc: 84.21%, Val Acc: 89.81%, LR: 0.000058
 Epoch [13/20], Loss: 0.5594, Train Acc: 85.42%, Val Acc: 88.59%, LR: 0.000051
 Epoch [14/20], Loss: 0.5535, Train Acc: 85.29%, Val Acc: 91.75%, LR: 0.000043
 Epoch [15/20], Loss: 0.5260, Train Acc: 87.36%, Val Acc: 91.99%, LR: 0.000035
 Epoch [16/20], Loss: 0.5141, Train Acc: 88.44%, Val Acc: 92.96%, LR: 0.000028
 Epoch [17/20], Loss: 0.5080, Train Acc: 88.66%, Val Acc: 93.69%, LR: 0.000021
 Epoch [18/20], Loss: 0.4961, Train Acc: 89.68%, Val Acc: 94.42%, LR: 0.000015
 Epoch [19/20], Loss: 0.4986, Train Acc: 88.95%, Val Acc: 93.93%, LR: 0.000010
 Epoch [20/20], Loss: 0.4888, Train Acc: 89.44%, Val Acc: 94.42%, LR: 0.000006

Starting phase 2 training (full fine-tuning)...

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
warnings.warn(
```

Epoch [1/15], Loss: 0.4808, Train Acc: 90.00%, Val Acc: 94.66%, LR: 0.000005
 Epoch [2/15], Loss: 0.4756, Train Acc: 90.30%, Val Acc: 94.42%, LR: 0.000005
 Epoch [3/15], Loss: 0.4680, Train Acc: 91.11%, Val Acc: 94.90%, LR: 0.000005
 Epoch [4/15], Loss: 0.4596, Train Acc: 91.40%, Val Acc: 94.17%, LR: 0.000005
 Epoch [5/15], Loss: 0.4544, Train Acc: 91.70%, Val Acc: 94.66%, LR: 0.000005
 Epoch [6/15], Loss: 0.4551, Train Acc: 91.54%, Val Acc: 95.15%, LR: 0.000005

Epoch [7/15], Loss: 0.4544, Train Acc: 92.08%, Val Acc: 94.66%, LR: 0.000005
 Epoch [8/15], Loss: 0.4382, Train Acc: 92.72%, Val Acc: 95.39%, LR: 0.000005
 Epoch [9/15], Loss: 0.4389, Train Acc: 92.05%, Val Acc: 95.15%, LR: 0.000005
 Epoch [10/15], Loss: 0.4336, Train Acc: 92.75%, Val Acc: 95.63%, LR: 0.000005
 Epoch [11/15], Loss: 0.4275, Train Acc: 93.18%, Val Acc: 96.60%, LR: 0.000005
 Epoch [12/15], Loss: 0.4310, Train Acc: 92.91%, Val Acc: 95.87%, LR: 0.000005
 Epoch [13/15], Loss: 0.4268, Train Acc: 93.24%, Val Acc: 96.36%, LR: 0.000005
 Epoch [14/15], Loss: 0.4243, Train Acc: 93.29%, Val Acc: 96.12%, LR: 0.000005
 Epoch [15/15], Loss: 0.4160, Train Acc: 93.91%, Val Acc: 96.60%, LR: 0.000005
 Testing model...

```
==== split_80_20 Results ===
Accuracy: 0.9505
Precision: 0.9504
Recall: 0.9505
F1-Score: 0.9504
```

Classification Report:

	precision	recall	f1-score	support
AD	0.93	0.91	0.92	225
CN	0.95	0.94	0.95	288
MCI	0.96	0.97	0.97	518
accuracy			0.95	1031
macro avg	0.95	0.94	0.94	1031
weighted avg	0.95	0.95	0.95	1031

```
=====
Processing: split_90_10
=====
Train samples: 4175
Val samples: 464
Test samples: 515
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0491, Train Acc: 49.44%, Val Acc: 51.29%, LR: 0.000033
Epoch [2/20], Loss: 1.0222, Train Acc: 50.83%, Val Acc: 56.90%, LR: 0.000067
Epoch [3/20], Loss: 0.9654, Train Acc: 55.93%, Val Acc: 64.66%, LR: 0.000100
Epoch [4/20], Loss: 0.8875, Train Acc: 61.99%, Val Acc: 68.75%, LR: 0.000099
Epoch [5/20], Loss: 0.8243, Train Acc: 67.64%, Val Acc: 75.22%, LR: 0.000098
Epoch [6/20], Loss: 0.7677, Train Acc: 71.07%, Val Acc: 74.35%, LR: 0.000095
Epoch [7/20], Loss: 0.7118, Train Acc: 75.59%, Val Acc: 83.84%, LR: 0.000091
Epoch [8/20], Loss: 0.6802, Train Acc: 77.08%, Val Acc: 82.97%, LR: 0.000086
Epoch [9/20], Loss: 0.6453, Train Acc: 79.19%, Val Acc: 86.64%, LR: 0.000080
Epoch [10/20], Loss: 0.6024, Train Acc: 82.37%, Val Acc: 85.13%, LR: 0.000073
Epoch [11/20], Loss: 0.5921, Train Acc: 83.26%, Val Acc: 90.30%, LR: 0.000066
Epoch [12/20], Loss: 0.5655, Train Acc: 85.03%, Val Acc: 88.36%, LR: 0.000058
```

```
Epoch [13/20], Loss: 0.5430, Train Acc: 86.47%, Val Acc: 89.66%, LR: 0.000051
Epoch [14/20], Loss: 0.5295, Train Acc: 87.21%, Val Acc: 90.52%, LR: 0.000043
Epoch [15/20], Loss: 0.5138, Train Acc: 88.10%, Val Acc: 92.67%, LR: 0.000035
Epoch [16/20], Loss: 0.5047, Train Acc: 88.79%, Val Acc: 93.32%, LR: 0.000028
Epoch [17/20], Loss: 0.4846, Train Acc: 89.53%, Val Acc: 93.10%, LR: 0.000021
Epoch [18/20], Loss: 0.4827, Train Acc: 90.11%, Val Acc: 92.67%, LR: 0.000015
Epoch [19/20], Loss: 0.4757, Train Acc: 90.20%, Val Acc: 93.75%, LR: 0.000010
Epoch [20/20], Loss: 0.4735, Train Acc: 90.85%, Val Acc: 93.75%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
warnings.warn(
```

```
Epoch [1/15], Loss: 0.4669, Train Acc: 91.19%, Val Acc: 93.75%, LR: 0.000005
Epoch [2/15], Loss: 0.4554, Train Acc: 90.97%, Val Acc: 93.97%, LR: 0.000005
Epoch [3/15], Loss: 0.4528, Train Acc: 91.78%, Val Acc: 93.53%, LR: 0.000005
Epoch [4/15], Loss: 0.4539, Train Acc: 91.74%, Val Acc: 94.18%, LR: 0.000005
Epoch [5/15], Loss: 0.4394, Train Acc: 92.36%, Val Acc: 94.61%, LR: 0.000005
Epoch [6/15], Loss: 0.4360, Train Acc: 92.48%, Val Acc: 94.83%, LR: 0.000005
Epoch [7/15], Loss: 0.4388, Train Acc: 92.67%, Val Acc: 95.47%, LR: 0.000005
Epoch [8/15], Loss: 0.4300, Train Acc: 92.69%, Val Acc: 95.69%, LR: 0.000005
Epoch [9/15], Loss: 0.4138, Train Acc: 94.49%, Val Acc: 95.04%, LR: 0.000005
Epoch [10/15], Loss: 0.4150, Train Acc: 94.13%, Val Acc: 95.26%, LR: 0.000005
Epoch [11/15], Loss: 0.4171, Train Acc: 93.92%, Val Acc: 95.69%, LR: 0.000005
Epoch [12/15], Loss: 0.4122, Train Acc: 94.40%, Val Acc: 96.12%, LR: 0.000005
Epoch [13/15], Loss: 0.4093, Train Acc: 94.11%, Val Acc: 96.12%, LR: 0.000005
Epoch [14/15], Loss: 0.4090, Train Acc: 94.49%, Val Acc: 96.34%, LR: 0.000005
Epoch [15/15], Loss: 0.4075, Train Acc: 94.56%, Val Acc: 96.55%, LR: 0.000005
Testing model...
```

```
==== split_90_10 Results ===
```

```
Accuracy: 0.9592
```

```
Precision: 0.9593
```

```
Recall: 0.9592
```

```
F1-Score: 0.9590
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.95	0.91	0.93	112
CN	0.97	0.96	0.97	144
MCI	0.95	0.98	0.97	259
accuracy			0.96	515
macro avg	0.96	0.95	0.95	515
weighted avg	0.96	0.96	0.96	515

```
=====
EfficientNetV2-S - SUMMARY OF ALL SPLITS
=====
```

	split	accuracy	precision	recall	f1_score	training_time
	split_10_90	0.575555	0.585557	0.575555	0.579376	209.364409
	split_20_80	0.707980	0.701371	0.707980	0.702363	601.484343
	split_30_70	0.800443	0.798541	0.800443	0.798868	867.035979
	split_40_60	0.848965	0.849452	0.848965	0.848979	1099.555901
	split_50_50	0.870004	0.872320	0.870004	0.868957	1194.387525
	split_60_40	0.923375	0.923386	0.923375	0.922852	1984.304198
	split_70_30	0.928202	0.928495	0.928202	0.927773	2231.632533
	split_80_20	0.950533	0.950384	0.950533	0.950398	2518.564639
	split_90_10	0.959223	0.959273	0.959223	0.959043	2807.190982

Detailed results saved to: /kaggle/working/efficientnetv2_s_results.csv