

mer-detection-pre-procesing-model2

October 28, 2025

```
[1]: import os
import matplotlib.pyplot as plt
from PIL import Image
import numpy as np

# Keep your original directories and subfolders
test_dir = "/kaggle/input/alzheimer-disease/test-20251010T094612Z-1-001/test"
train_dir = "/kaggle/input/alzheimer-disease/train-20251010T183510Z-1-001/train"
subfolders = ["AD", "CN", "MCI"]

# Function to show 3 images from each subfolder (unchanged signature)
def show_images_from_dir(path, n=3):
    # List only image files (same behavior; expanded to be robust if needed)
    files = [f for f in os.listdir(path) if f.lower().endswith('.png', '.jpg', '.jpeg', '.bmp', '.tif', '.tiff')]]
    files = files[:n] # Take only the first n images

    plt.figure(figsize=(15, 5))
    for i, file in enumerate(files):
        img_path = os.path.join(path, file)

        # Open with PIL
        img = Image.open(img_path)

        # --- CRITICAL FIX ---
        # Force grayscale to ensure single-channel display (no colorization)
        # Even if the source is already grayscale, this guarantees mode 'L'
        img = img.convert('L')

        # Convert to numpy for safe imshow with explicit bounds
        arr = np.asarray(img)

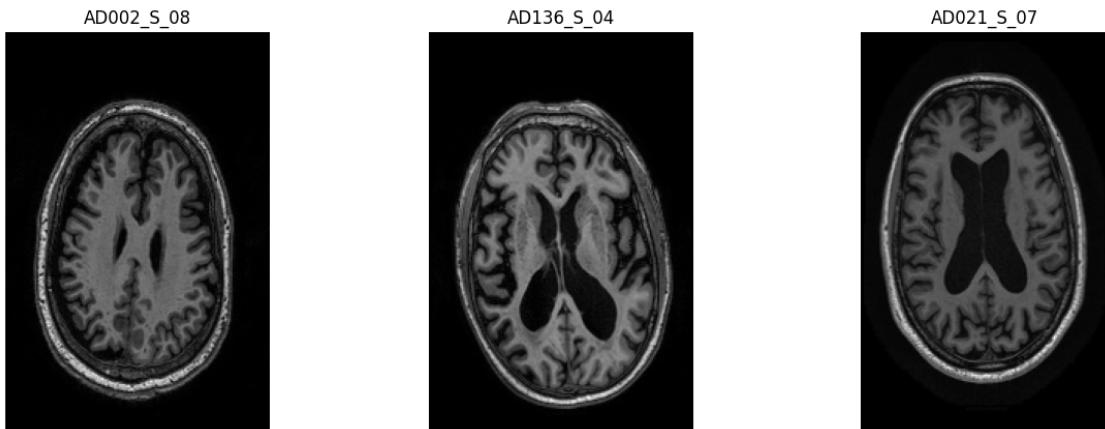
        # Display as true grayscale with fixed intensity bounds
        plt.subplot(1, len(files), i + 1)
        plt.imshow(arr, cmap='gray', vmin=0, vmax=255)
        plt.axis("off")
        plt.title(file[:10])
```

```
plt.show()

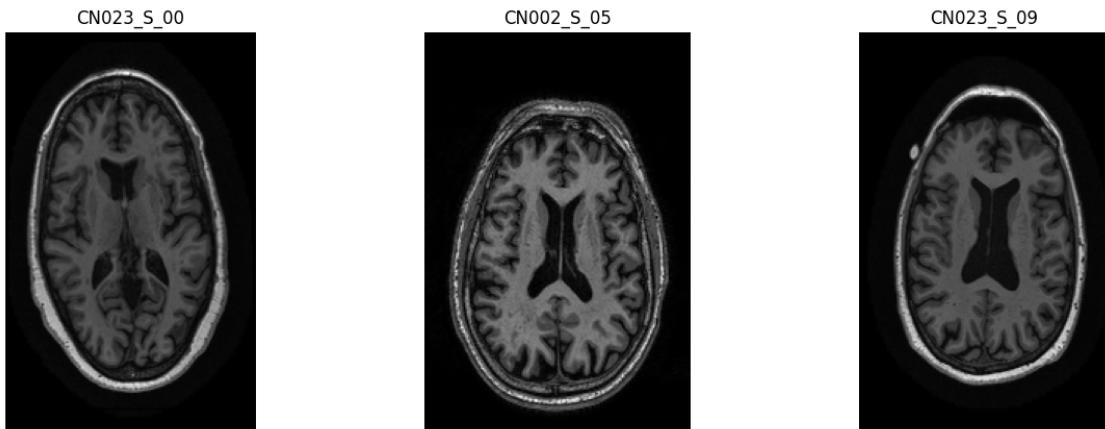
# Show images for Test set
for subfolder in subfolders:
    print(f"{subfolder} (Test)")
    show_images_from_dir(os.path.join(test_dir, subfolder))

# Show images for Train set
for subfolder in subfolders:
    print(f"{subfolder} (Train)")
    show_images_from_dir(os.path.join(train_dir, subfolder))
```

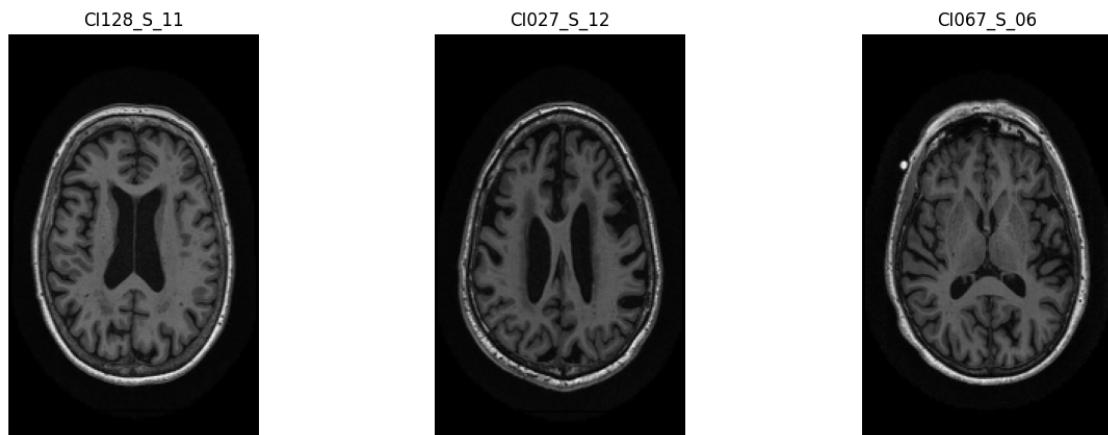
AD (Test)



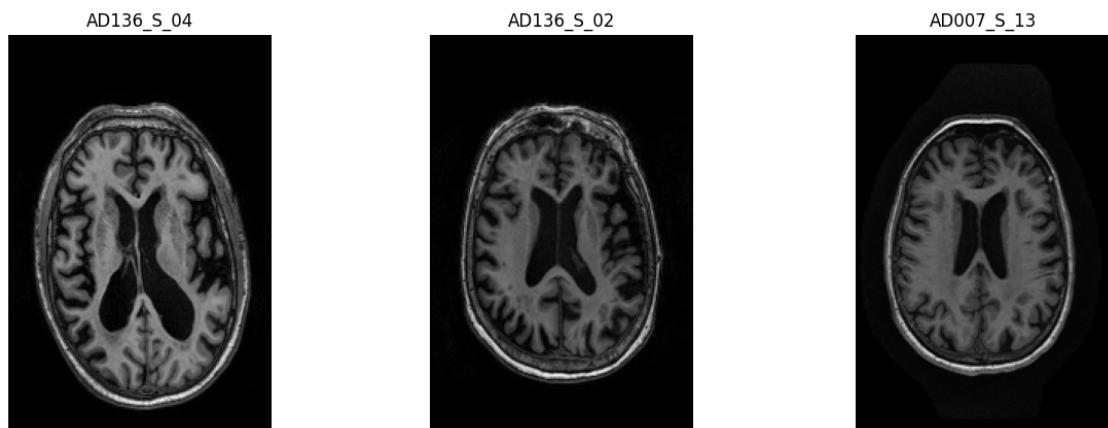
CN (Test)



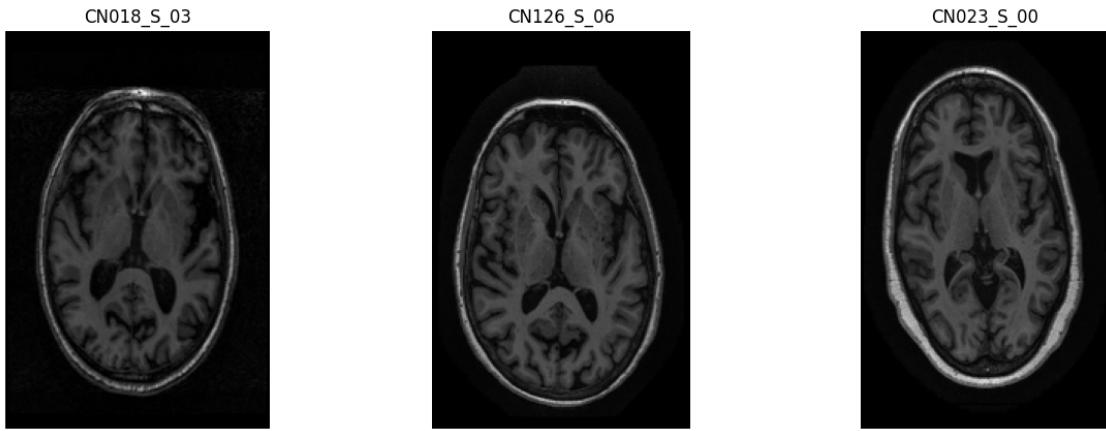
MCI (Test)



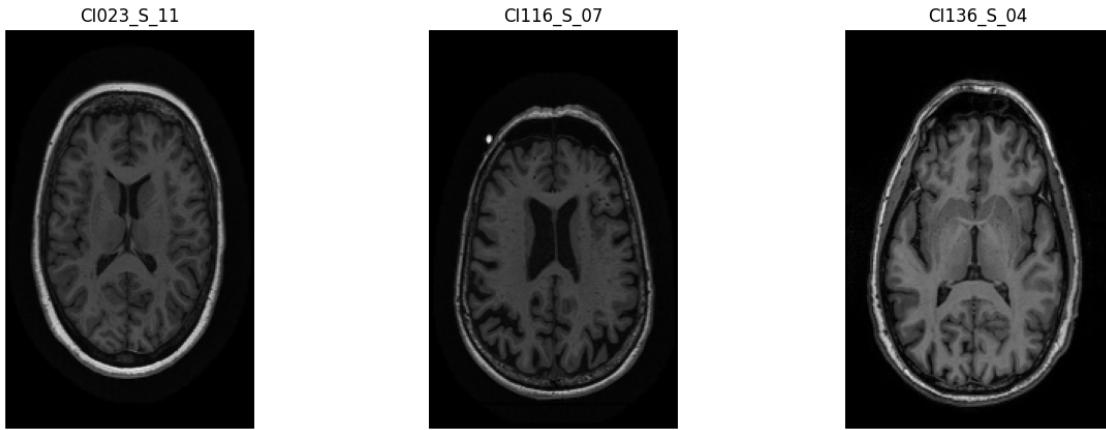
AD (Train)



CN (Train)



MCI (Train)



```
[2]: import os
import sys
import numpy as np
from PIL import Image, UnidentifiedImageError
import matplotlib.pyplot as plt

# SciPy for morphology / connected components (Kaggle preinstalled)
from scipy.ndimage import (
    gaussian_filter,
    binary_opening,
    binary_closing,
    binary_fill_holes,
    label
)
```

```

# Try skimage for CLAHE; fall back gracefully if missing
try:
    from skimage.exposure import equalize_adapthist
    _HAS_SKIMAGE = True
except Exception:
    _HAS_SKIMAGE = False

# GPU via PyTorch (for homomorphic filtering)

def _ensure(pkg):
    import importlib
    try:
        importlib.import_module(pkg)
    except Exception:
        import subprocess
        subprocess.check_call([sys.executable, "-m", "pip", "install", pkg, ↴"--quiet"])

_ensure("torch")
import torch
import torch.nn.functional as F

def torch_device():
    return torch.device("cuda" if torch.cuda.is_available() else "cpu")

def gaussian_kernel_2d(sigma: float, radius_factor: float = 3.0, device=None, ↴
    dtype=torch.float32):
    """Create a 2D Gaussian kernel tensor for conv2d (normalized)."""
    device = device or torch_device()
    rad = max(1, int(radius_factor * sigma))
    xs = torch.arange(-rad, rad + 1, device=device, dtype=dtype)
    g1 = torch.exp(-0.5 * (xs / sigma) ** 2)
    g1 = g1 / g1.sum()
    g2 = g1[:, None] @ g1[None, :]
    g2 = g2 / g2.sum()
    return g2 # (K, K)

def homomorphic_filter_gpu_u8(arr_u8: np.ndarray, sigma: float = 50.0) -> np. ↴
ndarray:
    """
    GPU homomorphic filtering (log + low-pass via conv2d + exp + robust ↴
    rescale).

    Input uint8 [0..255], output uint8 [0..255].
    """
    dev = torch_device()
    # [B=1,C=1,H,W] float32

```

```

    img = torch.from_numpy(arr_u8.astype(np.float32)).to(dev) + 1.0 # avoid underflow at log(0)
    img = img[None, None, :, :] # NCHW
    loga = torch.log(img)
    # 2D Gaussian low-pass via conv2d (reflect padding to avoid border artifacts)
    k = gaussian_kernel_2d(sigma=sigma, device=dev)
    k = k[None, None, :, :] # (out_c,in_c,H,W)
    pad = k.shape[-1] // 2
    low = F.conv2d(F.pad(loga, (pad, pad, pad, pad), mode="reflect"), k)
    high = loga - low
    corr = torch.exp(high) - 1.0 # back to linear domain

    # Robust percentile rescale to uint8
    a = corr.squeeze(0).squeeze(0) # HxW
    # Compute percentiles on CPU for simplicity
    a_np = a.detach().cpu().numpy().astype(np.float32)
    lo, hi = np.percentile(a_np, [0.5, 99.5])
    if hi - lo < 1e-6:
        a_np = (a_np - a_np.min()) / (a_np.ptp() + 1e-8)
    else:
        a_np = np.clip((a_np - lo) / (hi - lo), 0.0, 1.0)
    out_u8 = (a_np * 255.0 + 0.5).astype(np.uint8)
    return out_u8

test_dir = "/kaggle/input/alzheimer-disease/test-20251010T094612Z-1-001/test"
train_dir = "/kaggle/input/alzheimer-disease/train-20251010T183510Z-1-001/train"
subfolders = ["AD", "CN", "MCI"]

# Output (mirrors the structure)

preprocessed_root = "/kaggle/working/alzheimer-preprocessed"
preprocessed_test = os.path.join(preprocessed_root, "test")
preprocessed_train = os.path.join(preprocessed_root, "train")

IMG_EXT = (".png", ".jpg", ".jpeg", ".bmp", ".tif", ".tiff")

def ensure_dir(p: str):
    os.makedirs(p, exist_ok=True)

# Utilities

def load_gray(path: str) -> np.ndarray:
    """Load as uint8 grayscale [0..255]."""

```

```

    return np.asarray(Image.open(path).convert("L"), dtype=np.uint8)

def save_gray(arr_u8: np.ndarray, path: str):
    Image.fromarray(arr_u8, mode="L").save(path)

def percentile_rescale_u8(arr: np.ndarray, p_low=1.0, p_high=99.0) -> np.
    ndarray:
    """Robust contrast stretching to uint8."""
    a = arr.astype(np.float32)
    lo, hi = np.percentile(a, [p_low, p_high])
    if hi - lo < 1e-6:
        a = (a - a.min()) / (a.ptp() + 1e-8)
    else:
        a = np.clip((a - lo) / (hi - lo), 0.0, 1.0)
    return (a * 255.0 + 0.5).astype(np.uint8)

def otsu_threshold_u8(arr_u8: np.ndarray) -> int:
    """Pure NumPy Otsu threshold (returns 0..255)."""
    hist = np.bincount(arr_u8.ravel(), minlength=256).astype(np.float64)
    prob = hist / (arr_u8.size + 1e-12)
    omega = np.cumsum(prob)
    mu = np.cumsum(prob * np.arange(256))
    mu_t = mu[-1]
    sigma_b2 = (mu_t * omega - mu) ** 2 / (omega * (1 - omega) + 1e-12)
    sigma_b2[~np.isfinite(sigma_b2)] = 0.0
    return int(np.argmax(sigma_b2))

def largest_cc(mask: np.ndarray) -> np.ndarray:
    """Keep only largest connected component of a binary mask."""
    lbl, n = label(mask.astype(np.uint8))
    if n <= 1:
        return mask.astype(bool)
    counts = np.bincount(lbl.ravel())
    counts[0] = 0 # background
    keep = counts.argmax()
    return (lbl == keep)

def clahe_u8(arr_u8: np.ndarray) -> np.ndarray:
    """CLAHE if skimage is available; otherwise identity."""
    if _HAS_SKIMAGE:
        arr01 = arr_u8.astype(np.float32) / 255.0
        out01 = equalize_adapthist(arr01, clip_limit=0.01)
        return (np.clip(out01, 0.0, 1.0) * 255.0 + 0.5).astype(np.uint8)
    else:
        return arr_u8

```

```

# Core: one-image preprocessing

def preprocess_single_image(
    arr_u8: np.ndarray,
    do_skull_strip: bool = True,
    sigma_homomorphic: float = 50.0,
    min_area: int = 500
):
    """
    Steps:
        1) Robust pre-stretch (percentile)
        2) GPU homomorphic filter (PyTorch)
        3) (Optional) 2D skull/background stripping via Otsu + morphology + LCC + hole fill
        4) Gentle CLAHE (if available), else percentile rescale
    Returns processed uint8 image, plus an optional mask (uint8).
    """

    # 1) Robust stretch
    a1 = percentile_rescale_u8(arr_u8, 1.0, 99.0)

    # 2) Homomorphic filtering on GPU
    a2 = homomorphic_filter_gpu_u8(a1, sigma=sigma_homomorphic)

    brain_mask = None
    if do_skull_strip:
        # 3) Otsu thresholding (foreground bright)
        th = otsu_threshold_u8(a2)
        mask = (a2 >= th).astype(np.uint8)

        # Morphological clean-up
        mask = binary_opening(mask, structure=np.ones((3,3), dtype=np.uint8))
        mask = binary_closing(mask, structure=np.ones((5,5), dtype=np.uint8))
        mask = binary_fill_holes(mask)
        mask = largest_cc(mask)

        # Remove tiny masks (safety)
        if mask.sum() < min_area:
            mask = np.ones_like(mask, dtype=bool) # fallback: keep as-is

        a3 = (a2 * mask).astype(np.uint8)
        brain_mask = (mask.astype(np.uint8) * 255)
    else:
        a3 = a2

    # 4) CLAHE or robust stretch
    a4 = clahe_u8(a3)
    a4 = percentile_rescale_u8(a4, 0.5, 99.5)

```

```

    return a4, brain_mask

# Dataset-level processing & visualization

def preprocess_dataset(src_root: str, dst_root: str, n_preview: int = 3,
                      ↪do_skull_strip=True):
    """
    Applies the pipeline to all images under src_root/{AD/CN/MCI}
    and writes to dst_root/{AD/CN/MCI}. Silent on per-file issues.
    """
    ensure_dir(dst_root)
    summary = {}

    # Report GPU/CPU
    dev = torch_device()
    print(f"Device for homomorphic filtering: {dev}")

    for cls in subfolders:
        src_cls = os.path.join(src_root, cls)
        dst_cls = os.path.join(dst_root, cls)
        ensure_dir(dst_cls)

        processed = skipped = 0
        if not os.path.isdir(src_cls):
            summary[cls] = (0, 0)
            continue

        files = sorted([f for f in os.listdir(src_cls) if f.lower().
        ↪endswith(IMG_EXT)])
        # Process & save
        for fname in files:
            spath = os.path.join(src_cls, fname)
            dpath = os.path.join(dst_cls, fname)
            try:
                arr = load_gray(spath)
                proc, _ = preprocess_single_image(arr, ↪
                    ↪do_skull_strip=do_skull_strip)
                save_gray(proc, dpath)
                processed += 1
            except (UnidentifiedImageError, OSError, RuntimeError, ValueError):
                skipped += 1
                continue

        summary[cls] = (processed, skipped)

```

```

# Preview few examples
preview = files[:n_preview]
if preview:
    fig, axs = plt.subplots(len(preview), 3, figsize=(10, 3*len(preview)))
    if len(preview) == 1:
        axs = np.expand_dims(axs, 0)
    for i, fname in enumerate(preview):
        sp = os.path.join(src_cls, fname)
        dp = os.path.join(dst_cls, fname)
        try:
            orig = load_gray(sp)
            proc = load_gray(dp)
            diff = np.abs(proc.astype(np.int16) - orig.astype(np.int16)).astype(np.uint8)

            axs[i, 0].imshow(orig, cmap="gray", vmin=0, vmax=255); axs[i, 0].axis("off"); axs[i, 0].set_title(f"{cls}: Original")
            axs[i, 1].imshow(proc, cmap="gray", vmin=0, vmax=255); axs[i, 1].axis("off"); axs[i, 1].set_title("Processed")
            axs[i, 2].imshow(diff, cmap="gray", vmin=0, vmax=255); axs[i, 2].axis("off"); axs[i, 2].set_title("|Diff|")
        except Exception:
            continue
    plt.tight_layout()
    plt.show()

return summary

print("== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TEST ==")
sum_test = preprocess_dataset(test_dir, preprocessed_test, n_preview=3, do_skull_strip=True)

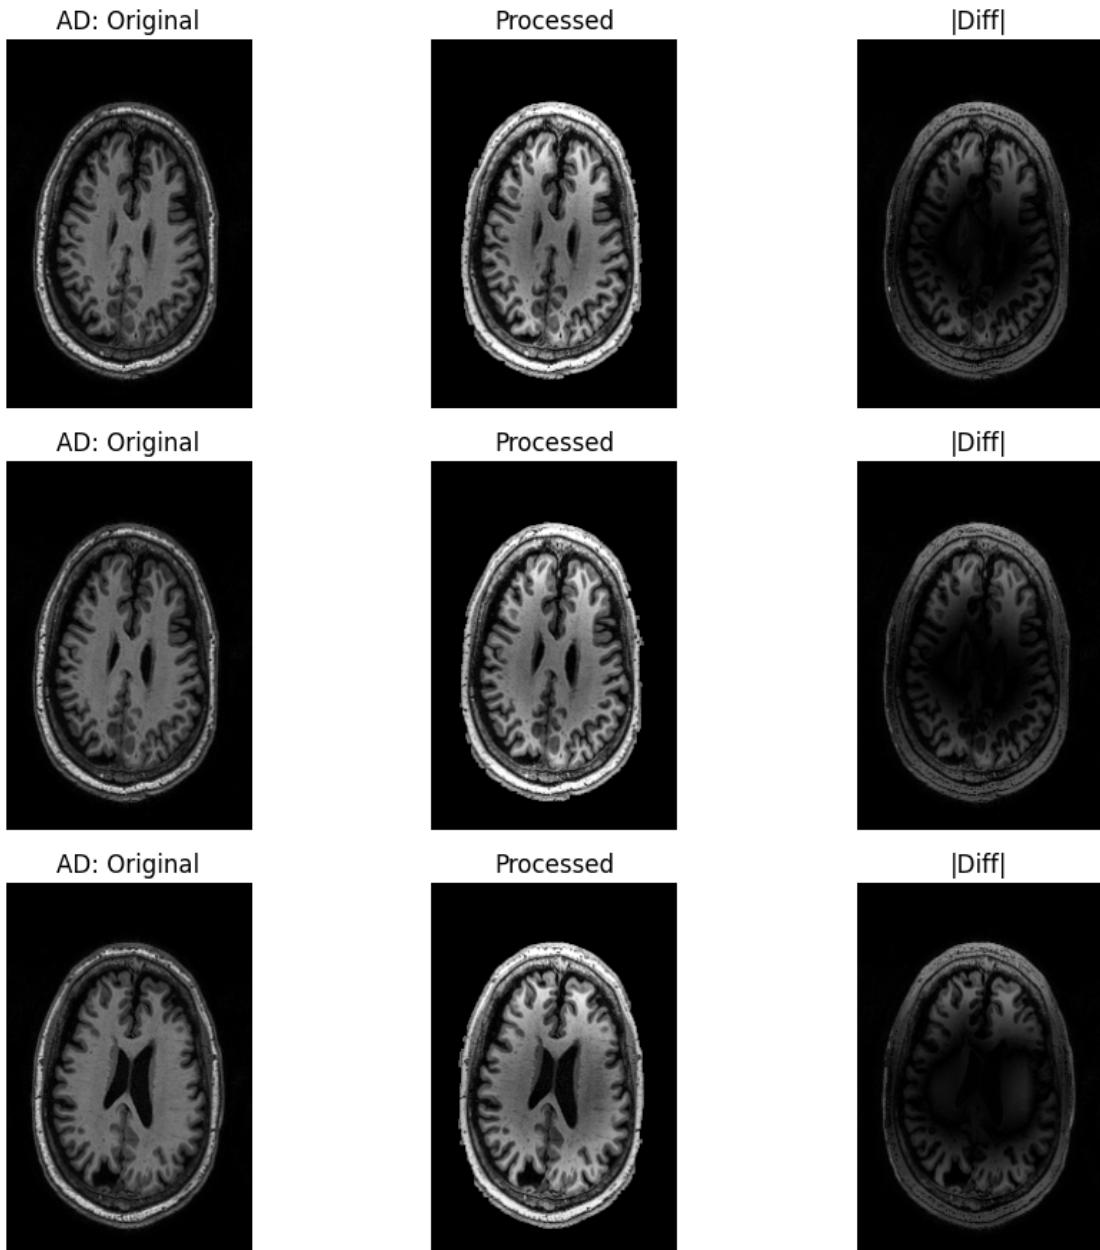
print("== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TRAIN ==")
sum_train = preprocess_dataset(train_dir, preprocessed_train, n_preview=3, do_skull_strip=True)

def _fmt(s): return ", ".join([f"{k}: {v[0]} ok / {v[1]} skipped" for k, v in s.items()])
print("Preprocessed images saved to:")
print(f" • Test : {preprocessed_test}")
print(f" • Train: {preprocessed_train}")
print("Summary TEST :", _fmt(sum_test))
print("Summary TRAIN:", _fmt(sum_train))

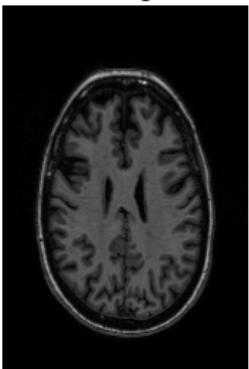
```

```
==== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TEST ====  
Device for homomorphic filtering: cuda
```

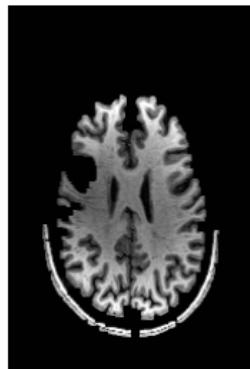
```
/tmp/ipykernel_19/2569239085.py:107: DeprecationWarning: 'mode' parameter is  
deprecated and will be removed in Pillow 13 (2026-10-15)  
Image.fromarray(arr_u8, mode="L").save(path)
```



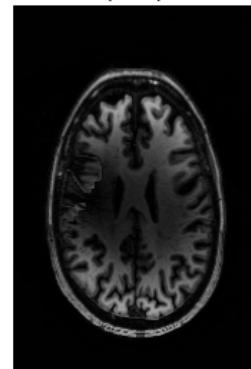
CN: Original



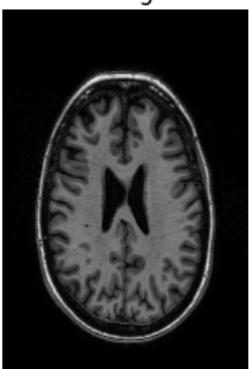
Processed



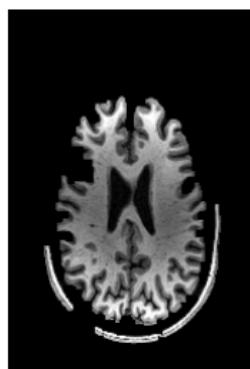
|Diff|



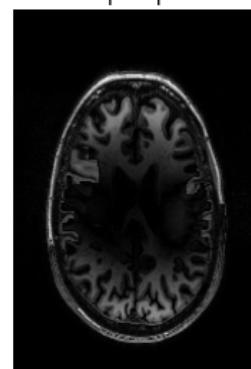
CN: Original



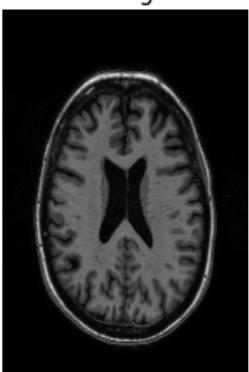
Processed



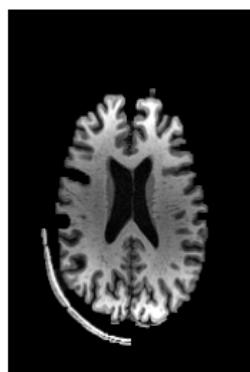
|Diff|



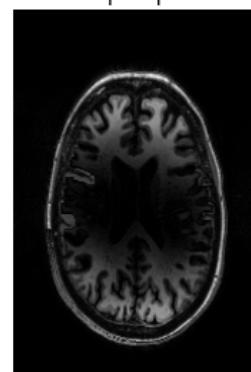
CN: Original

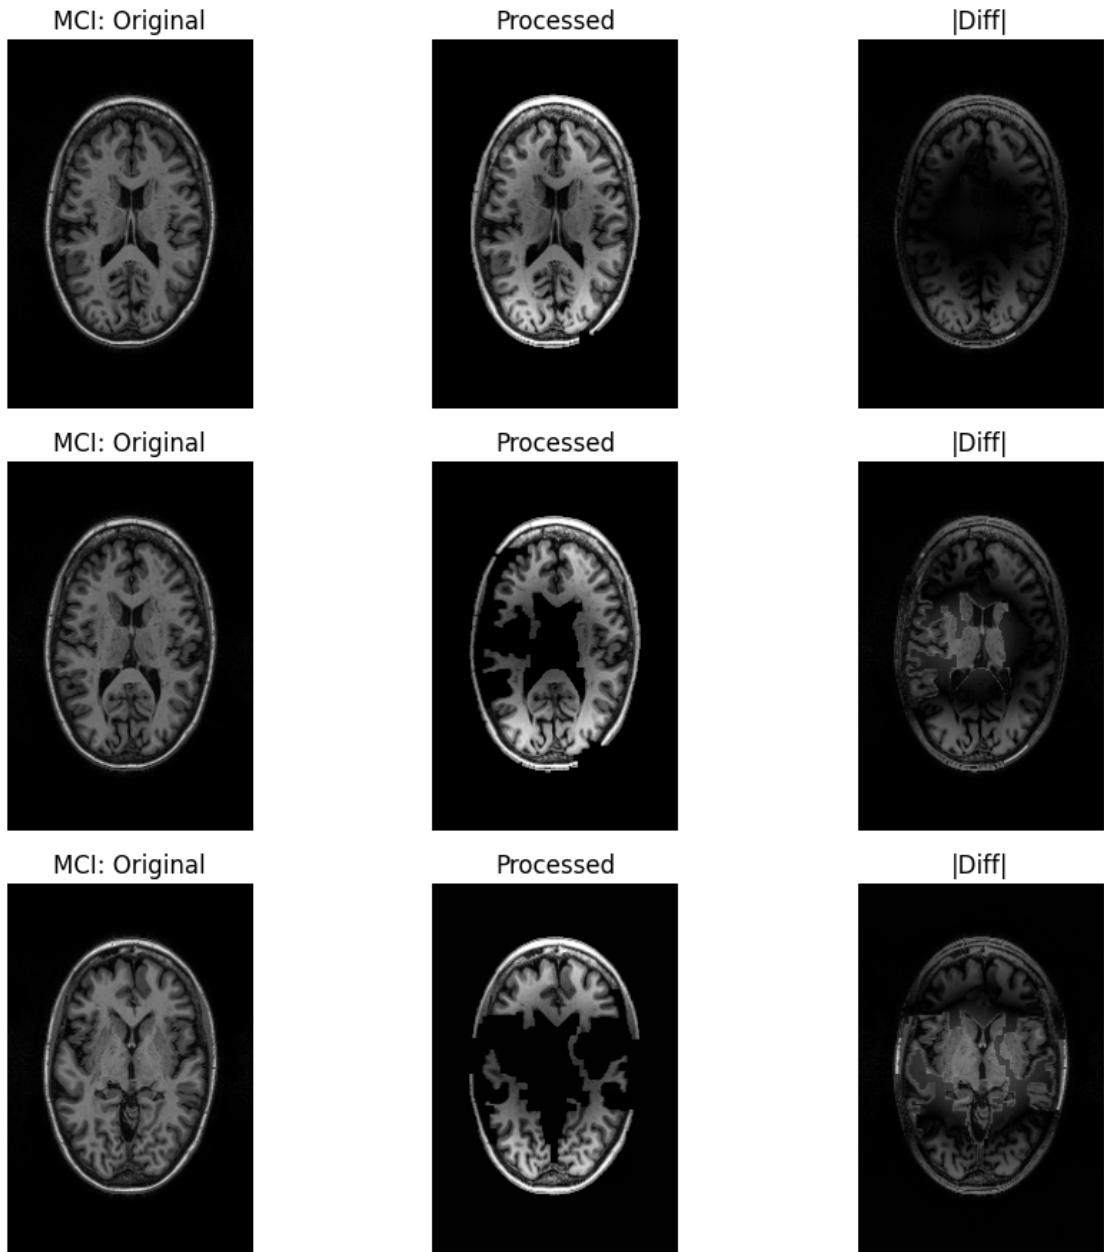


Processed



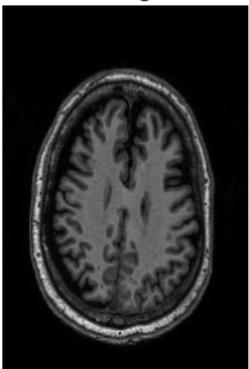
|Diff|



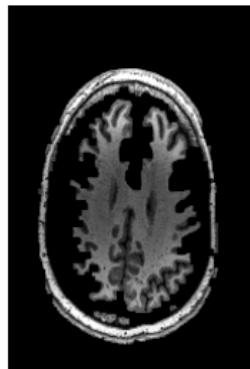


==== 2D MRI Preprocessing (GPU homomorphic + 2D skull-strip): TRAIN ====
Device for homomorphic filtering: cuda

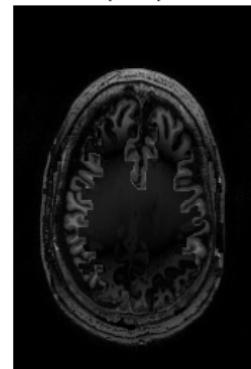
AD: Original



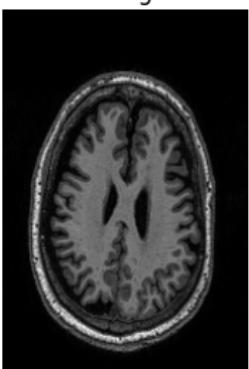
Processed



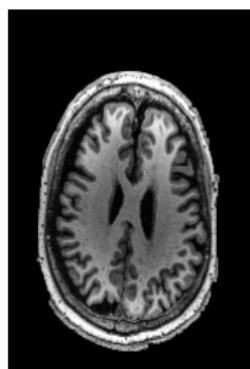
|Diff|



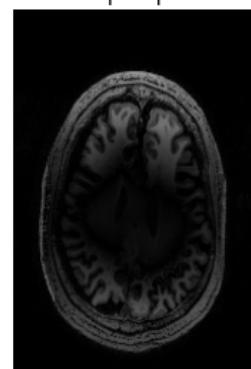
AD: Original



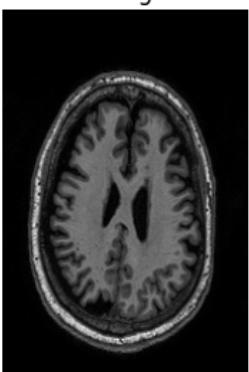
Processed



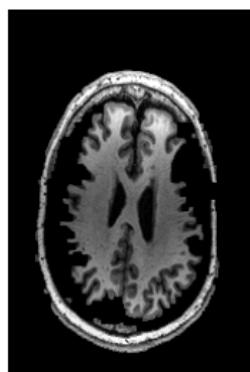
|Diff|



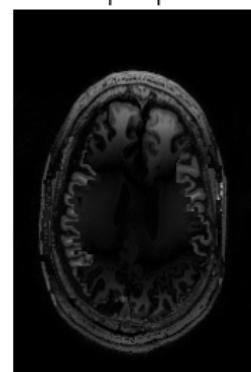
AD: Original



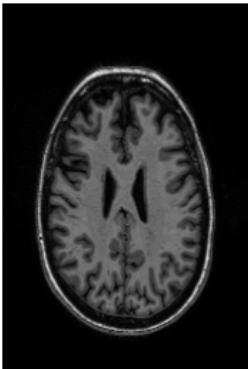
Processed



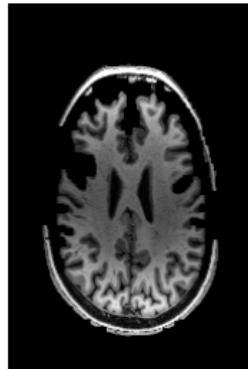
|Diff|



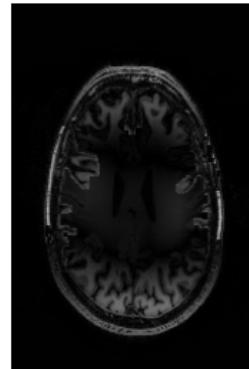
CN: Original



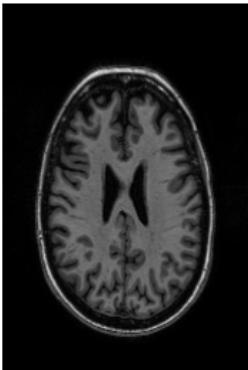
Processed



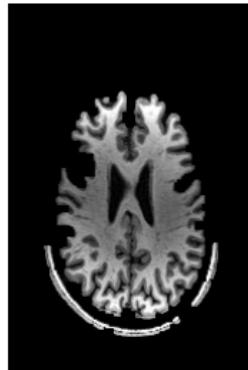
|Diff|



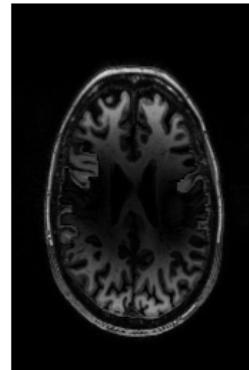
CN: Original



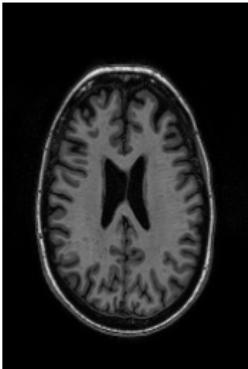
Processed



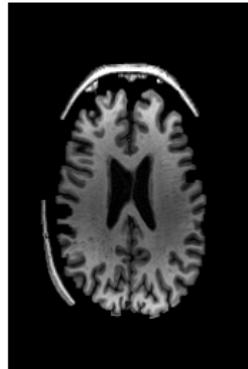
|Diff|



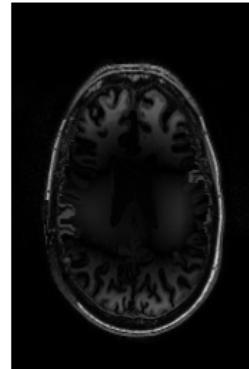
CN: Original

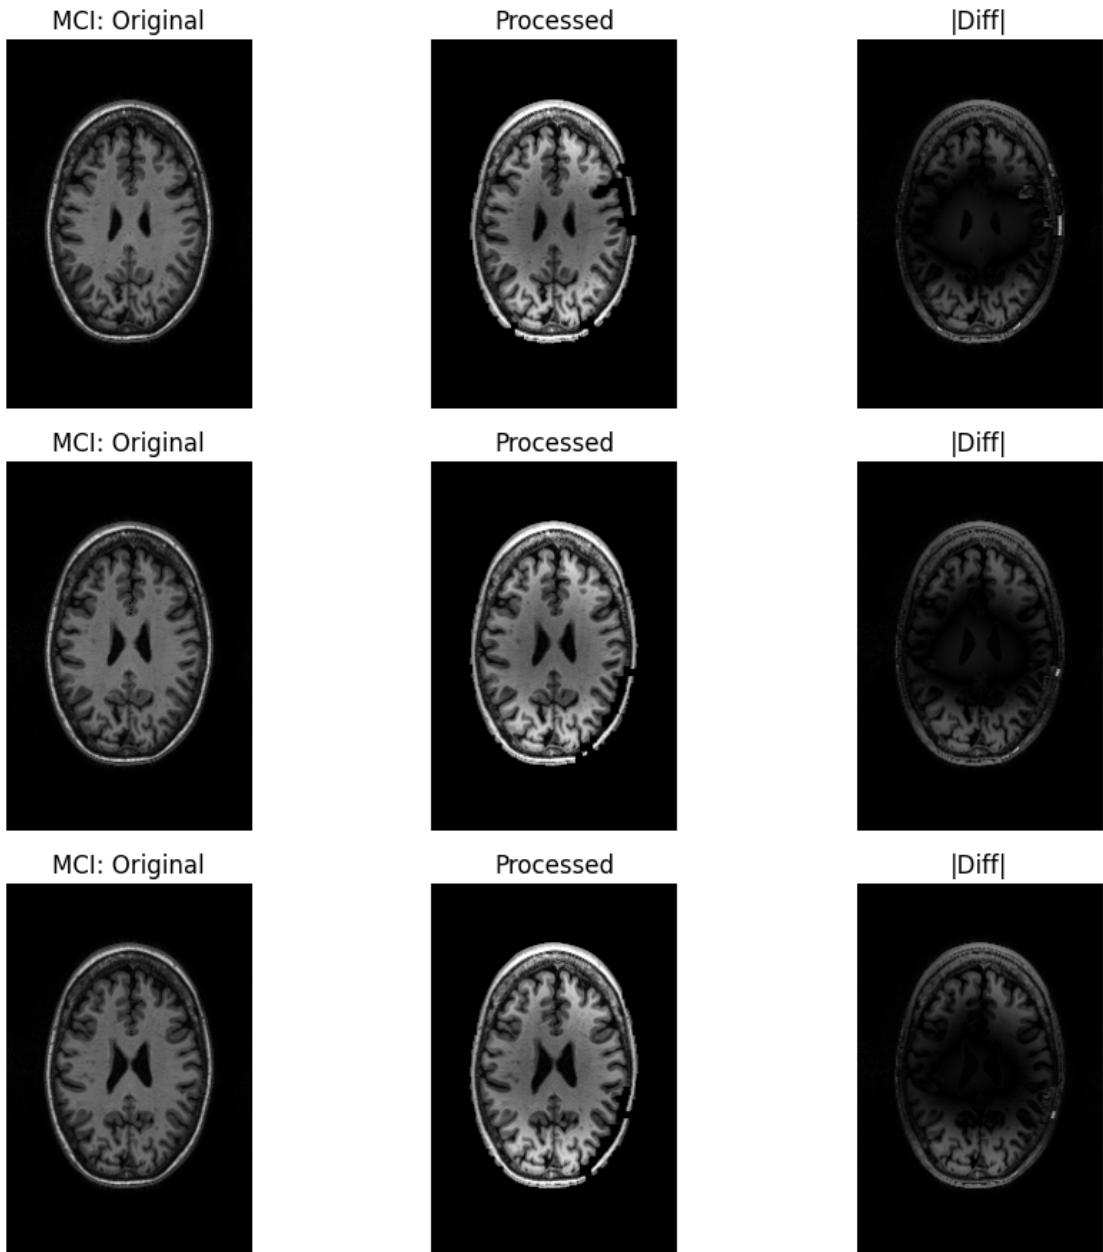


Processed



|Diff|





Preprocessed images saved to:

- Test : /kaggle/working/alzheimer-preprocessed/test
- Train: /kaggle/working/alzheimer-preprocessed/train

Summary TEST : AD: 225 ok / 0 skipped, CN: 288 ok / 0 skipped, MCI: 518 ok / 0 skipped

Summary TRAIN: AD: 899 ok / 0 skipped, CN: 1152 ok / 0 skipped, MCI: 2072 ok / 0 skipped

```
[3]: import os
import numpy as np
from PIL import Image, UnidentifiedImageError
import matplotlib.pyplot as plt

test_dir = "/kaggle/input/alzheimer-disease/test-20251010T094612Z-1-001/test"
train_dir = "/kaggle/input/alzheimer-disease/train-20251010T183510Z-1-001/train"
subfolders = ["AD", "CN", "MCI"]

preprocessed_test = "/kaggle/working/alzheimer-preprocessed/test"
preprocessed_train = "/kaggle/working/alzheimer-preprocessed/train"
resized_root = "/kaggle/working/alzheimer-resized-224"
resized_test = os.path.join(resized_root, "test")
resized_train = os.path.join(resized_root, "train")

IMG_EXT = (".png", ".jpg", ".jpeg", ".bmp", ".tif", ".tiff")
TARGET_SIZE = (224, 224)

def ensure_dir(p: str):
    os.makedirs(p, exist_ok=True)

def load_gray(path: str) -> np.ndarray:
    return np.asarray(Image.open(path).convert("L"), dtype=np.uint8)

def resize_and_save_gray(arr_u8: np.ndarray, path: str):
    img = Image.fromarray(arr_u8, mode="L")
    img_resized = img.resize(TARGET_SIZE, Image.Resampling.LANCZOS)
    img_resized.save(path)

def resize_dataset(src_root: str, dst_root: str, n_preview: int = 3):
    ensure_dir(dst_root)
    summary = {}

    for cls in subfolders:
        src_cls = os.path.join(src_root, cls)
        dst_cls = os.path.join(dst_root, cls)
        ensure_dir(dst_cls)

        processed = skipped = 0
        if not os.path.isdir(src_cls):
            summary[cls] = (0, 0)
            continue

        files = sorted([f for f in os.listdir(src_cls) if f.lower().
        ↪endswith(IMG_EXT)])
        for fname in files:
```

```

        spath = os.path.join(src_cls, fname)
        dpath = os.path.join(dst_cls, fname)
        try:
            arr = load_gray(spath)
            resize_and_save_gray(arr, dpath)
            processed += 1
        except (UnidentifiedImageError, OSError, RuntimeError, ValueError):
            skipped += 1
            continue

    summary[cls] = (processed, skipped)

    preview = files[:n_preview]
    if preview:
        fig, axs = plt.subplots(len(preview), 2, figsize=(8, 4*len(preview)))
        if len(preview) == 1:
            axs = np.expand_dims(axs, 0)
        for i, fname in enumerate(preview):
            sp = os.path.join(src_cls, fname)
            dp = os.path.join(dst_cls, fname)
            try:
                orig = load_gray(sp)
                resized = load_gray(dp)

                axs[i, 0].imshow(orig, cmap="gray", vmin=0, vmax=255)
                axs[i, 0].axis("off")
                axs[i, 0].set_title(f"{cls}: Original {orig.shape}")

                axs[i, 1].imshow(resized, cmap="gray", vmin=0, vmax=255)
                axs[i, 1].axis("off")
                axs[i, 1].set_title(f"Resized {resized.shape}")
            except Exception:
                continue
        plt.tight_layout()
        plt.show()

    return summary

print("== Resizing Preprocessed Images: TEST ==")
sum_test = resize_dataset(preprocessed_test, resized_test, n_preview=3)

print("== Resizing Preprocessed Images: TRAIN ==")
sum_train = resize_dataset(preprocessed_train, resized_train, n_preview=3)

def _fmt(s): return ", ".join([f"{k}: {v[0]} ok / {v[1]} skipped" for k, v in s.items()])

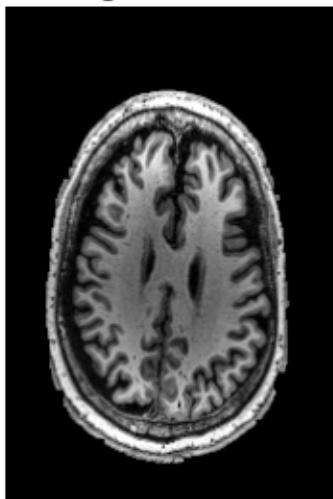
```

```
print("Resized images saved to:")
print(f" • Resized Test : {resized_test}")
print(f" • Resized Train: {resized_train}")
print("Summary TEST :", _fmt(sum_test))
print("Summary TRAIN:", _fmt(sum_train))
```

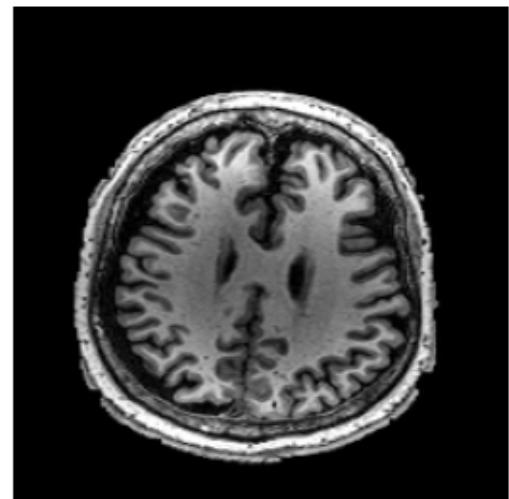
==== Resizing Preprocessed Images: TEST ===

```
/tmp/ipykernel_19/502029755.py:26: DeprecationWarning: 'mode' parameter is
deprecated and will be removed in Pillow 13 (2026-10-15)
    img = Image.fromarray(arr_u8, mode="L")
```

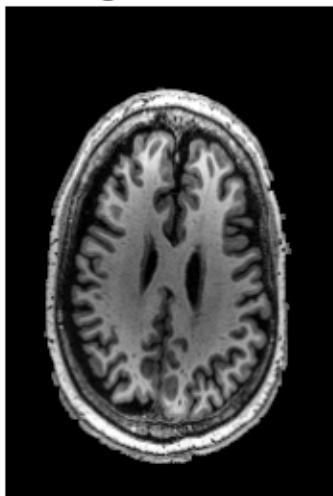
AD: Original (256, 170)



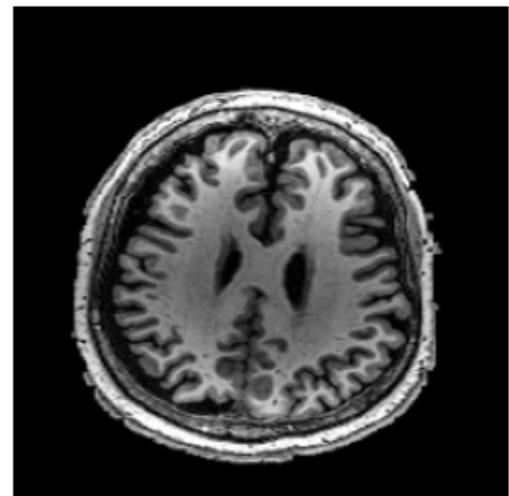
Resized (224, 224)



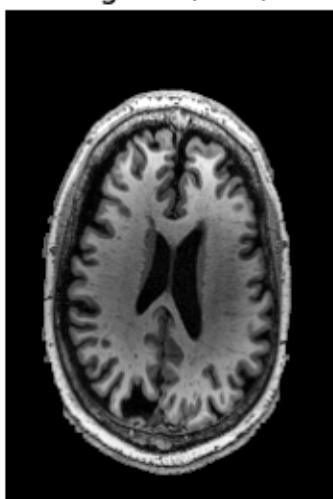
AD: Original (256, 170)



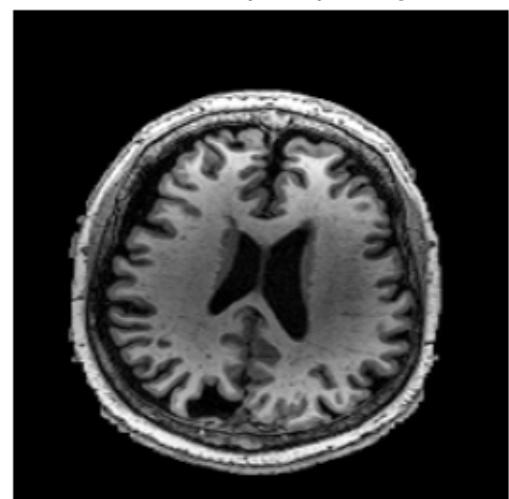
Resized (224, 224)



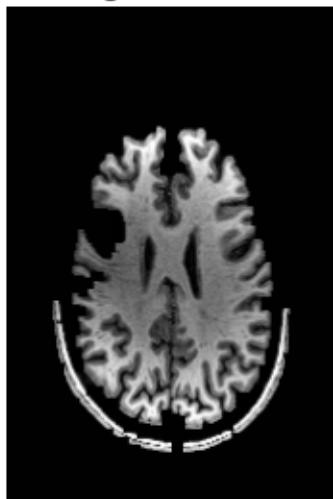
AD: Original (256, 170)



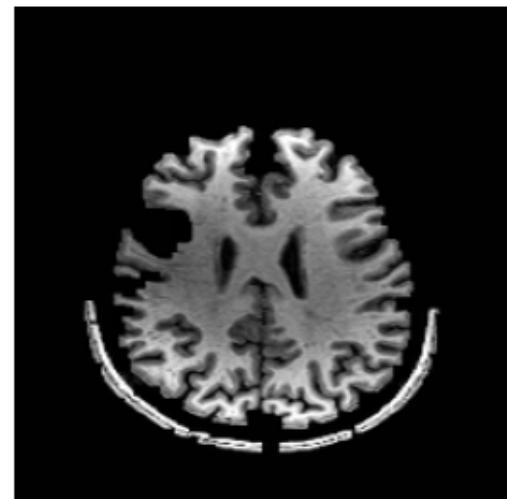
Resized (224, 224)



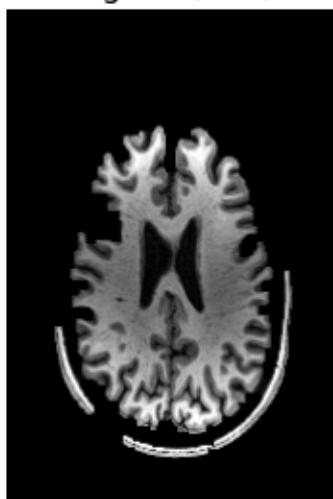
CN: Original (256, 170)



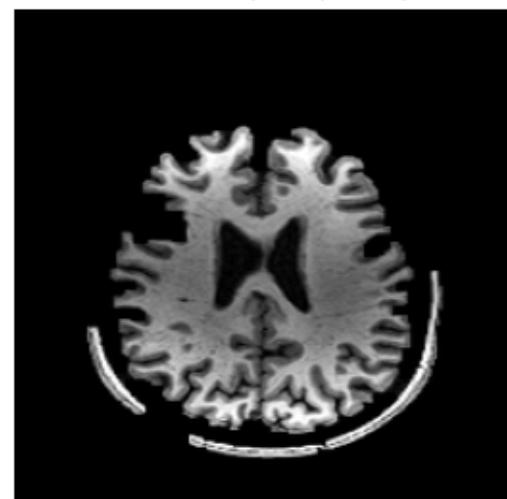
Resized (224, 224)



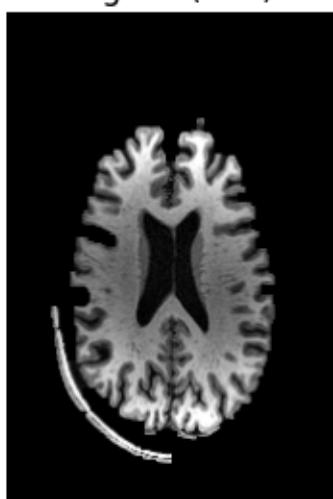
CN: Original (256, 170)



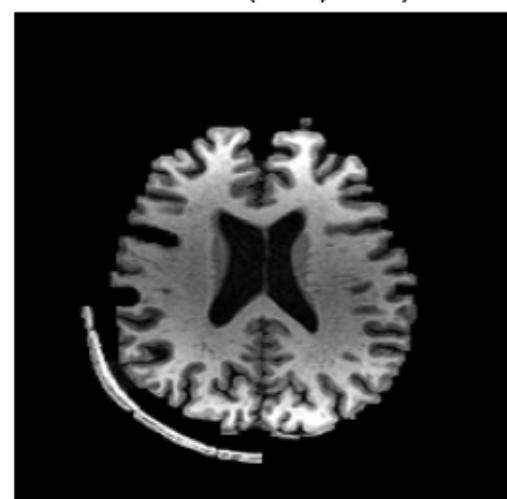
Resized (224, 224)



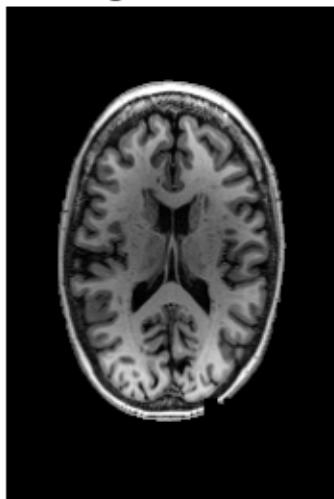
CN: Original (256, 170)



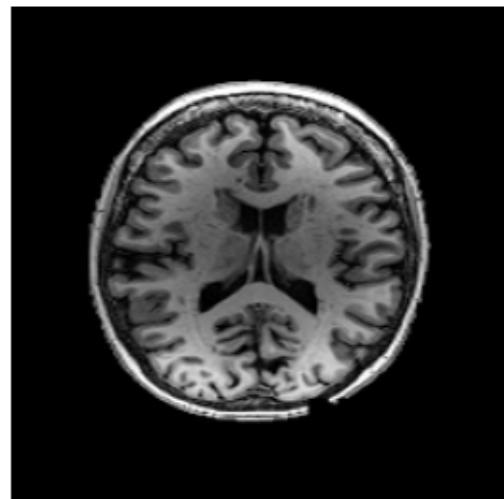
Resized (224, 224)



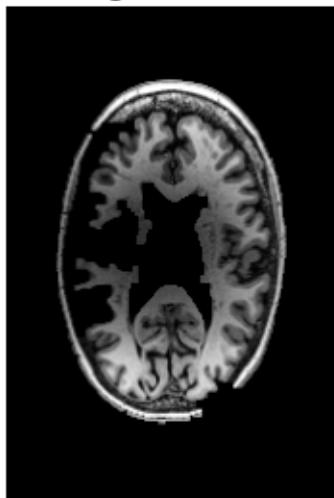
MCI: Original (256, 170)



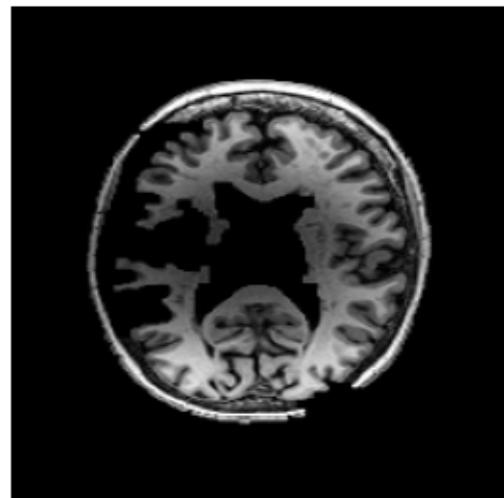
Resized (224, 224)



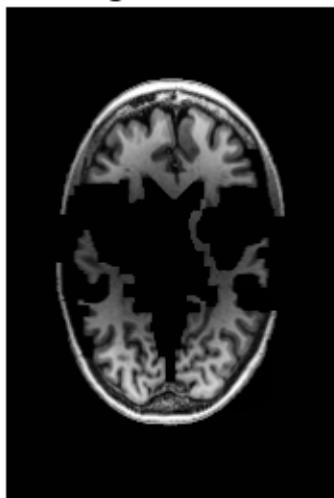
MCI: Original (256, 170)



Resized (224, 224)



MCI: Original (256, 170)

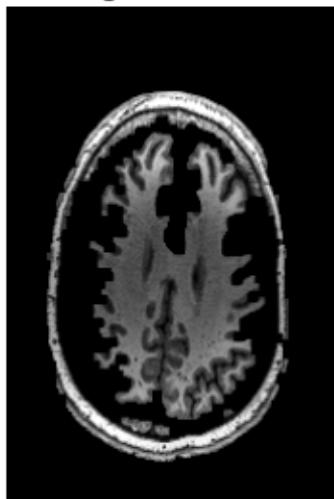


Resized (224, 224)

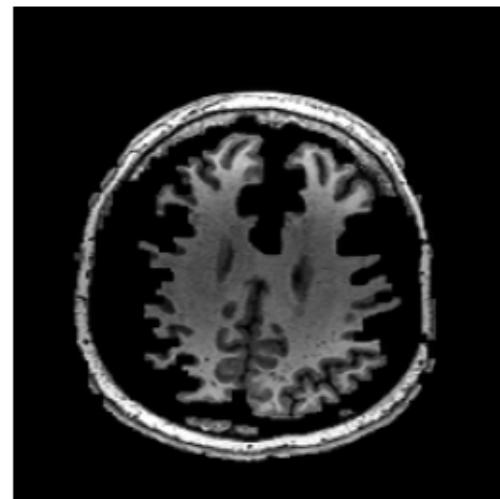


==== Resizing Preprocessed Images: TRAIN ===

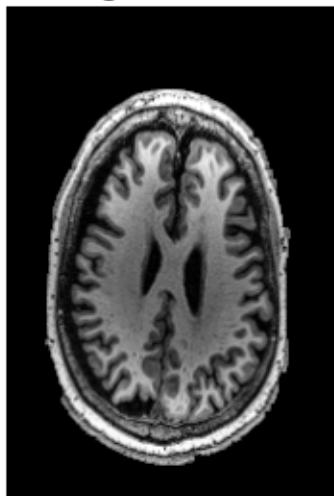
AD: Original (256, 170)



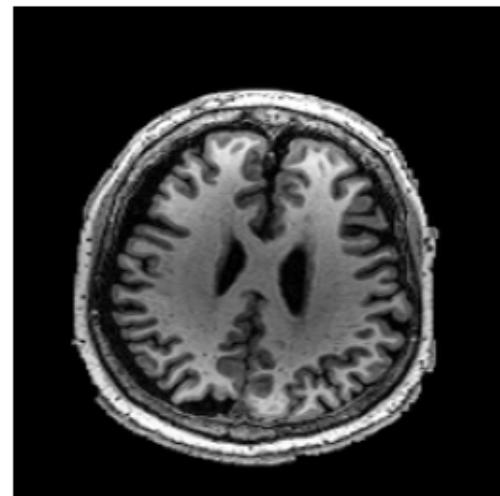
Resized (224, 224)



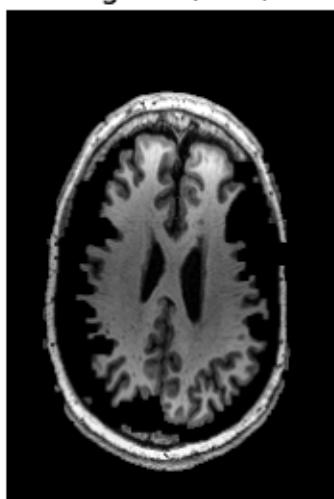
AD: Original (256, 170)



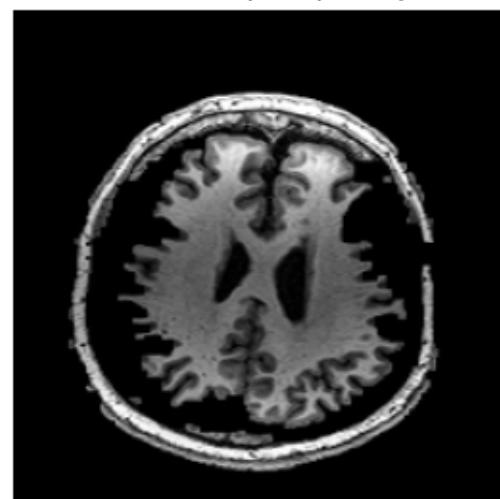
Resized (224, 224)



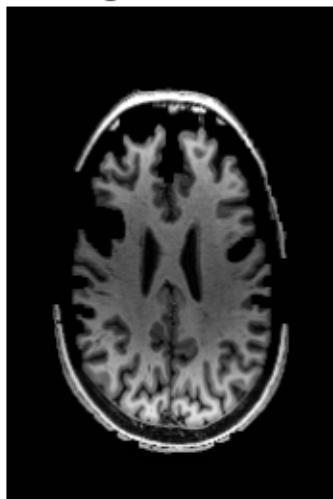
AD: Original (256, 170)



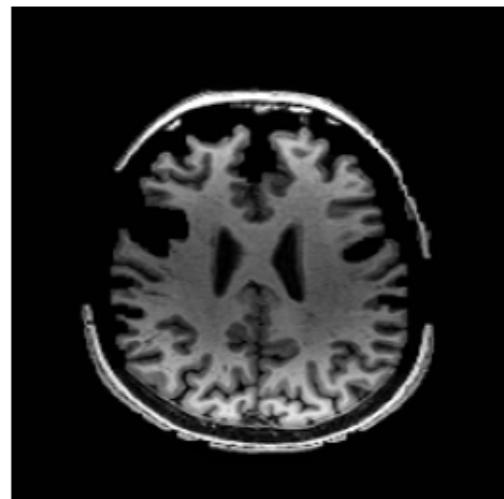
Resized (224, 224)



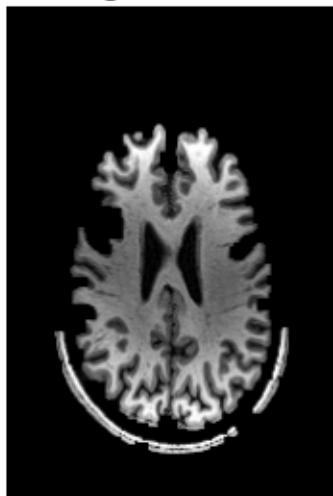
CN: Original (256, 170)



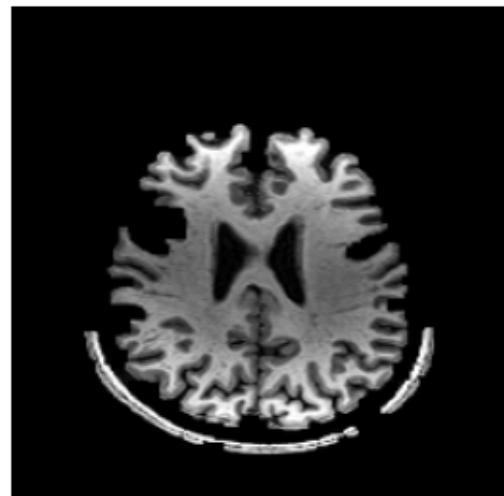
Resized (224, 224)



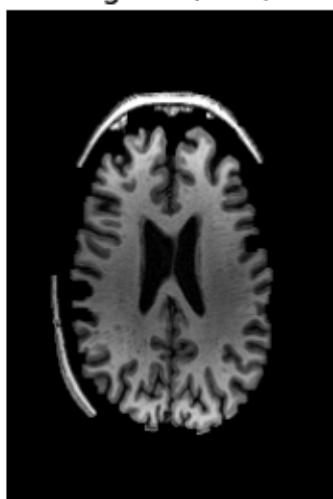
CN: Original (256, 170)



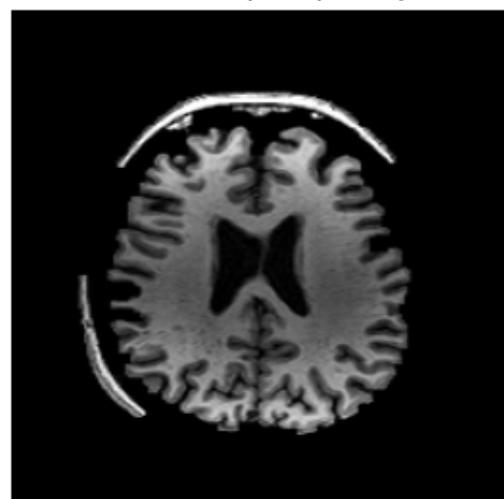
Resized (224, 224)



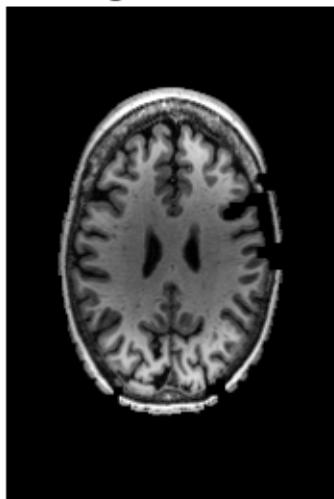
CN: Original (256, 170)



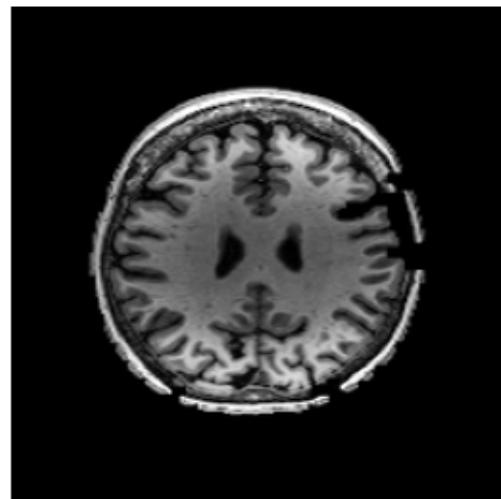
Resized (224, 224)



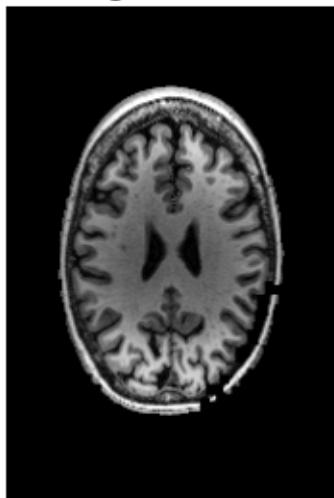
MCI: Original (256, 170)



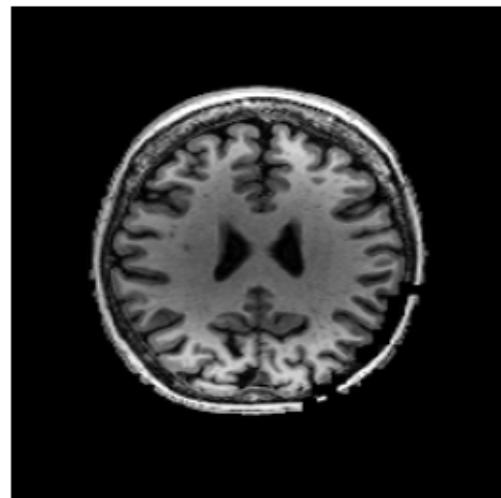
Resized (224, 224)



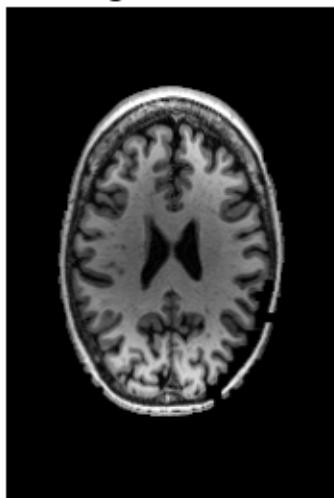
MCI: Original (256, 170)



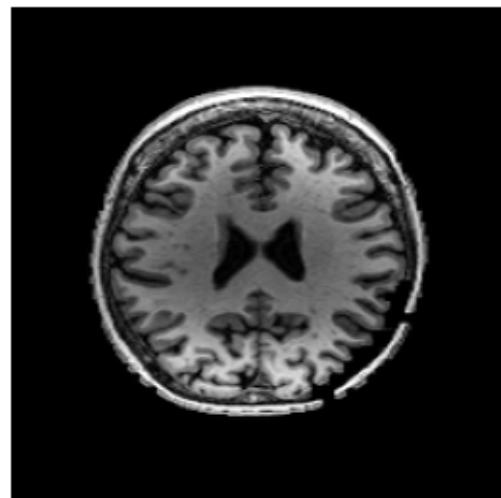
Resized (224, 224)



MCI: Original (256, 170)



Resized (224, 224)



```

Resized images saved to:
  • Resized Test : /kaggle/working/alzheimer-resized-224/test
  • Resized Train: /kaggle/working/alzheimer-resized-224/train
Summary TEST : AD: 225 ok / 0 skipped, CN: 288 ok / 0 skipped, MCI: 518 ok / 0
skipped
Summary TRAIN: AD: 899 ok / 0 skipped, CN: 1152 ok / 0 skipped, MCI: 2072 ok / 0
skipped

```

Train And Test Split

```
[4]: import os
import sys
import csv
import random
import shutil
from pathlib import Path
from typing import Dict, List, Tuple

# -----
# Configuration (modified for resized images)
# -----
# Source resized dataset (224x224)
RESIZED_ROOT = os.environ.get("RESIZED_ROOT", "/kaggle/working/
    ↪alzheimer-resized-224")
SRC_TRAIN = os.path.join(RESIZED_ROOT, "train")
SRC_TEST = os.path.join(RESIZED_ROOT, "test")

# Classes (same as before)
CLASSES = ["AD", "CN", "MCI"]

# Where to write splits
SPLITS_ROOT = "/kaggle/working/alzheimer-resized-224_splits"
os.makedirs(SPLITS_ROOT, exist_ok=True)

# The split ratios (train:test) you requested
RATIO_LIST = [
    (0.90, 0.10),
    (0.80, 0.20),
    (0.70, 0.30),
    (0.60, 0.40),
    (0.50, 0.50),
    (0.40, 0.60),
    (0.30, 0.70),
    (0.20, 0.80),
    (0.10, 0.90)
```

```

]

# Validation share taken from the training portion
VAL_FRACTION = 0.10

# Base random seed (deterministic builds). Each ratio derives its own seed.
BASE_SEED = 2025

# -----
# Helpers
# -----
IMG_EXTS = (".png", ".jpg", ".jpeg", ".bmp", ".tif", ".tiff")

def list_images_in_class(class_dir: str) -> List[str]:
    """Return absolute paths of all images in a class directory."""
    if not os.path.isdir(class_dir):
        return []
    files = sorted([
        str(Path(class_dir) / f) for f in os.listdir(class_dir)
        if f.lower().endswith(IMG_EXTS)
    ])
    return files

def gather_all_images() -> Dict[str, List[str]]:
    """
    Gather all resized images per class, from both 'train' and 'test'
    to form a single full pool for stratified splitting by ratio.
    """
    all_by_class = {c: [] for c in CLASSES}
    for c in CLASSES:
        # From resized train
        all_by_class[c].extend(list_images_in_class(os.path.join(SRC_TRAIN, c)))
        # From resized test
        all_by_class[c].extend(list_images_in_class(os.path.join(SRC_TEST, c)))
    return all_by_class

def ensure_dirs(*paths: str):
    for p in paths:
        os.makedirs(p, exist_ok=True)

def link_or_copy(src: str, dst: str):
    """Create a symlink; if not permitted, copy the file."""
    try:
        # Remove dst if exists
        if os.path.lexists(dst):
            os.unlink(dst)
        os.symlink(src, dst)
    except OSError:
        shutil.copy(src, dst)

```

```

except OSError:
    shutil.copy2(src, dst)

def write_manifest(split_dir: str, split_name: str, rows: List[Tuple[str, str, str, str]]):
    """
    Write both CSV manifest and plain path list.
    rows: list of (split, cls, filename, src_path)
    """
    # CSV
    csv_path = os.path.join(split_dir, f"{split_name}.csv")
    with open(csv_path, "w", newline="") as f:
        w = csv.writer(f)
        w.writerow(["split", "class", "filename", "path"])
        for r in rows:
            w.writerow(r)

    # TXT list
    txt_path = os.path.join(split_dir, f"{split_name}.txt")
    with open(txt_path, "w") as f:
        for _, _, _, p in rows:
            f.write(p + "\n")

def summarize_counts(counts: Dict[str, Dict[str, int]]):
    """
    Print counts per split (train/val/test) and per class for quick sanity
    check.
    """
    for split in ["train", "val", "test"]:
        info = counts.get(split, {})
        total = sum(info.values())
        detail = ", ".join([f"{k}: {v}" for k, v in info.items()])
        print(f"{split.capitalize():5s} => total {total:5d} | {detail}")

# -----
# Split builder
# -----
def build_splits():
    # 1) Pool all images by class (union of resized/train and resized/test)
    all_by_class = gather_all_images()

    # Optional: quick report of total availability
    print("Total images by class (union of resized/train + resized/test):")
    for c in CLASSES:
        print(f" {c}: {len(all_by_class[c])}")

    # 2) For each (train_ratio, test_ratio), create a split folder and populate

```

```

for tr_ratio, te_ratio in RATIO_LIST:
    # sanity: ratios sum approx 1
    assert abs(tr_ratio + te_ratio - 1.0) < 1e-6, "Train+Test ratio must sum to 1"

    # deterministic seed per ratio
    seed = BASE_SEED + int(round(te_ratio * 100))
    rng = random.Random(seed)

    # Split name and dirs
    split_name = f"split_{int(round(tr_ratio*100))}_{int(round(te_ratio*100))}"
    split_root = os.path.join(SPLITS_ROOT, split_name)
    train_root = os.path.join(split_root, "train")
    val_root = os.path.join(split_root, "val")
    test_root = os.path.join(split_root, "test")
    ensure_dirs(split_root, train_root, val_root, test_root)
    for c in CLASSES:
        ensure_dirs(os.path.join(train_root, c), os.path.join(val_root, c), os.path.join(test_root, c))

    rows_train, rows_val, rows_test = [], [], []
    counts = {"train": {}, "val": {}, "test": {}}

    # 3) Per-class stratified splitting
    for c in CLASSES:
        full_list = list(all_by_class[c]) # copy
        rng.shuffle(full_list)           # deterministic shuffle

        n_total = len(full_list)
        n_test = max(0, int(round(n_total * te_ratio)))
        n_test = min(n_test, n_total) # guard

        test_list = full_list[:n_test]
        train_pool = full_list[n_test:]

        # Validation from training portion (10%)
        n_val = max(0, int(round(len(train_pool) * VAL_FRACTION)))
        val_list = train_pool[:n_val]
        train_list = train_pool[n_val:]

    # 4) Materialize (symlink/copy) into folders and write manifests
    # test
    for src_path in test_list:
        fname = os.path.basename(src_path)
        dst_path = os.path.join(test_root, c, fname)
        link_or_copy(src_path, dst_path)

```

```

        rows_test.append(("test", c, fname, dst_path))

    # val
    for src_path in val_list:
        fname = os.path.basename(src_path)
        dst_path = os.path.join(val_root, c, fname)
        link_or_copy(src_path, dst_path)
        rows_val.append(("val", c, fname, dst_path))

    # train
    for src_path in train_list:
        fname = os.path.basename(src_path)
        dst_path = os.path.join(train_root, c, fname)
        link_or_copy(src_path, dst_path)
        rows_train.append(("train", c, fname, dst_path))

    # counts
    counts["test"][c] = len(test_list)
    counts["val"][c] = len(val_list)
    counts["train"][c] = len(train_list)

    # 5) Write manifest files for this split
    write_manifest(split_root, "train", rows_train)
    write_manifest(split_root, "val", rows_val)
    write_manifest(split_root, "test", rows_test)

    # 6) Summary printout
    print(f"\n==== {split_name} ====")
    summarize_counts(counts)
    print(f"Paths:\n  Train: {train_root}\n  Val : {val_root}\n  Test : {test_root}\n")

    # 7) Export env var for convenient access in later cells
    os.environ["RESIZED_SPLITS_ROOT"] = SPLITS_ROOT
    print(f"All splits created under: {SPLITS_ROOT}")

# -----
# Execute
# -----
build_splits()

```

Total images by class (union of resized/train + resized/test):

AD: 1124
CN: 1440
MCI: 2590

==== split_90_10 ===

```

Train => total 4175 | AD: 911, CN: 1166, MCI: 2098
Val   => total 464 | AD: 101, CN: 130, MCI: 233
Test  => total 515 | AD: 112, CN: 144, MCI: 259
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_90_10/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_90_10/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_90_10/test

==== split_80_20 ====
Train => total 3711 | AD: 809, CN: 1037, MCI: 1865
Val   => total 412 | AD: 90, CN: 115, MCI: 207
Test  => total 1031 | AD: 225, CN: 288, MCI: 518
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_80_20/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_80_20/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_80_20/test

==== split_70_30 ====
Train => total 3247 | AD: 708, CN: 907, MCI: 1632
Val   => total 361 | AD: 79, CN: 101, MCI: 181
Test  => total 1546 | AD: 337, CN: 432, MCI: 777
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_70_30/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_70_30/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_70_30/test

==== split_60_40 ====
Train => total 2784 | AD: 607, CN: 778, MCI: 1399
Val   => total 308 | AD: 67, CN: 86, MCI: 155
Test  => total 2062 | AD: 450, CN: 576, MCI: 1036
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_60_40/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_60_40/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_60_40/test

==== split_50_50 ====
Train => total 2319 | AD: 506, CN: 648, MCI: 1165
Val   => total 258 | AD: 56, CN: 72, MCI: 130
Test  => total 2577 | AD: 562, CN: 720, MCI: 1295
Paths:
  Train: /kaggle/working/alzheimer-resized-224_splits/split_50_50/train
  Val  : /kaggle/working/alzheimer-resized-224_splits/split_50_50/val
  Test : /kaggle/working/alzheimer-resized-224_splits/split_50_50/test

```

```

==== split_40_60 ====
Train => total 1855 | AD: 405, CN: 518, MCI: 932
Val   => total 207 | AD: 45, CN: 58, MCI: 104
Test  => total 3092 | AD: 674, CN: 864, MCI: 1554
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_40_60/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_40_60/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_40_60/test

==== split_30_70 ====
Train => total 1391 | AD: 303, CN: 389, MCI: 699
Val   => total 155 | AD: 34, CN: 43, MCI: 78
Test  => total 3608 | AD: 787, CN: 1008, MCI: 1813
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_30_70/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_30_70/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_30_70/test

==== split_20_80 ====
Train => total 928 | AD: 203, CN: 259, MCI: 466
Val   => total 103 | AD: 22, CN: 29, MCI: 52
Test  => total 4123 | AD: 899, CN: 1152, MCI: 2072
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_20_80/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_20_80/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_20_80/test

==== split_10_90 ====
Train => total 464 | AD: 101, CN: 130, MCI: 233
Val   => total 51 | AD: 11, CN: 14, MCI: 26
Test  => total 4639 | AD: 1012, CN: 1296, MCI: 2331
Paths:
    Train: /kaggle/working/alzheimer-resized-224_splits/split_10_90/train
    Val : /kaggle/working/alzheimer-resized-224_splits/split_10_90/val
    Test : /kaggle/working/alzheimer-resized-224_splits/split_10_90/test

All splits created under: /kaggle/working/alzheimer-resized-224_splits

```

```
[5]: #MobileNetV3-Large
import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
```

```

from torchvision import transforms, models
import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import ReduceLROnPlateau

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)

        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label

    def get_data_transforms():
        train_transform = transforms.Compose([
            transforms.Resize((224, 224)),
            transforms.RandomHorizontalFlip(p=0.5),
            transforms.RandomRotation(10),
            transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
            transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        ])

```

```

    ])

    val_transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
        ↪225])
    ])

    return train_transform, val_transform

def train_model(model, train_loader, val_loader, criterion, optimizer, ↪
    ↪scheduler, num_epochs, device):
    best_val_acc = 0
    patience = 5
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        epoch_loss = running_loss / len(train_loader)
        epoch_acc = 100 * correct / total

        val_acc = evaluate_model(model, val_loader, device)

        if scheduler:
            scheduler.step(val_acc)

        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train ↪
            ↪Acc: {epoch_acc:.2f}%, Val Acc: {val_acc:.2f}%')

```

```

        if val_acc > best_val_acc:
            best_val_acc = val_acc
            patience_counter = 0
        else:
            patience_counter += 1

        if patience_counter >= patience:
            print(f"Early stopping at epoch {epoch+1}")
            break

    def evaluate_model(model, data_loader, device):
        model.eval()
        correct = 0
        total = 0

        with torch.no_grad():
            for images, labels in data_loader:
                images, labels = images.to(device), labels.to(device)
                outputs = model(images)
                _, predicted = torch.max(outputs.data, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

        accuracy = 100 * correct / total
        return accuracy

    def test_model(model, test_loader, device):
        model.eval()
        all_preds = []
        all_labels = []

        with torch.no_grad():
            for images, labels in test_loader:
                images, labels = images.to(device), labels.to(device)
                outputs = model(images)
                _, predicted = torch.max(outputs.data, 1)

                all_preds.extend(predicted.cpu().numpy())
                all_labels.extend(labels.cpu().numpy())

        return all_preds, all_labels

    def calculate_metrics(y_true, y_pred, split_name):
        accuracy = accuracy_score(y_true, y_pred)
        precision = precision_score(y_true, y_pred, average='weighted', zero_division=0)

```

```

recall = recall_score(y_true, y_pred, average='weighted', zero_division=0)
f1 = f1_score(y_true, y_pred, average='weighted', zero_division=0)

print(f"\n==== {split_name} Results ===")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print("\nClassification Report:")
print(classification_report(y_true, y_pred, target_names=['AD', 'CN',
    ↪'MCI'], zero_division=0))

return {
    'split': split_name,
    'accuracy': accuracy,
    'precision': precision,
    'recall': recall,
    'f1_score': f1
}

def run_mobilenetv3_large_on_splits():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    splits_root = "/kaggle/working/alzheimer-resized-224_splits"
    split_folders = [f for f in os.listdir(splits_root) if f.
        ↪startswith('split_')]
    split_folders.sort()

    results = []

    train_transform, val_transform = get_data_transforms()

    for split_folder in split_folders:
        print(f"\n{'='*60}")
        print(f"Processing: {split_folder}")
        print(f"{'='*60}")

        split_path = os.path.join(splits_root, split_folder)

        train_dataset = AlzheimerDataset(split_path, 'train', train_transform)
        val_dataset = AlzheimerDataset(split_path, 'val', val_transform)
        test_dataset = AlzheimerDataset(split_path, 'test', val_transform)

        train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True,
            ↪num_workers=2)

```

```

    val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False,□
    ↵num_workers=2)
    test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False,□
    ↵num_workers=2)

    print(f"Train samples: {len(train_dataset)}")
    print(f"Val samples: {len(val_dataset)}")
    print(f"Test samples: {len(test_dataset)}")

    model = models.mobilenet_v3_large(weights=models.□
    ↵MobileNet_V3_Large_Weights.IMGNET1K_V2)
    num_ftrs = model.classifier[3].in_features
    model.classifier[3] = nn.Linear(num_ftrs, 3)
    model = model.to(device)

    criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
    optimizer = torch.optim.AdamW(model.parameters(), lr=0.0001,□
    ↵weight_decay=1e-4)
    scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5,□
    ↵patience=3, verbose=True)

    print("Starting training...")
    start_time = time.time()
    train_model(model, train_loader, val_loader, criterion, optimizer,□
    ↵scheduler, num_epochs=30, device=device)
    training_time = time.time() - start_time

    print("Testing model...")
    test_preds, test_labels = test_model(model, test_loader, device)

    split_results = calculate_metrics(test_labels, test_preds, split_folder)
    split_results['training_time'] = training_time
    results.append(split_results)

    torch.cuda.empty_cache()

results_df = pd.DataFrame(results)
print(f"\n{'='*80}")
print("MobileNetV3-Large - SUMMARY OF ALL SPLITS")
print(f"\n{'='*80}")
print(results_df.to_string(index=False))

results_csv_path = "/kaggle/working/mobilenetv3_large_results.csv"
results_df.to_csv(results_csv_path, index=False)
print(f"\nDetailed results saved to: {results_csv_path}")

```

```

    return results_df

if __name__ == "__main__":
    results = run_mobilenetv3_large_on_splits()

```

Using device: cuda

```

=====
Processing: split_10_90
=====
Train samples: 464
Val samples: 51
Test samples: 4639

Downloading:
"https://download.pytorch.org/models/mobilenet_v3_large-5c1a4163.pth" to
/root/.cache/torch/hub/checkpoints/mobilenet_v3_large-5c1a4163.pth
100%|      | 21.1M/21.1M [00:00<00:00, 232MB/s]

Starting training...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/30], Loss: 1.0622, Train Acc: 48.49%, Val Acc: 52.94%
Epoch [2/30], Loss: 1.0318, Train Acc: 50.43%, Val Acc: 54.90%
Epoch [3/30], Loss: 1.0103, Train Acc: 51.94%, Val Acc: 60.78%
Epoch [4/30], Loss: 0.9779, Train Acc: 54.09%, Val Acc: 58.82%
Epoch [5/30], Loss: 0.9336, Train Acc: 58.84%, Val Acc: 60.78%
Epoch [6/30], Loss: 0.9025, Train Acc: 62.72%, Val Acc: 60.78%
Epoch [7/30], Loss: 0.8577, Train Acc: 66.59%, Val Acc: 56.86%
Epoch [8/30], Loss: 0.7784, Train Acc: 71.98%, Val Acc: 52.94%
Early stopping at epoch 8
Testing model...

==== split_10_90 Results ===
Accuracy: 0.5115
Precision: 0.4727
Recall: 0.5115
F1-Score: 0.4502

Classification Report:
      precision    recall   f1-score   support
      AD        0.37      0.08      0.14      1012
      CN        0.43      0.26      0.33      1296
```

MCI	0.54	0.83	0.65	2331
accuracy			0.51	4639
macro avg	0.45	0.39	0.37	4639
weighted avg	0.47	0.51	0.45	4639

```
=====
Processing: split_20_80
=====
Train samples: 928
Val samples: 103
Test samples: 4123
Starting training...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/30], Loss: 1.0566, Train Acc: 48.92%, Val Acc: 50.49%
Epoch [2/30], Loss: 1.0221, Train Acc: 51.51%, Val Acc: 52.43%
Epoch [3/30], Loss: 0.9887, Train Acc: 52.16%, Val Acc: 55.34%
Epoch [4/30], Loss: 0.9414, Train Acc: 57.22%, Val Acc: 50.49%
Epoch [5/30], Loss: 0.9113, Train Acc: 59.91%, Val Acc: 52.43%
Epoch [6/30], Loss: 0.8799, Train Acc: 62.18%, Val Acc: 48.54%
Epoch [7/30], Loss: 0.8156, Train Acc: 68.00%, Val Acc: 60.19%
Epoch [8/30], Loss: 0.7624, Train Acc: 72.09%, Val Acc: 57.28%
Epoch [9/30], Loss: 0.7450, Train Acc: 73.17%, Val Acc: 56.31%
Epoch [10/30], Loss: 0.7020, Train Acc: 76.62%, Val Acc: 68.93%
Epoch [11/30], Loss: 0.7098, Train Acc: 74.35%, Val Acc: 69.90%
Epoch [12/30], Loss: 0.6342, Train Acc: 79.53%, Val Acc: 66.99%
Epoch [13/30], Loss: 0.6226, Train Acc: 80.82%, Val Acc: 55.34%
Epoch [14/30], Loss: 0.6073, Train Acc: 81.57%, Val Acc: 69.90%
Epoch [15/30], Loss: 0.5641, Train Acc: 84.81%, Val Acc: 77.67%
Epoch [16/30], Loss: 0.5719, Train Acc: 84.38%, Val Acc: 77.67%
Epoch [17/30], Loss: 0.5659, Train Acc: 84.48%, Val Acc: 64.08%
Epoch [18/30], Loss: 0.5448, Train Acc: 86.85%, Val Acc: 70.87%
Epoch [19/30], Loss: 0.5255, Train Acc: 87.72%, Val Acc: 79.61%
Epoch [20/30], Loss: 0.4976, Train Acc: 90.19%, Val Acc: 79.61%
Epoch [21/30], Loss: 0.5258, Train Acc: 87.18%, Val Acc: 78.64%
Epoch [22/30], Loss: 0.5110, Train Acc: 89.44%, Val Acc: 75.73%
Epoch [23/30], Loss: 0.4918, Train Acc: 90.41%, Val Acc: 80.58%
Epoch [24/30], Loss: 0.4656, Train Acc: 91.92%, Val Acc: 79.61%
Epoch [25/30], Loss: 0.4562, Train Acc: 91.81%, Val Acc: 67.96%
Epoch [26/30], Loss: 0.4681, Train Acc: 90.19%, Val Acc: 83.50%
Epoch [27/30], Loss: 0.4442, Train Acc: 93.21%, Val Acc: 86.41%
Epoch [28/30], Loss: 0.4237, Train Acc: 93.75%, Val Acc: 83.50%
Epoch [29/30], Loss: 0.4395, Train Acc: 93.32%, Val Acc: 84.47%
```

```
Epoch [30/30], Loss: 0.4294, Train Acc: 94.29%, Val Acc: 84.47%
Testing model...
```

```
==== split_20_80 Results ====
Accuracy: 0.7924
Precision: 0.8025
Recall: 0.7924
F1-Score: 0.7944
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.64	0.78	0.71	899
CN	0.86	0.73	0.79	1152
MCI	0.84	0.83	0.83	2072
accuracy			0.79	4123
macro avg	0.78	0.78	0.78	4123
weighted avg	0.80	0.79	0.79	4123

```
=====
Processing: split_30_70
=====
```

```
Train samples: 1391
Val samples: 155
Test samples: 3608
Starting training...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
```

```
Epoch [1/30], Loss: 1.0405, Train Acc: 50.04%, Val Acc: 52.90%
Epoch [2/30], Loss: 0.9956, Train Acc: 54.13%, Val Acc: 51.61%
Epoch [3/30], Loss: 0.9541, Train Acc: 57.15%, Val Acc: 50.32%
Epoch [4/30], Loss: 0.8840, Train Acc: 62.83%, Val Acc: 54.19%
Epoch [5/30], Loss: 0.8246, Train Acc: 66.28%, Val Acc: 58.71%
Epoch [6/30], Loss: 0.7949, Train Acc: 70.24%, Val Acc: 61.94%
Epoch [7/30], Loss: 0.7261, Train Acc: 75.84%, Val Acc: 62.58%
Epoch [8/30], Loss: 0.7022, Train Acc: 75.20%, Val Acc: 65.16%
Epoch [9/30], Loss: 0.6920, Train Acc: 75.70%, Val Acc: 74.19%
Epoch [10/30], Loss: 0.6347, Train Acc: 79.87%, Val Acc: 72.90%
Epoch [11/30], Loss: 0.6058, Train Acc: 83.32%, Val Acc: 67.74%
Epoch [12/30], Loss: 0.5822, Train Acc: 83.82%, Val Acc: 69.03%
Epoch [13/30], Loss: 0.5747, Train Acc: 84.90%, Val Acc: 76.77%
Epoch [14/30], Loss: 0.5610, Train Acc: 83.97%, Val Acc: 78.71%
Epoch [15/30], Loss: 0.5419, Train Acc: 87.71%, Val Acc: 78.71%
```

```
Epoch [16/30], Loss: 0.5313, Train Acc: 86.99%, Val Acc: 83.87%
Epoch [17/30], Loss: 0.5131, Train Acc: 88.71%, Val Acc: 80.65%
Epoch [18/30], Loss: 0.4967, Train Acc: 90.22%, Val Acc: 81.94%
Epoch [19/30], Loss: 0.4819, Train Acc: 90.73%, Val Acc: 83.23%
Epoch [20/30], Loss: 0.4839, Train Acc: 90.73%, Val Acc: 78.71%
Epoch [21/30], Loss: 0.4401, Train Acc: 93.89%, Val Acc: 83.87%
Early stopping at epoch 21
Testing model...
```

```
==== split_30_70 Results ===
```

```
Accuracy: 0.8428
Precision: 0.8439
Recall: 0.8428
F1-Score: 0.8400
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.81	0.70	0.75	787
CN	0.88	0.78	0.83	1008
MCI	0.83	0.94	0.88	1813
accuracy			0.84	3608
macro avg	0.84	0.81	0.82	3608
weighted avg	0.84	0.84	0.84	3608

```
=====
Processing: split_40_60
=====
```

```
Train samples: 1855
Val samples: 207
Test samples: 3092
Starting training...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
```

```
Epoch [1/30], Loss: 1.0254, Train Acc: 51.05%, Val Acc: 47.34%
Epoch [2/30], Loss: 0.9638, Train Acc: 57.68%, Val Acc: 49.76%
Epoch [3/30], Loss: 0.8876, Train Acc: 62.80%, Val Acc: 51.21%
Epoch [4/30], Loss: 0.8347, Train Acc: 67.12%, Val Acc: 61.35%
Epoch [5/30], Loss: 0.7704, Train Acc: 71.86%, Val Acc: 67.63%
Epoch [6/30], Loss: 0.7101, Train Acc: 74.72%, Val Acc: 72.95%
Epoch [7/30], Loss: 0.6867, Train Acc: 77.79%, Val Acc: 80.19%
Epoch [8/30], Loss: 0.6256, Train Acc: 81.99%, Val Acc: 74.88%
Epoch [9/30], Loss: 0.6156, Train Acc: 82.16%, Val Acc: 77.29%
```

```
Epoch [10/30], Loss: 0.5918, Train Acc: 83.18%, Val Acc: 75.85%
Epoch [11/30], Loss: 0.5921, Train Acc: 83.02%, Val Acc: 75.85%
Epoch [12/30], Loss: 0.5159, Train Acc: 87.76%, Val Acc: 83.09%
Epoch [13/30], Loss: 0.5231, Train Acc: 88.63%, Val Acc: 83.09%
Epoch [14/30], Loss: 0.5018, Train Acc: 88.73%, Val Acc: 82.61%
Epoch [15/30], Loss: 0.4832, Train Acc: 90.24%, Val Acc: 81.64%
Epoch [16/30], Loss: 0.4802, Train Acc: 90.03%, Val Acc: 83.09%
Epoch [17/30], Loss: 0.4637, Train Acc: 92.40%, Val Acc: 85.02%
Epoch [18/30], Loss: 0.4555, Train Acc: 92.78%, Val Acc: 86.47%
Epoch [19/30], Loss: 0.4557, Train Acc: 92.29%, Val Acc: 85.99%
Epoch [20/30], Loss: 0.4453, Train Acc: 92.72%, Val Acc: 85.51%
Epoch [21/30], Loss: 0.4485, Train Acc: 92.35%, Val Acc: 85.51%
Epoch [22/30], Loss: 0.4460, Train Acc: 93.10%, Val Acc: 85.99%
Epoch [23/30], Loss: 0.4323, Train Acc: 93.48%, Val Acc: 83.57%
Early stopping at epoch 23
Testing model...
```

```
==== split_40_60 Results ====
Accuracy: 0.8852
Precision: 0.8877
Recall: 0.8852
F1-Score: 0.8854
```

```
Classification Report:
precision    recall    f1-score    support
AD          0.80      0.86      0.83      674
CN          0.93      0.84      0.88      864
MCI         0.90      0.92      0.91     1554

accuracy                           0.89      3092
macro avg       0.88      0.87      0.87      3092
weighted avg    0.89      0.89      0.89      3092
```

```
=====
Processing: split_50_50
=====
Train samples: 2319
Val samples: 258
Test samples: 2577
Starting training...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/30], Loss: 1.0314, Train Acc: 50.37%, Val Acc: 46.12%
```

```
Epoch [2/30], Loss: 0.9639, Train Acc: 56.53%, Val Acc: 54.26%
Epoch [3/30], Loss: 0.9005, Train Acc: 61.23%, Val Acc: 52.33%
Epoch [4/30], Loss: 0.8187, Train Acc: 67.31%, Val Acc: 65.12%
Epoch [5/30], Loss: 0.7909, Train Acc: 69.64%, Val Acc: 71.71%
Epoch [6/30], Loss: 0.7151, Train Acc: 75.33%, Val Acc: 70.54%
Epoch [7/30], Loss: 0.6820, Train Acc: 77.19%, Val Acc: 70.93%
Epoch [8/30], Loss: 0.6321, Train Acc: 80.85%, Val Acc: 79.84%
Epoch [9/30], Loss: 0.6110, Train Acc: 82.92%, Val Acc: 79.07%
Epoch [10/30], Loss: 0.5798, Train Acc: 84.26%, Val Acc: 81.40%
Epoch [11/30], Loss: 0.5572, Train Acc: 85.68%, Val Acc: 76.36%
Epoch [12/30], Loss: 0.5387, Train Acc: 87.67%, Val Acc: 84.50%
Epoch [13/30], Loss: 0.5120, Train Acc: 88.27%, Val Acc: 90.70%
Epoch [14/30], Loss: 0.5145, Train Acc: 88.79%, Val Acc: 87.21%
Epoch [15/30], Loss: 0.4898, Train Acc: 90.21%, Val Acc: 86.82%
Epoch [16/30], Loss: 0.4773, Train Acc: 90.90%, Val Acc: 89.15%
Epoch [17/30], Loss: 0.4600, Train Acc: 91.89%, Val Acc: 87.21%
Epoch [18/30], Loss: 0.4137, Train Acc: 95.08%, Val Acc: 90.70%
Early stopping at epoch 18
Testing model...
```

```
==== split_50_50 Results ====
Accuracy: 0.9166
Precision: 0.9174
Recall: 0.9166
F1-Score: 0.9169
```

```
Classification Report:
precision    recall    f1-score    support
AD          0.86      0.89      0.87      562
CN          0.91      0.91      0.91      720
MCI         0.95      0.93      0.94     1295

accuracy                           0.92      2577
macro avg       0.91      0.91      0.91      2577
weighted avg    0.92      0.92      0.92      2577
```

```
=====
Processing: split_60_40
=====
Train samples: 2784
Val samples: 308
Test samples: 2062
Starting training...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
```

```

access the learning rate.
warnings.warn(
Epoch [1/30], Loss: 1.0283, Train Acc: 51.54%, Val Acc: 52.27%
Epoch [2/30], Loss: 0.9560, Train Acc: 56.97%, Val Acc: 50.97%
Epoch [3/30], Loss: 0.8718, Train Acc: 64.91%, Val Acc: 53.57%
Epoch [4/30], Loss: 0.7890, Train Acc: 69.47%, Val Acc: 62.66%
Epoch [5/30], Loss: 0.7321, Train Acc: 75.04%, Val Acc: 70.13%
Epoch [6/30], Loss: 0.6875, Train Acc: 77.33%, Val Acc: 66.56%
Epoch [7/30], Loss: 0.6512, Train Acc: 79.06%, Val Acc: 69.81%
Epoch [8/30], Loss: 0.6078, Train Acc: 83.08%, Val Acc: 76.95%
Epoch [9/30], Loss: 0.5888, Train Acc: 83.26%, Val Acc: 82.14%
Epoch [10/30], Loss: 0.5626, Train Acc: 84.95%, Val Acc: 80.52%
Epoch [11/30], Loss: 0.5253, Train Acc: 87.57%, Val Acc: 83.12%
Epoch [12/30], Loss: 0.5163, Train Acc: 88.07%, Val Acc: 86.69%
Epoch [13/30], Loss: 0.4899, Train Acc: 89.51%, Val Acc: 86.04%
Epoch [14/30], Loss: 0.4775, Train Acc: 90.62%, Val Acc: 89.29%
Epoch [15/30], Loss: 0.4583, Train Acc: 91.56%, Val Acc: 89.29%
Epoch [16/30], Loss: 0.4610, Train Acc: 92.06%, Val Acc: 85.71%
Epoch [17/30], Loss: 0.4245, Train Acc: 93.43%, Val Acc: 87.66%
Epoch [18/30], Loss: 0.4257, Train Acc: 94.00%, Val Acc: 88.31%
Epoch [19/30], Loss: 0.4078, Train Acc: 95.47%, Val Acc: 91.23%
Epoch [20/30], Loss: 0.3900, Train Acc: 95.65%, Val Acc: 90.58%
Epoch [21/30], Loss: 0.3851, Train Acc: 96.48%, Val Acc: 90.91%
Epoch [22/30], Loss: 0.3765, Train Acc: 96.62%, Val Acc: 91.88%
Epoch [23/30], Loss: 0.3724, Train Acc: 96.80%, Val Acc: 92.21%
Epoch [24/30], Loss: 0.3626, Train Acc: 97.52%, Val Acc: 93.51%
Epoch [25/30], Loss: 0.3707, Train Acc: 96.98%, Val Acc: 93.18%
Epoch [26/30], Loss: 0.3571, Train Acc: 97.49%, Val Acc: 93.18%
Epoch [27/30], Loss: 0.3680, Train Acc: 96.88%, Val Acc: 93.51%
Epoch [28/30], Loss: 0.3582, Train Acc: 97.56%, Val Acc: 93.51%
Epoch [29/30], Loss: 0.3488, Train Acc: 97.81%, Val Acc: 94.48%
Epoch [30/30], Loss: 0.3432, Train Acc: 98.38%, Val Acc: 95.13%
Testing model...

```

```

==== split_60_40 Results ===
Accuracy: 0.9685
Precision: 0.9685
Recall: 0.9685
F1-Score: 0.9685

```

Classification Report:

	precision	recall	f1-score	support
AD	0.96	0.96	0.96	450
CN	0.96	0.96	0.96	576
MCI	0.98	0.98	0.98	1036

accuracy			0.97	2062
macro avg	0.97	0.97	0.97	2062
weighted avg	0.97	0.97	0.97	2062

=====

Processing: split_70_30

=====

Train samples: 3247

Val samples: 361

Test samples: 1546

Starting training...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:

UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

 warnings.warn(

Epoch [1/30], Loss: 1.0272, Train Acc: 51.09%, Val Acc: 46.54%
 Epoch [2/30], Loss: 0.9280, Train Acc: 60.18%, Val Acc: 55.68%
 Epoch [3/30], Loss: 0.8179, Train Acc: 68.37%, Val Acc: 55.68%
 Epoch [4/30], Loss: 0.7353, Train Acc: 73.51%, Val Acc: 70.08%
 Epoch [5/30], Loss: 0.6736, Train Acc: 78.10%, Val Acc: 79.78%
 Epoch [6/30], Loss: 0.6368, Train Acc: 80.66%, Val Acc: 69.53%
 Epoch [7/30], Loss: 0.6031, Train Acc: 82.41%, Val Acc: 86.15%
 Epoch [8/30], Loss: 0.5682, Train Acc: 84.94%, Val Acc: 87.81%
 Epoch [9/30], Loss: 0.5265, Train Acc: 88.30%, Val Acc: 90.86%
 Epoch [10/30], Loss: 0.5201, Train Acc: 88.14%, Val Acc: 87.81%
 Epoch [11/30], Loss: 0.4897, Train Acc: 90.27%, Val Acc: 89.47%
 Epoch [12/30], Loss: 0.4798, Train Acc: 90.08%, Val Acc: 92.52%
 Epoch [13/30], Loss: 0.4575, Train Acc: 91.56%, Val Acc: 86.43%
 Epoch [14/30], Loss: 0.4449, Train Acc: 92.67%, Val Acc: 90.86%
 Epoch [15/30], Loss: 0.4288, Train Acc: 93.69%, Val Acc: 94.18%
 Epoch [16/30], Loss: 0.4170, Train Acc: 94.18%, Val Acc: 93.91%
 Epoch [17/30], Loss: 0.4178, Train Acc: 93.93%, Val Acc: 94.18%
 Epoch [18/30], Loss: 0.3965, Train Acc: 95.60%, Val Acc: 93.35%
 Epoch [19/30], Loss: 0.3900, Train Acc: 95.78%, Val Acc: 93.63%
 Epoch [20/30], Loss: 0.3663, Train Acc: 97.14%, Val Acc: 94.18%

Early stopping at epoch 20

Testing model...

==== split_70_30 Results ===

Accuracy: 0.9605

Precision: 0.9605

Recall: 0.9605

F1-Score: 0.9605

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

AD	0.95	0.94	0.95	337
CN	0.96	0.95	0.95	432
MCI	0.96	0.98	0.97	777
accuracy			0.96	1546
macro avg	0.96	0.95	0.96	1546
weighted avg	0.96	0.96	0.96	1546

```
=====
Processing: split_80_20
=====
Train samples: 3711
Val samples: 412
Test samples: 1031
Starting training...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/30], Loss: 1.0171, Train Acc: 51.31%, Val Acc: 49.51%
Epoch [2/30], Loss: 0.9045, Train Acc: 61.79%, Val Acc: 55.58%
Epoch [3/30], Loss: 0.8193, Train Acc: 67.64%, Val Acc: 67.96%
Epoch [4/30], Loss: 0.7369, Train Acc: 73.30%, Val Acc: 75.97%
Epoch [5/30], Loss: 0.6688, Train Acc: 77.90%, Val Acc: 82.28%
Epoch [6/30], Loss: 0.6242, Train Acc: 81.60%, Val Acc: 81.31%
Epoch [7/30], Loss: 0.5848, Train Acc: 83.35%, Val Acc: 81.55%
Epoch [8/30], Loss: 0.5497, Train Acc: 86.15%, Val Acc: 89.32%
Epoch [9/30], Loss: 0.5305, Train Acc: 86.82%, Val Acc: 84.47%
Epoch [10/30], Loss: 0.4962, Train Acc: 89.60%, Val Acc: 93.20%
Epoch [11/30], Loss: 0.4708, Train Acc: 91.22%, Val Acc: 95.15%
Epoch [12/30], Loss: 0.4611, Train Acc: 91.84%, Val Acc: 91.99%
Epoch [13/30], Loss: 0.4465, Train Acc: 92.45%, Val Acc: 94.17%
Epoch [14/30], Loss: 0.4262, Train Acc: 93.83%, Val Acc: 96.12%
Epoch [15/30], Loss: 0.4195, Train Acc: 93.88%, Val Acc: 96.36%
Epoch [16/30], Loss: 0.4064, Train Acc: 94.53%, Val Acc: 93.69%
Epoch [17/30], Loss: 0.3923, Train Acc: 95.39%, Val Acc: 95.39%
Epoch [18/30], Loss: 0.3887, Train Acc: 96.09%, Val Acc: 93.69%
Epoch [19/30], Loss: 0.3842, Train Acc: 95.88%, Val Acc: 96.60%
Epoch [20/30], Loss: 0.3712, Train Acc: 96.82%, Val Acc: 96.36%
Epoch [21/30], Loss: 0.3680, Train Acc: 96.71%, Val Acc: 96.84%
Epoch [22/30], Loss: 0.3608, Train Acc: 97.36%, Val Acc: 96.12%
Epoch [23/30], Loss: 0.3659, Train Acc: 96.85%, Val Acc: 97.57%
Epoch [24/30], Loss: 0.3537, Train Acc: 97.36%, Val Acc: 95.87%
Epoch [25/30], Loss: 0.3450, Train Acc: 97.93%, Val Acc: 96.84%
Epoch [26/30], Loss: 0.3464, Train Acc: 97.98%, Val Acc: 96.84%
```

```
Epoch [27/30], Loss: 0.3397, Train Acc: 98.09%, Val Acc: 96.36%
Epoch [28/30], Loss: 0.3310, Train Acc: 98.46%, Val Acc: 98.54%
Epoch [29/30], Loss: 0.3241, Train Acc: 98.81%, Val Acc: 98.79%
Epoch [30/30], Loss: 0.3236, Train Acc: 98.90%, Val Acc: 98.06%
Testing model...
```

```
==== split_80_20 Results ===
Accuracy: 0.9748
Precision: 0.9749
Recall: 0.9748
F1-Score: 0.9748
```

Classification Report:

	precision	recall	f1-score	support
AD	0.96	0.97	0.97	225
CN	0.97	0.99	0.98	288
MCI	0.98	0.97	0.98	518
accuracy			0.97	1031
macro avg	0.97	0.98	0.97	1031
weighted avg	0.97	0.97	0.97	1031

```
=====
Processing: split_90_10
=====
```

```
Train samples: 4175
Val samples: 464
Test samples: 515
Starting training...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/30], Loss: 1.0039, Train Acc: 53.41%, Val Acc: 52.16%
Epoch [2/30], Loss: 0.8793, Train Acc: 63.23%, Val Acc: 65.95%
Epoch [3/30], Loss: 0.7851, Train Acc: 69.37%, Val Acc: 75.86%
Epoch [4/30], Loss: 0.7055, Train Acc: 75.47%, Val Acc: 76.94%
Epoch [5/30], Loss: 0.6317, Train Acc: 80.81%, Val Acc: 83.62%
Epoch [6/30], Loss: 0.5990, Train Acc: 83.07%, Val Acc: 84.70%
Epoch [7/30], Loss: 0.5555, Train Acc: 86.18%, Val Acc: 84.70%
Epoch [8/30], Loss: 0.5252, Train Acc: 87.74%, Val Acc: 86.21%
Epoch [9/30], Loss: 0.4964, Train Acc: 89.70%, Val Acc: 90.09%
Epoch [10/30], Loss: 0.4775, Train Acc: 91.11%, Val Acc: 92.24%
Epoch [11/30], Loss: 0.4542, Train Acc: 92.26%, Val Acc: 93.53%
Epoch [12/30], Loss: 0.4328, Train Acc: 93.46%, Val Acc: 90.95%
```

Epoch [13/30], Loss: 0.4222, Train Acc: 93.94%, Val Acc: 93.32%
 Epoch [14/30], Loss: 0.4200, Train Acc: 93.70%, Val Acc: 93.53%
 Epoch [15/30], Loss: 0.3965, Train Acc: 95.23%, Val Acc: 96.34%
 Epoch [16/30], Loss: 0.3815, Train Acc: 96.10%, Val Acc: 93.75%
 Epoch [17/30], Loss: 0.3806, Train Acc: 96.14%, Val Acc: 96.12%
 Epoch [18/30], Loss: 0.3681, Train Acc: 96.81%, Val Acc: 95.26%
 Epoch [19/30], Loss: 0.3640, Train Acc: 97.05%, Val Acc: 95.47%
 Epoch [20/30], Loss: 0.3469, Train Acc: 97.89%, Val Acc: 96.98%
 Epoch [21/30], Loss: 0.3412, Train Acc: 98.01%, Val Acc: 96.55%
 Epoch [22/30], Loss: 0.3375, Train Acc: 98.20%, Val Acc: 96.55%
 Epoch [23/30], Loss: 0.3342, Train Acc: 98.44%, Val Acc: 96.77%
 Epoch [24/30], Loss: 0.3276, Train Acc: 99.02%, Val Acc: 95.69%
 Epoch [25/30], Loss: 0.3257, Train Acc: 98.92%, Val Acc: 97.41%
 Epoch [26/30], Loss: 0.3205, Train Acc: 99.33%, Val Acc: 96.34%
 Epoch [27/30], Loss: 0.3215, Train Acc: 99.04%, Val Acc: 97.41%
 Epoch [28/30], Loss: 0.3206, Train Acc: 99.07%, Val Acc: 97.41%
 Epoch [29/30], Loss: 0.3194, Train Acc: 99.02%, Val Acc: 97.63%
 Epoch [30/30], Loss: 0.3171, Train Acc: 99.28%, Val Acc: 97.84%
 Testing model...

==== split_90_10 Results ===

Accuracy: 0.9825

Precision: 0.9825

Recall: 0.9825

F1-Score: 0.9825

Classification Report:

	precision	recall	f1-score	support
AD	0.97	0.96	0.97	112
CN	0.98	0.99	0.98	144
MCI	0.99	0.99	0.99	259
accuracy			0.98	515
macro avg	0.98	0.98	0.98	515
weighted avg	0.98	0.98	0.98	515

=====

MobileNetV3-Large - SUMMARY OF ALL SPLITS

=====

split	accuracy	precision	recall	f1_score	training_time
split_10_90	0.511533	0.472666	0.511533	0.450184	15.468586
split_20_80	0.792384	0.802541	0.792384	0.794430	101.326068
split_30_70	0.842849	0.843855	0.842849	0.839961	104.212670
split_40_60	0.885188	0.887740	0.885188	0.885444	150.704233
split_50_50	0.916570	0.917441	0.916570	0.916873	144.666528
split_60_40	0.968477	0.968474	0.968477	0.968474	289.130308

```

split_70_30 0.960543 0.960498 0.960543 0.960462 223.756592
split_80_20 0.974782 0.974926 0.974782 0.974791 382.935752
split_90_10 0.982524 0.982512 0.982524 0.982511 429.871090

```

Detailed results saved to: /kaggle/working/mobilenetv3_large_results.csv

```
[6]: #ConvNeXt-Small
import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, models
import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score,
                           accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import ReduceLROnPlateau

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)

        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label
```

```

def get_data_transforms():
    train_transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.RandomHorizontalFlip(p=0.5),
        transforms.RandomRotation(10),
        transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
        transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
        transforms.RandomGrayscale(p=0.1),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    val_transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    return train_transform, val_transform

def train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs, device):
    best_val_acc = 0
    patience = 7
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()

            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
            optimizer.step()

            running_loss += loss.item()

```

```

    _, predicted = torch.max(outputs.data, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

    epoch_loss = running_loss / len(train_loader)
    epoch_acc = 100 * correct / total

    val_acc = evaluate_model(model, val_loader, device)

    if scheduler:
        scheduler.step(val_acc)

    current_lr = optimizer.param_groups[0]['lr']
    print(f'Epoch {epoch+1}/{num_epochs}, Loss: {epoch_loss:.4f}, Train\u2192Acc: {epoch_acc:.2f}%, Val Acc: {val_acc:.2f}%, LR: {current_lr:.6f}')

    if val_acc > best_val_acc:
        best_val_acc = val_acc
        patience_counter = 0
    else:
        patience_counter += 1

    if patience_counter >= patience:
        print("Early stopping at epoch", epoch+1)
        break

def evaluate_model(model, data_loader, device):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in data_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    return accuracy

def test_model(model, test_loader, device):
    model.eval()
    all_preds = []
    all_labels = []

```

```

    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    return all_preds, all_labels

def calculate_metrics(y_true, y_pred, split_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred, average='weighted', zero_division=0)
    recall = recall_score(y_true, y_pred, average='weighted', zero_division=0)
    f1 = f1_score(y_true, y_pred, average='weighted', zero_division=0)

    print(f"\n==== {split_name} Results ===")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-Score: {f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=['AD', 'CN', 'MCI'], zero_division=0))

    return {
        'split': split_name,
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1_score': f1
    }

def run_convnext_small_on_splits():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    splits_root = "/kaggle/working/alzheimer-resized-224_splits"
    split_folders = [f for f in os.listdir(splits_root) if f.startswith('split_')]
    split_folders.sort()

    results = []

    train_transform, val_transform = get_data_transforms()

```

```

for split_folder in split_folders:
    print(f"\n{'='*60}")
    print(f"Processing: {split_folder}")
    print(f"{'='*60}")

    split_path = os.path.join(splits_root, split_folder)

    train_dataset = AlzheimerDataset(split_path, 'train', train_transform)
    val_dataset = AlzheimerDataset(split_path, 'val', val_transform)
    test_dataset = AlzheimerDataset(split_path, 'test', val_transform)

    train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True, □
    ↪num_workers=2, pin_memory=True)
    val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False, □
    ↪num_workers=2, pin_memory=True)
    test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False, □
    ↪num_workers=2, pin_memory=True)

    print(f"Train samples: {len(train_dataset)}")
    print(f"Val samples: {len(val_dataset)}")
    print(f"Test samples: {len(test_dataset)}")

    model = models.convnext_small(weights=models.ConvNeXt_Small_Weights. \
    ↪IMAGENET1K_V1)
    num_ftrs = model.classifier[2].in_features
    model.classifier[2] = nn.Linear(num_ftrs, 3)

    for param in model.parameters():
        param.requires_grad = False

    for param in model.classifier.parameters():
        param.requires_grad = True

    for param in model.features[-2:].parameters():
        param.requires_grad = True

    model = model.to(device)

    criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
    optimizer = torch.optim.AdamW([
        {'params': model.classifier.parameters(), 'lr': 0.0001},
        {'params': model.features[-2:].parameters(), 'lr': 0.00005}
    ], weight_decay=1e-4)

    scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5, □
    ↪patience=3, verbose=True)

```

```

print("Starting training...")
start_time = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs=25, device=device)
training_time = time.time() - start_time

print("Fine-tuning all layers...")
for param in model.parameters():
    param.requires_grad = True

optimizer = torch.optim.AdamW(model.parameters(), lr=0.00001, weight_decay=1e-5)
scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5, patience=2, verbose=True)

start_time_finetune = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs=10, device=device)
training_time += time.time() - start_time_finetune

print("Testing model...")
test_preds, test_labels = test_model(model, test_loader, device)

split_results = calculate_metrics(test_labels, test_preds, split_folder)
split_results['training_time'] = training_time
results.append(split_results)

torch.cuda.empty_cache()

results_df = pd.DataFrame(results)
print(f"\n{'='*80}")
print("ConvNeXt-Small - SUMMARY OF ALL SPLITS")
print(f"\n{'='*80}")
print(results_df.to_string(index=False))

results_csv_path = "/kaggle/working/convnext_small_results.csv"
results_df.to_csv(results_csv_path, index=False)
print(f"\nDetailed results saved to: {results_csv_path}")

return results_df

if __name__ == "__main__":
    results = run_convnext_small_on_splits()

```

Using device: cuda

```
=====
Processing: split_10_90
=====
Train samples: 464
Val samples: 51
Test samples: 4639

Downloading: "https://download.pytorch.org/models/convnext_small-0c510722.pth"
to /root/.cache/torch/hub/checkpoints/convnext_small-0c510722.pth
100%|      | 192M/192M [00:00<00:00, 236MB/s]

Starting training...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/25], Loss: 1.0681, Train Acc: 48.71%, Val Acc: 49.02%, LR: 0.000100
Epoch [2/25], Loss: 1.0323, Train Acc: 51.72%, Val Acc: 49.02%, LR: 0.000100
Epoch [3/25], Loss: 1.0150, Train Acc: 50.86%, Val Acc: 52.94%, LR: 0.000100
Epoch [4/25], Loss: 1.0005, Train Acc: 55.82%, Val Acc: 50.98%, LR: 0.000100
Epoch [5/25], Loss: 0.9822, Train Acc: 56.68%, Val Acc: 47.06%, LR: 0.000100
Epoch [6/25], Loss: 0.9772, Train Acc: 56.47%, Val Acc: 52.94%, LR: 0.000100
Epoch [7/25], Loss: 0.9536, Train Acc: 57.33%, Val Acc: 60.78%, LR: 0.000100
Epoch [8/25], Loss: 0.9344, Train Acc: 57.97%, Val Acc: 56.86%, LR: 0.000100
Epoch [9/25], Loss: 0.9275, Train Acc: 60.13%, Val Acc: 60.78%, LR: 0.000100
Epoch [10/25], Loss: 0.8961, Train Acc: 60.78%, Val Acc: 60.78%, LR: 0.000100
Epoch [11/25], Loss: 0.8971, Train Acc: 62.72%, Val Acc: 64.71%, LR: 0.000100
Epoch [12/25], Loss: 0.8812, Train Acc: 62.28%, Val Acc: 64.71%, LR: 0.000100
Epoch [13/25], Loss: 0.8660, Train Acc: 63.79%, Val Acc: 60.78%, LR: 0.000100
Epoch [14/25], Loss: 0.8758, Train Acc: 67.46%, Val Acc: 66.67%, LR: 0.000100
Epoch [15/25], Loss: 0.8813, Train Acc: 63.36%, Val Acc: 64.71%, LR: 0.000100
Epoch [16/25], Loss: 0.8532, Train Acc: 65.09%, Val Acc: 60.78%, LR: 0.000100
Epoch [17/25], Loss: 0.8254, Train Acc: 67.24%, Val Acc: 60.78%, LR: 0.000100
Epoch [18/25], Loss: 0.8297, Train Acc: 67.24%, Val Acc: 49.02%, LR: 0.000050
Epoch [19/25], Loss: 0.8382, Train Acc: 67.03%, Val Acc: 66.67%, LR: 0.000050
Epoch [20/25], Loss: 0.7728, Train Acc: 71.98%, Val Acc: 66.67%, LR: 0.000050
Epoch [21/25], Loss: 0.7844, Train Acc: 71.77%, Val Acc: 68.63%, LR: 0.000050
Epoch [22/25], Loss: 0.7747, Train Acc: 71.34%, Val Acc: 66.67%, LR: 0.000050
Epoch [23/25], Loss: 0.7463, Train Acc: 75.65%, Val Acc: 64.71%, LR: 0.000050
Epoch [24/25], Loss: 0.7365, Train Acc: 73.92%, Val Acc: 66.67%, LR: 0.000050
Epoch [25/25], Loss: 0.7622, Train Acc: 71.98%, Val Acc: 68.63%, LR: 0.000025
Fine-tuning all layers...
Epoch [1/10], Loss: 0.7570, Train Acc: 73.49%, Val Acc: 62.75%, LR: 0.000010
Epoch [2/10], Loss: 0.7495, Train Acc: 73.49%, Val Acc: 64.71%, LR: 0.000010
Epoch [3/10], Loss: 0.7223, Train Acc: 75.00%, Val Acc: 66.67%, LR: 0.000010
Epoch [4/10], Loss: 0.6997, Train Acc: 78.23%, Val Acc: 68.63%, LR: 0.000010
Epoch [5/10], Loss: 0.6812, Train Acc: 78.66%, Val Acc: 68.63%, LR: 0.000010
Epoch [6/10], Loss: 0.6611, Train Acc: 79.53%, Val Acc: 64.71%, LR: 0.000010
```

```
Epoch [7/10], Loss: 0.6459, Train Acc: 80.39%, Val Acc: 72.55%, LR: 0.000010
Epoch [8/10], Loss: 0.6435, Train Acc: 81.68%, Val Acc: 66.67%, LR: 0.000010
Epoch [9/10], Loss: 0.6452, Train Acc: 78.88%, Val Acc: 66.67%, LR: 0.000010
Epoch [10/10], Loss: 0.6031, Train Acc: 83.84%, Val Acc: 66.67%, LR: 0.000005
Testing model...
```

```
==== split_10_90 Results ===
Accuracy: 0.6368
Precision: 0.6345
Recall: 0.6368
F1-Score: 0.6353
```

Classification Report:

	precision	recall	f1-score	support
AD	0.56	0.50	0.53	1012
CN	0.58	0.60	0.59	1296
MCI	0.70	0.72	0.71	2331
accuracy			0.64	4639
macro avg	0.61	0.61	0.61	4639
weighted avg	0.63	0.64	0.64	4639

```
=====
Processing: split_20_80
=====
```

```
Train samples: 928
Val samples: 103
Test samples: 4123
Starting training...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
Epoch [1/25], Loss: 1.0770, Train Acc: 46.88%, Val Acc: 50.49%, LR: 0.000100
Epoch [2/25], Loss: 1.0466, Train Acc: 48.60%, Val Acc: 53.40%, LR: 0.000100
Epoch [3/25], Loss: 1.0394, Train Acc: 51.72%, Val Acc: 56.31%, LR: 0.000100
Epoch [4/25], Loss: 1.0163, Train Acc: 51.51%, Val Acc: 55.34%, LR: 0.000100
Epoch [5/25], Loss: 0.9938, Train Acc: 55.06%, Val Acc: 52.43%, LR: 0.000100
Epoch [6/25], Loss: 0.9932, Train Acc: 55.71%, Val Acc: 57.28%, LR: 0.000100
Epoch [7/25], Loss: 0.9716, Train Acc: 54.31%, Val Acc: 58.25%, LR: 0.000100
Epoch [8/25], Loss: 0.9499, Train Acc: 57.65%, Val Acc: 55.34%, LR: 0.000100
Epoch [9/25], Loss: 0.9364, Train Acc: 58.19%, Val Acc: 53.40%, LR: 0.000100
Epoch [10/25], Loss: 0.9243, Train Acc: 59.27%, Val Acc: 53.40%, LR: 0.000100
Epoch [11/25], Loss: 0.9085, Train Acc: 59.70%, Val Acc: 58.25%, LR: 0.000050
Epoch [12/25], Loss: 0.9009, Train Acc: 61.31%, Val Acc: 58.25%, LR: 0.000050
```

```
Epoch [13/25], Loss: 0.8596, Train Acc: 63.90%, Val Acc: 58.25%, LR: 0.000050
Epoch [14/25], Loss: 0.8654, Train Acc: 64.55%, Val Acc: 60.19%, LR: 0.000050
Epoch [15/25], Loss: 0.8488, Train Acc: 65.30%, Val Acc: 61.17%, LR: 0.000050
Epoch [16/25], Loss: 0.8434, Train Acc: 65.73%, Val Acc: 57.28%, LR: 0.000050
Epoch [17/25], Loss: 0.8389, Train Acc: 68.10%, Val Acc: 58.25%, LR: 0.000050
Epoch [18/25], Loss: 0.8310, Train Acc: 66.49%, Val Acc: 60.19%, LR: 0.000050
Epoch [19/25], Loss: 0.8286, Train Acc: 67.24%, Val Acc: 60.19%, LR: 0.000025
Epoch [20/25], Loss: 0.7880, Train Acc: 70.91%, Val Acc: 60.19%, LR: 0.000025
Epoch [21/25], Loss: 0.8068, Train Acc: 67.89%, Val Acc: 58.25%, LR: 0.000025
Epoch [22/25], Loss: 0.8067, Train Acc: 69.07%, Val Acc: 58.25%, LR: 0.000025
Early stopping at epoch 22
Fine-tuning all layers...
Epoch [1/10], Loss: 0.8197, Train Acc: 67.46%, Val Acc: 64.08%, LR: 0.000010
Epoch [2/10], Loss: 0.7970, Train Acc: 68.53%, Val Acc: 65.05%, LR: 0.000010
Epoch [3/10], Loss: 0.7780, Train Acc: 70.80%, Val Acc: 70.87%, LR: 0.000010
Epoch [4/10], Loss: 0.7244, Train Acc: 76.19%, Val Acc: 68.93%, LR: 0.000010
Epoch [5/10], Loss: 0.7381, Train Acc: 73.38%, Val Acc: 72.82%, LR: 0.000010
Epoch [6/10], Loss: 0.7076, Train Acc: 76.40%, Val Acc: 71.84%, LR: 0.000010
Epoch [7/10], Loss: 0.7036, Train Acc: 76.94%, Val Acc: 71.84%, LR: 0.000010
Epoch [8/10], Loss: 0.6710, Train Acc: 79.74%, Val Acc: 69.90%, LR: 0.000005
Epoch [9/10], Loss: 0.6562, Train Acc: 79.53%, Val Acc: 73.79%, LR: 0.000005
Epoch [10/10], Loss: 0.6352, Train Acc: 82.65%, Val Acc: 75.73%, LR: 0.000005
Testing model...
```

==== split_20_80 Results ===

Accuracy: 0.7735

Precision: 0.7702

Recall: 0.7735

F1-Score: 0.7678

Classification Report:

	precision	recall	f1-score	support
AD	0.73	0.56	0.63	899
CN	0.77	0.74	0.75	1152
MCI	0.79	0.89	0.83	2072
accuracy			0.77	4123
macro avg	0.76	0.73	0.74	4123
weighted avg	0.77	0.77	0.77	4123

=====

Processing: split_30_70

=====

Train samples: 1391

Val samples: 155

Test samples: 3608

Starting training...

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:  
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to  
access the learning rate.  
    warnings.warn(  
  
Epoch [1/25], Loss: 1.0607, Train Acc: 47.81%, Val Acc: 50.32%, LR: 0.000100  
Epoch [2/25], Loss: 1.0246, Train Acc: 51.69%, Val Acc: 52.26%, LR: 0.000100  
Epoch [3/25], Loss: 1.0146, Train Acc: 51.55%, Val Acc: 54.19%, LR: 0.000100  
Epoch [4/25], Loss: 0.9884, Train Acc: 54.42%, Val Acc: 59.35%, LR: 0.000100  
Epoch [5/25], Loss: 0.9673, Train Acc: 56.51%, Val Acc: 54.84%, LR: 0.000100  
Epoch [6/25], Loss: 0.9459, Train Acc: 57.58%, Val Acc: 56.13%, LR: 0.000100  
Epoch [7/25], Loss: 0.9335, Train Acc: 59.74%, Val Acc: 66.45%, LR: 0.000100  
Epoch [8/25], Loss: 0.9010, Train Acc: 62.62%, Val Acc: 67.10%, LR: 0.000100  
Epoch [9/25], Loss: 0.8810, Train Acc: 64.13%, Val Acc: 66.45%, LR: 0.000100  
Epoch [10/25], Loss: 0.8694, Train Acc: 64.41%, Val Acc: 63.23%, LR: 0.000100  
Epoch [11/25], Loss: 0.8496, Train Acc: 64.92%, Val Acc: 64.52%, LR: 0.000100  
Epoch [12/25], Loss: 0.8387, Train Acc: 67.51%, Val Acc: 69.68%, LR: 0.000100  
Epoch [13/25], Loss: 0.8230, Train Acc: 67.22%, Val Acc: 69.03%, LR: 0.000100  
Epoch [14/25], Loss: 0.8019, Train Acc: 69.16%, Val Acc: 76.77%, LR: 0.000100  
Epoch [15/25], Loss: 0.8014, Train Acc: 68.66%, Val Acc: 74.19%, LR: 0.000100  
Epoch [16/25], Loss: 0.7953, Train Acc: 67.72%, Val Acc: 73.55%, LR: 0.000100  
Epoch [17/25], Loss: 0.7672, Train Acc: 71.39%, Val Acc: 76.13%, LR: 0.000100  
Epoch [18/25], Loss: 0.7480, Train Acc: 73.18%, Val Acc: 72.26%, LR: 0.000050  
Epoch [19/25], Loss: 0.7298, Train Acc: 74.48%, Val Acc: 73.55%, LR: 0.000050  
Epoch [20/25], Loss: 0.7019, Train Acc: 75.92%, Val Acc: 76.77%, LR: 0.000050  
Epoch [21/25], Loss: 0.7222, Train Acc: 74.69%, Val Acc: 77.42%, LR: 0.000050  
Epoch [22/25], Loss: 0.7059, Train Acc: 74.91%, Val Acc: 77.42%, LR: 0.000050  
Epoch [23/25], Loss: 0.6917, Train Acc: 77.64%, Val Acc: 80.65%, LR: 0.000050  
Epoch [24/25], Loss: 0.7001, Train Acc: 76.78%, Val Acc: 79.35%, LR: 0.000050  
Epoch [25/25], Loss: 0.6771, Train Acc: 78.00%, Val Acc: 77.42%, LR: 0.000050  
Fine-tuning all layers...  
Epoch [1/10], Loss: 0.7022, Train Acc: 74.84%, Val Acc: 81.29%, LR: 0.000010  
Epoch [2/10], Loss: 0.6616, Train Acc: 78.79%, Val Acc: 76.13%, LR: 0.000010  
Epoch [3/10], Loss: 0.6489, Train Acc: 79.51%, Val Acc: 78.71%, LR: 0.000010  
Epoch [4/10], Loss: 0.6318, Train Acc: 79.87%, Val Acc: 80.65%, LR: 0.000005  
Epoch [5/10], Loss: 0.5935, Train Acc: 83.54%, Val Acc: 83.87%, LR: 0.000005  
Epoch [6/10], Loss: 0.5697, Train Acc: 85.26%, Val Acc: 81.29%, LR: 0.000005  
Epoch [7/10], Loss: 0.5803, Train Acc: 83.61%, Val Acc: 83.87%, LR: 0.000005  
Epoch [8/10], Loss: 0.5600, Train Acc: 85.84%, Val Acc: 83.87%, LR: 0.000003  
Epoch [9/10], Loss: 0.5484, Train Acc: 87.63%, Val Acc: 82.58%, LR: 0.000003  
Epoch [10/10], Loss: 0.5473, Train Acc: 86.77%, Val Acc: 80.65%, LR: 0.000003  
Testing model...
```

```
==== split_30_70 Results ====  
Accuracy: 0.8196  
Precision: 0.8254  
Recall: 0.8196
```

F1-Score: 0.8156

Classification Report:

	precision	recall	f1-score	support
AD	0.83	0.66	0.73	787
CN	0.89	0.74	0.80	1008
MCI	0.79	0.94	0.86	1813
accuracy			0.82	3608
macro avg	0.83	0.78	0.80	3608
weighted avg	0.83	0.82	0.82	3608

=====

Processing: split_40_60

=====

Train samples: 1855
Val samples: 207
Test samples: 3092
Starting training...

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:  
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to  
access the learning rate.  
    warnings.warn(  
  
Epoch [1/25], Loss: 1.0505, Train Acc: 50.35%, Val Acc: 51.21%, LR: 0.000100  
Epoch [2/25], Loss: 1.0175, Train Acc: 52.45%, Val Acc: 55.07%, LR: 0.000100  
Epoch [3/25], Loss: 0.9818, Train Acc: 55.09%, Val Acc: 56.04%, LR: 0.000100  
Epoch [4/25], Loss: 0.9570, Train Acc: 56.55%, Val Acc: 57.97%, LR: 0.000100  
Epoch [5/25], Loss: 0.9302, Train Acc: 60.32%, Val Acc: 63.29%, LR: 0.000100  
Epoch [6/25], Loss: 0.8922, Train Acc: 62.96%, Val Acc: 59.90%, LR: 0.000100  
Epoch [7/25], Loss: 0.8802, Train Acc: 64.37%, Val Acc: 62.80%, LR: 0.000100  
Epoch [8/25], Loss: 0.8641, Train Acc: 63.50%, Val Acc: 64.25%, LR: 0.000100  
Epoch [9/25], Loss: 0.8422, Train Acc: 67.01%, Val Acc: 64.73%, LR: 0.000100  
Epoch [10/25], Loss: 0.8115, Train Acc: 68.03%, Val Acc: 66.18%, LR: 0.000100  
Epoch [11/25], Loss: 0.8005, Train Acc: 68.84%, Val Acc: 74.40%, LR: 0.000100  
Epoch [12/25], Loss: 0.7898, Train Acc: 70.51%, Val Acc: 72.46%, LR: 0.000100  
Epoch [13/25], Loss: 0.7693, Train Acc: 70.57%, Val Acc: 72.95%, LR: 0.000100  
Epoch [14/25], Loss: 0.7662, Train Acc: 72.18%, Val Acc: 75.36%, LR: 0.000100  
Epoch [15/25], Loss: 0.7373, Train Acc: 73.37%, Val Acc: 76.81%, LR: 0.000100  
Epoch [16/25], Loss: 0.7318, Train Acc: 73.53%, Val Acc: 73.43%, LR: 0.000100  
Epoch [17/25], Loss: 0.7064, Train Acc: 75.74%, Val Acc: 73.91%, LR: 0.000100  
Epoch [18/25], Loss: 0.7050, Train Acc: 75.63%, Val Acc: 75.36%, LR: 0.000100  
Epoch [19/25], Loss: 0.7208, Train Acc: 74.61%, Val Acc: 73.43%, LR: 0.000050  
Epoch [20/25], Loss: 0.6739, Train Acc: 78.44%, Val Acc: 78.26%, LR: 0.000050  
Epoch [21/25], Loss: 0.6735, Train Acc: 78.11%, Val Acc: 78.74%, LR: 0.000050  
Epoch [22/25], Loss: 0.6679, Train Acc: 78.27%, Val Acc: 77.29%, LR: 0.000050
```

```
Epoch [23/25], Loss: 0.6469, Train Acc: 79.62%, Val Acc: 79.71%, LR: 0.000050
Epoch [24/25], Loss: 0.6552, Train Acc: 79.08%, Val Acc: 78.26%, LR: 0.000050
Epoch [25/25], Loss: 0.6372, Train Acc: 80.75%, Val Acc: 80.19%, LR: 0.000050
Fine-tuning all layers...
Epoch [1/10], Loss: 0.6354, Train Acc: 80.65%, Val Acc: 79.71%, LR: 0.000010
Epoch [2/10], Loss: 0.6161, Train Acc: 82.59%, Val Acc: 80.68%, LR: 0.000010
Epoch [3/10], Loss: 0.5946, Train Acc: 83.18%, Val Acc: 80.68%, LR: 0.000010
Epoch [4/10], Loss: 0.5739, Train Acc: 85.66%, Val Acc: 84.06%, LR: 0.000010
Epoch [5/10], Loss: 0.5558, Train Acc: 86.47%, Val Acc: 81.64%, LR: 0.000010
Epoch [6/10], Loss: 0.5600, Train Acc: 86.15%, Val Acc: 85.99%, LR: 0.000010
Epoch [7/10], Loss: 0.5234, Train Acc: 88.52%, Val Acc: 84.06%, LR: 0.000010
Epoch [8/10], Loss: 0.5259, Train Acc: 87.92%, Val Acc: 84.54%, LR: 0.000010
Epoch [9/10], Loss: 0.5114, Train Acc: 88.89%, Val Acc: 87.44%, LR: 0.000010
Epoch [10/10], Loss: 0.4995, Train Acc: 90.08%, Val Acc: 86.47%, LR: 0.000010
Testing model...
```

```
==== split_40_60 Results ===
```

```
Accuracy: 0.8765
Precision: 0.8775
Recall: 0.8765
F1-Score: 0.8750
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.90	0.77	0.83	674
CN	0.88	0.84	0.86	864
MCI	0.87	0.94	0.90	1554
accuracy			0.88	3092
macro avg	0.88	0.85	0.86	3092
weighted avg	0.88	0.88	0.87	3092

```
=====
Processing: split_50_50
=====
```

```
Train samples: 2319
Val samples: 258
Test samples: 2577
Starting training...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
```

```
Epoch [1/25], Loss: 1.0428, Train Acc: 49.42%, Val Acc: 53.88%, LR: 0.000100
Epoch [2/25], Loss: 1.0036, Train Acc: 53.34%, Val Acc: 60.08%, LR: 0.000100
```

Epoch [3/25], Loss: 0.9718, Train Acc: 55.63%, Val Acc: 55.04%, LR: 0.000100
 Epoch [4/25], Loss: 0.9515, Train Acc: 56.75%, Val Acc: 63.95%, LR: 0.000100
 Epoch [5/25], Loss: 0.9153, Train Acc: 59.21%, Val Acc: 66.67%, LR: 0.000100
 Epoch [6/25], Loss: 0.8877, Train Acc: 62.35%, Val Acc: 62.02%, LR: 0.000100
 Epoch [7/25], Loss: 0.8653, Train Acc: 64.47%, Val Acc: 68.22%, LR: 0.000100
 Epoch [8/25], Loss: 0.8541, Train Acc: 66.02%, Val Acc: 74.03%, LR: 0.000100
 Epoch [9/25], Loss: 0.8346, Train Acc: 67.01%, Val Acc: 74.42%, LR: 0.000100
 Epoch [10/25], Loss: 0.8078, Train Acc: 67.79%, Val Acc: 74.81%, LR: 0.000100
 Epoch [11/25], Loss: 0.7812, Train Acc: 70.85%, Val Acc: 76.36%, LR: 0.000100
 Epoch [12/25], Loss: 0.7759, Train Acc: 70.94%, Val Acc: 74.03%, LR: 0.000100
 Epoch [13/25], Loss: 0.7548, Train Acc: 72.62%, Val Acc: 75.97%, LR: 0.000100
 Epoch [14/25], Loss: 0.7495, Train Acc: 72.53%, Val Acc: 79.07%, LR: 0.000100
 Epoch [15/25], Loss: 0.7411, Train Acc: 74.00%, Val Acc: 76.74%, LR: 0.000100
 Epoch [16/25], Loss: 0.7266, Train Acc: 73.74%, Val Acc: 83.33%, LR: 0.000100
 Epoch [17/25], Loss: 0.7325, Train Acc: 73.82%, Val Acc: 82.95%, LR: 0.000100
 Epoch [18/25], Loss: 0.7079, Train Acc: 76.50%, Val Acc: 79.84%, LR: 0.000100
 Epoch [19/25], Loss: 0.6982, Train Acc: 77.62%, Val Acc: 84.11%, LR: 0.000100
 Epoch [20/25], Loss: 0.6697, Train Acc: 79.04%, Val Acc: 81.78%, LR: 0.000100
 Epoch [21/25], Loss: 0.6700, Train Acc: 78.22%, Val Acc: 82.56%, LR: 0.000100
 Epoch [22/25], Loss: 0.6530, Train Acc: 79.60%, Val Acc: 84.11%, LR: 0.000100
 Epoch [23/25], Loss: 0.6657, Train Acc: 78.53%, Val Acc: 84.88%, LR: 0.000100
 Epoch [24/25], Loss: 0.6519, Train Acc: 80.38%, Val Acc: 84.88%, LR: 0.000100
 Epoch [25/25], Loss: 0.6325, Train Acc: 80.77%, Val Acc: 84.88%, LR: 0.000100
 Fine-tuning all layers...
 Epoch [1/10], Loss: 0.6238, Train Acc: 81.85%, Val Acc: 87.21%, LR: 0.000010
 Epoch [2/10], Loss: 0.5907, Train Acc: 83.31%, Val Acc: 88.76%, LR: 0.000010
 Epoch [3/10], Loss: 0.5659, Train Acc: 85.99%, Val Acc: 87.21%, LR: 0.000010
 Epoch [4/10], Loss: 0.5519, Train Acc: 86.42%, Val Acc: 88.37%, LR: 0.000010
 Epoch [5/10], Loss: 0.5234, Train Acc: 88.53%, Val Acc: 88.76%, LR: 0.000005
 Epoch [6/10], Loss: 0.4994, Train Acc: 89.56%, Val Acc: 90.31%, LR: 0.000005
 Epoch [7/10], Loss: 0.5071, Train Acc: 89.95%, Val Acc: 91.47%, LR: 0.000005
 Epoch [8/10], Loss: 0.4901, Train Acc: 89.87%, Val Acc: 91.47%, LR: 0.000005
 Epoch [9/10], Loss: 0.4800, Train Acc: 91.12%, Val Acc: 91.09%, LR: 0.000005
 Epoch [10/10], Loss: 0.4814, Train Acc: 91.03%, Val Acc: 91.09%, LR: 0.000003
 Testing model...

==== split_50_50 Results ===

Accuracy: 0.9065
 Precision: 0.9065
 Recall: 0.9065
 F1-Score: 0.9060

Classification Report:

	precision	recall	f1-score	support
AD	0.91	0.83	0.87	562
CN	0.89	0.90	0.90	720
MCI	0.91	0.94	0.93	1295

accuracy			0.91	2577
macro avg	0.91	0.89	0.90	2577
weighted avg	0.91	0.91	0.91	2577

=====

Processing: split_60_40

=====

Train samples: 2784

Val samples: 308

Test samples: 2062

Starting training...

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/25], Loss: 1.0498, Train Acc: 48.64%, Val Acc: 50.32%, LR: 0.000100
Epoch [2/25], Loss: 0.9921, Train Acc: 54.13%, Val Acc: 54.55%, LR: 0.000100
Epoch [3/25], Loss: 0.9572, Train Acc: 57.00%, Val Acc: 62.99%, LR: 0.000100
Epoch [4/25], Loss: 0.9264, Train Acc: 60.20%, Val Acc: 61.36%, LR: 0.000100
Epoch [5/25], Loss: 0.9095, Train Acc: 60.92%, Val Acc: 66.88%, LR: 0.000100
Epoch [6/25], Loss: 0.8699, Train Acc: 63.47%, Val Acc: 69.48%, LR: 0.000100
Epoch [7/25], Loss: 0.8511, Train Acc: 66.06%, Val Acc: 73.05%, LR: 0.000100
Epoch [8/25], Loss: 0.8261, Train Acc: 67.42%, Val Acc: 73.05%, LR: 0.000100
Epoch [9/25], Loss: 0.8098, Train Acc: 68.71%, Val Acc: 74.35%, LR: 0.000100
Epoch [10/25], Loss: 0.7839, Train Acc: 69.11%, Val Acc: 76.95%, LR: 0.000100
Epoch [11/25], Loss: 0.7573, Train Acc: 71.66%, Val Acc: 77.27%, LR: 0.000100
Epoch [12/25], Loss: 0.7377, Train Acc: 72.67%, Val Acc: 78.57%, LR: 0.000100
Epoch [13/25], Loss: 0.7467, Train Acc: 72.77%, Val Acc: 77.60%, LR: 0.000100
Epoch [14/25], Loss: 0.7132, Train Acc: 75.32%, Val Acc: 80.52%, LR: 0.000100
Epoch [15/25], Loss: 0.7070, Train Acc: 74.96%, Val Acc: 81.82%, LR: 0.000100
Epoch [16/25], Loss: 0.6920, Train Acc: 76.83%, Val Acc: 83.44%, LR: 0.000100
Epoch [17/25], Loss: 0.6819, Train Acc: 77.05%, Val Acc: 82.14%, LR: 0.000100
Epoch [18/25], Loss: 0.6757, Train Acc: 77.77%, Val Acc: 82.14%, LR: 0.000100
Epoch [19/25], Loss: 0.6577, Train Acc: 78.88%, Val Acc: 84.09%, LR: 0.000100
Epoch [20/25], Loss: 0.6474, Train Acc: 79.42%, Val Acc: 85.06%, LR: 0.000100
Epoch [21/25], Loss: 0.6443, Train Acc: 80.14%, Val Acc: 85.71%, LR: 0.000100
Epoch [22/25], Loss: 0.6295, Train Acc: 81.29%, Val Acc: 85.06%, LR: 0.000100
Epoch [23/25], Loss: 0.6290, Train Acc: 80.60%, Val Acc: 85.39%, LR: 0.000100
Epoch [24/25], Loss: 0.6095, Train Acc: 81.72%, Val Acc: 85.71%, LR: 0.000100
Epoch [25/25], Loss: 0.6192, Train Acc: 81.54%, Val Acc: 85.39%, LR: 0.000050
Fine-tuning all layers...
Epoch [1/10], Loss: 0.5952, Train Acc: 82.79%, Val Acc: 86.36%, LR: 0.000010
Epoch [2/10], Loss: 0.5624, Train Acc: 85.74%, Val Acc: 88.64%, LR: 0.000010
Epoch [3/10], Loss: 0.5432, Train Acc: 85.96%, Val Acc: 89.29%, LR: 0.000010
Epoch [4/10], Loss: 0.5354, Train Acc: 87.00%, Val Acc: 89.29%, LR: 0.000010
```

```
Epoch [5/10], Loss: 0.5081, Train Acc: 89.55%, Val Acc: 93.18%, LR: 0.000010
Epoch [6/10], Loss: 0.4906, Train Acc: 90.59%, Val Acc: 88.64%, LR: 0.000010
Epoch [7/10], Loss: 0.4918, Train Acc: 90.19%, Val Acc: 93.18%, LR: 0.000010
Epoch [8/10], Loss: 0.4859, Train Acc: 90.84%, Val Acc: 92.53%, LR: 0.000005
Epoch [9/10], Loss: 0.4595, Train Acc: 92.39%, Val Acc: 92.86%, LR: 0.000005
Epoch [10/10], Loss: 0.4545, Train Acc: 92.42%, Val Acc: 93.51%, LR: 0.000005
Testing model...
```

```
==== split_60_40 Results ===
```

```
Accuracy: 0.9365
```

```
Precision: 0.9365
```

```
Recall: 0.9365
```

```
F1-Score: 0.9362
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.93	0.89	0.91	450
CN	0.94	0.93	0.93	576
MCI	0.94	0.96	0.95	1036
accuracy			0.94	2062
macro avg	0.94	0.93	0.93	2062
weighted avg	0.94	0.94	0.94	2062

```
=====
Processing: split_70_30
=====
```

```
Train samples: 3247
```

```
Val samples: 361
```

```
Test samples: 1546
```

```
Starting training...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/25], Loss: 1.0570, Train Acc: 48.35%, Val Acc: 56.79%, LR: 0.000100
Epoch [2/25], Loss: 1.0049, Train Acc: 52.63%, Val Acc: 60.39%, LR: 0.000100
Epoch [3/25], Loss: 0.9595, Train Acc: 57.35%, Val Acc: 64.82%, LR: 0.000100
Epoch [4/25], Loss: 0.9230, Train Acc: 61.01%, Val Acc: 60.66%, LR: 0.000100
Epoch [5/25], Loss: 0.8942, Train Acc: 61.47%, Val Acc: 69.81%, LR: 0.000100
Epoch [6/25], Loss: 0.8650, Train Acc: 64.95%, Val Acc: 68.98%, LR: 0.000100
Epoch [7/25], Loss: 0.8447, Train Acc: 66.55%, Val Acc: 70.91%, LR: 0.000100
Epoch [8/25], Loss: 0.8156, Train Acc: 68.31%, Val Acc: 74.79%, LR: 0.000100
Epoch [9/25], Loss: 0.7937, Train Acc: 69.94%, Val Acc: 76.18%, LR: 0.000100
Epoch [10/25], Loss: 0.7766, Train Acc: 71.02%, Val Acc: 75.90%, LR: 0.000100
```

Epoch [11/25], Loss: 0.7537, Train Acc: 72.07%, Val Acc: 77.01%, LR: 0.000100
 Epoch [12/25], Loss: 0.7343, Train Acc: 73.54%, Val Acc: 76.45%, LR: 0.000100
 Epoch [13/25], Loss: 0.7277, Train Acc: 74.04%, Val Acc: 77.01%, LR: 0.000100
 Epoch [14/25], Loss: 0.7078, Train Acc: 76.41%, Val Acc: 82.55%, LR: 0.000100
 Epoch [15/25], Loss: 0.6866, Train Acc: 76.87%, Val Acc: 82.55%, LR: 0.000100
 Epoch [16/25], Loss: 0.6785, Train Acc: 77.92%, Val Acc: 82.55%, LR: 0.000100
 Epoch [17/25], Loss: 0.6727, Train Acc: 78.35%, Val Acc: 83.10%, LR: 0.000100
 Epoch [18/25], Loss: 0.6500, Train Acc: 79.33%, Val Acc: 83.10%, LR: 0.000100
 Epoch [19/25], Loss: 0.6372, Train Acc: 80.14%, Val Acc: 83.66%, LR: 0.000100
 Epoch [20/25], Loss: 0.6436, Train Acc: 79.64%, Val Acc: 84.76%, LR: 0.000100
 Epoch [21/25], Loss: 0.6291, Train Acc: 80.54%, Val Acc: 85.32%, LR: 0.000100
 Epoch [22/25], Loss: 0.6251, Train Acc: 81.71%, Val Acc: 86.15%, LR: 0.000100
 Epoch [23/25], Loss: 0.6071, Train Acc: 82.57%, Val Acc: 87.81%, LR: 0.000100
 Epoch [24/25], Loss: 0.6087, Train Acc: 82.17%, Val Acc: 85.87%, LR: 0.000100
 Epoch [25/25], Loss: 0.5872, Train Acc: 83.71%, Val Acc: 89.75%, LR: 0.000100
 Fine-tuning all layers...
 Epoch [1/10], Loss: 0.5874, Train Acc: 83.43%, Val Acc: 88.92%, LR: 0.000010
 Epoch [2/10], Loss: 0.5522, Train Acc: 86.20%, Val Acc: 88.09%, LR: 0.000010
 Epoch [3/10], Loss: 0.5325, Train Acc: 87.71%, Val Acc: 90.58%, LR: 0.000010
 Epoch [4/10], Loss: 0.5191, Train Acc: 88.64%, Val Acc: 91.97%, LR: 0.000010
 Epoch [5/10], Loss: 0.5009, Train Acc: 88.97%, Val Acc: 92.80%, LR: 0.000010
 Epoch [6/10], Loss: 0.4853, Train Acc: 89.74%, Val Acc: 92.80%, LR: 0.000010
 Epoch [7/10], Loss: 0.4729, Train Acc: 91.65%, Val Acc: 91.41%, LR: 0.000010
 Epoch [8/10], Loss: 0.4675, Train Acc: 91.84%, Val Acc: 93.91%, LR: 0.000010
 Epoch [9/10], Loss: 0.4593, Train Acc: 91.87%, Val Acc: 94.74%, LR: 0.000010
 Epoch [10/10], Loss: 0.4460, Train Acc: 93.04%, Val Acc: 93.91%, LR: 0.000010
 Testing model...

=== split_70_30 Results ===
 Accuracy: 0.9483
 Precision: 0.9484
 Recall: 0.9483
 F1-Score: 0.9483

Classification Report:

	precision	recall	f1-score	support
AD	0.94	0.92	0.93	337
CN	0.93	0.94	0.93	432
MCI	0.97	0.96	0.96	777
accuracy			0.95	1546
macro avg	0.94	0.94	0.94	1546
weighted avg	0.95	0.95	0.95	1546

=====

Processing: split_80_20

```
=====
Train samples: 3711
Val samples: 412
Test samples: 1031
Starting training...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.

    warnings.warn(
Epoch [1/25], Loss: 1.0393, Train Acc: 49.80%, Val Acc: 54.13%, LR: 0.000100
Epoch [2/25], Loss: 0.9873, Train Acc: 55.03%, Val Acc: 56.31%, LR: 0.000100
Epoch [3/25], Loss: 0.9428, Train Acc: 58.18%, Val Acc: 61.89%, LR: 0.000100
Epoch [4/25], Loss: 0.9112, Train Acc: 59.55%, Val Acc: 63.35%, LR: 0.000100
Epoch [5/25], Loss: 0.8735, Train Acc: 63.78%, Val Acc: 66.75%, LR: 0.000100
Epoch [6/25], Loss: 0.8496, Train Acc: 64.62%, Val Acc: 71.12%, LR: 0.000100
Epoch [7/25], Loss: 0.8290, Train Acc: 66.29%, Val Acc: 73.79%, LR: 0.000100
Epoch [8/25], Loss: 0.7940, Train Acc: 69.42%, Val Acc: 73.30%, LR: 0.000100
Epoch [9/25], Loss: 0.7774, Train Acc: 70.57%, Val Acc: 78.16%, LR: 0.000100
Epoch [10/25], Loss: 0.7604, Train Acc: 71.76%, Val Acc: 78.88%, LR: 0.000100
Epoch [11/25], Loss: 0.7295, Train Acc: 73.92%, Val Acc: 79.85%, LR: 0.000100
Epoch [12/25], Loss: 0.7266, Train Acc: 74.89%, Val Acc: 81.80%, LR: 0.000100
Epoch [13/25], Loss: 0.7113, Train Acc: 75.88%, Val Acc: 82.77%, LR: 0.000100
Epoch [14/25], Loss: 0.6897, Train Acc: 76.80%, Val Acc: 85.19%, LR: 0.000100
Epoch [15/25], Loss: 0.6781, Train Acc: 77.66%, Val Acc: 85.68%, LR: 0.000100
Epoch [16/25], Loss: 0.6697, Train Acc: 78.42%, Val Acc: 86.17%, LR: 0.000100
Epoch [17/25], Loss: 0.6507, Train Acc: 79.33%, Val Acc: 86.17%, LR: 0.000100
Epoch [18/25], Loss: 0.6455, Train Acc: 79.90%, Val Acc: 86.65%, LR: 0.000100
Epoch [19/25], Loss: 0.6302, Train Acc: 81.24%, Val Acc: 86.89%, LR: 0.000100
Epoch [20/25], Loss: 0.6214, Train Acc: 81.78%, Val Acc: 89.32%, LR: 0.000100
Epoch [21/25], Loss: 0.6153, Train Acc: 82.27%, Val Acc: 89.08%, LR: 0.000100
Epoch [22/25], Loss: 0.6012, Train Acc: 82.54%, Val Acc: 89.81%, LR: 0.000100
Epoch [23/25], Loss: 0.5919, Train Acc: 83.45%, Val Acc: 88.83%, LR: 0.000100
Epoch [24/25], Loss: 0.5903, Train Acc: 83.56%, Val Acc: 89.08%, LR: 0.000100
Epoch [25/25], Loss: 0.5817, Train Acc: 84.51%, Val Acc: 90.53%, LR: 0.000100
Fine-tuning all layers...
Epoch [1/10], Loss: 0.5686, Train Acc: 84.99%, Val Acc: 89.08%, LR: 0.000010
Epoch [2/10], Loss: 0.5354, Train Acc: 87.42%, Val Acc: 91.26%, LR: 0.000010
Epoch [3/10], Loss: 0.5145, Train Acc: 88.52%, Val Acc: 92.72%, LR: 0.000010
Epoch [4/10], Loss: 0.4966, Train Acc: 89.30%, Val Acc: 93.93%, LR: 0.000010
Epoch [5/10], Loss: 0.4826, Train Acc: 90.54%, Val Acc: 93.93%, LR: 0.000010
Epoch [6/10], Loss: 0.4702, Train Acc: 91.27%, Val Acc: 94.66%, LR: 0.000010
Epoch [7/10], Loss: 0.4603, Train Acc: 91.48%, Val Acc: 95.87%, LR: 0.000010
Epoch [8/10], Loss: 0.4528, Train Acc: 92.51%, Val Acc: 96.36%, LR: 0.000010
Epoch [9/10], Loss: 0.4438, Train Acc: 92.97%, Val Acc: 96.12%, LR: 0.000010
Epoch [10/10], Loss: 0.4247, Train Acc: 94.04%, Val Acc: 97.33%, LR: 0.000010
Testing model...
```

```
==== split_80_20 Results ====
Accuracy: 0.9525
Precision: 0.9525
Recall: 0.9525
F1-Score: 0.9524
```

Classification Report:

	precision	recall	f1-score	support
AD	0.93	0.92	0.93	225
CN	0.96	0.93	0.95	288
MCI	0.96	0.97	0.97	518
accuracy			0.95	1031
macro avg	0.95	0.94	0.95	1031
weighted avg	0.95	0.95	0.95	1031

```
=====
Processing: split_90_10
=====
```

Train samples: 4175

Val samples: 464

Test samples: 515

Starting training...

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/25], Loss: 1.0334, Train Acc: 50.97%, Val Acc: 53.23%, LR: 0.000100
Epoch [2/25], Loss: 0.9822, Train Acc: 54.42%, Val Acc: 58.41%, LR: 0.000100
Epoch [3/25], Loss: 0.9334, Train Acc: 58.68%, Val Acc: 63.58%, LR: 0.000100
Epoch [4/25], Loss: 0.8858, Train Acc: 62.42%, Val Acc: 68.10%, LR: 0.000100
Epoch [5/25], Loss: 0.8598, Train Acc: 65.15%, Val Acc: 71.12%, LR: 0.000100
Epoch [6/25], Loss: 0.8349, Train Acc: 66.30%, Val Acc: 73.28%, LR: 0.000100
Epoch [7/25], Loss: 0.8039, Train Acc: 69.27%, Val Acc: 76.08%, LR: 0.000100
Epoch [8/25], Loss: 0.7747, Train Acc: 71.43%, Val Acc: 80.82%, LR: 0.000100
Epoch [9/25], Loss: 0.7514, Train Acc: 72.74%, Val Acc: 79.53%, LR: 0.000100
Epoch [10/25], Loss: 0.7246, Train Acc: 74.92%, Val Acc: 80.60%, LR: 0.000100
Epoch [11/25], Loss: 0.7185, Train Acc: 75.09%, Val Acc: 83.41%, LR: 0.000100
Epoch [12/25], Loss: 0.7001, Train Acc: 75.93%, Val Acc: 83.19%, LR: 0.000100
Epoch [13/25], Loss: 0.6847, Train Acc: 77.44%, Val Acc: 86.42%, LR: 0.000100
Epoch [14/25], Loss: 0.6748, Train Acc: 77.96%, Val Acc: 84.91%, LR: 0.000100
Epoch [15/25], Loss: 0.6594, Train Acc: 78.68%, Val Acc: 85.56%, LR: 0.000100
Epoch [16/25], Loss: 0.6405, Train Acc: 80.12%, Val Acc: 86.21%, LR: 0.000100
Epoch [17/25], Loss: 0.6244, Train Acc: 80.86%, Val Acc: 87.28%, LR: 0.000100
Epoch [18/25], Loss: 0.6163, Train Acc: 81.89%, Val Acc: 90.09%, LR: 0.000100
```

Epoch [19/25], Loss: 0.6046, Train Acc: 83.66%, Val Acc: 89.22%, LR: 0.000100
 Epoch [20/25], Loss: 0.6099, Train Acc: 82.28%, Val Acc: 89.66%, LR: 0.000100
 Epoch [21/25], Loss: 0.5941, Train Acc: 83.07%, Val Acc: 88.79%, LR: 0.000100
 Epoch [22/25], Loss: 0.5851, Train Acc: 83.69%, Val Acc: 88.36%, LR: 0.000050
 Epoch [23/25], Loss: 0.5540, Train Acc: 85.92%, Val Acc: 90.73%, LR: 0.000050
 Epoch [24/25], Loss: 0.5536, Train Acc: 85.89%, Val Acc: 90.73%, LR: 0.000050
 Epoch [25/25], Loss: 0.5393, Train Acc: 86.80%, Val Acc: 92.24%, LR: 0.000050
 Fine-tuning all layers...

Epoch [1/10], Loss: 0.5384, Train Acc: 87.09%, Val Acc: 91.59%, LR: 0.000010
 Epoch [2/10], Loss: 0.5223, Train Acc: 88.07%, Val Acc: 92.89%, LR: 0.000010
 Epoch [3/10], Loss: 0.4962, Train Acc: 90.08%, Val Acc: 93.32%, LR: 0.000010
 Epoch [4/10], Loss: 0.4780, Train Acc: 91.07%, Val Acc: 94.40%, LR: 0.000010
 Epoch [5/10], Loss: 0.4673, Train Acc: 92.00%, Val Acc: 93.75%, LR: 0.000010
 Epoch [6/10], Loss: 0.4593, Train Acc: 92.00%, Val Acc: 94.40%, LR: 0.000010
 Epoch [7/10], Loss: 0.4448, Train Acc: 93.03%, Val Acc: 96.55%, LR: 0.000010
 Epoch [8/10], Loss: 0.4292, Train Acc: 93.84%, Val Acc: 95.69%, LR: 0.000010
 Epoch [9/10], Loss: 0.4239, Train Acc: 94.42%, Val Acc: 96.12%, LR: 0.000010
 Epoch [10/10], Loss: 0.4195, Train Acc: 94.23%, Val Acc: 95.91%, LR: 0.000005
 Testing model...

```
==== split_90_10 Results ===
Accuracy: 0.9417
Precision: 0.9423
Recall: 0.9417
F1-Score: 0.9413
```

Classification Report:

	precision	recall	f1-score	support
AD	0.95	0.87	0.91	112
CN	0.96	0.94	0.95	144
MCI	0.93	0.97	0.95	259
accuracy			0.94	515
macro avg	0.95	0.93	0.94	515
weighted avg	0.94	0.94	0.94	515

===== ConvNeXt-Small - SUMMARY OF ALL SPLITS =====

split	accuracy	precision	recall	f1_score	training_time
split_10_90	0.636775	0.634546	0.636775	0.635251	290.025709
split_20_80	0.773466	0.770169	0.773466	0.767802	542.540642
split_30_70	0.819568	0.825429	0.819568	0.815577	844.342106
split_40_60	0.876455	0.877528	0.876455	0.874956	1121.625662
split_50_50	0.906480	0.906529	0.906480	0.905963	1395.717910
split_60_40	0.936469	0.936461	0.936469	0.936248	1671.911056

```

split_70_30 0.948254 0.948387 0.948254 0.948274 1949.841236
split_80_20 0.952473 0.952461 0.952473 0.952366 2227.222875
split_90_10 0.941748 0.942265 0.941748 0.941317 2503.637483

```

Detailed results saved to: /kaggle/working/convnext_small_results.csv

```

[7]: import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, models
import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import CosineAnnealingLR, ReduceLROnPlateau

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)

        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        try:
            image = Image.open(img_path).convert('RGB')

            if self.transform:
                image = self.transform(image)

        return image, label
    except Exception as e:

```

```

        print(f"Error loading image {img_path}: {e}")
        # Return a dummy image and label
        dummy_image = torch.zeros(3, 224, 224)
        return dummy_image, label

def get_data_transforms():
    train_transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.RandomHorizontalFlip(p=0.5),
        transforms.RandomRotation(10),
        transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
        transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
        transforms.RandomGrayscale(p=0.1),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    val_transform = transforms.Compose([
        transforms.Resize((224, 224)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    return train_transform, val_transform

def train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs, device, warmup_epochs=3):
    best_val_acc = 0
    patience = 8
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        # Learning rate warmup
        if epoch < warmup_epochs:
            lr_scale = min(1.0, float(epoch + 1) / warmup_epochs)
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['initial_lr'] * lr_scale

        for batch_idx, (images, labels) in enumerate(train_loader):

```

```

        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()

        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    epoch_loss = running_loss / len(train_loader)
    epoch_acc = 100 * correct / total

    val_acc = evaluate_model(model, val_loader, device)

    if scheduler and epoch >= warmup_epochs:
        if isinstance(scheduler, CosineAnnealingLR):
            scheduler.step()
        else:
            scheduler.step(val_acc)

    current_lr = optimizer.param_groups[0]['lr']
    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train_Acc: {epoch_acc:.2f}%, Val Acc: {val_acc:.2f}%, LR: {current_lr:.6f}')

    if val_acc > best_val_acc:
        best_val_acc = val_acc
        patience_counter = 0
    else:
        patience_counter += 1

    if patience_counter >= patience:
        print("Early stopping at epoch {epoch+1}")
        break

def evaluate_model(model, data_loader, device):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in data_loader:

```

```

        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    return accuracy

def test_model(model, test_loader, device):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    return all_preds, all_labels

def calculate_metrics(y_true, y_pred, split_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred, average='weighted', zero_division=0)
    recall = recall_score(y_true, y_pred, average='weighted', zero_division=0)
    f1 = f1_score(y_true, y_pred, average='weighted', zero_division=0)

    print(f"\n==== {split_name} Results ===")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-Score: {f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=['AD', 'CN', 'MCI'], zero_division=0))

    return {
        'split': split_name,
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1_score': f1
    }

```

```

}

def run_densenet121_on_splits():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    splits_root = "/kaggle/working/alzheimer-resized-224_splits"
    split_folders = [f for f in os.listdir(splits_root) if f.
    ↪startswith('split_')]
    split_folders.sort()

    results = []

    train_transform, val_transform = get_data_transforms()

    for split_folder in split_folders:
        print(f"\n{'='*60}")
        print(f"Processing: {split_folder}")
        print(f"{'='*60}")

        split_path = os.path.join(splits_root, split_folder)

        train_dataset = AlzheimerDataset(split_path, 'train', train_transform)
        val_dataset = AlzheimerDataset(split_path, 'val', val_transform)
        test_dataset = AlzheimerDataset(split_path, 'test', val_transform)

        train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True,
        ↪num_workers=2, pin_memory=True)
        val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False,
        ↪num_workers=2, pin_memory=True)
        test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False,
        ↪num_workers=2, pin_memory=True)

        print(f"Train samples: {len(train_dataset)}")
        print(f"Val samples: {len(val_dataset)}")
        print(f"Test samples: {len(test_dataset)}")

        model = models.densenet121(weights=models.DenseNet121_Weights.
        ↪IMAGENET1K_V1)

        # Freeze all layers initially
        for param in model.parameters():
            param.requires_grad = False

        # Unfreeze classifier
        num_ftrs = model.classifier.in_features
        model.classifier = nn.Linear(num_ftrs, 3)

```

```

# Unfreeze last dense block and transition layer
for param in model.features.denseblock4.parameters():
    param.requires_grad = True
for param in model.features.norm5.parameters():
    param.requires_grad = True

model = model.to(device)

# Layer-wise learning rates for DenseNet
optimizer = torch.optim.AdamW([
    {'params': model.classifier.parameters(), 'lr': 0.0001, ↴
    'initial_lr': 0.0001},
    {'params': model.features.denseblock4.parameters(), 'lr': 0.00005, ↴
    'initial_lr': 0.00005},
    {'params': model.features.norm5.parameters(), 'lr': 0.00005, ↴
    'initial_lr': 0.00005},
    {'params': model.features.denseblock3.parameters(), 'lr': 0.00001, ↴
    'initial_lr': 0.00001}
], weight_decay=1e-4)

criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
scheduler = CosineAnnealingLR(optimizer, T_max=20, eta_min=1e-6)

print("Starting phase 1 training (partial unfreeze)...") 
start_time = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, ↴
scheduler, num_epochs=20, device=device)
phase1_time = time.time() - start_time

print("Starting phase 2 training (full fine-tuning)...") 
# Unfreeze all parameters for final fine-tuning
for param in model.parameters():
    param.requires_grad = True

optimizer = torch.optim.AdamW(model.parameters(), lr=0.00001, ↴
weight_decay=1e-5)
scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5, ↴
patience=3, verbose=True)

start_time_phase2 = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, ↴
scheduler, num_epochs=15, device=device, warmup_epochs=0)
total_time = phase1_time + (time.time() - start_time_phase2)

print("Testing model...")

```

```

    test_preds, test_labels = test_model(model, test_loader, device)

    split_results = calculate_metrics(test_labels, test_preds, split_folder)
    split_results['training_time'] = total_time
    results.append(split_results)

    torch.cuda.empty_cache()

    results_df = pd.DataFrame(results)
    print(f"\n{'='*80}")
    print("DenseNet-121 - SUMMARY OF ALL SPLITS")
    print(f"\n{'='*80}")
    print(results_df.to_string(index=False))

    results_csv_path = "/kaggle/working/densenet121_results.csv"
    results_df.to_csv(results_csv_path, index=False)
    print(f"\nDetailed results saved to: {results_csv_path}")

    return results_df

if __name__ == "__main__":
    results = run_densenet121_on_splits()

```

Downloading: "https://download.pytorch.org/models/densenet121-a639ec97.pth" to /root/.cache/torch/hub/checkpoints/densenet121-a639ec97.pth

Using device: cuda

```

=====
Processing: split_10_90
=====
Train samples: 464
Val samples: 51
Test samples: 4639

100% | 30.8M/30.8M [00:00<00:00, 226MB/s]

Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.1429, Train Acc: 32.11%, Val Acc: 35.29%, LR: 0.000033
Epoch [2/20], Loss: 1.0483, Train Acc: 49.35%, Val Acc: 50.98%, LR: 0.000067
Epoch [3/20], Loss: 1.0144, Train Acc: 53.23%, Val Acc: 49.02%, LR: 0.000100
Epoch [4/20], Loss: 0.9889, Train Acc: 53.66%, Val Acc: 49.02%, LR: 0.000099
Epoch [5/20], Loss: 0.9693, Train Acc: 54.31%, Val Acc: 54.90%, LR: 0.000098
Epoch [6/20], Loss: 0.9488, Train Acc: 58.84%, Val Acc: 54.90%, LR: 0.000095
Epoch [7/20], Loss: 0.9013, Train Acc: 61.42%, Val Acc: 52.94%, LR: 0.000091
Epoch [8/20], Loss: 0.8932, Train Acc: 63.79%, Val Acc: 52.94%, LR: 0.000086
Epoch [9/20], Loss: 0.8880, Train Acc: 62.07%, Val Acc: 52.94%, LR: 0.000080
Epoch [10/20], Loss: 0.8617, Train Acc: 63.79%, Val Acc: 56.86%, LR: 0.000073
Epoch [11/20], Loss: 0.8390, Train Acc: 66.16%, Val Acc: 58.82%, LR: 0.000066
```

```
Epoch [12/20], Loss: 0.8354, Train Acc: 68.32%, Val Acc: 58.82%, LR: 0.000058
Epoch [13/20], Loss: 0.8368, Train Acc: 66.38%, Val Acc: 56.86%, LR: 0.000051
Epoch [14/20], Loss: 0.8008, Train Acc: 70.91%, Val Acc: 64.71%, LR: 0.000043
Epoch [15/20], Loss: 0.8073, Train Acc: 71.34%, Val Acc: 64.71%, LR: 0.000035
Epoch [16/20], Loss: 0.7787, Train Acc: 71.55%, Val Acc: 58.82%, LR: 0.000028
Epoch [17/20], Loss: 0.7723, Train Acc: 74.35%, Val Acc: 60.78%, LR: 0.000021
Epoch [18/20], Loss: 0.7786, Train Acc: 73.28%, Val Acc: 64.71%, LR: 0.000015
Epoch [19/20], Loss: 0.7631, Train Acc: 73.49%, Val Acc: 62.75%, LR: 0.000010
Epoch [20/20], Loss: 0.7657, Train Acc: 74.57%, Val Acc: 60.78%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
warnings.warn(
```

```
Epoch [1/15], Loss: 0.7700, Train Acc: 73.06%, Val Acc: 68.63%, LR: 0.000010
Epoch [2/15], Loss: 0.7560, Train Acc: 74.57%, Val Acc: 64.71%, LR: 0.000010
Epoch [3/15], Loss: 0.7474, Train Acc: 74.57%, Val Acc: 64.71%, LR: 0.000010
Epoch [4/15], Loss: 0.7282, Train Acc: 76.08%, Val Acc: 70.59%, LR: 0.000010
Epoch [5/15], Loss: 0.7140, Train Acc: 75.86%, Val Acc: 70.59%, LR: 0.000010
Epoch [6/15], Loss: 0.6882, Train Acc: 81.03%, Val Acc: 72.55%, LR: 0.000010
Epoch [7/15], Loss: 0.6854, Train Acc: 79.96%, Val Acc: 70.59%, LR: 0.000010
Epoch [8/15], Loss: 0.6737, Train Acc: 80.82%, Val Acc: 66.67%, LR: 0.000010
Epoch [9/15], Loss: 0.6551, Train Acc: 82.11%, Val Acc: 68.63%, LR: 0.000010
Epoch [10/15], Loss: 0.6580, Train Acc: 82.33%, Val Acc: 68.63%, LR: 0.000005
Epoch [11/15], Loss: 0.6253, Train Acc: 83.41%, Val Acc: 68.63%, LR: 0.000005
Epoch [12/15], Loss: 0.6566, Train Acc: 81.68%, Val Acc: 70.59%, LR: 0.000005
Epoch [13/15], Loss: 0.6378, Train Acc: 83.62%, Val Acc: 70.59%, LR: 0.000005
Epoch [14/15], Loss: 0.6239, Train Acc: 82.54%, Val Acc: 70.59%, LR: 0.000003
Early stopping at epoch 14
```

```
Testing model...
```

```
== split_10_90 Results ==
```

```
Accuracy: 0.6264
```

```
Precision: 0.6384
```

```
Recall: 0.6264
```

```
F1-Score: 0.6157
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.49	0.56	0.52	1012
CN	0.71	0.39	0.50	1296
MCI	0.66	0.79	0.72	2331
accuracy			0.63	4639
macro avg	0.62	0.58	0.58	4639
weighted avg	0.64	0.63	0.62	4639

```
=====
Processing: split_20_80
=====
Train samples: 928
Val samples: 103
Test samples: 4123
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.1140, Train Acc: 38.47%, Val Acc: 48.54%, LR: 0.000033
Epoch [2/20], Loss: 1.0469, Train Acc: 49.03%, Val Acc: 49.51%, LR: 0.000067
Epoch [3/20], Loss: 1.0261, Train Acc: 52.26%, Val Acc: 51.46%, LR: 0.000100
Epoch [4/20], Loss: 0.9732, Train Acc: 55.60%, Val Acc: 52.43%, LR: 0.000099
Epoch [5/20], Loss: 0.9655, Train Acc: 56.57%, Val Acc: 54.37%, LR: 0.000098
Epoch [6/20], Loss: 0.9349, Train Acc: 58.94%, Val Acc: 53.40%, LR: 0.000095
Epoch [7/20], Loss: 0.9143, Train Acc: 59.59%, Val Acc: 57.28%, LR: 0.000091
Epoch [8/20], Loss: 0.8819, Train Acc: 62.93%, Val Acc: 60.19%, LR: 0.000086
Epoch [9/20], Loss: 0.8796, Train Acc: 63.36%, Val Acc: 60.19%, LR: 0.000080
Epoch [10/20], Loss: 0.8598, Train Acc: 64.01%, Val Acc: 62.14%, LR: 0.000073
Epoch [11/20], Loss: 0.8257, Train Acc: 67.89%, Val Acc: 62.14%, LR: 0.000066
Epoch [12/20], Loss: 0.8253, Train Acc: 66.38%, Val Acc: 63.11%, LR: 0.000058
Epoch [13/20], Loss: 0.8184, Train Acc: 68.75%, Val Acc: 66.99%, LR: 0.000051
Epoch [14/20], Loss: 0.8072, Train Acc: 68.64%, Val Acc: 63.11%, LR: 0.000043
Epoch [15/20], Loss: 0.7943, Train Acc: 69.40%, Val Acc: 68.93%, LR: 0.000035
Epoch [16/20], Loss: 0.7800, Train Acc: 71.23%, Val Acc: 66.99%, LR: 0.000028
Epoch [17/20], Loss: 0.7642, Train Acc: 71.66%, Val Acc: 66.02%, LR: 0.000021
Epoch [18/20], Loss: 0.7651, Train Acc: 72.74%, Val Acc: 68.93%, LR: 0.000015
Epoch [19/20], Loss: 0.7674, Train Acc: 71.12%, Val Acc: 65.05%, LR: 0.000010
Epoch [20/20], Loss: 0.7620, Train Acc: 73.71%, Val Acc: 66.99%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/15], Loss: 0.7667, Train Acc: 72.31%, Val Acc: 65.05%, LR: 0.000010
Epoch [2/15], Loss: 0.7454, Train Acc: 72.09%, Val Acc: 69.90%, LR: 0.000010
Epoch [3/15], Loss: 0.7389, Train Acc: 74.35%, Val Acc: 69.90%, LR: 0.000010
Epoch [4/15], Loss: 0.7178, Train Acc: 75.43%, Val Acc: 73.79%, LR: 0.000010
Epoch [5/15], Loss: 0.7081, Train Acc: 77.26%, Val Acc: 72.82%, LR: 0.000010
Epoch [6/15], Loss: 0.6806, Train Acc: 77.48%, Val Acc: 71.84%, LR: 0.000010
Epoch [7/15], Loss: 0.6982, Train Acc: 76.51%, Val Acc: 73.79%, LR: 0.000010
Epoch [8/15], Loss: 0.6496, Train Acc: 78.77%, Val Acc: 73.79%, LR: 0.000005
Epoch [9/15], Loss: 0.6598, Train Acc: 79.53%, Val Acc: 76.70%, LR: 0.000005
Epoch [10/15], Loss: 0.6292, Train Acc: 84.70%, Val Acc: 75.73%, LR: 0.000005
Epoch [11/15], Loss: 0.6509, Train Acc: 78.99%, Val Acc: 76.70%, LR: 0.000005
Epoch [12/15], Loss: 0.6334, Train Acc: 82.11%, Val Acc: 73.79%, LR: 0.000005
Epoch [13/15], Loss: 0.6226, Train Acc: 81.79%, Val Acc: 77.67%, LR: 0.000005
```

```
Epoch [14/15], Loss: 0.6041, Train Acc: 83.94%, Val Acc: 78.64%, LR: 0.000005
Epoch [15/15], Loss: 0.5866, Train Acc: 85.78%, Val Acc: 80.58%, LR: 0.000005
Testing model...
```

```
==== split_20_80 Results ===
```

```
Accuracy: 0.7317
```

```
Precision: 0.7282
```

```
Recall: 0.7317
```

```
F1-Score: 0.7261
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.65	0.54	0.59	899
CN	0.75	0.66	0.70	1152
MCI	0.75	0.86	0.80	2072
accuracy			0.73	4123
macro avg	0.72	0.68	0.70	4123
weighted avg	0.73	0.73	0.73	4123

```
=====
Processing: split_30_70
=====
```

```
Train samples: 1391
```

```
Val samples: 155
```

```
Test samples: 3608
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/20], Loss: 1.0540, Train Acc: 47.52%, Val Acc: 53.55%, LR: 0.000033
Epoch [2/20], Loss: 1.0246, Train Acc: 51.11%, Val Acc: 56.13%, LR: 0.000067
Epoch [3/20], Loss: 0.9905, Train Acc: 53.70%, Val Acc: 60.65%, LR: 0.000100
Epoch [4/20], Loss: 0.9559, Train Acc: 56.36%, Val Acc: 61.94%, LR: 0.000099
Epoch [5/20], Loss: 0.9310, Train Acc: 58.88%, Val Acc: 63.87%, LR: 0.000098
Epoch [6/20], Loss: 0.8967, Train Acc: 62.11%, Val Acc: 63.87%, LR: 0.000095
Epoch [7/20], Loss: 0.8896, Train Acc: 61.61%, Val Acc: 65.16%, LR: 0.000091
Epoch [8/20], Loss: 0.8423, Train Acc: 66.57%, Val Acc: 65.81%, LR: 0.000086
Epoch [9/20], Loss: 0.8457, Train Acc: 65.78%, Val Acc: 67.74%, LR: 0.000080
Epoch [10/20], Loss: 0.8197, Train Acc: 66.50%, Val Acc: 64.52%, LR: 0.000073
Epoch [11/20], Loss: 0.8035, Train Acc: 69.23%, Val Acc: 68.39%, LR: 0.000066
Epoch [12/20], Loss: 0.7875, Train Acc: 70.02%, Val Acc: 67.10%, LR: 0.000058
Epoch [13/20], Loss: 0.7737, Train Acc: 71.89%, Val Acc: 69.03%, LR: 0.000051
Epoch [14/20], Loss: 0.7664, Train Acc: 71.96%, Val Acc: 67.10%, LR: 0.000043
Epoch [15/20], Loss: 0.7706, Train Acc: 70.81%, Val Acc: 68.39%, LR: 0.000035
Epoch [16/20], Loss: 0.7323, Train Acc: 73.18%, Val Acc: 66.45%, LR: 0.000028
Epoch [17/20], Loss: 0.7143, Train Acc: 75.56%, Val Acc: 67.74%, LR: 0.000021
Epoch [18/20], Loss: 0.7145, Train Acc: 75.49%, Val Acc: 69.03%, LR: 0.000015
Epoch [19/20], Loss: 0.7188, Train Acc: 76.28%, Val Acc: 69.68%, LR: 0.000010
```

```
Epoch [20/20], Loss: 0.7126, Train Acc: 75.77%, Val Acc: 70.97%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.7177, Train Acc: 74.91%, Val Acc: 69.03%, LR: 0.000010
Epoch [2/15], Loss: 0.7031, Train Acc: 76.71%, Val Acc: 69.68%, LR: 0.000010
Epoch [3/15], Loss: 0.6930, Train Acc: 75.92%, Val Acc: 72.90%, LR: 0.000010
Epoch [4/15], Loss: 0.6771, Train Acc: 78.00%, Val Acc: 69.68%, LR: 0.000010
Epoch [5/15], Loss: 0.6608, Train Acc: 78.94%, Val Acc: 70.97%, LR: 0.000010
Epoch [6/15], Loss: 0.6213, Train Acc: 82.31%, Val Acc: 74.84%, LR: 0.000010
Epoch [7/15], Loss: 0.6317, Train Acc: 80.95%, Val Acc: 69.03%, LR: 0.000010
Epoch [8/15], Loss: 0.6125, Train Acc: 83.54%, Val Acc: 70.32%, LR: 0.000010
Epoch [9/15], Loss: 0.6018, Train Acc: 82.75%, Val Acc: 75.48%, LR: 0.000010
Epoch [10/15], Loss: 0.5915, Train Acc: 84.47%, Val Acc: 72.90%, LR: 0.000010
Epoch [11/15], Loss: 0.5831, Train Acc: 85.12%, Val Acc: 72.90%, LR: 0.000010
Epoch [12/15], Loss: 0.5761, Train Acc: 84.97%, Val Acc: 71.61%, LR: 0.000010
Epoch [13/15], Loss: 0.5423, Train Acc: 87.28%, Val Acc: 74.84%, LR: 0.000005
Epoch [14/15], Loss: 0.5420, Train Acc: 87.42%, Val Acc: 74.19%, LR: 0.000005
Epoch [15/15], Loss: 0.5399, Train Acc: 87.56%, Val Acc: 74.19%, LR: 0.000005
Testing model...
```

```
==== split_30_70 Results ===
```

```
Accuracy: 0.8071
```

```
Precision: 0.8058
```

```
Recall: 0.8071
```

```
F1-Score: 0.8056
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.76	0.70	0.72	787
CN	0.81	0.77	0.79	1008
MCI	0.82	0.88	0.85	1813
accuracy			0.81	3608
macro avg	0.80	0.78	0.79	3608
weighted avg	0.81	0.81	0.81	3608

```
=====
Processing: split_40_60
=====
```

```
Train samples: 1855
```

```
Val samples: 207
```

```
Test samples: 3092
```

```
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0519, Train Acc: 49.65%, Val Acc: 53.14%, LR: 0.000033
Epoch [2/20], Loss: 1.0122, Train Acc: 52.29%, Val Acc: 54.11%, LR: 0.000067
Epoch [3/20], Loss: 0.9748, Train Acc: 55.09%, Val Acc: 56.04%, LR: 0.000100
Epoch [4/20], Loss: 0.9387, Train Acc: 58.22%, Val Acc: 55.56%, LR: 0.000099
Epoch [5/20], Loss: 0.8999, Train Acc: 61.62%, Val Acc: 62.80%, LR: 0.000098
Epoch [6/20], Loss: 0.8769, Train Acc: 64.74%, Val Acc: 61.84%, LR: 0.000095
Epoch [7/20], Loss: 0.8388, Train Acc: 66.25%, Val Acc: 57.97%, LR: 0.000091
Epoch [8/20], Loss: 0.8208, Train Acc: 67.39%, Val Acc: 57.00%, LR: 0.000086
Epoch [9/20], Loss: 0.7928, Train Acc: 70.35%, Val Acc: 64.25%, LR: 0.000080
Epoch [10/20], Loss: 0.7676, Train Acc: 71.86%, Val Acc: 61.84%, LR: 0.000073
Epoch [11/20], Loss: 0.7444, Train Acc: 74.02%, Val Acc: 66.67%, LR: 0.000066
Epoch [12/20], Loss: 0.7452, Train Acc: 73.48%, Val Acc: 65.70%, LR: 0.000058
Epoch [13/20], Loss: 0.7327, Train Acc: 74.77%, Val Acc: 68.60%, LR: 0.000051
Epoch [14/20], Loss: 0.7052, Train Acc: 76.17%, Val Acc: 61.84%, LR: 0.000043
Epoch [15/20], Loss: 0.7107, Train Acc: 75.09%, Val Acc: 67.15%, LR: 0.000035
Epoch [16/20], Loss: 0.6822, Train Acc: 77.30%, Val Acc: 68.12%, LR: 0.000028
Epoch [17/20], Loss: 0.6831, Train Acc: 77.04%, Val Acc: 67.63%, LR: 0.000021
Epoch [18/20], Loss: 0.6669, Train Acc: 79.14%, Val Acc: 70.05%, LR: 0.000015
Epoch [19/20], Loss: 0.6658, Train Acc: 78.76%, Val Acc: 68.60%, LR: 0.000010
Epoch [20/20], Loss: 0.6826, Train Acc: 77.68%, Val Acc: 69.57%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.6779, Train Acc: 78.49%, Val Acc: 68.60%, LR: 0.000010
Epoch [2/15], Loss: 0.6554, Train Acc: 79.68%, Val Acc: 73.43%, LR: 0.000010
Epoch [3/15], Loss: 0.6293, Train Acc: 81.78%, Val Acc: 71.50%, LR: 0.000010
Epoch [4/15], Loss: 0.6201, Train Acc: 81.67%, Val Acc: 72.46%, LR: 0.000010
Epoch [5/15], Loss: 0.6024, Train Acc: 83.07%, Val Acc: 76.33%, LR: 0.000010
Epoch [6/15], Loss: 0.5865, Train Acc: 84.10%, Val Acc: 78.74%, LR: 0.000010
Epoch [7/15], Loss: 0.5897, Train Acc: 84.74%, Val Acc: 78.74%, LR: 0.000010
Epoch [8/15], Loss: 0.5706, Train Acc: 85.44%, Val Acc: 82.13%, LR: 0.000010
Epoch [9/15], Loss: 0.5521, Train Acc: 87.22%, Val Acc: 79.23%, LR: 0.000010
Epoch [10/15], Loss: 0.5446, Train Acc: 86.85%, Val Acc: 78.26%, LR: 0.000010
Epoch [11/15], Loss: 0.5328, Train Acc: 88.09%, Val Acc: 82.61%, LR: 0.000010
Epoch [12/15], Loss: 0.5221, Train Acc: 88.36%, Val Acc: 79.23%, LR: 0.000010
Epoch [13/15], Loss: 0.5229, Train Acc: 88.19%, Val Acc: 82.61%, LR: 0.000010
Epoch [14/15], Loss: 0.5103, Train Acc: 89.81%, Val Acc: 81.16%, LR: 0.000010
Epoch [15/15], Loss: 0.4857, Train Acc: 91.37%, Val Acc: 80.68%, LR: 0.000005
Testing model...
```

```
==== split_40_60 Results ===
```

```
Accuracy: 0.8195
```

```
Precision: 0.8268
```

```
Recall: 0.8195
```

F1-Score: 0.8207

Classification Report:

	precision	recall	f1-score	support
AD	0.69	0.81	0.74	674
CN	0.88	0.75	0.81	864
MCI	0.86	0.87	0.86	1554
accuracy			0.82	3092
macro avg	0.81	0.81	0.80	3092
weighted avg	0.83	0.82	0.82	3092

=====

Processing: split_50_50

=====

Train samples: 2319

Val samples: 258

Test samples: 2577

Starting phase 1 training (partial unfreeze)...

Epoch [1/20], Loss: 1.0487, Train Acc: 47.99%, Val Acc: 51.94%, LR: 0.000033
Epoch [2/20], Loss: 1.0032, Train Acc: 53.39%, Val Acc: 58.91%, LR: 0.000067
Epoch [3/20], Loss: 0.9598, Train Acc: 56.71%, Val Acc: 55.04%, LR: 0.000100
Epoch [4/20], Loss: 0.9081, Train Acc: 60.54%, Val Acc: 56.98%, LR: 0.000099
Epoch [5/20], Loss: 0.8767, Train Acc: 63.13%, Val Acc: 58.91%, LR: 0.000098
Epoch [6/20], Loss: 0.8543, Train Acc: 65.93%, Val Acc: 62.02%, LR: 0.000095
Epoch [7/20], Loss: 0.8176, Train Acc: 68.39%, Val Acc: 65.50%, LR: 0.000091
Epoch [8/20], Loss: 0.7861, Train Acc: 70.25%, Val Acc: 63.57%, LR: 0.000086
Epoch [9/20], Loss: 0.7786, Train Acc: 71.45%, Val Acc: 63.57%, LR: 0.000080
Epoch [10/20], Loss: 0.7593, Train Acc: 72.75%, Val Acc: 67.05%, LR: 0.000073
Epoch [11/20], Loss: 0.7521, Train Acc: 73.01%, Val Acc: 70.16%, LR: 0.000066
Epoch [12/20], Loss: 0.7236, Train Acc: 74.51%, Val Acc: 69.38%, LR: 0.000058
Epoch [13/20], Loss: 0.7104, Train Acc: 74.60%, Val Acc: 69.77%, LR: 0.000051
Epoch [14/20], Loss: 0.6980, Train Acc: 76.89%, Val Acc: 69.38%, LR: 0.000043
Epoch [15/20], Loss: 0.6845, Train Acc: 78.01%, Val Acc: 70.16%, LR: 0.000035
Epoch [16/20], Loss: 0.6713, Train Acc: 78.70%, Val Acc: 70.93%, LR: 0.000028
Epoch [17/20], Loss: 0.6743, Train Acc: 78.83%, Val Acc: 75.58%, LR: 0.000021
Epoch [18/20], Loss: 0.6533, Train Acc: 79.95%, Val Acc: 72.48%, LR: 0.000015
Epoch [19/20], Loss: 0.6567, Train Acc: 79.17%, Val Acc: 74.03%, LR: 0.000010
Epoch [20/20], Loss: 0.6605, Train Acc: 78.53%, Val Acc: 74.03%, LR: 0.000006

Starting phase 2 training (full fine-tuning)...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
warnings.warn(
Epoch [1/15], Loss: 0.6585, Train Acc: 79.26%, Val Acc: 76.74%, LR: 0.000010

```
Epoch [2/15], Loss: 0.6389, Train Acc: 81.46%, Val Acc: 74.81%, LR: 0.000010
Epoch [3/15], Loss: 0.6233, Train Acc: 81.72%, Val Acc: 75.58%, LR: 0.000010
Epoch [4/15], Loss: 0.6130, Train Acc: 82.66%, Val Acc: 75.19%, LR: 0.000010
Epoch [5/15], Loss: 0.6014, Train Acc: 83.35%, Val Acc: 82.17%, LR: 0.000010
Epoch [6/15], Loss: 0.5791, Train Acc: 84.56%, Val Acc: 81.01%, LR: 0.000010
Epoch [7/15], Loss: 0.5678, Train Acc: 85.12%, Val Acc: 81.78%, LR: 0.000010
Epoch [8/15], Loss: 0.5487, Train Acc: 87.02%, Val Acc: 81.40%, LR: 0.000010
Epoch [9/15], Loss: 0.5316, Train Acc: 87.54%, Val Acc: 82.56%, LR: 0.000010
Epoch [10/15], Loss: 0.5181, Train Acc: 88.70%, Val Acc: 86.05%, LR: 0.000010
Epoch [11/15], Loss: 0.5296, Train Acc: 88.18%, Val Acc: 85.66%, LR: 0.000010
Epoch [12/15], Loss: 0.5142, Train Acc: 89.18%, Val Acc: 85.27%, LR: 0.000010
Epoch [13/15], Loss: 0.5033, Train Acc: 90.38%, Val Acc: 83.72%, LR: 0.000010
Epoch [14/15], Loss: 0.5059, Train Acc: 89.00%, Val Acc: 86.43%, LR: 0.000010
Epoch [15/15], Loss: 0.4979, Train Acc: 89.87%, Val Acc: 86.82%, LR: 0.000010
Testing model...
```

==== split_50_50 Results ===

```
Accuracy: 0.8716
Precision: 0.8751
Recall: 0.8716
F1-Score: 0.8725
```

Classification Report:

	precision	recall	f1-score	support
AD	0.77	0.87	0.82	562
CN	0.88	0.84	0.86	720
MCI	0.92	0.89	0.90	1295
accuracy			0.87	2577
macro avg	0.86	0.87	0.86	2577
weighted avg	0.88	0.87	0.87	2577

=====

Processing: split_60_40

=====

Train samples: 2784

Val samples: 308

Test samples: 2062

Starting phase 1 training (partial unfreeze)...

```
Epoch [1/20], Loss: 1.0475, Train Acc: 48.92%, Val Acc: 48.70%, LR: 0.000033
Epoch [2/20], Loss: 0.9968, Train Acc: 53.92%, Val Acc: 50.97%, LR: 0.000067
Epoch [3/20], Loss: 0.9575, Train Acc: 56.86%, Val Acc: 59.74%, LR: 0.000100
Epoch [4/20], Loss: 0.9086, Train Acc: 61.89%, Val Acc: 61.36%, LR: 0.000099
Epoch [5/20], Loss: 0.8655, Train Acc: 64.51%, Val Acc: 62.66%, LR: 0.000098
Epoch [6/20], Loss: 0.8505, Train Acc: 66.06%, Val Acc: 64.29%, LR: 0.000095
Epoch [7/20], Loss: 0.8170, Train Acc: 68.35%, Val Acc: 65.26%, LR: 0.000091
```

```
Epoch [8/20], Loss: 0.7786, Train Acc: 70.58%, Val Acc: 69.81%, LR: 0.000086
Epoch [9/20], Loss: 0.7693, Train Acc: 71.19%, Val Acc: 69.48%, LR: 0.000080
Epoch [10/20], Loss: 0.7483, Train Acc: 73.64%, Val Acc: 69.48%, LR: 0.000073
Epoch [11/20], Loss: 0.7227, Train Acc: 75.04%, Val Acc: 72.08%, LR: 0.000066
Epoch [12/20], Loss: 0.7084, Train Acc: 75.11%, Val Acc: 74.03%, LR: 0.000058
Epoch [13/20], Loss: 0.6960, Train Acc: 76.08%, Val Acc: 72.73%, LR: 0.000051
Epoch [14/20], Loss: 0.6846, Train Acc: 77.19%, Val Acc: 75.65%, LR: 0.000043
Epoch [15/20], Loss: 0.6648, Train Acc: 78.81%, Val Acc: 75.00%, LR: 0.000035
Epoch [16/20], Loss: 0.6668, Train Acc: 78.59%, Val Acc: 74.03%, LR: 0.000028
Epoch [17/20], Loss: 0.6599, Train Acc: 78.34%, Val Acc: 74.03%, LR: 0.000021
Epoch [18/20], Loss: 0.6539, Train Acc: 80.14%, Val Acc: 74.03%, LR: 0.000015
Epoch [19/20], Loss: 0.6418, Train Acc: 80.24%, Val Acc: 72.40%, LR: 0.000010
Epoch [20/20], Loss: 0.6403, Train Acc: 80.10%, Val Acc: 74.68%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.6539, Train Acc: 79.71%, Val Acc: 74.35%, LR: 0.000010
Epoch [2/15], Loss: 0.6256, Train Acc: 81.50%, Val Acc: 76.62%, LR: 0.000010
Epoch [3/15], Loss: 0.6100, Train Acc: 82.58%, Val Acc: 77.92%, LR: 0.000010
Epoch [4/15], Loss: 0.5918, Train Acc: 83.66%, Val Acc: 77.27%, LR: 0.000010
Epoch [5/15], Loss: 0.5785, Train Acc: 84.30%, Val Acc: 80.19%, LR: 0.000010
Epoch [6/15], Loss: 0.5573, Train Acc: 86.35%, Val Acc: 82.79%, LR: 0.000010
Epoch [7/15], Loss: 0.5599, Train Acc: 86.03%, Val Acc: 82.79%, LR: 0.000010
Epoch [8/15], Loss: 0.5346, Train Acc: 87.57%, Val Acc: 87.01%, LR: 0.000010
Epoch [9/15], Loss: 0.5199, Train Acc: 88.36%, Val Acc: 86.36%, LR: 0.000010
Epoch [10/15], Loss: 0.5160, Train Acc: 88.76%, Val Acc: 87.66%, LR: 0.000010
Epoch [11/15], Loss: 0.5110, Train Acc: 88.58%, Val Acc: 86.69%, LR: 0.000010
Epoch [12/15], Loss: 0.5020, Train Acc: 89.80%, Val Acc: 86.36%, LR: 0.000010
Epoch [13/15], Loss: 0.4948, Train Acc: 90.55%, Val Acc: 89.29%, LR: 0.000010
Epoch [14/15], Loss: 0.4858, Train Acc: 90.70%, Val Acc: 88.96%, LR: 0.000010
Epoch [15/15], Loss: 0.4771, Train Acc: 91.16%, Val Acc: 88.64%, LR: 0.000010
Testing model...
```

```
==== split_60_40 Results ====
Accuracy: 0.9035
Precision: 0.9042
Recall: 0.9035
F1-Score: 0.9024
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.91	0.79	0.85	450
CN	0.92	0.89	0.91	576
MCI	0.89	0.96	0.92	1036

accuracy			0.90	2062
macro avg	0.91	0.88	0.89	2062
weighted avg	0.90	0.90	0.90	2062

```
=====
Processing: split_70_30
=====
```

Train samples: 3247

Val samples: 361

Test samples: 1546

Starting phase 1 training (partial unfreeze)...

Epoch [1/20], Loss: 1.0488, Train Acc: 48.41%, Val Acc: 52.08%, LR: 0.000033
 Epoch [2/20], Loss: 0.9982, Train Acc: 53.83%, Val Acc: 52.63%, LR: 0.000067
 Epoch [3/20], Loss: 0.9473, Train Acc: 57.99%, Val Acc: 56.51%, LR: 0.000100
 Epoch [4/20], Loss: 0.9046, Train Acc: 61.38%, Val Acc: 63.43%, LR: 0.000099
 Epoch [5/20], Loss: 0.8608, Train Acc: 65.85%, Val Acc: 63.71%, LR: 0.000098
 Epoch [6/20], Loss: 0.8240, Train Acc: 67.48%, Val Acc: 67.87%, LR: 0.000095
 Epoch [7/20], Loss: 0.7917, Train Acc: 69.73%, Val Acc: 63.71%, LR: 0.000091
 Epoch [8/20], Loss: 0.7693, Train Acc: 72.16%, Val Acc: 71.75%, LR: 0.000086
 Epoch [9/20], Loss: 0.7383, Train Acc: 72.96%, Val Acc: 73.41%, LR: 0.000080
 Epoch [10/20], Loss: 0.7190, Train Acc: 75.08%, Val Acc: 76.73%, LR: 0.000073
 Epoch [11/20], Loss: 0.7013, Train Acc: 76.47%, Val Acc: 75.07%, LR: 0.000066
 Epoch [12/20], Loss: 0.6813, Train Acc: 77.24%, Val Acc: 77.01%, LR: 0.000058
 Epoch [13/20], Loss: 0.6808, Train Acc: 78.72%, Val Acc: 76.18%, LR: 0.000051
 Epoch [14/20], Loss: 0.6672, Train Acc: 78.44%, Val Acc: 76.45%, LR: 0.000043
 Epoch [15/20], Loss: 0.6609, Train Acc: 79.12%, Val Acc: 76.18%, LR: 0.000035
 Epoch [16/20], Loss: 0.6491, Train Acc: 79.67%, Val Acc: 75.90%, LR: 0.000028
 Epoch [17/20], Loss: 0.6297, Train Acc: 81.06%, Val Acc: 78.12%, LR: 0.000021
 Epoch [18/20], Loss: 0.6318, Train Acc: 81.71%, Val Acc: 76.73%, LR: 0.000015
 Epoch [19/20], Loss: 0.6314, Train Acc: 81.52%, Val Acc: 78.67%, LR: 0.000010
 Epoch [20/20], Loss: 0.6229, Train Acc: 81.52%, Val Acc: 76.73%, LR: 0.000006

Starting phase 2 training (full fine-tuning)...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:

UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

warnings.warn(

Epoch [1/15], Loss: 0.6357, Train Acc: 80.04%, Val Acc: 78.39%, LR: 0.000010
 Epoch [2/15], Loss: 0.6110, Train Acc: 82.23%, Val Acc: 80.89%, LR: 0.000010
 Epoch [3/15], Loss: 0.5947, Train Acc: 83.22%, Val Acc: 82.55%, LR: 0.000010
 Epoch [4/15], Loss: 0.5676, Train Acc: 85.34%, Val Acc: 83.93%, LR: 0.000010
 Epoch [5/15], Loss: 0.5550, Train Acc: 86.60%, Val Acc: 84.49%, LR: 0.000010
 Epoch [6/15], Loss: 0.5521, Train Acc: 86.02%, Val Acc: 87.81%, LR: 0.000010
 Epoch [7/15], Loss: 0.5252, Train Acc: 87.99%, Val Acc: 84.21%, LR: 0.000010
 Epoch [8/15], Loss: 0.5191, Train Acc: 88.51%, Val Acc: 88.37%, LR: 0.000010
 Epoch [9/15], Loss: 0.5121, Train Acc: 88.45%, Val Acc: 86.98%, LR: 0.000010

```
Epoch [10/15], Loss: 0.5021, Train Acc: 89.59%, Val Acc: 88.37%, LR: 0.000010
Epoch [11/15], Loss: 0.4959, Train Acc: 89.93%, Val Acc: 89.20%, LR: 0.000010
Epoch [12/15], Loss: 0.4786, Train Acc: 90.82%, Val Acc: 90.03%, LR: 0.000010
Epoch [13/15], Loss: 0.4793, Train Acc: 91.22%, Val Acc: 90.86%, LR: 0.000010
Epoch [14/15], Loss: 0.4633, Train Acc: 91.87%, Val Acc: 89.75%, LR: 0.000010
Epoch [15/15], Loss: 0.4578, Train Acc: 92.52%, Val Acc: 89.47%, LR: 0.000010
Testing model...
```

```
==== split_70_30 Results ===
```

```
Accuracy: 0.9069
```

```
Precision: 0.9075
```

```
Recall: 0.9069
```

```
F1-Score: 0.9068
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.87	0.83	0.85	337
CN	0.87	0.93	0.90	432
MCI	0.94	0.93	0.94	777
accuracy			0.91	1546
macro avg	0.90	0.90	0.90	1546
weighted avg	0.91	0.91	0.91	1546

```
=====
Processing: split_80_20
=====
```

```
Train samples: 3711
```

```
Val samples: 412
```

```
Test samples: 1031
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/20], Loss: 1.0421, Train Acc: 49.02%, Val Acc: 53.64%, LR: 0.000033
Epoch [2/20], Loss: 0.9811, Train Acc: 54.43%, Val Acc: 57.52%, LR: 0.000067
Epoch [3/20], Loss: 0.9220, Train Acc: 60.36%, Val Acc: 59.47%, LR: 0.000100
Epoch [4/20], Loss: 0.8775, Train Acc: 64.21%, Val Acc: 67.48%, LR: 0.000099
Epoch [5/20], Loss: 0.8368, Train Acc: 67.21%, Val Acc: 68.45%, LR: 0.000098
Epoch [6/20], Loss: 0.8026, Train Acc: 69.23%, Val Acc: 71.60%, LR: 0.000095
Epoch [7/20], Loss: 0.7642, Train Acc: 72.00%, Val Acc: 68.93%, LR: 0.000091
Epoch [8/20], Loss: 0.7429, Train Acc: 73.43%, Val Acc: 73.30%, LR: 0.000086
Epoch [9/20], Loss: 0.7226, Train Acc: 74.75%, Val Acc: 74.51%, LR: 0.000080
Epoch [10/20], Loss: 0.6962, Train Acc: 76.26%, Val Acc: 74.03%, LR: 0.000073
Epoch [11/20], Loss: 0.7031, Train Acc: 75.69%, Val Acc: 75.97%, LR: 0.000066
Epoch [12/20], Loss: 0.6900, Train Acc: 76.21%, Val Acc: 76.70%, LR: 0.000058
Epoch [13/20], Loss: 0.6599, Train Acc: 78.63%, Val Acc: 75.49%, LR: 0.000051
Epoch [14/20], Loss: 0.6466, Train Acc: 80.54%, Val Acc: 77.91%, LR: 0.000043
Epoch [15/20], Loss: 0.6390, Train Acc: 80.19%, Val Acc: 78.88%, LR: 0.000035
```

```
Epoch [16/20], Loss: 0.6314, Train Acc: 81.03%, Val Acc: 79.85%, LR: 0.000028
Epoch [17/20], Loss: 0.6286, Train Acc: 81.16%, Val Acc: 79.37%, LR: 0.000021
Epoch [18/20], Loss: 0.6317, Train Acc: 80.01%, Val Acc: 81.07%, LR: 0.000015
Epoch [19/20], Loss: 0.6182, Train Acc: 82.27%, Val Acc: 79.85%, LR: 0.000010
Epoch [20/20], Loss: 0.6115, Train Acc: 82.54%, Val Acc: 81.31%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.6239, Train Acc: 81.81%, Val Acc: 80.58%, LR: 0.000010
Epoch [2/15], Loss: 0.5982, Train Acc: 82.86%, Val Acc: 82.77%, LR: 0.000010
Epoch [3/15], Loss: 0.5829, Train Acc: 84.10%, Val Acc: 84.22%, LR: 0.000010
Epoch [4/15], Loss: 0.5781, Train Acc: 84.29%, Val Acc: 80.58%, LR: 0.000010
Epoch [5/15], Loss: 0.5615, Train Acc: 85.93%, Val Acc: 85.92%, LR: 0.000010
Epoch [6/15], Loss: 0.5474, Train Acc: 86.58%, Val Acc: 88.59%, LR: 0.000010
Epoch [7/15], Loss: 0.5226, Train Acc: 87.98%, Val Acc: 87.86%, LR: 0.000010
Epoch [8/15], Loss: 0.5064, Train Acc: 89.46%, Val Acc: 87.62%, LR: 0.000010
Epoch [9/15], Loss: 0.5001, Train Acc: 90.14%, Val Acc: 88.59%, LR: 0.000010
Epoch [10/15], Loss: 0.4914, Train Acc: 90.62%, Val Acc: 87.86%, LR: 0.000005
Epoch [11/15], Loss: 0.4749, Train Acc: 91.70%, Val Acc: 90.05%, LR: 0.000005
Epoch [12/15], Loss: 0.4557, Train Acc: 92.54%, Val Acc: 90.05%, LR: 0.000005
Epoch [13/15], Loss: 0.4616, Train Acc: 92.24%, Val Acc: 92.23%, LR: 0.000005
Epoch [14/15], Loss: 0.4552, Train Acc: 93.05%, Val Acc: 91.50%, LR: 0.000005
Epoch [15/15], Loss: 0.4602, Train Acc: 92.29%, Val Acc: 91.99%, LR: 0.000005
Testing model...
```

```
==== split_80_20 Results ===
```

```
Accuracy: 0.9166
Precision: 0.9166
Recall: 0.9166
F1-Score: 0.9165
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.87	0.87	0.87	225
CN	0.94	0.92	0.93	288
MCI	0.92	0.94	0.93	518
accuracy			0.92	1031
macro avg	0.91	0.91	0.91	1031
weighted avg	0.92	0.92	0.92	1031

```
=====
```

```
Processing: split_90_10
```

```
=====
Train samples: 4175
Val samples: 464
Test samples: 515
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0574, Train Acc: 46.35%, Val Acc: 51.29%, LR: 0.000033
Epoch [2/20], Loss: 0.9790, Train Acc: 55.47%, Val Acc: 56.68%, LR: 0.000067
Epoch [3/20], Loss: 0.9216, Train Acc: 60.74%, Val Acc: 62.07%, LR: 0.000100
Epoch [4/20], Loss: 0.8730, Train Acc: 63.95%, Val Acc: 66.81%, LR: 0.000099
Epoch [5/20], Loss: 0.8264, Train Acc: 67.50%, Val Acc: 68.10%, LR: 0.000098
Epoch [6/20], Loss: 0.7824, Train Acc: 70.44%, Val Acc: 68.10%, LR: 0.000095
Epoch [7/20], Loss: 0.7588, Train Acc: 71.90%, Val Acc: 73.92%, LR: 0.000091
Epoch [8/20], Loss: 0.7268, Train Acc: 74.44%, Val Acc: 72.84%, LR: 0.000086
Epoch [9/20], Loss: 0.7088, Train Acc: 75.69%, Val Acc: 75.86%, LR: 0.000080
Epoch [10/20], Loss: 0.6920, Train Acc: 76.72%, Val Acc: 76.94%, LR: 0.000073
Epoch [11/20], Loss: 0.6703, Train Acc: 78.08%, Val Acc: 75.86%, LR: 0.000066
Epoch [12/20], Loss: 0.6656, Train Acc: 78.80%, Val Acc: 79.96%, LR: 0.000058
Epoch [13/20], Loss: 0.6386, Train Acc: 80.53%, Val Acc: 77.16%, LR: 0.000051
Epoch [14/20], Loss: 0.6336, Train Acc: 80.38%, Val Acc: 81.25%, LR: 0.000043
Epoch [15/20], Loss: 0.6372, Train Acc: 80.41%, Val Acc: 80.39%, LR: 0.000035
Epoch [16/20], Loss: 0.6110, Train Acc: 82.44%, Val Acc: 82.11%, LR: 0.000028
Epoch [17/20], Loss: 0.6039, Train Acc: 82.30%, Val Acc: 82.33%, LR: 0.000021
Epoch [18/20], Loss: 0.5964, Train Acc: 83.57%, Val Acc: 82.11%, LR: 0.000015
Epoch [19/20], Loss: 0.5953, Train Acc: 83.07%, Val Acc: 84.05%, LR: 0.000010
Epoch [20/20], Loss: 0.5896, Train Acc: 83.43%, Val Acc: 85.34%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/15], Loss: 0.6086, Train Acc: 82.63%, Val Acc: 79.96%, LR: 0.000010
Epoch [2/15], Loss: 0.5839, Train Acc: 83.98%, Val Acc: 84.27%, LR: 0.000010
Epoch [3/15], Loss: 0.5504, Train Acc: 86.66%, Val Acc: 84.91%, LR: 0.000010
Epoch [4/15], Loss: 0.5460, Train Acc: 86.66%, Val Acc: 86.42%, LR: 0.000010
Epoch [5/15], Loss: 0.5340, Train Acc: 87.33%, Val Acc: 86.85%, LR: 0.000010
Epoch [6/15], Loss: 0.5237, Train Acc: 87.86%, Val Acc: 88.15%, LR: 0.000010
Epoch [7/15], Loss: 0.5029, Train Acc: 89.53%, Val Acc: 89.66%, LR: 0.000010
Epoch [8/15], Loss: 0.4896, Train Acc: 90.97%, Val Acc: 90.73%, LR: 0.000010
Epoch [9/15], Loss: 0.4873, Train Acc: 90.99%, Val Acc: 87.28%, LR: 0.000010
Epoch [10/15], Loss: 0.4752, Train Acc: 91.62%, Val Acc: 89.87%, LR: 0.000010
Epoch [11/15], Loss: 0.4688, Train Acc: 91.88%, Val Acc: 91.59%, LR: 0.000010
Epoch [12/15], Loss: 0.4547, Train Acc: 92.41%, Val Acc: 92.89%, LR: 0.000010
Epoch [13/15], Loss: 0.4492, Train Acc: 93.20%, Val Acc: 92.24%, LR: 0.000010
Epoch [14/15], Loss: 0.4454, Train Acc: 93.51%, Val Acc: 93.32%, LR: 0.000010
Epoch [15/15], Loss: 0.4379, Train Acc: 93.75%, Val Acc: 92.89%, LR: 0.000010
Testing model...
```

```
==== split_90_10 Results ====
Accuracy: 0.9107
Precision: 0.9128
Recall: 0.9107
F1-Score: 0.9097
```

Classification Report:

	precision	recall	f1-score	support
AD	0.89	0.85	0.87	112
CN	0.96	0.84	0.90	144
MCI	0.90	0.98	0.94	259
accuracy			0.91	515
macro avg	0.92	0.89	0.90	515
weighted avg	0.91	0.91	0.91	515

DenseNet-121 - SUMMARY OF ALL SPLITS

split	accuracy	precision	recall	f1_score	training_time
split_10_90	0.626428	0.638367	0.626428	0.615734	88.385802
split_20_80	0.731749	0.728188	0.731749	0.726099	169.700626
split_30_70	0.807095	0.805800	0.807095	0.805617	248.054758
split_40_60	0.819534	0.826809	0.819534	0.820689	333.181419
split_50_50	0.871556	0.875118	0.871556	0.872503	414.224900
split_60_40	0.903492	0.904215	0.903492	0.902374	491.754236
split_70_30	0.906856	0.907481	0.906856	0.906768	562.409564
split_80_20	0.916586	0.916611	0.916586	0.916521	641.465663
split_90_10	0.910680	0.912797	0.910680	0.909669	720.117219

Detailed results saved to: /kaggle/working/densenet121_results.csv

```
[8]: #DenseNet-121
import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, models
import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import CosineAnnealingLR, ReduceLROnPlateau
```

```

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)

        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label

    def get_data_transforms():
        train_transform = transforms.Compose([
            transforms.Resize((224, 224)),
            transforms.RandomHorizontalFlip(p=0.5),
            transforms.RandomRotation(10),
            transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
            transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
            transforms.RandomGrayscale(p=0.1),
            transforms.GaussianBlur(kernel_size=3, sigma=(0.1, 2.0)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        ])

        val_transform = transforms.Compose([
            transforms.Resize((224, 224)),
            transforms.ToTensor(),

```

```

        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
˓→225])
    ])

    return train_transform, val_transform

def train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs, device, warmup_epochs=3):
    best_val_acc = 0
    patience = 8
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        # Learning rate warmup
        if epoch < warmup_epochs:
            lr_scale = min(1.0, float(epoch + 1) / warmup_epochs)
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['initial_lr'] * lr_scale

        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()

            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
            optimizer.step()

            running_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        epoch_loss = running_loss / len(train_loader)
        epoch_acc = 100 * correct / total

        val_acc = evaluate_model(model, val_loader, device)

        if scheduler and epoch >= warmup_epochs:
            if isinstance(scheduler, CosineAnnealingLR):

```

```

        scheduler.step()
    else:
        scheduler.step(val_acc)

    current_lr = optimizer.param_groups[0]['lr']
    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train_U
→Acc: {epoch_acc:.2f}%, Val Acc: {val_acc:.2f}%, LR: {current_lr:.6f}')
```

```

    if val_acc > best_val_acc:
        best_val_acc = val_acc
        patience_counter = 0
    else:
        patience_counter += 1

    if patience_counter >= patience:
        print("Early stopping at epoch {epoch+1}")
        break
```

```

def evaluate_model(model, data_loader, device):
    model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in data_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    return accuracy
```

```

def test_model(model, test_loader, device):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
```

```

    return all_preds, all_labels

def calculate_metrics(y_true, y_pred, split_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred, average='weighted', ▾
    ↪zero_division=0)
    recall = recall_score(y_true, y_pred, average='weighted', zero_division=0)
    f1 = f1_score(y_true, y_pred, average='weighted', zero_division=0)

    print(f"\n==== {split_name} Results ===")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-Score: {f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=['AD', 'CN', ▾
    ↪'MCI'], zero_division=0))

    return {
        'split': split_name,
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1_score': f1
    }

def run_densenet121_on_splits():
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Using device: {device}")

    splits_root = "/kaggle/working/alzheimer-resized-224_splits"
    split_folders = [f for f in os.listdir(splits_root) if f.
    ↪startswith('split_')]
    split_folders.sort()

    results = []

    train_transform, val_transform = get_data_transforms()

    for split_folder in split_folders:
        print(f"\n{'='*60}")
        print(f"Processing: {split_folder}")
        print(f"{'='*60}")

        split_path = os.path.join(splits_root, split_folder)

```

```

train_dataset = AlzheimerDataset(split_path, 'train', train_transform)
val_dataset = AlzheimerDataset(split_path, 'val', val_transform)
test_dataset = AlzheimerDataset(split_path, 'test', val_transform)

train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True, □
↪num_workers=2, pin_memory=True)
val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False, □
↪num_workers=2, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False, □
↪num_workers=2, pin_memory=True)

print(f"Train samples: {len(train_dataset)}")
print(f"Val samples: {len(val_dataset)}")
print(f"Test samples: {len(test_dataset)}")

model = models.densenet121(weights=models.DenseNet121_Weights. □
↪IMAGENET1K_V1)

# Freeze all layers initially
for param in model.parameters():
    param.requires_grad = False

# Unfreeze classifier
num_ftrs = model.classifier.in_features
model.classifier = nn.Linear(num_ftrs, 3)

# Unfreeze last dense block and transition layer
for param in model.features.denseblock4.parameters():
    param.requires_grad = True
for param in model.features.norm5.parameters():
    param.requires_grad = True

model = model.to(device)

# Layer-wise learning rates for DenseNet
optimizer = torch.optim.AdamW([
    {'params': model.classifier.parameters(), 'lr': 0.0001, □
↪'initial_lr': 0.0001},
    {'params': model.features.denseblock4.parameters(), 'lr': 0.00005, □
↪'initial_lr': 0.00005},
    {'params': model.features.norm5.parameters(), 'lr': 0.00005, □
↪'initial_lr': 0.00005},
    {'params': model.features.denseblock3.parameters(), 'lr': 0.00001, □
↪'initial_lr': 0.00001}
], weight_decay=1e-4)

```

```

criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
scheduler = CosineAnnealingLR(optimizer, T_max=20, eta_min=1e-6)

print("Starting phase 1 training (partial unfreeze)...")
start_time = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs=20, device=device)
phase1_time = time.time() - start_time

print("Starting phase 2 training (full fine-tuning)...")
# Unfreeze all parameters for final fine-tuning
for param in model.parameters():
    param.requires_grad = True

optimizer = torch.optim.AdamW(model.parameters(), lr=0.00001, weight_decay=1e-5)
scheduler = ReduceLROnPlateau(optimizer, mode='max', factor=0.5, patience=3, verbose=True)

start_time_phase2 = time.time()
train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs=15, device=device, warmup_epochs=0)
total_time = phase1_time + (time.time() - start_time_phase2)

print("Testing model...")
test_preds, test_labels = test_model(model, test_loader, device)

split_results = calculate_metrics(test_labels, test_preds, split_folder)
split_results['training_time'] = total_time
results.append(split_results)

torch.cuda.empty_cache()

results_df = pd.DataFrame(results)
print(f"\n{'='*80}")
print("DenseNet-121 - SUMMARY OF ALL SPLITS")
print(f"{'='*80}")
print(results_df.to_string(index=False))

results_csv_path = "/kaggle/working/densenet121_results.csv"
results_df.to_csv(results_csv_path, index=False)
print(f"\nDetailed results saved to: {results_csv_path}")

return results_df

if __name__ == "__main__":
    results = run_densenet121_on_splits()

```

```
Using device: cuda
```

```
=====
Processing: split_10_90
=====
```

```
Train samples: 464
```

```
Val samples: 51
```

```
Test samples: 4639
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/20], Loss: 1.0743, Train Acc: 43.97%, Val Acc: 45.10%, LR: 0.000033
```

```
Epoch [2/20], Loss: 1.0538, Train Acc: 48.71%, Val Acc: 49.02%, LR: 0.000067
```

```
Epoch [3/20], Loss: 1.0170, Train Acc: 49.57%, Val Acc: 54.90%, LR: 0.000100
```

```
Epoch [4/20], Loss: 0.9897, Train Acc: 53.45%, Val Acc: 58.82%, LR: 0.000099
```

```
Epoch [5/20], Loss: 0.9825, Train Acc: 54.74%, Val Acc: 60.78%, LR: 0.000098
```

```
Epoch [6/20], Loss: 0.9547, Train Acc: 57.33%, Val Acc: 56.86%, LR: 0.000095
```

```
Epoch [7/20], Loss: 0.9430, Train Acc: 57.54%, Val Acc: 60.78%, LR: 0.000091
```

```
Epoch [8/20], Loss: 0.9166, Train Acc: 62.72%, Val Acc: 62.75%, LR: 0.000086
```

```
Epoch [9/20], Loss: 0.8985, Train Acc: 63.58%, Val Acc: 58.82%, LR: 0.000080
```

```
Epoch [10/20], Loss: 0.8836, Train Acc: 64.44%, Val Acc: 60.78%, LR: 0.000073
```

```
Epoch [11/20], Loss: 0.8537, Train Acc: 66.38%, Val Acc: 60.78%, LR: 0.000066
```

```
Epoch [12/20], Loss: 0.8429, Train Acc: 67.24%, Val Acc: 60.78%, LR: 0.000058
```

```
Epoch [13/20], Loss: 0.8517, Train Acc: 64.87%, Val Acc: 60.78%, LR: 0.000051
```

```
Epoch [14/20], Loss: 0.8261, Train Acc: 69.61%, Val Acc: 60.78%, LR: 0.000043
```

```
Epoch [15/20], Loss: 0.8274, Train Acc: 69.18%, Val Acc: 64.71%, LR: 0.000035
```

```
Epoch [16/20], Loss: 0.8149, Train Acc: 69.83%, Val Acc: 62.75%, LR: 0.000028
```

```
Epoch [17/20], Loss: 0.8072, Train Acc: 70.69%, Val Acc: 66.67%, LR: 0.000021
```

```
Epoch [18/20], Loss: 0.8150, Train Acc: 68.75%, Val Acc: 62.75%, LR: 0.000015
```

```
Epoch [19/20], Loss: 0.7878, Train Acc: 73.71%, Val Acc: 64.71%, LR: 0.000010
```

```
Epoch [20/20], Loss: 0.7995, Train Acc: 71.34%, Val Acc: 66.67%, LR: 0.000006
```

```
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
```

```
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to  
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.7989, Train Acc: 70.26%, Val Acc: 64.71%, LR: 0.000010
```

```
Epoch [2/15], Loss: 0.7878, Train Acc: 72.20%, Val Acc: 66.67%, LR: 0.000010
```

```
Epoch [3/15], Loss: 0.7830, Train Acc: 72.20%, Val Acc: 64.71%, LR: 0.000010
```

```
Epoch [4/15], Loss: 0.7694, Train Acc: 74.14%, Val Acc: 66.67%, LR: 0.000010
```

```
Epoch [5/15], Loss: 0.7552, Train Acc: 75.86%, Val Acc: 64.71%, LR: 0.000010
```

```
Epoch [6/15], Loss: 0.7179, Train Acc: 78.02%, Val Acc: 66.67%, LR: 0.000005
```

```
Epoch [7/15], Loss: 0.7186, Train Acc: 78.45%, Val Acc: 64.71%, LR: 0.000005
```

```
Epoch [8/15], Loss: 0.7142, Train Acc: 76.51%, Val Acc: 64.71%, LR: 0.000005
```

```
Epoch [9/15], Loss: 0.7006, Train Acc: 79.74%, Val Acc: 62.75%, LR: 0.000005
```

```
Epoch [10/15], Loss: 0.6941, Train Acc: 77.80%, Val Acc: 62.75%, LR: 0.000003
```

```
Early stopping at epoch 10
```

```
Testing model...
```

```
==== split_10_90 Results ====
Accuracy: 0.5820
Precision: 0.5999
Recall: 0.5820
F1-Score: 0.5414
```

Classification Report:

	precision	recall	f1-score	support
AD	0.50	0.28	0.36	1012
CN	0.72	0.28	0.40	1296
MCI	0.58	0.88	0.70	2331
accuracy			0.58	4639
macro avg	0.60	0.48	0.49	4639
weighted avg	0.60	0.58	0.54	4639

```
=====
Processing: split_20_80
=====
```

Train samples: 928

Val samples: 103

Test samples: 4123

Starting phase 1 training (partial unfreeze)...

```
Epoch [1/20], Loss: 1.1661, Train Acc: 29.31%, Val Acc: 46.60%, LR: 0.000033
Epoch [2/20], Loss: 1.0541, Train Acc: 48.81%, Val Acc: 49.51%, LR: 0.000067
Epoch [3/20], Loss: 1.0195, Train Acc: 52.16%, Val Acc: 55.34%, LR: 0.000100
Epoch [4/20], Loss: 0.9949, Train Acc: 53.34%, Val Acc: 54.37%, LR: 0.000099
Epoch [5/20], Loss: 0.9750, Train Acc: 54.20%, Val Acc: 60.19%, LR: 0.000098
Epoch [6/20], Loss: 0.9554, Train Acc: 55.71%, Val Acc: 65.05%, LR: 0.000095
Epoch [7/20], Loss: 0.9225, Train Acc: 60.13%, Val Acc: 57.28%, LR: 0.000091
Epoch [8/20], Loss: 0.9095, Train Acc: 61.10%, Val Acc: 58.25%, LR: 0.000086
Epoch [9/20], Loss: 0.8998, Train Acc: 60.13%, Val Acc: 65.05%, LR: 0.000080
Epoch [10/20], Loss: 0.8857, Train Acc: 62.93%, Val Acc: 65.05%, LR: 0.000073
Epoch [11/20], Loss: 0.8636, Train Acc: 63.36%, Val Acc: 64.08%, LR: 0.000066
Epoch [12/20], Loss: 0.8444, Train Acc: 65.84%, Val Acc: 68.93%, LR: 0.000058
Epoch [13/20], Loss: 0.8476, Train Acc: 66.27%, Val Acc: 66.99%, LR: 0.000051
Epoch [14/20], Loss: 0.8270, Train Acc: 68.53%, Val Acc: 60.19%, LR: 0.000043
Epoch [15/20], Loss: 0.8186, Train Acc: 66.70%, Val Acc: 67.96%, LR: 0.000035
Epoch [16/20], Loss: 0.8111, Train Acc: 68.75%, Val Acc: 64.08%, LR: 0.000028
Epoch [17/20], Loss: 0.7918, Train Acc: 71.44%, Val Acc: 66.99%, LR: 0.000021
Epoch [18/20], Loss: 0.7931, Train Acc: 70.47%, Val Acc: 66.02%, LR: 0.000015
Epoch [19/20], Loss: 0.7957, Train Acc: 69.29%, Val Acc: 64.08%, LR: 0.000010
Epoch [20/20], Loss: 0.7771, Train Acc: 72.74%, Val Acc: 68.93%, LR: 0.000006
```

Early stopping at epoch 20

Starting phase 2 training (full fine-tuning)...

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.

    warnings.warn(
Epoch [1/15], Loss: 0.7744, Train Acc: 72.84%, Val Acc: 66.99%, LR: 0.000010
Epoch [2/15], Loss: 0.7830, Train Acc: 71.98%, Val Acc: 72.82%, LR: 0.000010
Epoch [3/15], Loss: 0.7452, Train Acc: 73.92%, Val Acc: 68.93%, LR: 0.000010
Epoch [4/15], Loss: 0.7505, Train Acc: 72.95%, Val Acc: 74.76%, LR: 0.000010
Epoch [5/15], Loss: 0.7365, Train Acc: 73.71%, Val Acc: 75.73%, LR: 0.000010
Epoch [6/15], Loss: 0.7232, Train Acc: 74.89%, Val Acc: 73.79%, LR: 0.000010
Epoch [7/15], Loss: 0.7134, Train Acc: 77.05%, Val Acc: 72.82%, LR: 0.000010
Epoch [8/15], Loss: 0.6932, Train Acc: 76.94%, Val Acc: 76.70%, LR: 0.000010
Epoch [9/15], Loss: 0.6786, Train Acc: 78.45%, Val Acc: 75.73%, LR: 0.000010
Epoch [10/15], Loss: 0.6602, Train Acc: 81.47%, Val Acc: 76.70%, LR: 0.000010
Epoch [11/15], Loss: 0.6616, Train Acc: 81.14%, Val Acc: 79.61%, LR: 0.000010
Epoch [12/15], Loss: 0.6180, Train Acc: 83.30%, Val Acc: 76.70%, LR: 0.000010
Epoch [13/15], Loss: 0.6222, Train Acc: 83.08%, Val Acc: 78.64%, LR: 0.000010
Epoch [14/15], Loss: 0.5980, Train Acc: 85.13%, Val Acc: 72.82%, LR: 0.000010
Epoch [15/15], Loss: 0.6094, Train Acc: 84.16%, Val Acc: 72.82%, LR: 0.000005
Testing model...

==== split_20_80 Results ===
Accuracy: 0.6949
Precision: 0.6943
Recall: 0.6949
F1-Score: 0.6910

Classification Report:
      precision    recall   f1-score   support
AD        0.57      0.55      0.56      899
CN        0.74      0.59      0.66     1152
MCI       0.72      0.82      0.77     2072
accuracy                           0.69      4123
macro avg       0.68      0.65      0.66      4123
weighted avg    0.69      0.69      0.69      4123

=====
Processing: split_30_70
=====
Train samples: 1391
Val samples: 155
Test samples: 3608
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0815, Train Acc: 43.49%, Val Acc: 49.03%, LR: 0.000033

```

```
Epoch [2/20], Loss: 1.0262, Train Acc: 53.13%, Val Acc: 53.55%, LR: 0.000067
Epoch [3/20], Loss: 0.9904, Train Acc: 54.06%, Val Acc: 54.84%, LR: 0.000100
Epoch [4/20], Loss: 0.9739, Train Acc: 55.00%, Val Acc: 58.71%, LR: 0.000099
Epoch [5/20], Loss: 0.9406, Train Acc: 57.66%, Val Acc: 58.06%, LR: 0.000098
Epoch [6/20], Loss: 0.9186, Train Acc: 61.61%, Val Acc: 60.00%, LR: 0.000095
Epoch [7/20], Loss: 0.8901, Train Acc: 61.75%, Val Acc: 62.58%, LR: 0.000091
Epoch [8/20], Loss: 0.8802, Train Acc: 63.26%, Val Acc: 63.87%, LR: 0.000086
Epoch [9/20], Loss: 0.8452, Train Acc: 66.00%, Val Acc: 64.52%, LR: 0.000080
Epoch [10/20], Loss: 0.8278, Train Acc: 67.58%, Val Acc: 69.03%, LR: 0.000073
Epoch [11/20], Loss: 0.8268, Train Acc: 67.36%, Val Acc: 70.97%, LR: 0.000066
Epoch [12/20], Loss: 0.8130, Train Acc: 68.73%, Val Acc: 65.81%, LR: 0.000058
Epoch [13/20], Loss: 0.8013, Train Acc: 68.51%, Val Acc: 67.10%, LR: 0.000051
Epoch [14/20], Loss: 0.7772, Train Acc: 70.17%, Val Acc: 67.10%, LR: 0.000043
Epoch [15/20], Loss: 0.7629, Train Acc: 71.89%, Val Acc: 63.87%, LR: 0.000035
Epoch [16/20], Loss: 0.7539, Train Acc: 73.11%, Val Acc: 65.81%, LR: 0.000028
Epoch [17/20], Loss: 0.7453, Train Acc: 73.40%, Val Acc: 63.23%, LR: 0.000021
Epoch [18/20], Loss: 0.7427, Train Acc: 73.62%, Val Acc: 67.74%, LR: 0.000015
Epoch [19/20], Loss: 0.7339, Train Acc: 73.69%, Val Acc: 66.45%, LR: 0.000010
Early stopping at epoch 19
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
warnings.warn(
Epoch [1/15], Loss: 0.7599, Train Acc: 73.11%, Val Acc: 67.10%, LR: 0.000010
Epoch [2/15], Loss: 0.7198, Train Acc: 75.13%, Val Acc: 68.39%, LR: 0.000010
Epoch [3/15], Loss: 0.7052, Train Acc: 75.70%, Val Acc: 68.39%, LR: 0.000010
Epoch [4/15], Loss: 0.6923, Train Acc: 77.35%, Val Acc: 69.68%, LR: 0.000010
Epoch [5/15], Loss: 0.6868, Train Acc: 77.86%, Val Acc: 72.26%, LR: 0.000010
Epoch [6/15], Loss: 0.6711, Train Acc: 79.51%, Val Acc: 67.74%, LR: 0.000010
Epoch [7/15], Loss: 0.6517, Train Acc: 81.24%, Val Acc: 67.10%, LR: 0.000010
Epoch [8/15], Loss: 0.6533, Train Acc: 79.65%, Val Acc: 73.55%, LR: 0.000010
Epoch [9/15], Loss: 0.6434, Train Acc: 80.52%, Val Acc: 68.39%, LR: 0.000010
Epoch [10/15], Loss: 0.6178, Train Acc: 82.60%, Val Acc: 71.61%, LR: 0.000010
Epoch [11/15], Loss: 0.6067, Train Acc: 83.11%, Val Acc: 74.19%, LR: 0.000010
Epoch [12/15], Loss: 0.5929, Train Acc: 84.47%, Val Acc: 73.55%, LR: 0.000010
Epoch [13/15], Loss: 0.5971, Train Acc: 82.96%, Val Acc: 75.48%, LR: 0.000010
Epoch [14/15], Loss: 0.5995, Train Acc: 82.60%, Val Acc: 74.19%, LR: 0.000010
Epoch [15/15], Loss: 0.5749, Train Acc: 85.12%, Val Acc: 74.84%, LR: 0.000010
Testing model...
```

```
==== split_30_70 Results ====
Accuracy: 0.7589
Precision: 0.7567
Recall: 0.7589
F1-Score: 0.7572
```

Classification Report:

	precision	recall	f1-score	support
AD	0.69	0.62	0.66	787
CN	0.73	0.74	0.74	1008
MCI	0.80	0.83	0.81	1813
accuracy			0.76	3608
macro avg	0.74	0.73	0.74	3608
weighted avg	0.76	0.76	0.76	3608

=====

Processing: split_40_60

=====

Train samples: 1855

Val samples: 207

Test samples: 3092

Starting phase 1 training (partial unfreeze)...

Epoch [1/20], Loss: 1.0465, Train Acc: 49.06%, Val Acc: 50.72%, LR: 0.000033
Epoch [2/20], Loss: 1.0159, Train Acc: 52.40%, Val Acc: 54.11%, LR: 0.000067
Epoch [3/20], Loss: 0.9818, Train Acc: 55.26%, Val Acc: 55.56%, LR: 0.000100
Epoch [4/20], Loss: 0.9465, Train Acc: 58.27%, Val Acc: 58.45%, LR: 0.000099
Epoch [5/20], Loss: 0.9204, Train Acc: 61.67%, Val Acc: 57.97%, LR: 0.000098
Epoch [6/20], Loss: 0.8865, Train Acc: 63.72%, Val Acc: 63.77%, LR: 0.000095
Epoch [7/20], Loss: 0.8647, Train Acc: 64.91%, Val Acc: 64.73%, LR: 0.000091
Epoch [8/20], Loss: 0.8320, Train Acc: 67.33%, Val Acc: 62.32%, LR: 0.000086
Epoch [9/20], Loss: 0.8188, Train Acc: 67.17%, Val Acc: 64.25%, LR: 0.000080
Epoch [10/20], Loss: 0.8056, Train Acc: 70.51%, Val Acc: 62.80%, LR: 0.000073
Epoch [11/20], Loss: 0.7829, Train Acc: 71.48%, Val Acc: 60.87%, LR: 0.000066
Epoch [12/20], Loss: 0.7635, Train Acc: 72.35%, Val Acc: 68.12%, LR: 0.000058
Epoch [13/20], Loss: 0.7400, Train Acc: 74.82%, Val Acc: 65.70%, LR: 0.000051
Epoch [14/20], Loss: 0.7355, Train Acc: 74.23%, Val Acc: 67.63%, LR: 0.000043
Epoch [15/20], Loss: 0.7291, Train Acc: 74.88%, Val Acc: 65.22%, LR: 0.000035
Epoch [16/20], Loss: 0.7127, Train Acc: 75.90%, Val Acc: 67.15%, LR: 0.000028
Epoch [17/20], Loss: 0.7118, Train Acc: 75.26%, Val Acc: 71.01%, LR: 0.000021
Epoch [18/20], Loss: 0.7052, Train Acc: 76.98%, Val Acc: 68.60%, LR: 0.000015
Epoch [19/20], Loss: 0.7066, Train Acc: 75.31%, Val Acc: 67.63%, LR: 0.000010
Epoch [20/20], Loss: 0.6898, Train Acc: 76.66%, Val Acc: 70.05%, LR: 0.000006

Starting phase 2 training (full fine-tuning)...

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.

 warnings.warn(

Epoch [1/15], Loss: 0.6922, Train Acc: 77.47%, Val Acc: 68.60%, LR: 0.000010
Epoch [2/15], Loss: 0.6847, Train Acc: 77.14%, Val Acc: 71.01%, LR: 0.000010
Epoch [3/15], Loss: 0.6681, Train Acc: 78.44%, Val Acc: 69.57%, LR: 0.000010

```
Epoch [4/15], Loss: 0.6678, Train Acc: 78.11%, Val Acc: 69.08%, LR: 0.000010
Epoch [5/15], Loss: 0.6351, Train Acc: 81.67%, Val Acc: 74.88%, LR: 0.000010
Epoch [6/15], Loss: 0.6343, Train Acc: 81.35%, Val Acc: 72.95%, LR: 0.000010
Epoch [7/15], Loss: 0.6118, Train Acc: 82.86%, Val Acc: 72.95%, LR: 0.000010
Epoch [8/15], Loss: 0.6041, Train Acc: 83.50%, Val Acc: 75.85%, LR: 0.000010
Epoch [9/15], Loss: 0.5956, Train Acc: 83.29%, Val Acc: 75.36%, LR: 0.000010
Epoch [10/15], Loss: 0.5676, Train Acc: 85.07%, Val Acc: 76.33%, LR: 0.000010
Epoch [11/15], Loss: 0.5663, Train Acc: 85.23%, Val Acc: 76.33%, LR: 0.000010
Epoch [12/15], Loss: 0.5450, Train Acc: 87.01%, Val Acc: 75.85%, LR: 0.000010
Epoch [13/15], Loss: 0.5445, Train Acc: 86.85%, Val Acc: 78.74%, LR: 0.000010
Epoch [14/15], Loss: 0.5397, Train Acc: 87.60%, Val Acc: 77.78%, LR: 0.000010
Epoch [15/15], Loss: 0.5275, Train Acc: 87.71%, Val Acc: 78.74%, LR: 0.000010
Testing model...
```

==== split_40_60 Results ===

Accuracy: 0.7791
Precision: 0.7799
Recall: 0.7791
F1-Score: 0.7740

Classification Report:

	precision	recall	f1-score	support
AD	0.80	0.61	0.69	674
CN	0.77	0.69	0.73	864
MCI	0.78	0.90	0.83	1554
accuracy			0.78	3092
macro avg	0.78	0.73	0.75	3092
weighted avg	0.78	0.78	0.77	3092

=====

Processing: split_50_50

=====

Train samples: 2319
Val samples: 258
Test samples: 2577

Starting phase 1 training (partial unfreeze)...

```
Epoch [1/20], Loss: 1.0504, Train Acc: 46.53%, Val Acc: 48.45%, LR: 0.000033
Epoch [2/20], Loss: 1.0066, Train Acc: 52.82%, Val Acc: 55.04%, LR: 0.000067
Epoch [3/20], Loss: 0.9709, Train Acc: 56.19%, Val Acc: 52.33%, LR: 0.000100
Epoch [4/20], Loss: 0.9278, Train Acc: 59.68%, Val Acc: 57.36%, LR: 0.000099
Epoch [5/20], Loss: 0.8943, Train Acc: 62.10%, Val Acc: 54.26%, LR: 0.000098
Epoch [6/20], Loss: 0.8647, Train Acc: 64.64%, Val Acc: 60.47%, LR: 0.000095
Epoch [7/20], Loss: 0.8509, Train Acc: 65.07%, Val Acc: 59.30%, LR: 0.000091
Epoch [8/20], Loss: 0.8309, Train Acc: 66.62%, Val Acc: 62.40%, LR: 0.000086
Epoch [9/20], Loss: 0.8144, Train Acc: 68.05%, Val Acc: 64.34%, LR: 0.000080
```

```
Epoch [10/20], Loss: 0.8034, Train Acc: 68.56%, Val Acc: 66.67%, LR: 0.000073
Epoch [11/20], Loss: 0.7776, Train Acc: 71.24%, Val Acc: 63.57%, LR: 0.000066
Epoch [12/20], Loss: 0.7515, Train Acc: 73.39%, Val Acc: 69.77%, LR: 0.000058
Epoch [13/20], Loss: 0.7579, Train Acc: 71.54%, Val Acc: 67.44%, LR: 0.000051
Epoch [14/20], Loss: 0.7325, Train Acc: 74.43%, Val Acc: 67.83%, LR: 0.000043
Epoch [15/20], Loss: 0.7173, Train Acc: 75.33%, Val Acc: 69.77%, LR: 0.000035
Epoch [16/20], Loss: 0.7226, Train Acc: 74.82%, Val Acc: 67.05%, LR: 0.000028
Epoch [17/20], Loss: 0.7209, Train Acc: 74.64%, Val Acc: 68.99%, LR: 0.000021
Epoch [18/20], Loss: 0.7087, Train Acc: 75.55%, Val Acc: 71.32%, LR: 0.000015
Epoch [19/20], Loss: 0.6890, Train Acc: 77.10%, Val Acc: 71.32%, LR: 0.000010
Epoch [20/20], Loss: 0.6951, Train Acc: 75.55%, Val Acc: 71.71%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.6993, Train Acc: 75.64%, Val Acc: 70.16%, LR: 0.000010
Epoch [2/15], Loss: 0.6647, Train Acc: 78.65%, Val Acc: 71.32%, LR: 0.000010
Epoch [3/15], Loss: 0.6515, Train Acc: 79.56%, Val Acc: 73.64%, LR: 0.000010
Epoch [4/15], Loss: 0.6437, Train Acc: 79.78%, Val Acc: 71.32%, LR: 0.000010
Epoch [5/15], Loss: 0.6396, Train Acc: 81.29%, Val Acc: 74.81%, LR: 0.000010
Epoch [6/15], Loss: 0.6082, Train Acc: 83.57%, Val Acc: 75.58%, LR: 0.000010
Epoch [7/15], Loss: 0.6049, Train Acc: 82.23%, Val Acc: 78.29%, LR: 0.000010
Epoch [8/15], Loss: 0.5906, Train Acc: 84.48%, Val Acc: 79.84%, LR: 0.000010
Epoch [9/15], Loss: 0.5713, Train Acc: 84.82%, Val Acc: 78.68%, LR: 0.000010
Epoch [10/15], Loss: 0.5581, Train Acc: 85.42%, Val Acc: 80.23%, LR: 0.000010
Epoch [11/15], Loss: 0.5435, Train Acc: 87.15%, Val Acc: 82.95%, LR: 0.000010
Epoch [12/15], Loss: 0.5404, Train Acc: 87.41%, Val Acc: 81.78%, LR: 0.000010
Epoch [13/15], Loss: 0.5387, Train Acc: 87.32%, Val Acc: 82.95%, LR: 0.000010
Epoch [14/15], Loss: 0.5334, Train Acc: 87.45%, Val Acc: 82.17%, LR: 0.000010
Epoch [15/15], Loss: 0.5237, Train Acc: 88.27%, Val Acc: 84.11%, LR: 0.000010
Testing model...
```

```
==== split_50_50 Results ===
```

```
Accuracy: 0.8184
Precision: 0.8176
Recall: 0.8184
F1-Score: 0.8168
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.80	0.69	0.74	562
CN	0.79	0.81	0.80	720
MCI	0.84	0.88	0.86	1295
accuracy		0.82		2577

macro avg	0.81	0.79	0.80	2577
weighted avg	0.82	0.82	0.82	2577

```
=====
Processing: split_60_40
=====

Train samples: 2784
Val samples: 308
Test samples: 2062
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0588, Train Acc: 47.13%, Val Acc: 51.62%, LR: 0.000033
Epoch [2/20], Loss: 1.0067, Train Acc: 52.44%, Val Acc: 52.60%, LR: 0.000067
Epoch [3/20], Loss: 0.9702, Train Acc: 56.43%, Val Acc: 52.92%, LR: 0.000100
Epoch [4/20], Loss: 0.9269, Train Acc: 59.66%, Val Acc: 59.42%, LR: 0.000099
Epoch [5/20], Loss: 0.8930, Train Acc: 61.85%, Val Acc: 58.12%, LR: 0.000098
Epoch [6/20], Loss: 0.8664, Train Acc: 63.90%, Val Acc: 62.99%, LR: 0.000095
Epoch [7/20], Loss: 0.8351, Train Acc: 66.81%, Val Acc: 61.04%, LR: 0.000091
Epoch [8/20], Loss: 0.8097, Train Acc: 68.32%, Val Acc: 63.96%, LR: 0.000086
Epoch [9/20], Loss: 0.7928, Train Acc: 69.97%, Val Acc: 64.61%, LR: 0.000080
Epoch [10/20], Loss: 0.7596, Train Acc: 70.87%, Val Acc: 63.96%, LR: 0.000073
Epoch [11/20], Loss: 0.7512, Train Acc: 72.49%, Val Acc: 67.21%, LR: 0.000066
Epoch [12/20], Loss: 0.7371, Train Acc: 73.85%, Val Acc: 67.21%, LR: 0.000058
Epoch [13/20], Loss: 0.7215, Train Acc: 75.14%, Val Acc: 66.56%, LR: 0.000051
Epoch [14/20], Loss: 0.7046, Train Acc: 75.90%, Val Acc: 67.86%, LR: 0.000043
Epoch [15/20], Loss: 0.6991, Train Acc: 76.29%, Val Acc: 69.81%, LR: 0.000035
Epoch [16/20], Loss: 0.6740, Train Acc: 78.45%, Val Acc: 69.81%, LR: 0.000028
Epoch [17/20], Loss: 0.6825, Train Acc: 76.69%, Val Acc: 71.10%, LR: 0.000021
Epoch [18/20], Loss: 0.6808, Train Acc: 77.12%, Val Acc: 70.13%, LR: 0.000015
Epoch [19/20], Loss: 0.6715, Train Acc: 77.77%, Val Acc: 71.43%, LR: 0.000010
Epoch [20/20], Loss: 0.6608, Train Acc: 79.06%, Val Acc: 72.73%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/15], Loss: 0.6749, Train Acc: 78.56%, Val Acc: 69.81%, LR: 0.000010
Epoch [2/15], Loss: 0.6590, Train Acc: 78.27%, Val Acc: 71.43%, LR: 0.000010
Epoch [3/15], Loss: 0.6376, Train Acc: 79.92%, Val Acc: 70.78%, LR: 0.000010
Epoch [4/15], Loss: 0.6136, Train Acc: 82.51%, Val Acc: 73.38%, LR: 0.000010
Epoch [5/15], Loss: 0.6050, Train Acc: 82.76%, Val Acc: 73.38%, LR: 0.000010
Epoch [6/15], Loss: 0.5847, Train Acc: 84.81%, Val Acc: 76.30%, LR: 0.000010
Epoch [7/15], Loss: 0.5904, Train Acc: 83.66%, Val Acc: 79.22%, LR: 0.000010
Epoch [8/15], Loss: 0.5762, Train Acc: 84.73%, Val Acc: 76.62%, LR: 0.000010
Epoch [9/15], Loss: 0.5534, Train Acc: 85.85%, Val Acc: 81.17%, LR: 0.000010
Epoch [10/15], Loss: 0.5486, Train Acc: 86.31%, Val Acc: 80.19%, LR: 0.000010
Epoch [11/15], Loss: 0.5342, Train Acc: 87.90%, Val Acc: 79.22%, LR: 0.000010
```

```
Epoch [12/15], Loss: 0.5138, Train Acc: 89.76%, Val Acc: 80.19%, LR: 0.000010
Epoch [13/15], Loss: 0.5190, Train Acc: 88.97%, Val Acc: 84.09%, LR: 0.000010
Epoch [14/15], Loss: 0.5004, Train Acc: 90.52%, Val Acc: 84.74%, LR: 0.000010
Epoch [15/15], Loss: 0.5046, Train Acc: 89.69%, Val Acc: 87.34%, LR: 0.000010
Testing model...
```

```
==== split_60_40 Results ===
Accuracy: 0.8657
Precision: 0.8658
Recall: 0.8657
F1-Score: 0.8646
```

Classification Report:

	precision	recall	f1-score	support
AD	0.85	0.79	0.82	450
CN	0.88	0.82	0.85	576
MCI	0.86	0.93	0.89	1036
accuracy			0.87	2062
macro avg	0.86	0.84	0.85	2062
weighted avg	0.87	0.87	0.86	2062

```
=====
Processing: split_70_30
=====
```

Train samples: 3247

Val samples: 361

Test samples: 1546

Starting phase 1 training (partial unfreeze)...

```
Epoch [1/20], Loss: 1.0714, Train Acc: 46.04%, Val Acc: 50.42%, LR: 0.000033
Epoch [2/20], Loss: 1.0015, Train Acc: 53.06%, Val Acc: 54.29%, LR: 0.000067
Epoch [3/20], Loss: 0.9632, Train Acc: 56.36%, Val Acc: 61.50%, LR: 0.000100
Epoch [4/20], Loss: 0.9169, Train Acc: 60.92%, Val Acc: 64.82%, LR: 0.000099
Epoch [5/20], Loss: 0.8870, Train Acc: 62.77%, Val Acc: 67.31%, LR: 0.000098
Epoch [6/20], Loss: 0.8434, Train Acc: 67.35%, Val Acc: 70.36%, LR: 0.000095
Epoch [7/20], Loss: 0.8163, Train Acc: 67.48%, Val Acc: 65.37%, LR: 0.000091
Epoch [8/20], Loss: 0.7938, Train Acc: 69.82%, Val Acc: 69.25%, LR: 0.000086
Epoch [9/20], Loss: 0.7573, Train Acc: 72.13%, Val Acc: 68.42%, LR: 0.000080
Epoch [10/20], Loss: 0.7399, Train Acc: 72.77%, Val Acc: 68.42%, LR: 0.000073
Epoch [11/20], Loss: 0.7327, Train Acc: 74.35%, Val Acc: 72.02%, LR: 0.000066
Epoch [12/20], Loss: 0.7078, Train Acc: 76.29%, Val Acc: 71.75%, LR: 0.000058
Epoch [13/20], Loss: 0.6998, Train Acc: 76.56%, Val Acc: 68.42%, LR: 0.000051
Epoch [14/20], Loss: 0.6928, Train Acc: 76.44%, Val Acc: 74.52%, LR: 0.000043
Epoch [15/20], Loss: 0.6935, Train Acc: 76.66%, Val Acc: 74.52%, LR: 0.000035
Epoch [16/20], Loss: 0.6674, Train Acc: 79.12%, Val Acc: 75.90%, LR: 0.000028
Epoch [17/20], Loss: 0.6784, Train Acc: 77.09%, Val Acc: 74.79%, LR: 0.000021
```

```
Epoch [18/20], Loss: 0.6527, Train Acc: 79.74%, Val Acc: 73.41%, LR: 0.000015
Epoch [19/20], Loss: 0.6602, Train Acc: 79.49%, Val Acc: 73.13%, LR: 0.000010
Epoch [20/20], Loss: 0.6565, Train Acc: 79.40%, Val Acc: 73.96%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/15], Loss: 0.6754, Train Acc: 77.98%, Val Acc: 74.79%, LR: 0.000010
Epoch [2/15], Loss: 0.6487, Train Acc: 79.15%, Val Acc: 76.18%, LR: 0.000010
Epoch [3/15], Loss: 0.6260, Train Acc: 81.24%, Val Acc: 78.12%, LR: 0.000010
Epoch [4/15], Loss: 0.6180, Train Acc: 81.92%, Val Acc: 79.50%, LR: 0.000010
Epoch [5/15], Loss: 0.5850, Train Acc: 83.95%, Val Acc: 81.44%, LR: 0.000010
Epoch [6/15], Loss: 0.5738, Train Acc: 85.06%, Val Acc: 79.78%, LR: 0.000010
Epoch [7/15], Loss: 0.5625, Train Acc: 85.59%, Val Acc: 82.55%, LR: 0.000010
Epoch [8/15], Loss: 0.5529, Train Acc: 86.08%, Val Acc: 83.38%, LR: 0.000010
Epoch [9/15], Loss: 0.5483, Train Acc: 86.70%, Val Acc: 82.83%, LR: 0.000010
Epoch [10/15], Loss: 0.5306, Train Acc: 88.08%, Val Acc: 82.83%, LR: 0.000010
Epoch [11/15], Loss: 0.5122, Train Acc: 89.28%, Val Acc: 86.15%, LR: 0.000010
Epoch [12/15], Loss: 0.5138, Train Acc: 88.82%, Val Acc: 86.43%, LR: 0.000010
Epoch [13/15], Loss: 0.4961, Train Acc: 89.87%, Val Acc: 87.26%, LR: 0.000010
Epoch [14/15], Loss: 0.5001, Train Acc: 89.31%, Val Acc: 87.53%, LR: 0.000010
Epoch [15/15], Loss: 0.4870, Train Acc: 91.01%, Val Acc: 90.58%, LR: 0.000010
Testing model...
```

```
==== split_70_30 Results ===
```

```
Accuracy: 0.8706
```

```
Precision: 0.8721
```

```
Recall: 0.8706
```

```
F1-Score: 0.8712
```

```
Classification Report:
```

	precision	recall	f1-score	support
AD	0.78	0.82	0.80	337
CN	0.86	0.87	0.86	432
MCI	0.92	0.89	0.91	777
accuracy			0.87	1546
macro avg	0.85	0.86	0.86	1546
weighted avg	0.87	0.87	0.87	1546

```
=====
Processing: split_80_20
=====
```

```
Train samples: 3711
```

```

Val samples: 412
Test samples: 1031
Starting phase 1 training (partial unfreeze)...
Epoch [1/20], Loss: 1.0402, Train Acc: 49.02%, Val Acc: 51.46%, LR: 0.000033
Epoch [2/20], Loss: 0.9893, Train Acc: 53.84%, Val Acc: 55.10%, LR: 0.000067
Epoch [3/20], Loss: 0.9490, Train Acc: 57.83%, Val Acc: 59.22%, LR: 0.000100
Epoch [4/20], Loss: 0.9028, Train Acc: 62.68%, Val Acc: 63.35%, LR: 0.000099
Epoch [5/20], Loss: 0.8606, Train Acc: 64.38%, Val Acc: 66.02%, LR: 0.000098
Epoch [6/20], Loss: 0.8253, Train Acc: 67.12%, Val Acc: 64.32%, LR: 0.000095
Epoch [7/20], Loss: 0.7969, Train Acc: 69.74%, Val Acc: 63.83%, LR: 0.000091
Epoch [8/20], Loss: 0.7780, Train Acc: 70.52%, Val Acc: 66.99%, LR: 0.000086
Epoch [9/20], Loss: 0.7658, Train Acc: 71.71%, Val Acc: 71.36%, LR: 0.000080
Epoch [10/20], Loss: 0.7432, Train Acc: 73.86%, Val Acc: 69.66%, LR: 0.000073
Epoch [11/20], Loss: 0.7200, Train Acc: 74.67%, Val Acc: 73.30%, LR: 0.000066
Epoch [12/20], Loss: 0.7037, Train Acc: 75.75%, Val Acc: 72.33%, LR: 0.000058
Epoch [13/20], Loss: 0.6967, Train Acc: 76.23%, Val Acc: 75.24%, LR: 0.000051
Epoch [14/20], Loss: 0.6806, Train Acc: 78.01%, Val Acc: 76.21%, LR: 0.000043
Epoch [15/20], Loss: 0.6654, Train Acc: 79.12%, Val Acc: 74.76%, LR: 0.000035
Epoch [16/20], Loss: 0.6700, Train Acc: 79.09%, Val Acc: 75.97%, LR: 0.000028
Epoch [17/20], Loss: 0.6528, Train Acc: 79.66%, Val Acc: 77.18%, LR: 0.000021
Epoch [18/20], Loss: 0.6507, Train Acc: 79.98%, Val Acc: 76.70%, LR: 0.000015
Epoch [19/20], Loss: 0.6329, Train Acc: 81.38%, Val Acc: 77.67%, LR: 0.000010
Epoch [20/20], Loss: 0.6392, Train Acc: 80.44%, Val Acc: 77.18%, LR: 0.000006
Starting phase 2 training (full fine-tuning)...
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to
access the learning rate.
    warnings.warn(
Epoch [1/15], Loss: 0.6487, Train Acc: 79.76%, Val Acc: 71.84%, LR: 0.000010
Epoch [2/15], Loss: 0.6375, Train Acc: 80.68%, Val Acc: 77.91%, LR: 0.000010
Epoch [3/15], Loss: 0.6113, Train Acc: 82.81%, Val Acc: 77.67%, LR: 0.000010
Epoch [4/15], Loss: 0.5991, Train Acc: 83.37%, Val Acc: 79.61%, LR: 0.000010
Epoch [5/15], Loss: 0.5686, Train Acc: 85.53%, Val Acc: 79.85%, LR: 0.000010
Epoch [6/15], Loss: 0.5631, Train Acc: 85.61%, Val Acc: 81.31%, LR: 0.000010
Epoch [7/15], Loss: 0.5366, Train Acc: 88.06%, Val Acc: 80.58%, LR: 0.000010
Epoch [8/15], Loss: 0.5396, Train Acc: 86.93%, Val Acc: 85.44%, LR: 0.000010
Epoch [9/15], Loss: 0.5231, Train Acc: 88.04%, Val Acc: 84.47%, LR: 0.000010
Epoch [10/15], Loss: 0.5043, Train Acc: 89.63%, Val Acc: 83.25%, LR: 0.000010
Epoch [11/15], Loss: 0.4964, Train Acc: 89.98%, Val Acc: 86.17%, LR: 0.000010
Epoch [12/15], Loss: 0.4958, Train Acc: 90.06%, Val Acc: 85.68%, LR: 0.000010
Epoch [13/15], Loss: 0.4798, Train Acc: 91.32%, Val Acc: 87.38%, LR: 0.000010
Epoch [14/15], Loss: 0.4739, Train Acc: 91.75%, Val Acc: 88.35%, LR: 0.000010
Epoch [15/15], Loss: 0.4748, Train Acc: 91.46%, Val Acc: 88.35%, LR: 0.000010
Testing model...

```

```

==== split_80_20 Results ====
Accuracy: 0.9108

```

```
Precision: 0.9122  
Recall: 0.9108  
F1-Score: 0.9106
```

Classification Report:

	precision	recall	f1-score	support
AD	0.85	0.86	0.86	225
CN	0.96	0.87	0.91	288
MCI	0.91	0.96	0.93	518
accuracy			0.91	1031
macro avg	0.91	0.89	0.90	1031
weighted avg	0.91	0.91	0.91	1031

```
=====
```

```
Processing: split_90_10
```

```
=====
```

```
Train samples: 4175
```

```
Val samples: 464
```

```
Test samples: 515
```

```
Starting phase 1 training (partial unfreeze)...
```

```
Epoch [1/20], Loss: 1.0352, Train Acc: 50.97%, Val Acc: 53.88%, LR: 0.000033  
Epoch [2/20], Loss: 0.9782, Train Acc: 54.23%, Val Acc: 59.27%, LR: 0.000067  
Epoch [3/20], Loss: 0.9329, Train Acc: 59.31%, Val Acc: 61.42%, LR: 0.000100  
Epoch [4/20], Loss: 0.8842, Train Acc: 62.80%, Val Acc: 61.64%, LR: 0.000099  
Epoch [5/20], Loss: 0.8483, Train Acc: 65.87%, Val Acc: 63.58%, LR: 0.000098  
Epoch [6/20], Loss: 0.8134, Train Acc: 68.50%, Val Acc: 65.73%, LR: 0.000095  
Epoch [7/20], Loss: 0.7847, Train Acc: 70.04%, Val Acc: 70.04%, LR: 0.000091  
Epoch [8/20], Loss: 0.7594, Train Acc: 72.22%, Val Acc: 71.98%, LR: 0.000086  
Epoch [9/20], Loss: 0.7370, Train Acc: 74.04%, Val Acc: 69.83%, LR: 0.000080  
Epoch [10/20], Loss: 0.7236, Train Acc: 74.32%, Val Acc: 72.20%, LR: 0.000073  
Epoch [11/20], Loss: 0.7016, Train Acc: 76.38%, Val Acc: 76.08%, LR: 0.000066  
Epoch [12/20], Loss: 0.6880, Train Acc: 76.69%, Val Acc: 72.41%, LR: 0.000058  
Epoch [13/20], Loss: 0.6760, Train Acc: 77.99%, Val Acc: 75.43%, LR: 0.000051  
Epoch [14/20], Loss: 0.6679, Train Acc: 78.73%, Val Acc: 74.57%, LR: 0.000043  
Epoch [15/20], Loss: 0.6626, Train Acc: 78.78%, Val Acc: 77.37%, LR: 0.000035  
Epoch [16/20], Loss: 0.6443, Train Acc: 79.64%, Val Acc: 76.08%, LR: 0.000028  
Epoch [17/20], Loss: 0.6409, Train Acc: 80.14%, Val Acc: 77.37%, LR: 0.000021  
Epoch [18/20], Loss: 0.6218, Train Acc: 82.42%, Val Acc: 80.17%, LR: 0.000015  
Epoch [19/20], Loss: 0.6308, Train Acc: 81.51%, Val Acc: 80.17%, LR: 0.000010  
Epoch [20/20], Loss: 0.6310, Train Acc: 81.05%, Val Acc: 78.66%, LR: 0.000006
```

```
Starting phase 2 training (full fine-tuning)...
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62:
```

```
UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to  
access the learning rate.
```

```

warnings.warn(
Epoch [1/15], Loss: 0.6338, Train Acc: 80.86%, Val Acc: 79.96%, LR: 0.000010
Epoch [2/15], Loss: 0.6126, Train Acc: 82.18%, Val Acc: 77.59%, LR: 0.000010
Epoch [3/15], Loss: 0.5917, Train Acc: 83.81%, Val Acc: 82.33%, LR: 0.000010
Epoch [4/15], Loss: 0.5748, Train Acc: 85.05%, Val Acc: 83.62%, LR: 0.000010
Epoch [5/15], Loss: 0.5602, Train Acc: 85.60%, Val Acc: 82.11%, LR: 0.000010
Epoch [6/15], Loss: 0.5351, Train Acc: 87.88%, Val Acc: 83.62%, LR: 0.000010
Epoch [7/15], Loss: 0.5376, Train Acc: 87.59%, Val Acc: 85.34%, LR: 0.000010
Epoch [8/15], Loss: 0.5165, Train Acc: 88.74%, Val Acc: 87.07%, LR: 0.000010
Epoch [9/15], Loss: 0.5101, Train Acc: 89.22%, Val Acc: 89.22%, LR: 0.000010
Epoch [10/15], Loss: 0.5036, Train Acc: 89.41%, Val Acc: 87.93%, LR: 0.000010
Epoch [11/15], Loss: 0.4931, Train Acc: 90.08%, Val Acc: 89.22%, LR: 0.000010
Epoch [12/15], Loss: 0.4762, Train Acc: 91.43%, Val Acc: 89.22%, LR: 0.000010
Epoch [13/15], Loss: 0.4777, Train Acc: 91.64%, Val Acc: 90.95%, LR: 0.000010
Epoch [14/15], Loss: 0.4621, Train Acc: 92.29%, Val Acc: 91.16%, LR: 0.000010
Epoch [15/15], Loss: 0.4552, Train Acc: 92.50%, Val Acc: 88.36%, LR: 0.000010
Testing model...

```

```

==== split_90_10 Results ===
Accuracy: 0.8854
Precision: 0.8857
Recall: 0.8854
F1-Score: 0.8849

```

Classification Report:

	precision	recall	f1-score	support
AD	0.82	0.83	0.82	112
CN	0.90	0.83	0.86	144
MCI	0.91	0.94	0.92	259
accuracy			0.89	515
macro avg	0.87	0.87	0.87	515
weighted avg	0.89	0.89	0.88	515

DenseNet-121 - SUMMARY OF ALL SPLITS

split	accuracy	precision	recall	f1_score	training_time
split_10_90	0.582022	0.599896	0.582022	0.541413	90.557603
split_20_80	0.694882	0.694307	0.694882	0.690960	204.277958
split_30_70	0.758869	0.756701	0.758869	0.757176	293.513171
split_40_60	0.779107	0.779911	0.779107	0.774039	398.096600
split_50_50	0.818393	0.817609	0.818393	0.816841	497.732667
split_60_40	0.865664	0.865824	0.865664	0.864625	585.891587
split_70_30	0.870634	0.872084	0.870634	0.871198	688.507176

```
split_80_20  0.910766  0.912176  0.910766  0.910570    782.058805
split_90_10  0.885437  0.885662  0.885437  0.884912    878.856004
```

Detailed results saved to: /kaggle/working/densenet121_results.csv

```
[9]: import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, models
import pandas as pd
import numpy as np
from PIL import Image
from sklearn.metrics import precision_score, recall_score, f1_score, □
    ↪accuracy_score, classification_report
import time
from torch.optim.lr_scheduler import CosineAnnealingLR, ReduceLROnPlateau

class AlzheimerDataset(Dataset):
    def __init__(self, split_dir, split_type, transform=None):
        self.split_dir = split_dir
        self.split_type = split_type
        self.transform = transform

        csv_path = os.path.join(split_dir, f"{split_type}.csv")
        self.df = pd.read_csv(csv_path)

        self.class_to_idx = {"AD": 0, "CN": 1, "MCI": 2}
        self.idx_to_class = {v: k for k, v in self.class_to_idx.items()}

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_path = row['path']
        label = self.class_to_idx[row['class']]

        image = Image.open(img_path).convert('RGB')

        if self.transform:
            image = self.transform(image)

        return image, label

    def get_data_transforms():
        train_transform = transforms.Compose([
```

```

        transforms.Resize((380, 380)),
        transforms.RandomHorizontalFlip(p=0.5),
        transforms.RandomRotation(10),
        transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1),
        transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)),
        transforms.RandomGrayscale(p=0.1),
        transforms.GaussianBlur(kernel_size=3, sigma=(0.1, 2.0)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
        transforms.RandomErasing(p=0.2, scale=(0.02, 0.2), ratio=(0.3, 3.3))
    ])

    val_transform = transforms.Compose([
        transforms.Resize((380, 380)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    return train_transform, val_transform

def train_model(model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs, device, warmup_epochs=3):
    best_val_acc = 0
    patience = 8
    patience_counter = 0

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0

        # Learning rate warmup
        if epoch < warmup_epochs:
            lr_scale = min(1.0, float(epoch + 1) / warmup_epochs)
            for param_group in optimizer.param_groups:
                param_group['lr'] = param_group['initial_lr'] * lr_scale

        for batch_idx, (images, labels) in enumerate(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)

```

```

        loss.backward()

        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        epoch_loss = running_loss / len(train_loader)
        epoch_acc = 100 * correct / total

        val_acc = evaluate_model(model, val_loader, device)

        if scheduler and epoch >= warmup_epochs:
            if isinstance(scheduler, CosineAnnealingLR):
                scheduler.step()
            else:
                scheduler.step(val_acc)

        current_lr = optimizer.param_groups[0]['lr']
        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train Acc: {train_acc:.4f}, Val Acc: {val_acc:.4f}')
    
```

```

File "/tmp/ipykernel_19/419255855.py", line 107
    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train Acc: {train_acc:.4f}, Val Acc: {val_acc:.4f}')
^
SyntaxError: unterminated string literal (detected at line 107)

```